




PERCONA[®]

Percona XtraBackup

Documentation

8.0.35–36 (June 04, 2026)

Table of Contents

[Home](#)

[Get help from Percona](#)

[Release notes](#)

[Release notes index](#)

[Percona XtraBackup 8.0.35-36 \(2026-06-04\)](#)

[Percona XtraBackup 8.0.35-35 \(2026-01-26\)](#)

[Percona XtraBackup 8.0.35-34 \(2025-08-11\)](#)

[Percona XtraBackup 8.0.35-33 \(2025-05-19\)](#)

[Percona XtraBackup 8.0.35-32 \(2025-01-09\)](#)

[Percona XtraBackup 8.0.35-31-Update \(2024-09-20\)](#)

[Percona XtraBackup 8.0.35-31 \(2024-06-20\)](#)

[Percona XtraBackup 8.0.35-30 \(2023-12-04\)](#)

[Percona XtraBackup 8.0.34-29 \(2023-08-21\)](#)

[Percona XtraBackup 8.0.33-28 \(2023-07-19\)](#)

[Percona XtraBackup 8.0.33-27 \(2023-05-25\)](#)

[Percona XtraBackup 8.0.32-26 \(2023-04-04\)](#)

[Percona XtraBackup 8.0.32-25 \(2023-02-27\)](#)

[Percona XtraBackup 8.0.31-24 \(2023-02-07\)](#)

[2022 \(8.0.30-23 - 8.0.23-16\)](#)

[Percona XtraBackup 8.0.30-23 \(2022-11-14\)](#)

[Percona XtraBackup 8.0.29-22 \(2022-07-19\)](#)

[Percona XtraBackup 8.0.28-21 \(2022-05-25\)](#)

[Percona XtraBackup 8.0.28-20](#)

[Percona XtraBackup 8.0.27-19](#)

[2021 \(8.0.26-18 - 8.0.23-16\)](#)

[Percona XtraBackup 8.0.26-18.0](#)

[Percona XtraBackup 8.0.25-17.0](#)

[Percona XtraBackup 8.0.23-16.0](#)

[2020 \(8.0.22-15 - 8.0.10\)](#)

[Percona XtraBackup 8.0.22-15.0](#)

[Percona XtraBackup 8.0.14](#)

[Percona XtraBackup 8.0.13](#)

[Percona XtraBackup 8.0.12](#)

[Percona XtraBackup 8.0.11](#)

[Percona XtraBackup 8.0.10](#)

[2019 \(8.0.9 - 8.0.5\)](#)

[Percona XtraBackup 8.0.9](#)

[Percona XtraBackup 8.0.8](#)

[Percona XtraBackup 8.0.7](#)

[Percona XtraBackup 8.0.6](#)

[Percona XtraBackup 8.0.5](#)

[2018 \(8.0.4 - 8.0.3-rc1\)](#)

[Percona XtraBackup 8.0.4](#)

[Percona XtraBackup 8.0-3-rc1](#)

[Quickstart Guide](#)

[Overview](#)

[Start a Docker container and take a backup](#)

[Restore the backup](#)

[Clean up](#)

[Next steps](#)

[How Percona XtraBackup works](#)

[About Percona XtraBackup](#)

[How Percona XtraBackup works](#)

[Percona XtraBackup Version Compatibility and Server Version Checks](#)

[Differences between a logical backup and Percona XtraBackup](#)

[Backup strategy and backup policy](#)

[Backup example](#)

[Features](#)

[Backup features](#)

[LRU dump backup](#)

[Throttling backups](#)

[Store backup history on the server](#)

[Take a streaming backup](#)

[Accelerate the backup process](#)

[Prepare features](#)

[Dictionary cache](#)

[Restore features](#)

[Point-in-time recovery](#)

[Encrypted InnoDB tablespace backups](#)

[FLUSH TABLES WITH READ LOCK option](#)

[Improved log statements](#)

[lock-ddl-per-table option improvements](#)

[InnoDB redo log archiving](#)

[Smart memory estimation](#)

[Work with binary logs](#)

[Analyze table statistics](#)

[xbcloud features](#)

[FIFO data sink](#)

[Exponential backoff](#)

[Binaries](#)

[Overview](#)

[The xtrabackup binary](#)

[The xtrabackup binary overview](#)

[xtrabackup implementation details](#)

[The xtrabackup command-line options](#)

[Configure xtrabackup](#)

[The xbcrypt binary](#)

[The xbcrypt binary overview](#)

[The xbcrypt command-line options](#)

[The xbstream binary](#)

[The xbstream binary overview](#)

[The xbstream command-line options](#)

[The xbcloud binary](#)

[The xbcloud binary overview](#)

[The xbcloud command-line options](#)

[Use the xbcloud binary](#)

[Use the xbcloud binary with Amazon S3](#)

[Use the xbccloud binary with an IAM instance profile](#)

[Use the xbccloud binary with Swift](#)

[Use the xbccloud binary with Google Cloud Storage](#)

[Use the xbccloud binary with Microsoft Azure Cloud Storage](#)

[Use the xbccloud binary with MinIO](#)

[Dependency compatibility and limitations](#)

[Install](#)

[Install Percona XtraBackup 8.0 overview](#)

[Before you start](#)

[Connection and privileges needed](#)

[Permissions needed](#)

[Percona Software Download instructions](#)

[Use APT](#)

[Percona release and APT](#)

[Files in DEB package](#)

[Download DEB packages](#)

[Pin packages](#)

[Work with AppArmor](#)

[Uninstall](#)

[Use YUM](#)

[Percona release and YUM](#)

[Files in RPM package](#)

[Download RPM packages](#)

[Work with SELinux](#)

[Uninstall](#)

[Use binary tarballs](#)

[Install with binary tarballs](#)

[Binary tarballs available](#)

[Compile from source](#)

[Compile and install Percona XtraBackup 8.0 from source](#)

[Docker](#)

[Run in a Docker container](#)

[How to](#)

[Take a full backup](#)

[Create a full backup](#)

[Prepare a full backup](#)

[Take an incremental backup](#)

[Create an incremental backup](#)

[Prepare an incremental backup](#)

[Take an incremental backup using page tracking](#)

[Take a compressed backup](#)

[Create a compressed backup](#)

[Decompress and prepare a backup](#)

[Take an encrypted backup](#)

[Encrypt backups](#)

[Take a partial backup](#)

[Create a partial backup](#)

[Prepare a partial backup](#)

[Back up individual partitions](#)

[Create an individual partitions backup](#)

[Prepare an individual partitions backup](#)

[Restore a backup](#)

[Restore full, incremental, compressed backups](#)

[Restore a partial backup](#)

[Restore an individual partitions backup](#)

[Restore individual tables](#)

[Replicate](#)

[How to set up a replica for replication in 6 simple steps with Percona XtraBackup](#)

[How to create a new \(or repair a broken\) GTID-based replica](#)

[Make backups in replication environments](#)

[Verify backups with replication and pt-checksum](#)

[Troubleshoot](#)

[Error Message: Found tables with row versions due to INSTANT ADD/DROP columns](#)

[xtrabackup exit codes](#)

[Reference](#)

[Topic index](#)

[Index of files created by Percona XtraBackup](#)

[Understand version numbers](#)

[Frequently asked questions](#)

[Glossary](#)

[Dump schema for all tables](#)

[Supported versions](#)

[Percona Toolkit version checking](#)

[Trademark policy](#)

[Copyright and licensing information](#)

[Topic Index](#)

Percona XtraBackup – Documentation

This documentation is for the latest release: [Percona XtraBackup 8.0.35-36](#) release notes.

Percona XtraBackup is an open source hot backup utility for MySQL-based servers, that keeps your database fully available during planned maintenance windows.

Whether it is a 24x7 highly loaded server or a low-transaction-volume Percona XtraBackup is designed to make backups seamless without disrupting the performance of the server in a production environment. Percona XtraBackup (PXB) is a 100% open source backup solution with [commercial support](#) available for organizations who want to benefit from comprehensive, responsive, and cost-flexible database support for MySQL.

Supported versions

Percona XtraBackup 8.0.35-36 can take backups of MySQL-compatible databases running earlier releases within the same major version, if specific conditions are met. [See the supported versions](#) for more details.

For more information, see [Percona XtraBackup Version Compatibility and Server Version Checks](#).

Taking your backup with Percona XtraBackup is easy. Follow our documentation guides, and you'll be set up in quickly.

Quickstart guide

Get started quickly with our Quickstart guide.

[Quickstart guide](#)

Installation guides

Find the best installation solution with our step-by-step installation instructions.

[Installation instructions](#)

Binaries

Learn about the Percona XtraBackup binaries: xtrabackup, xbcloud, xbcrypt, and xbstream.

Backup management

Learn about the different types of backups and how to take them.

Supported storage engines

Percona XtraBackup can back up data from *InnoDB*, *XtraDB*, *MyISAM*, and *MyRocks* tables on *MySQL* 8.0 servers as well as *Percona Server for MySQL* with *XtraDB*, [Percona Server for MySQL 8.0](#), and [Percona XtraDB Cluster 8.0](#).

Version updates

Version 8.0.6 and later supports the *MyRocks* storage engine. An incremental backup on the *MyRocks* storage engine does not determine if an earlier full or incremental backup contains the same files. **Percona XtraBackup** copies all *MyRocks* files each time it takes a backup. **Percona XtraBackup** does not support the *TokuDB* storage engine. Find more information in [Percona TokuBackup](#).

Limitations

Percona XtraBackup 8.0 does not support making backups of databases created in versions prior to 8.0 of *MySQL*, *Percona Server for MySQL* or *Percona XtraDB Cluster*. As the changes that *MySQL* 8.0 introduced in *data dictionaries*, *redo log* and *undo log* are incompatible with previous versions, it is currently impossible for *Percona XtraBackup* 8.0 to also support versions prior to 8.0.

Due to changes in *MySQL* 8.0.20 released by Oracle at the end of April 2020, *Percona XtraBackup* 8.0, up to version 8.0.11, is not compatible with *MySQL* version 8.0.20 or higher, or *Percona* products that are based on it: *Percona Server for MySQL* and *Percona XtraDB Cluster*.

For more information, see [Percona XtraBackup 8.x and MySQL 8.0.20](#)

See also

[Dependency compatibility and limitations](#)

Learn more about other Percona products

[Percona Distribution for MySQL 8.0](#)

Get help from Percona

Our documentation guides are packed with information, but they can't cover everything you need to know about Percona XtraBackup. They also won't cover every scenario you might come across. Don't be afraid to try things out and ask questions when you get stuck.

Percona's Community Forum

Be a part of a space where you can tap into a wealth of knowledge from other database enthusiasts and experts who work with Percona's software every day. While our service is entirely free, keep in mind that response times can vary depending on the complexity of the question. You are engaging with people who genuinely love solving database challenges.

We recommend visiting our [Community Forum](#). It's an excellent place for discussions, technical insights, and support around Percona database software. If you're new and feeling a bit unsure, our [FAQ](#) and [Guide for New Users](#) ease you in.

If you have thoughts, feedback, or ideas, the community team would like to hear from you at [Any ideas on how to make the forum better?](#). We're always excited to connect and improve everyone's experience.

Percona experts

Percona experts bring years of experience in tackling tough database performance issues and design challenges.

Talk to a Percona Expert

We understand your challenges when managing complex database environments. That's why we offer various services to help you simplify your operations and achieve your goals.

Service	Description
24/7 Expert Support	Our dedicated team of database experts is available 24/7 to assist you with any database issues. We provide flexible support plans tailored to your specific needs.
Hands-On Database Management	Our managed services team can take over the day-to-day management of your database infrastructure, freeing up your time to focus on other priorities.

Service	Description
Expert Consulting	Our experienced consultants provide guidance on database topics like architecture design, migration planning, performance optimization, and security best practices.
Comprehensive Training	Our training programs help your team develop skills to manage databases effectively, offering virtual and in-person courses.

We're here to help you every step of the way. Whether you need a quick fix or a long-term partnership, we're ready to provide your expertise and support.

Release notes

Percona XtraBackup 8.0 release notes index

- [Percona XtraBackup 8.0.35-36 \(2026-06-04\)](#).
- [Percona XtraBackup 8.0.35-35 \(2026-01-26\)](#).
- [Percona XtraBackup 8.0.35-34 \(2025-08-11\)](#).
- [Percona XtraBackup 8.0.35-33 \(2025-05-19\)](#).
- [Percona XtraBackup 8.0.35-32 \(2025-01-09\)](#).
- [Percona XtraBackup 8.0.35-31 Update \(2024-09-20\)](#).
- [Percona XtraBackup 8.0.35-31 \(2024-06-20\)](#).
- [Percona XtraBackup 8.0.35-30 \(2023-12-04\)](#).
- [Percona XtraBackup 8.0.34-29 \(2023-08-21\)](#).
- [Percona XtraBackup 8.0.33-28 \(2023-07-19\)](#).
- [Percona XtraBackup 8.0.33-27 \(2023-05-25\)](#).
- [Percona XtraBackup 8.0.32-26 \(2023-04-04\)](#).
- [Percona XtraBackup 8.0.32-25 \(2023-02-27\)](#).
- [Percona XtraBackup 8.0.31-24 \(2023-02-07\)](#).
- [Percona XtraBackup 8.0.30-23 \(2022-11-14\)](#).
- [Percona XtraBackup 8.0.29-22 \(2022-07-19\)](#).
- [Percona XtraBackup 8.0.28-21 \(2022-05-25\)](#).
- [Percona XtraBackup 8.0.28-20.0 \(2022-04-26\)](#).
- [Percona Xtrabackup 8.0.27-19.0 \(2022-02-02\)](#).
- [Percona XtraBackup 8.0.26-18.0 \(2021-09-02\)](#).
- [Percona Xtrabackup 8.0.25-17.0 \(2021-06-03\)](#).
- [Percona XtraBackup 8.0.23-16.0 \(2021-03-22\)](#).
- [Percona Xtrabackup 8.0.22-15.0 \(2020-12-14\)](#).
- [Percona XtraBackup 8.0.14 \(2020-08-31\)](#).
- [Percona Xtrabackup 8.0.13 \(2020-06-17\)](#).
- [Percona XtraBackup 8.0.12 \(2020-05-27\)](#).
- [Percona Xtrabackup 8.0.11 \(2020-04-13\)](#).
- [Percona XtraBackup 8.0.10 \(2020-03-16\)](#).

- [Percona Xtrabackup 8.0.9 \(2019-12-16\)](#)
- [Percona XtraBackup 8.0.8 \(2019-11-21\)](#)
- [Percona Xtrabackup 8.0.7 \(2019-08-07\)](#)
- [Percona Xtrabackup 8.0.6 \(2019-05-09\)](#)
- [Percona XtraBackup 8.0.5 \(2019-03-04\)](#)
- [Percona Xtrabackup 8.0.4 \(2018-12-10\)](#)
- [Percona XtraBackup 8.0-3-rc1 \(2018-10-31\)](#)

Percona XtraBackup 8.0.35-36 (2026-06-04)

Important

Percona XtraBackup 8.0.35-36 is the final release of the Percona XtraBackup 8.0 series.

We recommend that you either upgrade to MySQL 8.4 or stay on 8.0; we'll support you.

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Starting with Percona XtraBackup 8.0.34, you can take backups of Percona Server for MySQL and MySQL version 8.0.34 or later within the 8.0.x series.

Check the latest list of all [Supported versions](#).

Release highlights

Percona XtraBackup 8.0.35-36 release includes improvements and bug fixes.

- Percona XtraBackup adds support for Azure Storage accounts with [Hierarchical Namespace \(HNS\)](#) enabled in `xbcloud delete`. Percona XtraBackup now correctly removes backups from Azure Data Lake Storage Gen2 environments where deletion previously failed due to unsupported HNS behavior.
- Percona XtraBackup improves sparse file handling efficiency by removing redundant `fallocate(PUNCH_HOLE)` calls from `local_write_sparse`.

We recommend that you download the Percona XtraBackup for the same architecture as the MySQL-compatible server. For example, if your server is on an ARM64 platform, you should download and use the Percona XtraBackup for ARM64 for that operating system.

Improvements

- [PXB-3643](#): Adds support for Azure Storage accounts with [Hierarchical Namespace \(HNS\)](#) enabled in `xbcloud delete`. Percona XtraBackup now correctly removes backups from Azure Data Lake Storage Gen2 environments where deletion previously failed due to unsupported HNS behavior.

- [PXB-3658](#): Improves sparse file handling efficiency by removing redundant `fallocate(PUNCH_HOLE)` calls from `local_write_sparse`.

Bug fixes

- [PXB-3568](#): Fixed an issue where `xtrabackup` terminated unexpectedly on RHEL 9 when using `--compress=lz4`.
- [PXB-3759](#): Fixed an issue where `xtrabackup` terminated unexpectedly when redo log encryption was enabled and `innodb_log_buffer_size` exceeded 16 MB.

Packaging notes

- Percona XtraBackup 8.0 adds support for Ubuntu 26.04.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#)

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.35-35 (2026-01-26)

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Starting with Percona XtraBackup 8.0.34, you can take backups of Percona Server for MySQL and MySQL version 8.0.34 or later within the 8.0.x series.

Check the latest list of all [Supported versions](#).

Release highlights

Announcement – Discontinuation of the Percona PRO Program

Percona has consolidated its build offerings. The Pro builds are no longer required. All features from the Percona XtraBackup 8.0.x Pro build series have been merged into the main open source Community release. Percona XtraBackup 8.0.35-35 release build is available for Amazon Linux 2023 (AL2023). We support both AMD64 and ARM64 versions of Amazon Linux 2023.

We recommend that you download the Percona XtraBackup for the same architecture as the MySQL-compatible server. For example, if your server is on an ARM64 platform, you should download and use the Percona XtraBackup for ARM64 for that operating system.

Improvements

- [PXB-3458](#): Improves restore reliability by preventing failures caused by invalid `innodb_undo_directory` in the `--copy-back` mode.
- [PXB-2724](#): Adds `--move-back` support for the RocksDB storage engine.

Bug fixes

- [PXB-3543](#): Fixed a performance degradation that occurred during incremental backups when the base backup contained 10,000 or more unchanged `.ibd` files.

- [PXB-3624](#): Fixed an issue where `xbcloud` could exit because `libcurl` attempted to use a `libev` event loop after it had already been shut down. While `xbcloud` threw a warning with `libcurl` versions `8.11.x` and potentially exited due to [CVE-2025-0665](#), this behavior was resolved in `libcurl` 8.12.0 and later.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#)

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.35-34 (2025-08-11)

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Starting with Percona XtraBackup 8.0.34, you can take backups of Percona Server for MySQL 8.0.34, MySQL 8.0.34, and **higher versions within the 8.0 series**.

Check the latest list of all [Supported versions](#).

Release highlights

This release provides a bug fix.

We recommend that you download the Percona XtraBackup for the same platform as the MySQL-compatible server. For example, if your server is on an ARM64 platform, you should download and use the Percona XtraBackup for ARM64 for that operating system.

Bug fix

- [PXB-3470](#): When streaming a backup that has been compressed (using `lz4` or `zstd` with the `--compress` option), the decompression process, whether performed by `xbstream --decompress` or `xtrabackup --decompress`, requires a significant amount of time.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#).

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.35-33 (2025-05-19)

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Starting with Percona XtraBackup 8.0.34, you can take backups of Percona Server for MySQL 8.0.34, MySQL 8.0.34, and **higher versions within the 8.0 series**.

Check the latest list of all [Supported versions](#).

Release highlights

This release provides improvements and bug fixes.

We recommend that you download the Percona XtraBackup for the same platform as the MySQL-compatible server. For example, if your server is on an ARM64 platform, you should download and use the Percona XtraBackup for ARM64 for that operating system.

Improvements

- [PXB-3427](#): Percona XtraBackup now prepares incremental backups faster. The `--prepare` command directly applies the `.delta` files. To speed up this process, use the `--parallel=X` option, replacing `X` with the number of threads you want to use simultaneously. This option applies the delta files concurrently. For more information, see [Prepare an incremental backup](#).
- [PXB-3199](#): The `xbcloud put` operations were updated to include support for [ObjectLock-enabled AWS S3 buckets](#) (Thanks to volver for contributing the fix for this issue.).

Bug fix

- [PXB-3426](#): Fix xtrabackup crashes caused by Key Management Interoperability Protocol (KMIP) component keyring usage.
- [PXB-3490](#): Fixed an invalid URL for the Boost tarball.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#).

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.35-32 (2025-01-09)

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Starting with Percona XtraBackup 8.0.34, you can take backups of Percona Server for MySQL 8.0.34, MySQL 8.0.34, and **higher versions within the 8.0 series**.

Check the latest list of all [Supported versions](#).

Release highlights

This release provides bug fixes.

We recommend that you download the Percona XtraBackup for the same platform as the MySQL-compatible server. For example, if your server is on an ARM64 platform, you should download and use the Percona XtraBackup for ARM64 for that operating system.

Bug fix

- [PXB-3302](#): When a MySQL instance contained a large number of GTID sets, the “GTID of the last change” output in the Xtrabackup log was truncated compared to the complete output found in `xtrabackup_binlog_info` and `xtrabackup_info`. This truncation created problems for external tools that needed to get GTID coordinates from the log.
- [PXB-3336](#): A typographical error existed in the Full Text Search index pattern.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#)

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.35-31-Update (2024-09-20)

Packaging notes

This update to the release of Percona XtraBackup includes the addition of `arm64.deb` packages in [Percona Software Downloads](#) for the following platforms:

- Debian 12
- Debian 11
- Ubuntu 24.04
- Ubuntu 22.04
- Ubuntu 20.04

Percona XtraBackup 8.0.35–31 (2024-06-20)

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Starting with Percona XtraBackup 8.0.34, you can take backups of Percona Server for MySQL 8.0.34, MySQL 8.0.34, and **higher versions within the 8.0 series**.

Check the latest list of all [Supported versions](#).

Release highlights

This release adds Percona XtraBackup ARM64 packages and provides bug fixes.

We recommend that you download the Percona XtraBackup for the same platform as the MySQL-compatible server. For example, if your server is on an ARM64 platform, you should download and use the Percona XtraBackup for ARM64 for that operating system.

Bug fix

- [PXB-3283](#): If backup contains tables with [temporal datatypes](#) and their charset/collation changed (ALTER TABLE CONVERT CHARACTER SET) before the backup is taken, xtrabackup export couldn't generate the correct configuration (.cfg) files required by the import.
- [PXB-2797](#): When you import a table with a full-text index from an XtraBackup backup, you might encounter a schema mismatch. This happens when the table's structure in the backup does not match the structure expected by the database.
- [PXB-3210](#): Percona XtraBackup version 8.0.33-28 and later fail to build on MacOS because of the FIND_PROCP() function. This function causes issues during the build process, preventing the software from compiling successfully.
- [PXB-3130](#): When preparing 8.0.30 backups with newer xtrabackup versions, the prepare step failed. (Thank you to LiNing for the contribution!)

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#).

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.35-30 (2023-12-04)

Get started with [Quickstart Guide for Percona XtraBackup](#).

Supported versions

Check the latest list of all [Supported versions](#).

Release highlights

This release merges the MySQL 8.0.35 code base.

Bug fix

- [PXB-3003](#): Percona XtraBackup discovering redo logs to parse and the server purging redo logs simultaneously could cause a race condition.
- [PXB-3139](#): Fix for gcc-13 and clang-16 compilation issues.
- [PXB-3147](#): Percona XtraBackup failed to execute the `DO innodb_redo_log_consumer_register("PXB"); query`.
- [PXB-3181](#): The use of the `innodb-use-native-aio=true` option resulted in printing duplicate timestamps.
- [PXB-3173](#): The xbcloud was refactored to use multiple threads with a global list of files. This list did not allow users to download a subset of files.
- [PXB-3141](#): Adjusted folder creation on `ds_fifo.cc` and xbcloud to use proper permission.
- [PXB-3168](#): Under high write load, a backup could fail with a `log block numbers mismatch` error.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub repository](#).

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

Percona XtraBackup 8.0.34-29 (2023-08-21)

Get started with [Quickstart Guide for Percona XtraBackup](#).

[Percona XtraBackup \(PXB\)](#) is a 100% open source backup solution for all versions of Percona Server for MySQL and MySQL® that performs online non-blocking, tightly compressed, highly secure full backups on transactional systems. Maintain fully available applications during planned maintenance windows with Percona XtraBackup.

Release highlights

Starting with Percona XtraBackup 8.0.34 you can take backups of Percona Server for MySQL 8.0.34, MySQL 8.0.34, and **higher versions within the 8.0 series**.

Check the latest list of all [Supported versions](#).

Compression

The [ZSTD compression algorithm](#) is moved to [General Availability](#) and becomes the default compression algorithm for the `--compress` option.

The `QuickLZ` compression algorithm is removed and is not used for `--backup`. The `--decompress` option still uses `qpress` for backward compatibility to decompress backups taken by older versions of Percona XtraBackup.

New features

- [PXB-3059](#): Adds the `--redo-log-arch-dir=name` option to set the redo log archive destination directory. (Thanks to wangbinmss for providing the patch)

Improvements

- [PXB-3007](#): The [ZSTD compression algorithm](#) is moved to [General Availability](#) and becomes the default compression algorithm for the `--compress` option.

Bug fix

- [PXB-3073](#): An unresponsive host may cause `xbcloud` to become stuck. Added a `--timeout` option to `xbcloud`.
- [PXB-3088](#): Percona XtraBackup could not prepare the backup if the redo log size was more than 128G. (Thanks to Ahmed Et-tanany for providing the fix for this issue)
- [PXB-3079](#): The prepare phase skipped rollback on encrypted tables and was completed successfully if the keyring plugin was not loaded.
- [PXB-2860](#): Percona XtraBackup locked table even when the `--tables-exclude` and `--lock-ddl-per-table` options were used.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub location](#)

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

[Contribute to the documentation](#)

For [training](#), contact [Percona Training - Start learning now](#).

Percona XtraBackup 8.0.33-28 (2023-07-19)

[Percona XtraBackup \(PXB\)](#) is a 100% open source backup solution for all versions of Percona Server for MySQL and MySQL® that performs online non-blocking, tightly compressed, highly secure full backups on transactional systems. Maintain fully available applications during planned maintenance windows with Percona XtraBackup.

Release highlights

This release includes improvements, features, and bug fixes. Percona has implemented a two-stage release process for each version. The first release primarily ensures compatibility with the latest MySQL version, to help those customers who need an updated version of Percona XtraBackup as soon as possible. The second release contains additional bug fixes and any improvements or new features.

Percona XtraBackup 8.0.33-28 implements a new data sink that uses the first in, first out (FIFO) method. With the [FIFO data sink](#) feature, users with a streaming capacity of 10Gbps (typically on a Local Area Network (LAN)) can benefit from faster backups by streaming data in parallel to object storage. This feature is in [tech preview](#).

Percona XtraBackup 8.0.33-28 implements the [Dictionary cache](#) feature. Dictionary cache improves the performance of `xtrabackup --prepare` operation.

New feature

[PXB-2514](#): Implements a new [data sink that uses FIFO \(named pipes\)](#).

Improvements

- [PXB-2955](#): Implements [dictionary cache](#).
- [PXB-3033](#): Executes `STOP SLAVE` conditionally depending on `--[lock-ddl]`.
- [PXB-3066](#): Fixes compilation issues on Debian 12.
- [PXB-2993](#): Makes Percona XtraBackup compatible with `procps-4`.

Bug fix

- [PXB-2811](#): Fix for several Percona XtraBackup compilation warnings.

- [PXB-3075](#): Percona XtraBackup 8.0.33 could not prepare backups from 8.0.29 and previous versions.

Packaging notes

Percona XtraBackup 8.0.33-28 supports Debian 12.

Useful links

Install [Percona XtraBackup 8.0](#)

The [Percona XtraBackup GitHub location](#)

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

[Contribute to the documentation](#)

For [training](#), contact [Percona Training - Start learning now](#).

Percona XtraBackup 8.0.33-27 (2023-05-25)

Release date	May 25, 2023
Install instructions	Install Percona XtraBackup

[Percona XtraBackup \(PXB\)](#) is a 100% open source backup solution for all versions of Percona Server for MySQL and MySQL® that performs online non-blocking, tightly compressed, highly secure full backups on transactional systems. Maintain fully available applications during planned maintenance windows with Percona XtraBackup.

Release highlights

This release merges the MySQL 8.0.33 code base. Percona has implemented a two-stage release process for each version. The first release primarily ensures compatibility with the latest MySQL version, to help those customers who need an updated version of Percona XtraBackup as soon as possible. The second release contains additional bug fixes and any improvements or new features.

Bug fix

[PXB-3026](#): `DWITH_NDB=OFF` option did not turn off the `DWITH_NDB_CLUSTER_STORAGE_ENGINE` and compiled `NDB` files.

Useful links

The [Percona XtraBackup GitHub location](#)

Download product binaries, packages, and tarballs at [Percona Product Downloads](#)

[Contribute to the documentation](#)

For [training](#), contact [Percona Training - Start learning now](#).

Percona XtraBackup 8.0.32-26 (2023-04-04)

Release date	April 4, 2023
Install instructions	Install Percona XtraBackup

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries. Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

Release highlights

This release includes improvements and bug fixes. Percona has implemented a two-stage release process for each version. The first release primarily ensures compatibility with the latest MySQL version, to help those customers who need an updated version of Percona XtraBackup as soon as possible. The second release contains additional bug fixes and any improvements or new features.

This release fixes the security vulnerability [CVE-2022-25834](#) with [PXB-2977](#).

This version adds the [--estimate-memory](#) parameter to enable/disable the [Smart memory estimation](#) feature during the backup phase.

Improvements

[PXB-2882](#): Removes bitmap code and variables.

[PXB-2979](#): The Smart memory estimation feature parses metadata from the redo log to estimate the required memory to prepare a backup. On write-intensive workloads, this behavior caused the delay of the redo follow thread.

[PXB-2980](#): Adds the [--estimate-memory](#) parameter to enable/disable the Smart memory estimation feature during the backup phase.

Bug fixes

[PXB-2264](#): When a table with the next auto incremental value `2` was exported using Percona XtraBackup, the table got the next auto incremental value `1` instead of `2`.

[PXB-2954](#): During prepare phase Percona XtraBackup uses Serialized Dictionary information (SDI) to create the table metadata. If there was an orphan table, leftover from a upgrade, without SDI information, Percona XtraBackup failed to prepare a backup.

[PXB-2970](#): Removed unnecessary check for redo log files at the end of backup in case the redo archive was enabled.

[PXB-2972](#): An error after the redo apply phase caused an assertion failure instead of a graceful exit.

[PXB-2977](#): Fixed the security vulnerability [CVE-2022-25834](#).

[PXB-2998](#): Due to a server [bug](#) in the upgrade, the server leaves the dictionary tables in a temporary schema instead of the `mysql` schema. Percona Xtrabackup can now successfully prepare these backup directories.

[PXB-2999](#): Removed the warning message that MLOG_INDEX_LOAD redo log record was found.

[PXB-3022](#): Reduced the memory required for `--prepare` phase.

Useful links

The [Percona XtraBackup GitHub location](#)

[Contribute to the documentation](#)

For [training](#), contact [Percona Training - Start learning now](#).

Percona XtraBackup 8.0.32-25 (2023-02-27)

Release date	February 27, 2023
Install instructions	Install Percona XtraBackup

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries. Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

Release highlights

This release merges the MySQL 8.0.32 code base. Percona has implemented a two-stage release process for each version. The first release primarily ensures compatibility with the latest MySQL version, to help those customers who need an updated version of Percona XtraBackup as soon as possible. The second release contains additional bug fixes and any improvements or new features.

Starting with version 8.0.32-25, Percona XtraBackup allows instant columns in the backup for MySQL 8.0.32. This limitation still exists for backing up MySQL 8.0.31 and lower versions.

This limitation does not exist for backing up any version of Percona Server for MySQL.

Useful links

The [Percona XtraBackup GitHub location](#)

[Contribute to the documentation](#)

For [training](#), contact [Percona Training - Start learning now](#).

Percona XtraBackup 8.0.31-24 (2023-02-07)

Release date	February 7, 2023
Install instructions	Install Percona XtraBackup
Download this version	Percona XtraBackup 8.0

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries. Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

For [training](#), contact [Percona Training - Start learning now](#).

Release highlights

This release adds the ability to use an IAM instance profile when running xbcld from an EC2 instance. An IAM instance profile does not need the `--s3-secret-key` nor the `--s3-access-key` but the profile has a six-hour lifespan.

With this version, `qpress` / `QuickLZ` are deprecated.

New features

[PXB-2856](#): Use [IAM instance profile when running xbcld from an EC2 instance](#).

[PXB-2869](#): A warning is printed on backup that `qpress/QuickLZ` is deprecated.

Bug fixes

[PXB-2890](#): XtraBackup did not compress already compressed tables.

[PXB-2928](#): XtraBackup exits with signal 11 when taking a backup using page tracking.

[PXB-2854](#): Fixed a memory corruption issue in QuickLZ.

[PXB-2925](#): The PXB start_lsn, during the prepare phase, did not match the redo file start LSN when a wait was added in the backup phase after calling `recv_find_max_checkpoint` and before copying the checkpoint header.

[PXB-2964](#): XtraBackup required a very high open files limit.

[PXB-2959](#): XtraBackup now compiles on platforms with OpenSSL 3.0.

[PXB-2971](#): XtraBackup copied local manifest files which caused node failures when the files overrode settings on another node.

Deprecation and removal

Starting with Percona XtraBackup 8.0.31-24, using qpress/QuickLZ to compress backups is deprecated and may be removed in future versions. We recommend using either the `LZ4` or the Zstandard (`ZSTD`) compression algorithms.

Useful links

- The [Percona XtraBackup GitHub location](#)
- [Contribute to the documentation](#)

2022 (8.0.30-23 – 8.0.23-16)

Percona XtraBackup 8.0.30-23 (2022-11-14)

Release date	November 14, 2022
Install instructions	Install Percona XtraBackup
Download this version	Percona XtraBackup 8.0

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries. Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

For paid [support](#), [managed services](#) or [consulting services](#), contact [Percona Sales](#).

For [training](#), contact [Percona Training - Start learning now](#).

Release highlights

The [Take an incremental backup using page tracking](#) feature is [Generally Available \(GA\)](#).

The following features and improvements introduced in Percona XtraBackup 8.0.30-23 are available only in [tech preview](#):

- Percona XtraBackup adds support for [Zstandard \(ZSTD\)](#) compression algorithm. ZSTD is a fast lossless compression algorithm, targeting real-time compression scenarios at zlib-level and better compression ratios.
- Percona XtraBackup implements the [Smart memory estimation](#) feature to compute the memory required to `--prepare` a backup. Percona XtraBackup has extended the crash recovery logic to extract the formula used to allocate memory. Now, in the backup phase, while copying redo log entries, Percona XtraBackup computes the required memory for `--prepare`. Percona XtraBackup also takes into consideration the number of InnoDB pages to be fetched from the disk. Percona XtraBackup checks the server's available free memory and uses that memory up to the limit specified in the `--use-free-memory-pct` option to run `--prepare`.

New features

[PXB-2669](#): Implements support for [Zstandard \(ZSTD\)](#) compression algorithm.

Improvements

[PXB-2710](#): Implements support for `Smart Memory Estimation`.

[PXB-2844](#): Adjusts the logic to detect multi-threaded slave.

Bug fixes

[PXB-2681](#): Percona XtraBackup got stuck during incremental backup prepare of encrypted tables.

[PXB-2517](#): Check if files-from is successfully closed before rsync (Thanks to Garen Chan for the patch.)

[PXB-2702](#): Backup failed when undo truncation log is present (Thanks to Dmitry Smal for his contribution to the initial patch.)

[PXB-2729](#): Percona XtraBackup compilation failed if GTest packages are installed.

[PXB-2810](#): Percona XtraBackup crashed when AIO initialization failed (Thanks to Yan Huang for the patch.)

[PXB-2840](#): Xbcloud attempted to create a non-existing bucket.

[PXB-2874](#): When Percona XtraBackup generated a new master key on `-copy-back` or `-move-back`, the master key id was reset back to 1.

[PXB-2809](#): `wsrep_sync_wait<>0` caused `Lock wait timeout exceeded; try restarting transaction` error (Thanks to Frank Well for providing the initial patch.)

Platform support

Percona XtraBackup 8.0.30 supports Oracle Linux/Red Hat Enterprise Linux 9.

Useful links

- The [Percona XtraBackup GitHub location](#)
- [Contribute to the documentation](#)

Percona XtraBackup 8.0.29-22 (2022-07-19)

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries.

Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

Release Highlights

Percona XtraBackup 8.0.29-22 adds new redo log types to support the changes in the `INSTANT` algorithm behavior. *MySQL* 8.0.29 extended the support for `ALGORITHM=INSTANT` to allow columns to be added to any position in a table and column drops. Older versions of *Percona XtraBackup* are incompatible with *MySQL* 8.0.29 because of this new functionality.

Note

MySQL 8.0.29 extended the support of the `ALTER TABLE ... DROP COLUMNS` statement to use `ALGORITHM=INSTANT`. These operations only modify the metadata in the data directory and do not affect the physical structure. The `ALGORITHM=INSTANT` is the default setting for supported DDL operations. Prior to *MySQL* 8.0.29, an `ALGORITHM=INSTANT` column could only be added as the last table column. As of *MySQL* 8.0.29, the column can be added to any table position. The ability to add or drop a column in any position increases the complexity of crash recovery since the redo log does not have the logical positions from the data dictionary. The redo log contains the physical order for the table.

Percona Server for MySQL 8.0.29-21 contains fixes for these issues. *Percona XtraBackup* 8.0.29 can backup and restore tables that used the `INSTANT` algorithm in ADD/DROP in *Percona Server for MySQL*. However, *MySQL* 8.0.29, if the server contains tables modified by `INSTANT` ADD/DROP, is unsafe for backups at this time. *Percona XtraBackup* detects the tables modified by `INSTANT` ADD/DROP and generates an error. This error lists the affected tables and provides instructions to convert the modified tables to regular tables.

Useful Links

- The [Percona XtraBackup installation instructions](#)
- The [Percona XtraBackup downloads](#)
- The [Percona XtraBackup GitHub location](#)
- To contribute to the documentation, review the [Documentation Contribution Guide](#)

Percona XtraBackup 8.0.28-21 (2022-05-25)

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries.

Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

Release Highlights

Percona XtraBackup adds support for the [Amazon Key Management Service \(KMS\)](#) component.

New Features

- [PXB-2721](#): Implements support for the Amazon Key Management Service component in Percona XtraBackup.

Bugs Fixed

- [PXB-2761](#): Fix for a compilation error in GCC 11. (Thanks to user tcoyvwac for providing the patch to fix this issue.)
- [PXB-2422](#): The extraction of files failed when a file was located in another directory.

Useful Links

- The [Percona XtraBackup installation instructions](#)
- The [Percona XtraBackup downloads](#)
- The [Percona XtraBackup GitHub location](#)
- To contribute to the documentation, review the [Documentation Contribution Guide](#)

Percona XtraBackup 8.0.28-20

- **Date**

April 26, 2022

[Percona XtraBackup for MySQL Databases](#) enables MySQL backups without blocking user queries. Percona XtraBackup is ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, Percona XtraBackup drives down backup costs while providing unique features for MySQL backups.

Release Highlights

The log statements structure in Percona XtraBackup has been improved. Now, the error logs have a standard structure. The uniformity of the headers makes it easier to follow an operation's progress or review the log to diagnose issues. The improved log structure is displayed in the backup, prepare, move-back/copy-back error logs.

Each log statement has the following attributes:

- **Timestamp** - a timestamp for when the event occurred in a UTC format.
- **Severity** - the severity level of a statement indicates the importance of an event.
- **ID** - this identifier is currently not used but may be used in future versions.
- **Context** - the name of the module that issued the log statement, such as **XtraBackup**, **InnoDB**, or **Server**.
- **Message** - a description of the event generated by the module.

New Features

- [PXB-2670](#): Improves the error logging framework in Percona XtraBackup.

Bugs Fixed

- [PXB-1676](#): The error message after creating a backup was incorrect.
- [PXB-2716](#): Fix for a compilation error on systems without fallocation. (Thanks to user Bo98 for providing the patch to fix this issue.)

- [PXB-2506](#): During copy back and move back operations, added a check if the backup is fully prepared. Percona XtraBackup throws an error if the backup is not prepared or was prepared with `–apply-log-only`.
- [PXB-2662](#): Percona XtraBackup exited during prepare when creating an incremental backup using page tracking.
- [PXB-2714](#): Fix for a memory leak on the keyring component.
- [PXB-2697](#): Undefined symbols in macOS resulted in an error when linked with Percona XtraBackup 8.0.27.
- [PXB-2722](#): Fix for when via command line, a password, passed using the `-p` option, was written into the backup `tool_command` in `xtrabackup_info`.

Useful Links

- The [Percona XtraBackup installation instructions](#)
- The [Percona XtraBackup downloads](#)
- The [Percona XtraBackup GitHub location](#)
- To contribute to the documentation, review the [Documentation Contribution Guide](#)

Percona XtraBackup 8.0.27-19

- **Date**

February 2, 2022

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Release Highlights

The following list contains a number of the bug fixes for *MySQL 8.0.27*, provided by Oracle, and included in Percona Server for MySQL:

- The `default_authentication_plugin` is deprecated. Support for this plugin may be removed in future versions. Use the `authentication_policy` variable.
- The `binary` operator is deprecated. Support for this operator may be removed in future versions. Use `CAST(... AS BINARY)`.
- When a parent table updated or deleted a row, this operation initiated a cascading `SET NULL` operation on the child table. On the child table, a virtual column value could be set to NULL instead of the value derived from the base column value.

Find the full list of bug fixes and changes in the [MySQL 8.0.27 Release Notes](#).

New Features

- [PXB-1883](#): Added support for Microsoft Azure Cloud Storage in the xcloud binary. (Thanks to Ivan for reporting this issue)

Improvements

- [PXB-2434](#): Use MySQL page tracking for incremental backup

Bugs Fixed

- [PXB-1741](#): PS startup displays errors for databases excluded during partial backup
- [PXB-2614](#): The `xbstream` binary was enhanced to support sparse files on the XFS file system.
- [PXB-2608](#): Upgrade the Vault API to V2 (Thanks to Benedito Marques Magalhaes for reporting this issue)
- [PXB-2543](#): Fix that adds `IB_EXPORT_CFG_VERSION_V6` when exporting a `.cfg` file for a table (Thanks to Peng Gao for reporting this issue)
- [PXB-2465](#): Fix for querying the `ps.log_status` when group replication channels are items on that list. XtraBackup skips the name if it matches either `group_replication_applier` or `group_replication_recovery`. (Thanks to Galamb Gergő for reporting this issue)
- [PXB-2625](#): The default version of CURL in Debian Buster failed to identify a TLS HTTP2 connection as closed and attempted to reuse the connection. (Thanks to Johan Andersson for reporting this issue)

2021 (8.0.26-18 – 8.0.23-16)

Percona XtraBackup 8.0.26-18.0

- **Date**

September 2, 2021

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Improvements

- [PXB-2477](#): The xbcloud should retry on error and utilize incremental backoff (Thanks to Baptiste Mille-Mathias for reporting this issue)
- [PXB-2580](#): With the xbcloud binary, a chunk-upload on an SSL connect error to S3 was not retried (Thanks to Tim Vaillancourt for providing the patch)
- [PXB-2317](#): Remove the obsolete LOCKLESS binary log functionality since the `performance_schema.log_status` table is now used to get the log information on all the storages without locking them

Bugs Fixed

- [PXB-2101](#): The Prepare step fails due to duplicate Serialized Dictionary Information (SDI) (Thanks to Fungo Wang for reporting this issue)
- [PXB-1504](#): The FIND_GCRYPT macro is broken (Thanks to Maxim Bubliss for reporting this issue)
- [PXB-1961](#): The Prepare step of a full backup fails for encrypted tables with a generic error
- [PXB-2585](#): XtraBackup fails to backup archive RocksDB Write-Ahead Log (WAL) files

Percona XtraBackup 8.0.25–17.0

- **Date**

June 3, 2021

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Bugs Fixed

- [PXB-2442](#) : Backup cannot be decompressed using the AppArmor profile
- [PXB-2444](#) : The xbcloud binary fails to upload backup with enforcing SELinux mode
- [PXB-2443](#) : Version check fails with AppArmor profile
- [PXB-2445](#) : Initializing the libgcrypt in xbcloud
- [PXB-2455](#) : XtraBackup prepare fails if the checkpoint LSN is greater than the last LSN
- [PXB-2473](#) : SELinux errors in audit logs
- [PXB-1462](#) : Long gtid_executed breaks —history functionality
- [PXB-2369](#) : Issues with “Installing” page in XtraBackup documentation
- [PXB-2457](#) : Incorrect binlog names if binlog name contains periods
- [PXB-2106](#) : Copy-back creates wrong binlog.index

Percona XtraBackup 8.0.23-16.0

- **Date**

March 22, 2021

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

This release fixes the security vulnerability [CVE-2020-29488](#)

Improvements

- [PXB-2280](#): Provide SELinux and AppArmor default policies
- [PXB-1979](#): Enable `--lock-ddl` by default to prevent corruption of the backup

Bugs Fixed

- [PXB-2274](#): Correct restore processing when there are DML statements running during backup stage by writing the last_checkpoint and LSN from ps.log_status instead of the redo log (Thanks to user Li Biao for reporting this issue)
- [PXB-2430](#): Correct build failure by removing the dependency of no-server-version-check with version-check (Thanks to user agarner for reporting this issue)
- [PXB-2415](#): Add build dependencies to correct Debian/Ubuntu packages in docker (Thanks to user Matt Cole for reporting this issue)
- [PXB-2429](#): Correct Backup failure for encrypted tablespace by skipping the encryption reset operation during the redo scan phase
- [PXB-2418](#): Remove PROTOBUF_LITE_LIBRARY - it is not used and is not needed
- [PXB-2395](#): Update versions for xbstream and xbcrypt
- [PXB-2394](#): Correct spellings in xbcloud help

- [PXB-2357](#): Correct hang in backup with redo log archive by using block number instead of checksum (checksum of redo file and archive file can be different)
- [PXB-2180](#): Correct incremental prepare failure with logical redo by skipping the apply of logical redos (MLOG_TABLE_DYNAMIC_META) during the incremental prepare (except the last prepare).

2020 (8.0.22-15 – 8.0.10)

Percona XtraBackup 8.0.22–15.0

- **Date**

December 14, 2020

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

NOTE: The naming structure for the releases has changed. See the [aligning Percona XtraBackup versions with Percona Server for MySQL](#) blog post for more information.

New Features

- [PXB-2112](#): xbcld: support storage_class option with `–storage=s3` (Thanks to user rluisr for reporting this issue)
- [PXB-2242](#): Prevent back up if Source system (DB) version is higher then current PXB version

Improvements

- [PXB-2254](#): Redesign `–lock-ddl-per-table`

Bugs Fixed

- [PXB-793](#): Fix syntax error when executing `–lock-ddl-per-table` queries
- [PXB-953](#): Improve stdout for the end of usage of `–lock-ddl-per-table`
- [PXB-787](#): Modify scan to provide correct status when finding corrupted tablespace
- [PXB-2314](#): Handle Disk format changes introduced in MySQL 8.0.22 release
- [PXB-2279](#): Modify to look for prefixes in `xtrabackup_tablespaces*` without extensions (Thanks to user mrrmainnet for reporting this issue)
- [PXB-2336](#): Modify to check to see if first block is an all-zero block and process accordingly

- [PXB-2275](#): Modify backup processing to add validations if an encrypted table is created
- [PXB-2272](#): Modify regexp to consider all #sql as temporary tables
- [PXB-2257](#): Modify –lock-ddl-per-table to properly close database connection

Percona XtraBackup 8.0.14

- **Date**

August 31, 2020

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Percona XtraBackup 8.0.14 supports backup and restore processing for all versions of MySQL and has been tested with the latest MySQL 8.0.21.

Improvements

- [PXB-1605](#): Document how to use Percona Xtrabackup with Docker
- [PXB-2252](#): Introduce debug option to print redo log records scanned and applied

Bugs Fixed

- [PXB-2215](#): Modify Import tablespace process to correctly calculate pack_length
- [PXB-2114](#): Document when table is ignored by full backup it is not automatically ignored for incremental backup
- [PXB-2255](#): Modify processing to Error out if undo truncation during backup
- [PXB-2249](#): Verify perl binary exists before completing version check
- [PXB-2243](#): Correct processing when undo truncation happens between full backup and incremental
- [PXB-2238](#): Provide binary tarball with shared libs and glibc suffix & minimal tarballs
- [PXB-2202](#): Modify Xbcloud to display an error when xtrabackup fails to create a backup

Percona XtraBackup 8.0.13

- **Date**

June 17, 2020

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Percona XtraBackup 8.0.13 supports backup and restore processing for all versions of MySQL and has been tested with the latest MySQL 8.0.20.

Bugs Fixed

- [PXB-2165](#): Modify xbcloud to store backups using s3 access key parameters if AWS access key env variables are set
- [PXB-2164](#): Modify xbcloud to return the error when the backup doesn't exist in s3 bucket
- [PXB-2127](#): Modify xbcloud to upload backups with empty database to min.io storage
- [PXB-2198](#): Modify xbcloud delete to return the error when the backup doesn't exist in s3 bucket
- [PXB-2155](#): Correct corruption when redo logs are encrypted using xtrabackup --backup --compress=lz4
- [PXB-2166](#): Modify xbcloud to check for bucket existence and return `ok` status when domain name can be resolved

Percona XtraBackup 8.0.12

- **Date**

May 27, 2020

- **Installation**

[Installing Percona XtraBackup](#)

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Percona XtraBackup 8.0.12 now supports backup and restore processing for all versions of MySQL; previous versions of Percona XtraBackup will not work with MySQL 8.0.20 and higher.

Bugs Fixed

- [PXB-2133](#): Correct prtype stored for DECIMAL columns in .cfg file by `–export`
- [PXB-2146](#): Amazon S3 session token support added to xbcloud
- [PXB-2179](#): Modify `xtrabackup prepare` to shut down keyring plugin when run with `–generate-transition-key`
- [PXB-2157](#): Modify `xtrabackup –prepare` to generate cfg files for partition tables when `–export` is used

Percona XtraBackup 8.0.11

- **Date**

April 13, 2020

[Installing Percona XtraBackup 8.0](#)

Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

All *Percona* software is open-source and free.

This release fixes security vulnerability *CVE-2020-10997*.

Percona XtraBackup 8.0.10

- **Date**

March 16, 2020

[Installing Percona XtraBackup 8.0](#)

Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

All *Percona* software is open-source and free.

Bugs Fixed

- [PXB-1982](#): The history table showed a wrong value for `lock_time`.
- [PXB-2099](#): `copy-back` did not move undo files to the *data directory*.
- [PXB-2107](#): If the undo tablespace was created in a directory which is a symbolic link to another directory then the backup failed at backup stage.
- [PXB-2118](#): Undo log file was renamed incorrectly to undo tablespace name after restore

2019 (8.0.9 – 8.0.5)

Percona XtraBackup 8.0.9

Percona is glad to announce the release of *Percona XtraBackup* 8.0.9 on December 16, 2019. Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

All *Percona* software is open-source and free.

Bugs Fixed

- Sometime between December 3rd and December 10th, a change was introduced in that caused an incompatibility with our *Percona XtraBackup* `xbc1oud` utility. Bug fixed [PXB-1978](#).

Percona XtraBackup 8.0.8

Percona is glad to announce the release of Percona XtraBackup 8.0.8 on November 21, 2019. Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

All Percona software is open-source and free.

New Features and Improvements

- Support log archiving feature in PXB 8.0. More information in [PXB-1912](#) and in the [Redo Log section of MySQL documentation](#)
- For the MyRocks storage engine, support the creation of renewable checkpoints (controlled via `--rocksdb-checkpoint-max-age` and `--rocksdb-checkpoint-max-count`) to minimize the amount of binary logs to apply after the backup was completed. Using renewable checkpoints, Percona XtraBackup only copies the SST files that were created after the previous checkpoint. More information in [PXB-1915](#).
- Two options (`--backup-lock-timeout` and `--backup-lock-retry-count`) were added to enable the configuring of the timeout for acquiring metadata locks in `FLUSH TABLES WITH READ LOCK`, `LOCK TABLE FOR BACKUP`, and `LOCK BINLOG FOR BACKUP` statements. More information in [PXB-1914](#)

Bugs Fixed

- An encrypted table could not be restored when `ADD INDEX` or `DROP INDEX` commands had been run on the table. Bug fixed [PXB-1905](#)
- In some cases `xtrabackup --prepare` could fail to decrypt a table but reported that the operation completed ok. Bug fixed [PXB-1936](#)
- `xtrabackup --move-back` did not complete successfully when the encrypted binlog file. Bug fixed [PXB-1937](#).
- Percona XtraBackup could crash during the prepare stage when making an incremental

backup when a multi valued index was being added or dropped for JSON data. Bug fixed [PXB-1913](#).

Other bugs fixed: [PXB-1928](#), [PXB-1938](#), [PXB-1951](#), [PXB-1953](#), [PXB-1954](#).

Percona XtraBackup 8.0.7

Percona is glad to announce the release of *Percona XtraBackup* 8.0.7 on August 7, 2019. Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

In release 8.0.7, *Percona XtraBackup* enables making backups of databases that contain the [encrypted system tablespace](#). Encrypted mysql tablespace is now also supported.

Percona XtraBackup 8.0.7 implements the support of the `lz4` compression algorithm so that you could make compressed backups using lz4 (`--compress=lz4`) in addition to the default quicklz method.

All *Percona* software is open-source and free.

New Features and Improvements

- Add support of the system tablespace encryption. More information in [PXB-1649](#).
- Implemented the support of the lz4 compression algorithm. More information in [PXB-1857](#).

Bugs Fixed

- When the *encrypted tablespaces* feature was enabled, encrypted and compressed tables were not usable on the joiner node (Percona XtraDB Cluster) via SST (State Snapshot Transfer) with the xtrabackup-v2 method. Bug fixed [PXB-1867](#).
- `xbcloud` did not update date related fields of the HTTP header when retrying a request. Bug fixed [PXB-1874](#).
- `xbcloud` did not retry to send the request after receiving the HTTP 408 error (request timeout). Bug fixed [PXB-1875](#).
- `xtrabackup` did not accept decimal fractions as values of the `innodb_max_dirty_pages_pct` option. Bug fixed [PXB-1807](#).
- If the user tried to merge an already prepared incremental backup, a misleading error was produced without informing that incremental backups may not be used twice. Bug fixed [PXB-1862](#).

Other bugs fixed: [PXB-1493](#), [PXB-1557](#), [PXB-1887](#), [PXB-1870](#), [PXB-1879](#), [PXB-1901](#).

Percona XtraBackup 8.0.6

Percona is glad to announce the release of *Percona XtraBackup* 8.0.6 on May 9, 2019. Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

In version 8.0.6, *Percona XtraBackup* introduces the support of the MyRocks storage engine with Percona Server for MySQL version 8.0.15-6 or higher.

Percona XtraBackup 8.0.6 enables saving backups to an Amazon S3, MinIO, and Google Cloud Storage (using interoperability mode) when using xbcld. The following example demonstrates how to use an Amazon S3 storage to make a full backup:

```
$ xtrabackup --backup --stream=xbstream --extra-lsdir=/tmp --target-dir=/tmp | \
\
xbcloud put --storage=s3 \
--s3-endpoint='s3.amazonaws.com' \
--s3-access-key='YOUR-ACCESSKEYID' \
--s3-secret-key='YOUR-SECRETACCESSKEY' \
--s3-bucket='mysql_backups'
--parallel=10 \
${date -I}-full_backup
```

All *Percona* software is open-source and free

New Features

- The MyRocks storage engine is now supported with *Percona XtraBackup*. More information in [PXB-1754](#).
- Amazon S3 is now supported in xbcld. More information in [PXB-1813](#).

Bugs Fixed

- *Percona XtraBackup* could fail to restore the undo tablespace created during or before incremental backup. Bug fixed [PXB-1780](#).

- A backup could fail if `log_bin_index` was defined in `my.cnf`. Bug fixed [PXB-1801](#).
- When the row format was changed during the backup, xtrabackup could crash during the incremental prepare stage. Bug fixed [PXB-1824](#).
- During the `prepare` phase, *Percona XtraBackup* could freeze and never finish execution. Bug fixed [PXB-1819](#).
- *Percona XtraBackup* could crash during the `prepare` stage when making a backup of a host running MySQL Server v8.0.16. Bug fixed [PXB-1839](#).

Other bugs fixed: [PXB-1809](#), [PXB-1810](#), [PXB-1832](#), [PXB-1837](#).

Percona XtraBackup 8.0.5

Percona is glad to announce the release of *Percona XtraBackup* 8.0.5 on March 4, 2019. Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

Percona XtraBackup 8.0.5 introduces the support of undo tablespaces created using the new syntax (CREATE UNDO TABLESPACE) [available since MySQL 8.0.14](#). *Percona XtraBackup* also supports the binary log encryption introduced in MySQL 8.0.14.

Two new options were added to *xbstream*. Use the `--decompress` option with *xbstream* to decompress individual qpress files. With the `--decompress-threads` option, specify the number of threads to apply when decompressing. Thanks to [Rauli Ikonen](#) for this contribution.

This release of *Percona XtraBackup* is a **General Availability** release ready for use in a production environment.

All Percona software is open-source and free.

Please note the following about this release:

- The deprecated innobackupex has been removed. Use the xtrabackup command to back up your instances: `$ xtrabackup --backup --target-dir=/data/backup`
- When migrating from earlier database server versions, [backup and restore and using Percona XtraBackup 2.4](#) and then use `mysql_upgrade` from MySQL 8.0.x
- If using `yum` or `apt` repositories to install *Percona Xtrabackup* 8.0.5, ensure that you have enabled the new tools repository. You can do this with the **percona-release enable tools release** command and then install the `percona-xtrabackup-80` package.

New Features

- [PXB-1548](#): *Percona XtraBackup* enables updating the `ib_buffer_pool` file with the latest pages present in the buffer pool using the `--dump-innodb-buffer-pool` option. Thanks to Marcelo Altmann for contribution.
- [PXB-1768](#): Added support for undo tablespaces created with the new MySQL 8.0.14 syntax.

- [PXB-1781](#): Added support for binary log encryption introduced in MySQL 8.0.14.
- [PXB-1797](#): For *xbstream*, two new options were added. The `--decompress` option enables *xbstream* to decompress individual qpress files. The `--decompress-threads` option controls the number of threads to apply when decompressing. Thanks to [Rauli Ikonen](#) for this contribution.

Bugs Fixed

- Using `--lock-ddl-per-table` caused the server to scan all records of partitioned tables which could lead to the “out of memory” error. Bugs fixed [PXB-1691](#) and [PXB-1698](#).
- When *Percona XtraBackup* was started run with the `--slave-info`, incorrect coordinates were written to the `xtrabackup_slave_info` file.. Bug fixed [PXB-1737](#)
- *Percona XtraBackup* could crash at the prepare stage when making an incremental backup if the variable `innodb-rollback-segments` was changed after starting the MySQL Server. Bug fixed [PXB-1785](#).
- The full backup could fail when *Percona Server for MySQL* was started with the `--innodb-encrypt-tables` parameter. Bug fixed [PXB-1793](#).

Other bugs fixed: [PXB-1632](#), [PXB-1715](#), [PXB-1770](#), [PXB-1771](#), [PXB-1773](#).

2018 (8.0.4 – 8.0.3-rc1)

Percona XtraBackup 8.0.4

Percona is glad to announce the release of *Percona XtraBackup* 8.0.4 on December 10, 2018. Downloads are available from our [download site](#) and from apt and yum repositories.

Percona XtraBackup enables MySQL backups without blocking user queries, making it ideal for companies with large data sets and mission-critical applications that cannot tolerate long periods of downtime. Offered free as an open source solution, it drives down backup costs while providing unique features for MySQL backups.

This release of Percona Xtrabackup is a **General Availability** release ready for use in a production environment.

Please note the following about this release:

- The deprecated innobackupex has been removed. Use the xtrabackup command to back up your instances: `$ xtrabackup --backup --target-dir=/data/backup`
- When migrating from earlier database server versions, [backup and restore and using Percona XtraBackup 2.4](#) and then use `mysql_upgrade` from MySQL 8.0.x
- If using `yum` or `apt` repositories to install *Percona XtraBackup* 8.0.4, ensure that you have enabled the new tools repository. You can do this with the **percona-release enable tools release** command and then install the `percona-xtrabackup-80` package.

All Percona software is open-source and free. We are grateful to the community for the invaluable contributions to *Percona XtraBackup*. We would especially like to highlight the input of Alexey Kopytov who has been actively offering improvements and submitting bug reports for *Percona XtraBackup*.

New Features

- *Percona XtraBackup* 8.0.4 is based on MySQL 8.0.13 and fully supports Percona Server for MySQL 8.0 series and MySQL 8.0 series.

Bugs Fixed

- [PXB-1699](#): xtrabackup --prepare could fail on backups of MySQL 8.0.13 databases
- [PXB-1704](#): xtrabackup --prepare could hang while performing insert buffer merge

- [PXB-1668](#): When the `-throttle` option was used, the applied value was different from the one specified by the user (off by one error)
- [PXB-1679](#): PXB could crash when `ALTER TABLE ... TRUNCATE PARTITION` command was run during a backup without locking DDL

Percona XtraBackup 8.0-3-rc1

Percona is glad to announce the release of Percona XtraBackup 8.0-3-rc1 on October 31 2018. Downloads are available from our [download site](#) and from apt and yum repositories.

This is a **Release Candidate** quality release and it is not intended for production. If you want a high quality, Generally Available release, use the current Stable version (the most recent stable version at the time of writing is 2.4.12 in the 2.4 series).

Things to Note

- `innobackupex` was previously deprecated and has been removed
- Due to the new MySQL redo log and data dictionary formats the Percona XtraBackup 8.0.x versions will only be compatible with MySQL 8.0.x and the upcoming Percona Server for MySQL 8.0.x
- For experimental migrations from earlier database server versions, you will need to [backup and restore and using XtraBackup 2.4](#) and then use `mysql_upgrade` from MySQL 8.0.x

Improvements

- [PXB-1655](#): The `--lock-ddl` option is supported when backing up MySQL 8

Bugs Fixed

- [PXB-1678](#): Incremental backup prepare with the `--apply-log-only` option could roll back uncommitted transactions
- [PXB-1672](#): The MTS slave without GTID could be backed up when the `--safe-slave-backup` option was applied.

Quickstart Guide

Overview

Percona XtraBackup is a 100% open source backup solution for all versions of Percona Server for MySQL and MySQL® that performs online non-blocking, tightly compressed, highly secure full backups on transactional systems.

We recommend using Docker to quickly start with Percona XtraBackup. This Quickstart guide focuses on this installation method. You can explore alternative options in the [Install](#) section.

Why to start Percona XtraBackup in a Docker container

Docker containers are built from Docker images, which are snapshots of the configuration needed to run applications, such as the code and libraries. With Docker, you can build, deploy, run, update, and manage containers, which isolate applications from the host system. A Docker container lets you work with Percona XtraBackup without installing the product on a local drive.

Purpose of the Quickstart

This section provides the basic steps to start Percona XtraBackup in a Docker container, take a backup, prepare a backup, and restore a backup of Percona Server for MySQL database.

You are welcome to change the names of any items in commands. If you do, the results may vary from those we provide in this Quickstart guide.

In this section, container refers to the Docker container, and instance refers to the database server in the container.

Prerequisites

- [Install Docker](#) on your system.

To take a backup of Percona Server for MySQL, run Percona Server for MySQL in a Docker container and create a database and a table.

- [Start Percona Server in a Docker container](#)
- [Create a database and table in Percona Server](#)

Limitations

Percona XtraBackup 8.0 does not support making backups of databases created in versions before the 8.0 series of MySQL, Percona Server for MySQL, or Percona XtraDB Cluster or versions greater than 8.0, such as the 8.1 Innovation release.

Supported versions

Check the latest list of all [Supported versions](#).

Supported servers

Percona XtraBackup 8.0 can back up data stored on MySQL servers, Percona Server for MySQL, and Percona XtraDB Cluster.

Supported storage engines

Percona XtraBackup 8.0 can back up data from InnoDB, XtraDB, MyISAM, and MyRocks tables.

Steps to perform

In this Quickstart, you will learn how to:

- 1 [Start a Docker container and take a backup](#)
- 2 [Restore a backup](#)
- 3 [Clean up](#)
- 4 [Choose your next steps](#)

Next steps

[Start a Docker container and take a backup](#)

Start a Docker container and take a backup

In this scenario, Percona XtraBackup works in combination with a MySQL-compatible database instance. To use the tool, you must run Percona XtraBackup in a separate Docker container and then connect the Percona XtraBackup container directly to the Percona Server for MySQL container. This allows the backup tool access to the database for backup and restore operations.

The following steps create a Docker volume, start Percona XtraBackup in a Docker container, take and prepare a backup of Percona Server for MySQL database.

Create a Docker volume

The benefits of using a Docker volume are the following:

- You can keep your data safe and accessible across different containers. For example, if you want to upgrade your database image to a newer version, you can stop the old container and start a new one with the same volume attached without losing any data.
- You can quickly backup and restore your data using the `docker cp` command or mount the volume to another container.

```
$ docker volume create backupvol
```



Expected output



```
backupvol
```

Start Percona XtraBackup 8.0.34 in a container, take and prepare a backup

The `docker run` command creates and runs a new container from an image. The command modifies the container's behavior with options.

In our example, the command has the following options:

Option	Description
<code>--name</code>	Provides a name to the container. If you do not use this option, Docker adds a random name.
<code>--volumes-from</code>	Refers to Percona Server for MySQL and indicates that you intend to use the same data as the <code>psmysql</code> container.
<code>it</code>	Interacts with the container and be a pseudo-terminal.
<code>--user root</code>	Sets the user to root inside the Percona XtraBackup container. This option is required to access the MySQL data directory and run the <code>xtrabackup</code> command.
<code>backup_8034</code>	Indicates the directory inside the container where the backup files are stored.
<code>-v</code>	Mounts a volume from the host machine to volumes in another container that is being run. In our example, the <code>-v</code> option mounts a volume from the <code>psmysql</code> container to the <code>backupvol</code> volume on the host machine.
<code>backupvol</code>	Indicates the persistent storage for the database.
<code>psmysql</code>	Indicates the Percona Server for MySQL container.
<code>--password</code>	Prompts the user to enter the password for the root user. For convenience you can add the password for the root user to this option, for example, <code>--password=secret</code> . Then the password will be passed automatically when running the command. Note that if you specify the password with <code>--password=secret</code> , the password is visible in <code>docker ps</code> , <code>docker ps -a</code> (docker history) and regular <code>ps</code> command.
<code>8.0.34</code>	Use this tag to specify a specific version. Avoid using the <code>latest</code> tag. In our example, we use the <code>8.0.34</code> tag. In Docker, a tag is a label assigned to an image and is used to maintain different versions of an image.

If you do not add a tag, Docker uses `latest` as the default tag and downloads the newest image from [percona/percona-xtrabackup on the Docker Hub](#). This image can be in a different series or version from what you expect since the latest image changes over time. If you are using Percona XtraBackup version prior to 8.0.34, use tags to ensure that you use [compatible versions](#) of Percona Server for MySQL and Percona XtraBackup.

Starting with Percona XtraBackup 8.0.35-31, you can run the Docker ARM64 version of Percona XtraBackup. Use the `8.0.35-31-aarch64` tag instead of `8.0.35-31`.

Connect the Percona XtraBackup container to the Percona Server container

We recommend using the `--user root` option in the Docker command.

Run a Docker container example

```
$ docker run --name pxb --volumes-from psmysql -v backupvol:/backup_8034 -it --  
user root percona/percona-xtrabackup:8.0.34 /bin/bash -c "xtrabackup --backup --  
datadir=/var/lib/mysql/ --target-dir=/backup_8034 --user=root --password;  
xtrabackup --prepare --target-dir=/backup_8034"
```

You are prompted to enter the password. In our example, the password is `secret`.

```
2023-12-21T16:16:10.796716-00:00 0 [Note] [MY-011825] [Xtrabackup] recognized  
server arguments: --datadir=/var/lib/mysql/2023-12-21T16:16:10.799979-00:00 0  
[Note] [MY-011825] [Xtrabackup] recognized client arguments: --backup=1 --target-  
dir=/backup_8034 --user=root --password
```

Enter password:

In this example of expected output, we provide the first and last section of the backup log and use ellipses instead of the entire backup log.

Expected output

```
2023-12-14T15:45:16.277503-00:00 0 [Note] [MY-011825] [Xtrabackup] recognized server  
arguments: --datadir=/var/lib/mysql/  
  
2023-12-14T15:45:16.280208-00:00 0 [Note] [MY-011825] [Xtrabackup] recognized client  
arguments: --backup=1 --target-dir=/backup_80 --user=root --password=  
  
xtrabackup version 8.0.34-29 based on MySQL server 8.0.34 Linux (x86_64) (revision id:  
5ba706ee)  
  
2023-12-14T15:45:16.306670-00:00 0 [Note] [MY-011825] [Xtrabackup] perl binary not found.  
Skipping the version check  
  
2023-12-14T15:45:16.307026-00:00 0 [Note] [MY-011825] [Xtrabackup] Connecting to MySQL server  
host: localhost, user: root, password: set, port: not set, socket: not set  
  
2023-12-14T15:45:16.320214-00:00 0 [Note] [MY-011825] [Xtrabackup] Using server version  
8.0.34-26  
  
...  
  
2023-12-14T15:45:20.619892-00:00 0 [Note] [MY-011825] [Xtrabackup] completed OK!
```

The command runs a Docker container `pxb` from the `percona/percona-xtrabackup:8.0.34` image and mounts two volumes: one from another container named `psmysql`, which contains Percona Server data directory, and another named `backupvol`, which is where the backup files are stored. The command also sets the user to root and prompts the user to enter the password.

The command then executes two steps:

- Runs `xtrabackup` with the `--backup` option to copy the data files from `/var/lib/mysql/` to `/backup_8034`
- Runs `xtrabackup` with the `--prepare` option to apply the log files and make the backup consistent and ready for restore

You can check the Xtrabackup logs with the `docker container logs <container-name>` command.

For example:

```
$ docker container logs pxb
```



Next steps

[Restore the backup](#)

Restore the backup

The following steps describe how to restore your backup to another Percona Server container and check whether the data is available.

1 Create another Docker volume

```
$ docker volume create myvol2
```

Expected output

myvol2

2 Run another Percona Server container to restore the backup to

Option	Description
--name	Provides a name to the container. If you do not use this option, Docker adds a random name.
-d	Detaches the container. The container runs in the background.
-e	Adds an environmental variable. This example adds the MYSQL_ROOT_PASSWORD environment variable. The instance refuses to initialize if no environmental variable is added. Choose a more secure password, if needed.
myvol2	Indicates the persistent storage for the database.
8.0.34	Use this tag to specify a specific version.

You must provide at least one environment variable to access the database, such as MYSQL_ROOT_PASSWORD, MYSQL_DATABASE, MYSQL_USER, and MYSQL_PASSWORD or the instance refuses to initialize.

For this document, we are using the 8.0.34 tag. In Docker, a tag is a label assigned to an image. Tags are used to maintain different versions of an image. If we did not add a tag, Docker uses latest as the default tag and downloads the latest image from [percona/percona-server on the Docker Hub].

To use a Docker volume for persistent storage with a database, specify the path where the database stores its data inside the container, usually `/var/lib/mysql`.

→ Run a Docker container example

```
$ docker run -d -p 3307:3306 --name psmysql2 -e MYSQL_ROOT_PASSWORD=secret  
-v myvol2:/var/lib/mysql percona/percona-server:8.0.34
```

 Expected output

```
684464ce43d8880500f81abccb1e7440146a43f7d79c983a2a2a622b9a0d2080
```

3 Stop `psmysql2` container

```
$ docker stop psmysql2
```

 Expected output

```
psmysql2
```

4 Remove all created files in `myvol2` to allow xtrabackup restore data to `myvol2`

The `--rm` option automatically removes the temporary container created from `percona/percona-xtrabackup:8.0.34` image after the container exits.

```
$ docker run --volumes-from psmysql2 -v backupvol:/backup_8034 -it --rm --user  
root percona/percona-xtrabackup:8.0.34 /bin/bash -c "rm -rf /var/lib/mysql/*"
```

If the command executes successfully, the expected output is empty.

5 Restore backup of `psmysql` from `backupvol` to a new `psmysql2` instance.

```
$ docker run --volumes-from psmysql2 -v backupvol:/backup_8034 -it --rm --user  
root percona/percona-xtrabackup:8.0.34 /bin/bash -c "xtrabackup --copy-back --  
datadir=/var/lib/mysql/ --target-dir=/backup_8034"
```



Expected output



```
2023-12-14T16:06:39.845762-00:00 0 [Note] [MY-011825] [Xtrabackup] recognized server
arguments: --datadir=/var/lib/mysql/

2023-12-14T16:06:39.848408-00:00 0 [Note] [MY-011825] [Xtrabackup] recognized client
arguments: --copy-back=1 --target-dir=/backup_8034

xtrabackup version 8.0.34-29 based on MySQL server 8.0.34 Linux (x86_64) (revision id:
5ba706ee)

2023-12-14T16:06:39.849053-00:00 0 [Note] [MY-011825] [Xtrabackup] cd to /backup_8034/

...

2023-12-14T16:06:40.111577-00:00 1 [Note] [MY-011825] [Xtrabackup] Creating directory
./#innodb_redo

2023-12-14T16:06:40.111603-00:00 1 [Note] [MY-011825] [Xtrabackup] Done: creating
directory ./#innodb_redo

2023-12-14T16:06:40.111634-00:00 1 [Note] [MY-011825] [Xtrabackup] Copying ./ibtmp1 to
/var/lib/mysql/ibtmp1

2023-12-14T16:06:40.125971-00:00 1 [Note] [MY-011825] [Xtrabackup] Done: Copying
./ibtmp1 to /var/lib/mysql/ibtmp1

2023-12-14T16:06:40.144298-00:00 0 [Note] [MY-011825] [Xtrabackup] completed OK!
```

This command restores the backup files from the `backup_8034` volume to the `psmysql2` container. It uses the `percona/percona-xtrabackup:8.0.34` image to run a temporary container with root privileges. It executes the `xtrabackup` tool with the `--copy-back` option, which copies the files from the `backup_8034` volume to the `/var/lib/mysql/` directory in the `psmysql2` container. The command uses `-rm` option to delete the temporary container after it exits.

Validate the backup

This section describes the backup validation steps assuming that you backed up `mydb` database with the `employees` table as we have in our example. The validation steps may vary depending on your data set.

- 1 Change the owner of the files in `myvol2` from `root` to `mysql`

To avoid permission issues when running `psmysql2` container, you need to change the owner because the files were restored by `root` user and `psmysql2` will use `mysql` user.

```
$ docker run --volumes-from psmysql2 -v backupvol:/backup_8034 -it --rm --user  
root percona/percona-xtrabackup:8.0.34 /bin/bash -c "chown -R mysql:mysql  
/var/lib/mysql/"
```

If the command executes successfully, the expected output is empty.

2 Start the `psmysql2` container

```
$ docker start psmysql2
```

Expected output

```
psmysql2
```

3 Connect to the database instance.

For this example, we have the following options:

Option	Description
<code>it</code>	Interact with the container and be a pseudo-terminal
<code>psmysql2</code>	Running container name
<code>mysql</code>	Connects to a database instance
<code>-u</code>	Specifies the user account used to connect
<code>-p</code>	Use this password when connecting

→ Connect to the database instance example

```
$ docker exec -it psmysql2 mysql -uroot -p
```

You are prompted to enter the password, which is `secret`.

```
Enter password:
```

You should see the following result.



Expected output



```
Welcome to the MySQL monitor.  Commands end with ; or \g.

Your MySQL connection id is 31

Server version: 8.0.34-26 Percona Server (GPL), Release 26, Revision 0fe62c85

Copyright (c) 2009-2023 Percona LLC and/or its affiliates

Copyright (c) 2000, 2023, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
```

- 4 Check the list of databases to make sure, the backed up, `mydb` database is in the list.

```
mysql> show databases;
```



Expected output



```
+-----+
| Database          |
+-----+
| information_schema |
| mydb               |
| mysql              |
| performance_schema |
| sys                |
+-----+

5 rows in set (0.03 sec)
```

- 5 Use the `mydb` database.

```
mysql> use mydb;
```





Expected output



Database changed

6 Check that your table contains data

```
mysql> SELECT * FROM employees;
```



Expected output



```
+-----+-----+-----+-----+
| id | name                | email                | country |
+-----+-----+-----+-----+
|  1 | Erasmus Richardson | posuere.cubilia.curae@outlook.net | England |
|  2 | Jenna French       | rhoncus.donec@hotmail.couk        | Canada  |
|  3 | Alfred Dejesus     | interdum@aol.org                | Austria |
|  4 | Hamilton Puckett   | dapibus.quam@outlook.com          | Canada  |
|  5 | Michal Brzezinski  | magna@icloud.pl                  | Poland  |
|  6 | Zofia Lis          | zofial00@hotmail.pl               | Poland  |
|  7 | Aisha Yakubu       | ayakubu80@outlook.com             | Nigeria |
|  8 | Miguel Cardenas    | euismod@yahoo.com                 | Peru    |
|  9 | Luke Jansen        | nibh@hotmail.edu                  | Netherlands |
| 10 | Roger Pettersen    | nunc@protonmail.no                | Norway  |
+-----+-----+-----+-----+

10 rows in set (0.00 sec)
```

Next steps

Clean up

Clean up

In this document you'll learn how to:

- Exit the MySQL command client shell and the Docker container
- Remove the Docker container and the Docker image
- Remove the Docker volume.

The steps are as follows:

- 1 To exit the MySQL command client shell in `psmysql12` container, we use `exit`. You can also use the `\q` or `quit` commands. The execution of the statement also closes the connection.

```
mysql> exit
```

Expected output

```
Bye
```

- 2 Remove docker containers and images, if you no longer need docker containers / images or want to free up disk space. To remove a docker container, use the command `docker rm` followed by the container name or the container ID. To remove a docker image, use the command `docker rmi` followed by the image ID or name and the tag. The example uses the `-f` option to force the removal.

- Remove `psmysql12`, `psmysql1` and `pxb` Docker containers:

```
$ docker container rm psmysql12 -f
```

Expected output

```
psmysql12
```

```
$ docker container rm psmysql1 -f
```

 Expected output

psmysql

```
$ docker container rm pxb -f
```

 Expected output

pxb

- Remove `percona/percona-server:8.0.34` and `percona/percona-xtrabackup:8.0.34` Docker images

```
$ docker image rmi percona/percona-server:8.0.34
```

 Expected output

Untagged: percona/percona-server:8.0.34

```
$ docker image rmi percona/percona-xtrabackup:8.0.34
```

 Expected output

```
Untagged: percona/percona-server:8.0.34
Untagged: percona/percona-
server@sha256:4944f9b365e0dc88f41b3b704ff2a02d1459fd07763d7d1a444b263db8498e1f
Deleted: sha256:b2588da614b1f382468fc9f44600863e324067a9cae57c204a30a2105d61d9d9
Deleted: sha256:1ceaa6dc89e328281b426854a3b00509b5df13826a9618a09e819a830b752ebd
Deleted: sha256:77471692427a227eb16d06907357956c3bb43f0fdc3ecf6f8937e1acecae24fe
Deleted: sha256:8db06cc7b0430437edc7f118b139d2195cb65e2e8025f9a4517d16778f615384
Deleted: sha256:e5a57a2fafec4ab9482240f28927651d56545c19626e344aceb8be3704c3c397
Deleted: sha256:f86198f39b893674d44d424c958f23183bf919d2ced20e1f519714d0972d75ed
Deleted: sha256:db9672f7e12e374d5e9016b758a29d5444e8b0fd1246a6f1fc5c2b3c847dddcf
```

- 3 Remove docker volume if a container does not use the volume, and you no longer need it
Remove `backupvol` and `myvol2` Docker volumes:

```
$ docker volume rm backupvol
```



 Expected output



```
backupvol
```

```
$ docker volume rm myvol2
```



 Expected output



```
myvol2
```

Next step

[Next steps](#)

Next steps

After taking your first backup with Percona Xtrabackup, it's high time to expand your knowledge and skills in using Percona Xtrabackup.

Review the [Percona Xtrabackup documentation](#) for more information.

Learn more about other Percona products

For superior and optimized performance

Percona Server for MySQL (PS) is a freely available, fully compatible, enhanced, and open source drop-in replacement for any MySQL database. It provides superior and optimized performance, greater scalability and availability, enhanced backups, increased visibility, and instrumentation. Percona Server for MySQL is trusted by thousands of enterprises to provide better performance and concurrency for their most demanding workloads.

Install [Percona Server for MySQL](#).

For high availability

Percona XtraDB Cluster (PXC) is a 100% open source, enterprise-grade, highly available clustering solution for MySQL multi-master setups based on Galera. PXC helps enterprises minimize unexpected downtime and data loss, reduce costs, and improve performance and scalability of your database environments supporting your critical business applications in the most demanding public, private, and hybrid cloud environments.

Install [Percona XtraDB Cluster](#).

For Monitoring and Management

Percona Monitoring and Management (PMM) monitors and provides actionable performance data for MySQL variants, including Percona Server for MySQL, Percona XtraDB Cluster, Oracle MySQL Community Edition, Oracle MySQL Enterprise Edition, and MariaDB. PMM captures metrics and data for the InnoDB, XtraDB, and MyRocks storage engines, and has specialized dashboards for specific engine details.

[Install PMM and connect your MySQL instances to it.](#)

Advanced command-line tools

Percona Toolkit is a collection of advanced command-line tools used by the Percona support staff to perform a variety of MySQL, MongoDB, and system tasks that are complex or difficult to perform manually. These tools are ideal alternatives to “one-off” scripts because they are professionally developed, formally tested, and documented. Each tool is self-contained, so installation is quick and easy and does not install libraries.

[Percona Toolkit documentation](#)

How Percona XtraBackup works

About Percona XtraBackup

Percona XtraBackup is the world's only open source, free *MySQL* hot backup software that performs non-blocking backups for *InnoDB* and *XtraDB* databases.

With *Percona XtraBackup*, you can achieve the following benefits:

- Backups that complete quickly and reliably
- Uninterrupted transaction processing during backups
- Savings on disk space and network bandwidth
- Automatic backup verification
- Higher uptime due to faster restore time

Percona XtraBackup makes *MySQL* hot backups for all versions of *Percona Server for MySQL*, and *MySQL*. It performs streaming, compressed, and incremental *MySQL* backups.

Check the latest list of all [Supported versions](#).

Important

Percona XtraBackup 2.4 supports *MySQL* and *Percona Server for MySQL* 5.6 and 5.7 databases. Due to changes in the *MySQL* redo log and data dictionary formats, the **Percona XtraBackup 8.0.x** versions are only compatible with *MySQL* 8.0.x, *Percona Server for MySQL* 8.0.x, and compatible versions.

Percona's enterprise-grade commercial [MySQL Support](#) contracts include support for *Percona XtraBackup*. We recommend support for critical production deployments. *Percona XtraDB Backup* supports encryption.

Supported storage engines

Percona XtraBackup can back up data from *InnoDB*, *XtraDB*, *MyISAM*, and *MyRocks* tables on *MySQL* 8.0 servers as well as *Percona Server for MySQL* with *XtraDB*, [Percona Server for MySQL 8.0](#), and [Percona XtraDB Cluster 8.0](#).

Version updates

Version 8.0.6 and later supports the *MyRocks* storage engine. An incremental backup on the *MyRocks* storage engine does not determine if an earlier full or incremental backup contains the same files. **Percona XtraBackup** copies all *MyRocks* files each time it takes a backup. **Percona XtraBackup** does not support the *TokuDB* storage engine.

 See also

[Percona TokuBackup](#)

Limitations

Percona XtraBackup 8.0 does not support making backups of databases created in versions prior to 8.0 of *MySQL*, *Percona Server for MySQL* or *Percona XtraDB Cluster*. As the changes that *MySQL 8.0* introduced in *data dictionaries*, *redo log* and *undo log* are incompatible with previous versions, it is currently impossible for *Percona XtraBackup 8.0* to also support versions prior to 8.0.

Due to changes in *MySQL 8.0.20* released by Oracle at the end of April 2020, *Percona XtraBackup 8.0*, up to version 8.0.11, is not compatible with *MySQL* version 8.0.20 or higher, or *Percona* products that are based on it: *Percona Server for MySQL* and *Percona XtraDB Cluster*.

For more information, see [Percona XtraBackup 8.x and MySQL 8.0.20](#)

What are the features of Percona XtraBackup?

Here is a short list of the *Percona XtraBackup* features. See the documentation for more.

- Create hot InnoDB backups without pausing your database
- Make incremental backups of *MySQL*
- Stream compressed *MySQL* backups to another server
- Move tables between *MySQL* servers on-line
- Create new *MySQL* replication replicas easily
- Backup *MySQL* without adding load to the server
- *Percona XtraBackup* performs throttling based on the number of IO operations per second
- *Percona XtraBackup* can export individual tables even from a full backup, regardless of the InnoDB version
- Backup locks is a lightweight alternative to `FLUSH TABLES WITH READ LOCK` available in *Percona Server*. *Percona XtraBackup* uses them automatically to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.



See also

[How Percona XtraBackup works](#)

Additional information

The *InnoDB* tables are locked while copying non-*InnoDB* data.

How Percona XtraBackup works

Percona XtraBackup is based on InnoDB's crash-recovery functionality. It copies your InnoDB data files, which results in data that is internally inconsistent; but then it performs crash recovery on the files to make them a consistent, usable database again.

This works because InnoDB maintains a redo log, also called the transaction log. This contains a record of every change to InnoDB data. When InnoDB starts, it inspects the data files and the transaction log, and performs two steps. It applies committed transaction log entries to the data files, and it performs an undo operation on any transactions that modified data but did not commit.

[Percona XtraBackup 8.0.30-23](#) adds the `--register-redo-log-consumer` parameter. This parameter is disabled by default. When enabled, the parameter lets Percona XtraBackup register as a redo log consumer at the start of the backup. The server does not remove a redo log that Percona XtraBackup (the consumer) has not yet copied. The consumer reads the redo log and manually advances the log sequence number (LSN). The server blocks the writes during the process. Based on the redo log consumption, the server determines when it can purge the log.

Percona XtraBackup works by remembering the LSN when it starts, and then copying away the data files. It takes some time to do this, so if the files are changing, then they reflect the state of the database at different points in time. At the same time, Percona XtraBackup runs a background process that watches the transaction log files, and copies changes from it. Percona XtraBackup needs to do this continually because the transaction logs are written in a round-robin fashion, and can be reused after a while. Percona XtraBackup needs the transaction log records for every change to the data files since it began execution.

Percona XtraBackup uses [Backup locks](#)

where available as a lightweight alternative to `FLUSH TABLES WITH READ LOCK`. This feature is available in Percona Server for MySQL 5.6+. MySQL 8.0 allows acquiring an instance level backup lock via the `LOCK INSTANCE FOR BACKUP` statement.

Locking is only done for MyISAM and other non-InnoDB tables **after** Percona XtraBackup finishes backing up all InnoDB/XtraDB data and logs. Percona XtraBackup uses this automatically to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.

Important

The `BACKUP_ADMIN` privilege is required to query the `performance_schema_log_status` for either `LOCK INSTANCE FOR BACKUP` or `LOCK TABLES FOR BACKUP`.

xtrabackup tries to avoid backup locks and `FLUSH TABLES WITH READ LOCK` when the instance contains only InnoDB tables. In this case, xtrabackup obtains binary log coordinates from `performance_schema.log_status`. `FLUSH TABLES WITH READ LOCK` is still required in MySQL 8.0 when xtrabackup is started with the `--slave-info`. The `log_status` table in Percona Server for MySQL 8.0 is extended to include the relay log coordinates, so no locks are needed even with the `--slave-info` option.

 See also

[MySQL Documentation: LOCK INSTANCE FOR BACKUP](#)

When backup locks are supported by the server, xtrabackup first copies InnoDB data, runs the `LOCK TABLES FOR BACKUP` and then copies the MyISAM tables. Once this is done, the backup of the files will begin. It will backup `.frm`, `.MRG`, `.MYD`, `.MYI`, `.CSM`, `.CSV`, `.sdi` and `.par` files.

After that xtrabackup will use `LOCK BINLOG FOR BACKUP` to block all operations that might change either binary log position or `Exec_Master_Log_Pos` or `Exec_Gtid_Set` (i.e. source binary log coordinates corresponding to the current SQL thread state on a replication replica) as reported by `SHOW MASTER/SLAVE STATUS`. xtrabackup will then finish copying the REDO log files and fetch the binary log coordinates. After this is completed xtrabackup will unlock the binary log and tables.

Finally, the binary log position will be printed to `STDERR` and xtrabackup will exit returning 0 if all went OK.

Note that the `STDERR` of xtrabackup is not written in any file. You will have to redirect it to a file, e.g., `xtrabackup OPTIONS 2> backupout.log`.

It will also create the following files in the directory of the backup.

During the `prepare` phase, Percona XtraBackup performs crash recovery against the copied data files, using the copied transaction log file. After this is done, the database is ready to restore and use.

The backed-up MyISAM and InnoDB tables will be eventually consistent with each other, because after the `prepare` (recovery) process, InnoDB's data is rolled forward to the point at which the backup completed, not rolled back to the point at which it started. This point in time matches where the `FLUSH TABLES WITH READ LOCK` was taken, so the MyISAM data and the prepared InnoDB data are in sync.

The xtrabackup offers many features not mentioned in the preceding explanation. The functionality of each tool is explained in more detail further in this manual. In brief, though, the tools enable you to do operations such as streaming and incremental backups with various combinations of copying the data files, copying the log files, and applying the logs to the data.

Restoring a backup

To restore a backup with xtrabackup you can use the [--copy-back](#) or [--move-back](#) options.

xtrabackup will read from the `my.cnf` the variables `datadir`, `innodb_data_home_dir`, `innodb_data_file_path`, `innodb_log_group_home_dir` and check that the directories exist.

It will copy the MyISAM tables, indexes, etc. (.MRG, .MYD, .MYI, .CSM, .CSV, `.sdi`, and `par` files) first, InnoDB tables and indexes next and the log files at last. It will preserve file's attributes when copying them, you may have to change the files' ownership to `mysql` before starting the database server, as they will be owned by the user who created the backup.

Alternatively, the [--move-back](#) option may be used to restore a backup. This option is similar to [--copy-back](#) with the only difference that instead of copying files it moves them to their target locations. As this option removes backup files, it must be used with caution. It is useful in cases when there is not enough free disk space to hold both data files and their backup copies.

Percona XtraBackup Version Compatibility and Server Version Checks

When the server undergoes internal changes, such as modifying the structure of the redo log records, older versions of Percona XtraBackup may either fail or produce corrupted backups. To prevent these issues, XtraBackup checks the version of the source system before starting the backup and compares it to the Percona XtraBackup version.

Here are the possible scenarios and the actions taken:

Scenario	Actions Taken
Server version is the same as Percona XtraBackup version	The backup process continues without any issues. This is the ideal scenario because both the server and the backup tool are fully compatible.
Server version is older than Percona XtraBackup version	The backup process continues without any issues. Percona XtraBackup is designed to handle older versions of the server without problems.
Server version is newer than Percona XtraBackup version	The backup process stops with an error unless you explicitly override the version check. This happens because the backup tool may not support new features or changes introduced in the newer server version.
Server version is newer than Percona XtraBackup version, and the version check is overridden	The backup process continues, but there is a risk of failure or corruption. Overriding the version check can be risky because the backup tool might not be able to handle new features or changes correctly.

Until version 8.0.34, to ensure successful backups and restores, it is recommended to use a Percona XtraBackup version that matches or is newer than your server version. However, version 8.0.34 can back up newer server versions and older server versions if you set the `lock-ddl=ON` option.

Version updates

Percona XtraBackup 8.0.34: disabled the server-version-check.

Percona XtraBackup 8.0.22-15: introduced `--no-server-version-check`. This option allows you to override the version comparison process. This option is off by default, meaning the comparison occurs unless explicitly overridden.

Override Server Version check

You can add the `--no-server-version-check` option to your backup command to override the server version check and proceed with the backup. Here's an example of how to use this option:

```
$ xtrabackup --backup --no-server-version-check --target-dir=$mysql/backup1
```



Additional resources

For more details, refer to the documentation on [How XtraBackup Works](#).

For more information, refer to [Supported versions](#).

Differences between a logical backup and Percona XtraBackup

You'll notice several key differences when transitioning from a logical database to Percona XtraBackup.

Type	Logical backup	Physical backup
Definition	Contains logical data (tables, schemas, procedures).	Copies entire database directories and files.
Method	Exported as binary files using EXPORT/IMPORT tools.	Copies raw data files directly from the database.
Data contents	Includes data structure (schema, tables, procedures).	Contains complete data (structure, tables, transactions)
Instance Independence	Can be restored on any other machine or environment.	Tied to the specific database instance and environment.
OS-level restoration	Not suitable for OS-level restoration	Suitable for full disaster recovery scenarios.
Granularity	Granular recovery of specific components.	Restores the entire database.
Performance	Slower due to SQL parsing and execution	Faster due to direct file copying
Use cases	Data migration platform-independent copies	Full database restoration, disaster recovery
Security	Less secure	More secure

XtraBackup is designed to perform hot backups of MySQL, MariaDB, and Percona Server databases without interrupting database services. XtraBackup performs physical backups by copying the data files, ensuring minimal impact on server performance. XtraBackup includes advanced features such as backup compression and encryption, providing efficient and secure data protection. The tool also offers a streaming backup option, which is useful for backing up large databases or transferring backups to another server. XtraBackup integrates with Percona's additional tools like Percona Toolkit and Percona Monitoring and Management (PMM), offering a comprehensive database performance monitoring, management, and optimization solution. These features are not typically found in logical database backup solutions.

Another difference is the support for incremental backups with XtraBackup. You can back up only the data that has changed since the last full backup, significantly reducing the time and storage requirements compared to a full logical backup.

A logical backup contains copies of information about a database, such as tables, schemas, and procedures. Commonly, these backups are exported as binary files using EXPORT/IMPORT tools. Logical backups focus on the data structure and SQL statements, and you can back up a selected set of data or exclude specific data during the export, saving time and storage. They contain only structural information about the database, not instance-specific details. Logical backups are slower than physical backups because of the SQL parsing and execution. A logical backup does not preserve the file system-level security settings, user accounts, privileges, or authentication settings. During restoration, you must manually reconfigure these settings.

You can restore a logical backup on any other machine and in different environments. Logical backups provide granular recovery capabilities to restore specific components, such as individual tables or schemas. The recovery process can be more difficult because it involves recreating the entire data structure. Use logical backups when restoring or moving a database copy to another environment. They are an option for platform-independent data migration.

Unlike physical backups, logical backups are less secure because the data is unencrypted. Anyone with access to the backup could read sensitive information.

Backup strategy and backup policy

In conversation, 'backup strategy,' 'backup policy,' and 'backup plan' are used interchangeably, but the terms have distinct meanings.

A backup strategy is an overall approach or methodology for backing up the data in a database system. The strategy includes the decisions and considerations for what data to back up, how frequently backups should be taken, what backup methods and technologies to use, where to store the backup files, and what steps should be taken to ensure data integrity and the ability to recover.

A backup strategy typically involves business requirements, the Recovery Point Objective (RPO), the maximum acceptable data loss in case of system failure or disaster, and the Recovery Time Objective (RTO), the acceptable downtime for restoring the system after a failure, and compliance regulations.

A backup policy defines the guidelines, rules, or standards for specific procedures and practices for implementing the backup strategy.

The policy outlines the following:

- Detailed steps and parameters for performing backups
- Retention periods
- Backup types
- Backup verification methods
- Error handling procedures
- Security measures

A backup plan, based on the backup strategy and the backup policy, is a detailed operation document that specifies how to execute the backup strategy in actionable steps and defines the timelines for backup activities.

A backup plan can include the following details:

- Who performs the backup
- When backups are scheduled
- How backups are monitored and audited
- What actions to take in case of backup failures or emergencies

Types of backup

A [full backup](#) copies all data to a storage location. This type of backup creates a complete copy of the data set. The benefit of a full backup is a quick recovery since the data is stored in one place. However, this type of backup is also the most time-consuming and storage-intensive; it duplicates all data regardless of whether it has changed since the last backup.

An [incremental backup](#) only copies data that has changed since the last backup, whether that was a full or another incremental backup. This method is faster and more storage-efficient than a full backup. However, the recovery can be slower because you must restore each set of changes in sequence.

For more information, see the [example of using a full backup with incremental backups](#).

Backup policy defined

The backup policy sets the guidelines and objectives for data protection, while the backup strategy provides the roadmap for achieving those objectives.

A backup policy may state that all customer data must be backed up daily, with a full backup at the start of each week and incremental backups taken each business day. These definitions are based on the following: It also defines how long backups should be retained based on legal or regulatory requirements.

The difference between a backup policy and a backup strategy lies in their scope and detail. A backup policy is more about the following:

- ‘what’ — What should be backed up and may include information on compliance and governance
- ‘when’ - When should you do backups

The backup strategy contains the ‘how’ — how the backups are scheduled and includes information on the backup technologies and the restoration process.

The backup strategy details the use of specific backup software, the configuration of backup schedules, the process for monitoring backup success, and the necessary steps for restoring data.

Backup policy best practices

An effective backup policy is crucial for safeguarding an organization’s data integrity and availability. Here are some best practices to consider:

Backup Policy Best Practices | Markdown Table

Best Practice	Description
Regularly Update and Review	Ensure the backup policy is a living document, updated to reflect changes in technology, business processes, and regulations.
Clearly Define Roles and Responsibilities	Assign specific individuals or teams for executing backups, monitoring success, and managing restoration.
Adopt the 3-2-1 Backup Rule	Have 3 data copies, store them on 2 different media types, and keep 1 copy offsite.
Ensure Comprehensive Coverage	The policy should cover all critical data, applications, and systems to avoid data loss gaps.
Encrypt Backup Data	Protect backups with encryption, especially for offsite or cloud storage.
Test Backups Regularly	Regularly test backup processes to ensure data can be effectively restored.
Implement a Disaster Recovery Plan	Integrate the backup policy with a broader disaster recovery plan for data loss scenarios.
Use Reliable Backup Software and Hardware	Invest in proven technologies with robust backup and restoration capabilities.
Monitor Backup Procedures	Continuously monitor procedures to ensure they function as intended and identify issues promptly.
Train Staff	Educate relevant staff on the importance of backups and their role in the process.
Document the Backup Process	Maintain thorough documentation of the backup process, including schedules, technologies, and restoration procedures.
Define Service Level Agreements (SLAs)	Establish clear SLAs for backup and recovery times to set expectations and ensure accountability.
Comply with Legal and Regulatory Requirements	Ensure the policy complies with relevant laws and regulations regarding data retention and protection.

Best Practice	Description
Plan for Scalability	Design the policy to accommodate future data volume growth and infrastructure changes.
Conduct Periodic Risk Assessments	Regularly assess data risks and adjust the backup policy to mitigate those risks.

By adhering to these best practices, organizations can create a robust backup policy that protects their data and supports their business continuity and disaster recovery strategies. It’s important to remember that the effectiveness of a backup policy is not just in its creation but in its implementation and ongoing management. Regular reviews, updates, and staff training ensure the backup policy remains relevant and effective in evolving data landscapes and threats. The goal is to have a backup policy that is not only comprehensive and compliant but also practical and responsive to the organization’s needs.

Backup policy example

An example of a backup policy might include the following components:

Aspect	Details
Purpose	To define the procedures and responsibilities for backing up organizational data to ensure its availability in case of data loss.
Scope	This policy applies to all data, hardware, and software in the organization’s information systems.
Policy Statement	The organization shall maintain a regular backup schedule for all critical data. Full backups are to be performed weekly, with incremental backups occurring daily. All backups must be encrypted and stored in a secure, off-site location.
Roles and Responsibilities	The IT department is responsible for executing backups as scheduled. The department head will ensure that backup procedures are followed and reviewed regularly.
Backup Procedures	The IT department must use software and hardware. Data to be backed up includes, but is not limited to, customer information, financial records, and operational data.

Aspect	Details
Data Retention	Retain the backup data for five years or as required by law or industry regulations.
Disaster Recovery	In the event of data loss, the IT department shall follow the disaster recovery plan to restore data from the most recent backup.
Review and Testing	Review the backup policy and procedures annually. Test the backup and restoration procedures quarterly to ensure they are effective.

Test the backups

You must test the effectiveness of your backups and demonstrate that you can recover your data in the event of loss or corruption.

Here are several methods to assess the reliability of your backup system:

Aspect	Details
Perform Regular Test Restores	The most reliable way to test backups is to periodically restore the data to a separate system. This restoration verifies the backup's integrity and the restoration process's functionality.
Monitor Backup Reports	Regularly review backup reports and logs for any errors or failures during the backup process. This can help identify issues early before they become critical problems.
Check for Data Consistency	Compare the original data with the backed-up data to ensure consistency. This can be done by checking file sizes, using checksums, or employing other data validation techniques. One tool would be to use <code>[pt-table-checksum]</code> .
Simulate Disaster Scenarios	Conducting a simulated disaster recovery exercise can help determine how well your backup system will perform under stress and identify any gaps in your recovery plan.
Audit Backup Procedures	Regular audits of the backup procedures can ensure that the backups are being performed in accordance with the established policies and that the backup system is functioning as expected.

Aspect	Details
Validate Backup Configuration	Ensure that the backup configuration is correct and that all necessary data is being backed up. This includes checking backup schedules, retention policies, and the scope of the data backed up.
Verify Security Measures	Confirm that the backup data is secure and that encryption and other security measures are functioning correctly to protect against unauthorized access.
Evaluate Recovery Time Objectives	Measure how long it takes to restore data from backups and compare this with your recovery time objectives to ensure they align with your business requirements.
Assess Backup Media Reliability	Regularly check the physical and logical condition of the backup media, whether it's tape, disk, or cloud storage, to ensure it's in good working order.
Conduct Incremental Backup Tests	Test the restoration of incremental backups to ensure that they work in conjunction with full backups and that the data can be restored to a specific point in time.
Review Backup Software Updates	Keep the backup software updated and test after updates to ensure new versions or patches do not introduce issues with the backup or restoration processes.
Check for Compatibility Issues	Ensure that the backup system is compatible with all the software and data formats used in your organization to avoid restoration issues.
Document Test Results	Keep detailed records of all backup tests, including the scope, date, issues encountered, and resolutions. This documentation can be invaluable for troubleshooting and improving the backup process.

By implementing these methods, you are confident that your backup system protects your data and can recover your data when necessary. Regular testing and monitoring are key to maintaining an effective backup strategy and ensuring business continuity. Remember, the goal of testing is not just to prove that backups can be restored, but to identify and resolve any issues before an actual disaster occurs. This proactive approach can save valuable time and resources, and ultimately, safeguard your organization's data integrity.

Common challenges

Common backup challenges often include issues such as damaged backups, missing or failed backups, maintaining legacy systems, adhering to strict privacy and security regulations, and managing the costs associated with data backup. To overcome these challenges, organizations can adopt several strategies. Implementing a modern, reliable backup solution that includes automated data consistency checking can help protect against damaged backups. Regular recovery tests can ensure that data is recoverable and not corrupt.

For missing or failed backups, it’s essential to have a robust monitoring system that alerts IT staff to any issues with the backup process. This allows for immediate action to rectify any problems before they escalate. Maintaining legacy systems can be addressed by gradually phasing out older technology and replacing it with more current, supported solutions. This can help avoid the pitfalls of outdated systems that may not be compatible with newer technologies or that may no longer be supported by the vendor.

Adhering to privacy and security regulations requires a thorough understanding of the laws and standards that apply to an organization’s industry and geography. Implementing encryption and access controls can help ensure that backups are secure and compliant with regulations. Finally, managing the costs of data backup involves careful planning and budgeting. Organizations should consider the total cost of ownership of their backup solutions, including initial investment, ongoing maintenance, and potential costs associated with data recovery.

In addition to these strategies, it’s essential to follow best practices such as the 3-2-1 backup rule, which suggests having three copies of data on two different media, with one copy stored offsite. Regular testing and validation of backups, clear documentation of backup procedures, and continuous training for staff involved in the backup process are also crucial for overcoming common backup challenges. By staying vigilant and proactive, organizations can ensure that their data remains secure and recoverable, even in these challenges.

Disaster recovery plan

A disaster recovery plan (DRP) is a structured approach with detailed instructions for responding to unplanned incidents such as natural disasters, power outages, cyber-attacks, and other disruptive events.

This table organizes the different sections of the disaster recovery plan, along with brief descriptions of each section.

Section	Description
Introduction	This section outlines the purpose of the plan, its scope, and the types of disasters it covers. It also includes the plan’s objectives, such as minimizing downtime and data loss.

Section	Description
Roles and Responsibilities	This part defines the disaster recovery team's specific roles. It lists the contact information for each team member and outlines the chain of command.
Incident Response	This section provides the steps to detect, report, and assess damage caused by a disaster. It includes procedures for immediate response to limit the impact on operations.
Plan Activation	Here, the criteria for launching the disaster recovery plan are detailed. It specifies who has the authority to activate the plan and under what circumstances.
Communication Plan	This part details how communication should be handled during a disaster, including internal communication with employees and external communication with customers, suppliers, and the media.
Recovery Strategies	This section outlines the strategies for restoring hardware, software, and data that are critical to resuming business operations. It includes prioritized actions for recovery and the resources required.
Data Backup Procedures	Detailed instructions for restoring data from backups are provided here. It includes the locations of backups, the types of backups (full, incremental), and the backup schedule.
Site Relocation	If the primary business location is compromised, this part of the plan details the process for moving operations to an alternate site.
Plan Testing and Maintenance	This section describes the schedule for regular testing of the DRP to ensure its effectiveness. It also outlines the process for updating the plan as the business and technology environment changes.
Appendices	The appendices may include site maps, network diagrams, floor plans, and any other documentation that supports disaster recovery efforts.

It's important to note that a DRP should be tailored to the specific needs of the organization and should be reviewed and tested regularly to ensure it remains effective and up-to-date. The example provided here is a high-level overview and would need to be expanded with detailed, actionable steps specific to the organization's operational requirements. For more detailed templates and guidance on creating a disaster recovery plan, there are resources available online that can provide a starting point for developing a comprehensive plan.

Backup example

This example demonstrates how a beginner DBA can set up daily full backups and hourly incremental backups for a large organization using Percona XtraBackup 8.0 with a Percona Server for MySQL 8.0 database server.

Assumptions

- You have Percona XtraBackup 8.0 installed on your system.
- You have a Percona Server for MySQL 8.0 database server running.
- You have created dedicated directories for storing backups:
 - `/backups/mysql/full` - for daily full backups
 - `/backups/mysql/incr` - for hourly incremental backups

Daily Full Backup Script

This following script defines the server details like hostname, username, password, data directory, and backup directory. Replace these with your actual values.

- Defines the server credentials
- Generates a timestamp with the `date` command for the full backup filename
- Uses `xtrabackup` with the `--backup` flag for a full backup
- Specifies the directory to store the backup with `--target-dir`
- `--datadir` points to the location of your MySQL data files.
- Provides the credentials with `--user` and `--password` to access the database server.
- Echoes a message confirming the successful backup completion.

```
#!/bin/bash

# Set server details (replace with your actual values)
SERVER_HOST="your_server_hostname"
SERVER_USER="your_db_username"
SERVER_PASSWORD="your_db_password"
DATA_DIR="/var/lib/mysql" # Replace with your data directory path
BACKUP_DIR="/backups/mysql/full"

# Get current date for filename
TODAY=$(date +%Y-%m-%d)
FULL_BACKUP_FILENAME="${BACKUP_DIR}/${TODAY}"

# Run xtrabackup with full backup command
xtrabackup --backup \
  --target-dir="${FULL_BACKUP_FILENAME}" \
  --datadir="${DATA_DIR}" \
  --user="${SERVER_USER}" \
  --password="${SERVER_PASSWORD}"

echo "Daily full backup completed at $(date)"
```

Running the Daily Script

You can schedule this script to run automatically using cron. For example, to run it every day at midnight (00:00), add this line to your crontab:

```
$ 0 0 * * * /path/to/daily_full_backup.sh
```

Hourly Incremental Backup Script

This script uses the same server details as the daily script. It retrieves the current date and hour for the incremental backup filename. The core command uses `xtrabackup` with the `--backup` and `--incremental` flags for an incremental backup. `--target-dir` specifies the directory to store the incremental backup. `--base-dir` points to the latest full backup directory obtained using `ls` and `tail` commands. This ensures the incremental backup references the most recent full backup. Similar to the daily script, it echoes a confirmation message.

```
#!/bin/bash
```



```
# Set server details (same as daily script)
SERVER_HOST="your_server_hostname"
SERVER_USER="your_db_username"
SERVER_PASSWORD="your_db_password"
DATA_DIR="/var/lib/mysql" # Replace with your data directory path
BACKUP_DIR="/backups/mysql/incr"

# Get current date and hour for filename
TODAY=$(date +%Y-%m-%d)
HOUR=$(date +%H) # Get current hour (00-23)
INCR_BACKUP_FILENAME="${BACKUP_DIR}/${TODAY}_${HOUR}"

# Set the base directory for incremental backups (point to the latest full backup)
BASE_DIR="/backups/mysql/full/$(ls -tr /backups/mysql/full | tail -n 1)"

# Run xtrabackup with incremental backup command
xtrabackup --backup \
  --incremental \
  --target-dir="${INCR_BACKUP_FILENAME}" \
  --base-dir="${BASE_DIR}" \
  --user="${SERVER_USER}" \
  --password="${SERVER_PASSWORD}"

echo "Hourly incremental backup completed at $(date)"
```

Run the Hourly Script

Schedule this script to run every hour using cron.

```
$ 0 * * * * /path/to/your/script.sh
```




Features

Backup features

LRU dump backup

Percona XtraBackup includes a saved buffer pool dump into a backup to enable reducing the warm up time. It restores the buffer pool state from `ib_buffer_pool` file after restart. *Percona XtraBackup* discovers `ib_buffer_pool` and backs it up automatically.



If the buffer restore option is enabled in `my.cnf` buffer pool will be in the warm state after backup is restored.

Throttling backups

xtrabackup does not block database activity during a backup, but xtrabackup generates additional disk input/output (I/O) activity. On systems with limited I/O capacity, this can affect database performance.

Use the `--throttle` option to limit the backup rate.

How throttling works

Use the `--throttle` option to minimize backup impact on system performance. By limiting the number of **10 MB chunks** processed per second, you primarily limit data throughput, which in turn reduces overall I/O pressure on the system.

When used with the `--backup` option, `--throttle` limits the rate of read and write operations performed during the copy process. For incremental backups, the limit applies only to read operations.

For example:

- `--throttle=1` limits processing to approximately 10 MB/s
- `--throttle=10` limits processing to approximately 100 MB/s

By default, throttling is **disabled**, and xtrabackup runs at the maximum rate permitted by the system.

On write-heavy systems, overly aggressive throttling can cause backups to fall behind redo log generation, see [Redo log wraparound risk](#).

When to use throttling

Use the `--throttle` option in the following scenarios:

- **High-traffic production systems**

On latency-sensitive workloads, unthrottled backups can cause I/O contention. This can increase query latency as the database waits for disk operations to complete.

- **Limited storage throughput (cloud environments)**

On platforms such as AWS with provisioned IOPS or burst limits (for example, EBS volumes), unthrottled backups can exceed available I/O capacity. This may degrade overall system performance when limits are reached.

- **Systems with slower storage**

On HDDs or lower-performance SSDs, higher latency and seek times limit the number of concurrent I/O operations. Throttling prevents backups from saturating the storage subsystem.

- **Multi-tenant or co-located environments**

When multiple services share the same host, unthrottled backups can consume available I/O bandwidth and affect other applications.

Resource monitoring: finding your baseline

Before setting a throttle value, understand your storage subsystem limits. If your disk is already near saturation, an unthrottled backup can introduce significant I/O wait and increase query latency.

Use the following tools to observe disk utilization and validate your throttle settings.

iostat (sysstat)

Use:

```
iostat -xz 1
```



Key metrics:

- `aqu-sz` - average queue size, a better indicator of saturation. Values consistently above 1 suggest the device cannot keep up with demand.
- `await` — average time (ms) for I/O requests. Increasing values during backup indicate growing latency.
- `%util` — indicates how busy the device is (time spent servicing I/O), not saturation. A device can show 100% utilization and still have spare capacity.

pidstat (sysstat)

Use `pidstat` to monitor per-process I/O activity.

Note

`pidstat` requires root privileges (or appropriate capabilities) to display system-wide I/O activity. Without this, output may be incomplete.

```
pidstat -d 1
```



Key notes:

- Shows read/write throughput per process (including xtrabackup and mysqld).
- Helps distinguish I/O generated by xtrabackup from that of the MySQL server.
- Commonly available as part of the `sysstat` package.
- When run without elevated privileges, it only shows processes owned by the current user.

pv (Pipe Viewer)

For streaming backups, `pv` can display real-time throughput:

```
xtrabackup --backup --throttle=5 --stream=xbstream \  
| pv \  
| gzip > backup.xb.gz
```

This helps verify whether the effective throughput matches the configured throttle, although compression (`gzip`) may become a bottleneck on CPU-constrained systems.

Configuration guidelines

Select a throttle value based on available I/O capacity and workload requirements. Use the insights from [resource monitoring](#) to select and refine an appropriate throttle value.

Typical starting points:

- Busy systems: `--throttle=1-5`
- Moderate load: `--throttle=5-20`
- Dedicated backup window: no throttling

Note

If the throttle value is too low, the backup may not keep up with the rate of transaction log generation and may not complete.

Adjust the value based on observed system performance.

Redo log wraparound risk (important)

When using `--throttle`, backups may become too slow to keep up with redo log generation rate on busy systems.

Because the InnoDB redo log is a circular buffer with limited size, it can wrap around before xtrabackup has finished reading it. When this happens, the backup fails with an error similar to:

```
xtrabackup: error: expected log block no. ..., but got no. ...  
xtrabackup: error: it looks like InnoDB log has wrapped around before xtrabackup  
could process all records  
xtrabackup: Error: xtrabackup_copy_logfile() failed
```

This issue is more likely on write-heavy systems or when using very low throttle values.

There are two primary ways to prevent this issue:

Option 1: Increase redo log capacity (recommended)

```
SET GLOBAL innodb_redo_log_capacity = <larger value>;
```

Increasing redo log size:

- Reduces the risk of wraparound during throttled backups.
- Improves write performance under heavy load.
- May increase crash recovery time.

Option 2: Prevent redo log wraparound

Use the `--register-redo-log-consumer` option to prevent redo log wraparound by coordinating with the server.

However:

- It can block or significantly delay write operations if xtrabackup cannot keep up.
- It is best suited for replicas, where replication lag is acceptable.

Advanced tuning scenarios

On high-performance storage (for example, NVMe), consider using higher `--parallel` values (4–8) with little or no throttling.

On shared or limited environments (such as cloud volumes):

- Identify your provisioned IOPS or throughput limits.

- Start with approximately **50% of available I/O capacity**, then adjust based on observed latency and workload behavior.
- Use moderate `--parallel` values (2–4) to ensure consistent utilization without overwhelming the system.

Interaction with `--parallel`

The `--throttle` and `--parallel` options control different aspects of backup performance and can be used together. If `--parallel` is set too high without adjusting `--throttle`, threads may compete for the limited throughput, reducing efficiency.

How they work together

- `--throttle` sets a **global limit** on backup throughput.
- `--parallel` controls the number of worker threads copying data.

If both are used:

- All threads share the same throttle limit.
- Increasing `--parallel` does not increase total throughput beyond the throttle limit.

For example:

- `--throttle=10` (approximately 100 MB/s total)
- `--parallel=4`

Each thread will average approximately 25 MB/s to stay within the global limit.

Why use both

Using both options allows you to:

- Improve efficiency when backing up many small tables (via parallelism)
- Control overall I/O impact (via throttling)

This is especially useful on systems with mixed workloads.

Example usage

A) Streaming backup with throttling

Limit the backup speed to approximately **50 MB/s**:

```
xtrabackup --backup --throttle=5 --stream=xbstream \  
| pv \  
| gzip > backup.xb.gz
```



B) Testing throttling behavior

To measure the effective throughput of `--throttle` without interference from compression or disk writes:

```
xtrabackup --backup --throttle=2 --stream=xbstream 2>/dev/null \  
| pv \  
| cat > /dev/null
```



This setup ensures:

- No disk write bottleneck
- No CPU bottleneck from compression
- Clean pv output (stderr from xtrabackup is suppressed)

Use this method to validate that the configured throttle produces the expected throughput.

Local backup with throttling

Limit the backup to approximately **100 MB/s**:

```
xtrabackup --backup --target-dir=/backups/full --throttle=10
```



Expected output



```
xtrabackup: Starting backup  
xtrabackup: Throttling set to 10 chunks per second  
...  
xtrabackup: completed OK!
```



Incremental backup with throttling

For incremental backups, the limit applies only to read operations. Redo log copying continues independently and may become a bottleneck if the backup is throttled too aggressively.

```
xtrabackup \  
  --backup \  
  --target-dir=/backups/inc1 \  
  --incremental-basedir=/backups/full \  
  --throttle=5
```

Expected output

```
xtrabackup: Starting incremental backup  
xtrabackup: Throttling set to 5 chunks per second  
...  
xtrabackup: completed OK!
```

Store backup history on the server

Percona XtraBackup supports storing the backups history on the server to provide users with additional information about backups that are being taken. Backup history information is in the `PERCONA_SCHEMA.XTRABACKUP_HISTORY` table.

To use this feature the following options are available:

- `-history` = <name> : This option enables the history feature and allows the user to specify a backup series name that will be placed within the history record.
- `-incremental-history-name` =<name> : This option allows an incremental backup to be made based on a specific history series by name. *xtrabackup* will search the history table looking for the most recent (highest `to_lsn`) backup in the series and take the `to_lsn` value to use as it's starting lsn. This is mutually exclusive with `-incremental-history-uuid`, `-incremental-basedir` and `-incremental-lsn` options. If no valid LSN can be found (no series by that name) *xtrabackup* will return with an error.
- `-incremental-history-uuid` =<uuid> : Allows an incremental backup to be made based on a specific history record identified by UUID. *xtrabackup* will search the history table looking for the record matching UUID and take the `to_lsn` value to use as it's starting LSN. This options is mutually exclusive with `-incremental-basedir`, `-incremental-lsn` and `-incremental-history-name` options. If no valid LSN can be found (no record by that UUID or missing `to_lsn`), *xtrabackup* will return with an error.

Note

Backup that's currently being performed will **NOT** exist in the `xtrabackup_history` table within the resulting backup set as the record will not be added to that table until after the backup has been taken.

If you want access to backup history outside of your backup set in the case of some catastrophic event, you will need to either perform a `mysqldump`, partial backup or `SELECT *` on the history table after *xtrabackup* completes and store the results with you backup set.

For the necessary privileges, see Permissions and Privileges Needed.

PERCONA_SCHEMA.XTRABACKUP_HISTORY table

This table contains the information about the previous server backups. Information about the backups will only be written if the backup was taken with `-history` option.

Column Name	Description
uuid	Unique backup id
name	User provided name of backup series. There may be multiple entries with the same name used to identify related backups in a series.
tool_name	Name of tool used to take backup
tool_command	Exact command line given to the tool with _password and <code>[-encrypt_key]</code> obfuscated
tool_version	Version of tool used to take backup
ibbackup_version	Version of the xtrabackup binary used to take backup
server_version	Server version on which backup was taken
start_time	Time at the start of the backup
end_time	Time at the end of the backup
lock_time	Amount of time, in seconds, spent calling and holding locks for <code>FLUSH TABLES WITH READ LOCK</code>
binlog_pos	Binary log file and position at end of <code>FLUSH TABLES WITH READ LOCK</code>
innodb_from_lsn	LSN at beginning of backup which can be used to determine prior backups
innodb_to_lsn	LSN at end of backup which can be used as the starting lsn for the next incremental
partial	Is this a partial backup, if <code>N</code> that means that it's the full backup
incremental	Is this an incremental backup
format	Description of result format (<code>xbstream</code>)
compressed	Is this a compressed backup
encrypted	Is this an encrypted backup

Limitations

- [-history](#) option must be specified only on the command line and not within a configuration file in order to be effective.
- [-incremental-history-name](#) and [-incremental-history-uuid](#) options must be specified only on the command line and not within a configuration file in order to be effective.

Take a streaming backup

Percona XtraBackup supports streaming mode. Streaming mode sends a backup to `STDOUT` in the `xbstream` format instead of copying the files to the backup directory.

This method enables you to utilize other programs to filter the backup output, enhancing flexibility in backup storage. For instance, compression can be achieved by directing the output to a compression utility. One advantage of streaming backups and employing Unix pipes is that backups can be automatically encrypted.

Version changes

- Starting with Percona XtraBackup 8.0.33-28, the `--stream` option does not require an `xbstream` parameter. The `xbstream` parameter is optional.
- Using `--encrypt` might create larger backups than expected when used with InnoDB Page Compression.

To avoid this issue with compressed backups, use the `--compress` option with the `--xbstream` option in Percona XtraBackup 8.0.31-24 and later.

Use streaming

To utilize the streaming feature, you need to employ the `--stream` option, specifying the stream format (`xbstream`). Specifying the stream format (`xbstream`) is optional, starting with Percona XtraBackup 8.0.33-28.

```
$ xtrabackup --stream=xbstream
```



`xtrabackup` uses `xbstream` to stream all of the data files to `STDOUT`, in a special `xbstream` format. After all data is streamed, `xtrabackup` stops and also streams the saved transaction log.

When compression is enabled, `xtrabackup` compresses the output data, except for the metadata, using the specified compression algorithm. Read about the supported compression algorithms in the [Create a compressed backup](#) document.

With `xbstream`, backups can be copied and compressed simultaneously, significantly speeding up the process. However, if backups are both compressed and encrypted, they need to be decrypted before uncompressing.

Task	Command
Stream the backup into an archive named backup.xbstream	<code>xtrabackup --backup --stream=xbstream > backup.xbstream</code>
Stream the backup into a compressed archive named backup.xbstream	<code>xtrabackup --backup --stream=xbstream --compress > backup.xbstream</code>
Encrypt the backup	<code>xtrabackup --backup --stream=xbstream \ gzip \ openssl des3 -salt -k 'password' -out backup.xbstream.gz.des3</code>
Unpack the backup to the current directory	<code>xbstream -x < backup.xbstream</code>
Send the backup compressed directly to another host and unpack it	<code>xtrabackup --backup --compress --stream=xbstream ssh user@otherhost "xbstream -x"</code>
Send the backup to another server using netcat	On the destination host: <code>nc -l 9999 cat - > /data/backups/backup.xbstream</code> On the source host: <code>xtrabackup --backup --stream=xbstream nc desthost 9999</code>
Send the backup to another server using a one-liner	<code>ssh user@desthost "(nc -l 9999 > /data/backups/backup.xbstream &)" && xtrabackup --backup --stream=xbstream nc desthost 9999</code>
Throttle the throughput to 10MB/sec using the pipe viewer tool	<code>xtrabackup --backup --stream=xbstream pv -q -L10m ssh user@desthost "cat - > /data/backups/backup.xbstream"</code>
Checksum the backup during the streaming	On the destination host: <code>nc -l 9999 tee >(sha1sum > destination_checksum) > /data/backups/backup.xbstream</code> On the source host: <code>xtrabackup --backup --stream=xbstream tee >(sha1sum > source_checksum) nc desthost 9999</code> Compare the checksums on the source host: <code>cat source_checksum</code> 65e4f916a49c1f216e0887ce54cf59bf3934dbad

Task	Command
	Compare the checksums on the destination host: <code>cat destination_checksum</code> <code>65e4f916a49c1f216e0887ce54cf59bf3934dbad</code>
Parallel compression with parallel copying backup	<code>xtrabackup --backup --compress --compress-threads=8 --stream=xbstream --parallel=4 > backup.xbstream</code>

 Important

The streamed backup must be prepared before restoration. Streaming mode does not prepare the backup.

Accelerate the backup process

Copy with the `--parallel` and `--compress-threads` options

When making a local or streaming backup with *xbstream* option, multiple files can be copied at the same time when using the `--parallel` option. This option specifies the number of threads created by *xtrabackup* to copy data files.

To take advantage of this option either the multiple tablespaces option must be enabled (`innodb_file_per_table`) or the shared tablespace must be stored in multiple *ibdata* files with the `innodb_data_file_path` option. Having multiple files for the database (or splitting one into many) doesn't have a measurable impact on performance.

As this feature is implemented **at the file level**, concurrent file transfer can sometimes increase I/O throughput when doing a backup on highly fragmented data files, due to the overlap of a greater number of random read requests. You should consider tuning the filesystem also to obtain the maximum performance (e.g. checking fragmentation).

If the data is stored on a single file, this option will have no effect.

To use this feature, simply add the option to a local backup, for example:

```
xtrabackup --backup --parallel=4 --target-dir=/path/to/backup
```



By using the *xbstream* in streaming backups, you can additionally speed up the compression process with the `--compress-threads` option. This option specifies the number of threads created by *xtrabackup* for parallel data compression. The default value for this option is 1.

To use this feature, simply add the option to a local backup, for example:

```
xtrabackup --backup --stream=xbstream --compress --compress-threads=4 --target-dir=./ > backup.xbstream
```



Before applying logs, compressed files will need to be uncompressed.

The `--rsync` option

In order to speed up the backup process and to minimize the time `FLUSH TABLES WITH READ LOCK` is blocking the writes, the option `--rsync` should be used. When this option is specified, *xtrabackup* uses `rsync` to copy all non-InnoDB files instead of spawning a separate `cp` for each file, which can be much faster for servers with a large number of databases or tables. *xtrabackup* will call the `rsync` twice, once before the `FLUSH TABLES WITH READ LOCK` and once during to minimize the time the read lock is being held. During the second `rsync` call, it will only synchronize the changes to non-transactional data (if any) since the first call performed before the `FLUSH TABLES WITH READ LOCK`. Note that *Percona XtraBackup* will use [Backup locks](#) where available as a lightweight alternative to `FLUSH TABLES WITH READ LOCK`. This feature is available in *Percona Server for MySQL 5.6+*. *Percona XtraBackup* uses this automatically to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.



Note

This option cannot be used together with the `--stream` option.

Prepare features

Dictionary cache

Percona XtraBackup is based on how [crash recovery](#) works. Percona XtraBackup copies the InnoDB data files, which results in data that is internally inconsistent; then the [prepare](#) phase performs crash recovery on the files to make a consistent, usable database again.

The [_prepare](#) phase has the following operations:

- Applies the redo log
- Applies the undo log

As a physical operation, the changes from the redo log modifications are applied to a page. In the redo log operation, there is no concept of row or transaction. The redo apply operation does not make the database consistent with a transaction. An uncommitted transaction can be flushed or written to the redo log by the server. Percona XtraBackup applies changes recorded in the redo log.

Percona XtraBackup physically modifies a specific offset of a page within a tablespace (IBD file) when applying a redo log.

As a logical operation, Percona XtraBackup applies the undo log on a specific row. When the redo log completes, XtraBackup uses the undo log to roll back changes from uncommitted transactions during the backup.

Undo log

There are two types of undo log records:

- INSERT
- UPDATE

An undo log record contains a [table_id](#). Percona XtraBackup uses this [table_id](#) to find the table definition, and then uses that information to parse the records on an index page. The transaction rollback reads the undo log records and applies the changes.

After initializing the data dictionary engine and the data dictionary cache, the storage engine can request the [table_id](#) and uses this ID to fetch the table schema. An index search tuple (key) is created from the table schema and key fields from the undo log record. The server finds the record using the search tuple (key) and performs the undo operation.

For example, InnoDB uses the `table_id` (also known as the `se_private_id`) for a table definition. Percona XtraBackup does not behave like a server and does not have access to the data dictionary. XtraBackup initializes the InnoDB engine and uses the InnoDB `table` object (`dict_table_t`) when needed. XtraBackup relies on Serialized Dictionary Information (SDI) that is stored in the tablespace. SDI is a JSON representation of a table.

Version changes

Prior to Percona XtraBackup 8.0.33-28, XtraBackup loads the SDI from every `.IBD` file and all tables into cache as non-evictable. Making the tables non-evictable essentially disables the LRU cache. Every table remains in memory until the operation ends.

This method has the following issues:

- Unnecessary IO operations from reading SDI pages to load the tables that are not required for rollback
- All tables loaded into the cache increases the time required in the [--prepare](#) phase
- Unnecessary tables can lead to out-of-memory errors
- Huge number of tables and IBD files may cause an exit in the [--prepare](#) phase

Percona XtraBackup 8.0.33-28 implements a new design and tables are loaded as `evictable`. XtraBackup scans the B-tree index of the data dictionary tables `mysql.indexes` and `mysql.index_partitions` to establish a relationship between the `table_id` and the `tablespace(space_id)`. XtraBackup uses this relationship during transaction rollback. XtraBackup does not load user tables unless there is a transaction rollback on them.

A background thread or a Percona XtraBackup main thread handles the cache eviction when the cache size limit is reached.

This new design provides the following benefits during the [--prepare](#) phase:

- Uses less memory
- Uses less IO
- Improves the time taken in the [--prepare](#) phase
- Completes successfully, even if the [--prepare](#) phase has a huge number of tables
- Completes the Percona XtraDB Cluster SST process faster

Restore features

Point-in-time recovery

Recovering up to particular moment in database's history can be done with xtrabackup and the binary logs of the server.

Note that the binary log contains the operations that modified the database from a point in the past. You must have full data directory as a base. Apply a series of operations from the binary log to make the data match the point in time data.

```
$ xtrabackup --backup --target-dir=/path/to/backup  
$ xtrabackup --prepare --target-dir=/path/to/backup
```

For more details on these procedures, see [Create a full backup](#) and [Prepare a full backup](#).

Now, suppose that some time has passed, and you want to restore the database to a certain point in the past, having in mind that there is the constraint of the point where the snapshot was taken.

To find out what is the situation of binary logging in the server, execute the following queries:

```
mysql> SHOW BINARY LOGS;
```

Expected output

Log_name	File_size
mysql-bin.000001	126
mysql-bin.000002	1306
mysql-bin.000003	126
mysql-bin.000004	497

This query returns the binary log names and file size.

```
mysql> SHOW MASTER STATUS;
```

Expected output

File	Position	Binlog_Do_DB	Binlog_Ignore_DB
mysql-bin.000004	497		

The query returns which file is currently being used to record changes, and the current position within it. Typically, the files are stored in the `datadir`, unless, when starting the server, another location is specified with the `--log-bin=` option).

To find out the position of the snapshot taken, see the `xtrabackup_binlog_info` at the backup's directory:

```
$ cat /path/to/backup/xtrabackup_binlog_info
```

Expected output

mysql-bin.000003	57
------------------	----

This result tells you which file was used at the moment of the backup for the binary log and its position. That position will be the effective one when you restore the backup:

```
$ xtrabackup --copy-back --target-dir=/path/to/backup
```

As the restoration will not affect the binary log files (you may need to adjust file permissions, see [Restore a backup](#)), the next step is extracting the queries from the binary log with **mysqlbinlog** starting from the position of the snapshot and redirecting the results to a file.

If you have multiple files for the binary log, as in the example, you must extract the queries with one process.

```
$ mysqlbinlog /path/to/datadir/mysql-bin.000003 /path/to/datadir/mysql-bin.000004 \
--start-position=57 > mybinlog.sql
```

Inspect the file with the queries to determine which position or date corresponds to the wanted point-in-time. Once determined, pipe it to the server. For example, if the point is `11-12-25 01:00:00`, the command moves the database forward to that point-in-time.

```
$ mysqlbinlog /path/to/datadir/mysql-bin.000003 /path/to/datadir/mysql-bin.000004
\
--start-position=57 --stop-datetime="11-12-25 01:00:00" | mysql -u root -p
```

Encrypted InnoDB tablespace backups

InnoDB supports [data encryption for InnoDB tables](#) stored in file-per-table tablespaces. This feature provides an at-rest encryption for physical tablespace data files.

For an authenticated user or application to access an encrypted tablespace, InnoDB uses the master encryption key to decrypt the tablespace key. The master encryption key is stored in a keyring. *xtrabackup* supports two keyring plugins: `keyring_file`, and `keyring_vault`. These plugins are installed into the `plugin` directory.

Version updates

Percona XtraBackup 8.0.25-17 adds support for the `keyring_file` component, which is part of the component-based infrastructure MySQL which extends the server capabilities. The component is stored in the `plugin` directory.

See a [comparison of keyring components and keyring plugins](#) for more information.

Percona XtraBackup 8.0.27-19 adds support for the Key Management Interoperability Protocol (KMIP) which enables the communication between the key management system and encrypted database server.

Percona XtraBackup 8.0.28-21 adds support for the Amazon Key Management Service (AWS KMS). AWS KMS is cloud-based encryption and key management service. The keys and functionality can be used for other AWS services or your applications that use AWS. No configuration is required to back up a server with AWS KMS-enabled encryption.

Use `keyring_file` plugin

In order to back up and prepare a database containing encrypted InnoDB tablespaces, specify the path to a keyring file as the value of the `--keyring-file-data` option.

```
$ xtrabackup --backup --target-dir=/data/backup/ --user=root \  
--keyring-file-data=/var/lib/mysql-keyring/keyring
```



After *xtrabackup* takes the backup, the following message confirms the action:

Confirmation message

```
xtrabackup: Transaction log of lsn (5696709) to (5696718) was copied.  
160401 10:25:51 completed OK!
```

Warning

xtrabackup does not copy the keyring file into the backup directory. To prepare the backup, you must copy the keyring file manually.

Prepare the backup with the `keyring_file` plugin

To prepare the backup specify the `keyring-file-data`.

```
$ xtrabackup --prepare --target-dir=/data/backup \  
--keyring-file-data=/var/lib/mysql-keyring/keyring
```

After *xtrabackup* takes the backup, the following message confirms the action:

Confirmation message

```
InnoDB: Shutdown completed; log sequence number 5697064  
160401 10:34:28 completed OK!
```

The backup is now prepared and can be restored with the `--copy-back` option. In case the keyring has been rotated, you must restore the keyring used when the backup was taken and prepared.

Use `keyring_vault` plugin

How to use the keyring vault plugin is described [in this document](#).

Create a backup with the `keyring_vault` plugin

The following command creates a backup in the `/data/backup` directory:

```
$ xtrabackup --backup --target-dir=/data/backup --user=root
```

After *xtrabackup* completes the action, the following message confirms the action:

Confirmation message

```
xtrabackup: Transaction log of lsn (5696709) to (5696718) was copied.  
160401 10:25:51 completed OK!
```

Prepare the backup with the `keyring_vault` plugin

To prepare the backup, *xtrabackup* must access the keyring. *xtrabackup* does not communicate with the MySQL server or read the default `my.cnf` configuration file. Specify the keyring settings in the command line:

```
$ xtrabackup --prepare --target-dir=/data/backup \  
--keyring-vault-config=/etc/vault.cnf
```

How to use the keyring vault plugin is described [in this document](#).

After *xtrabackup* completes the action, the following message confirms the action:

Confirmation message

```
InnoDB: Shutdown completed; log sequence number 5697064  
160401 10:34:28 completed OK!
```

The backup is now prepared and can be restored with the `--copy-back` option:

```
$ xtrabackup --copy-back --target-dir=/data/backup --datadir=/data/mysql
```

Use the `keyring_file` component

A component is not loaded with the `--early_plugin_load` option. The server uses a manifest to load the component and the component has its own configuration file. See the [MySQL documentation on the component installation](#) for more information.

An example of a manifest and a configuration file follows:

An example of `./bin/mysqld.my`:


```
{
  "components": "file://component_keyring_file"
}
```

An example of `/lib/plugin/component_keyring_file.cnf`:

```
{
  "path": "/var/lib/mysql-keyring/keyring_file", "read_only": false
}
```

For more information, see [Keyring Component Installation](#) and [Using the keyring_file File-Based Keyring Plugin](#).

With the appropriate privilege, you can `SELECT` on the [performance_schema.keyring_component_status table](#) to view the attributes and status of the installed keyring component when in use.

The component has no special requirements for backing up a database that contains encrypted InnoDB tablespaces.

```
$ xtrabackup --backup --target-dir=/data/backup --user=root
```

After *xtrabackup* completes the action, the following message confirms the action:

Confirmation message

```
xtrabackup: Transaction log of lsn (5696709) to (5696718) was copied.
160401 10:25:51 completed OK!
```

Prepare the backup with `component-keyring-config`

You can copy the `component_keyring_file.cnf` file to the backup directory before starting the `prepare` step or add the `component_keyring_config` option. You must specify the path if the `component-keyring-config` is not in the default location.

```
$ xtrabackup --prepare --target-dir=/data/backup \
--component-keyring-config=/usr/lib64/mysql/plugin/component_keyring_file.cnf
```

After *xtrabackup* completes preparing the backup, the following message confirms the action:



Confirmation message



```
InnoDB: Shutdown completed; log sequence number 5697064  
160401 10:34:28 completed OK!
```

The backup is prepared. To restore the backup use the `--copy-back` option. If the keyring has been rotated, you must restore the specific keyring used to take and prepare the backup.

Incremental encrypted InnoDB tablespace backups with `keyring_file`

The process of taking incremental backups with InnoDB tablespace encryption is similar to taking the Incremental Backups with unencrypted tablespace. The `keyring-file` component should not be used in production or for regulatory compliance.

Create an incremental backup

To make an incremental backup, begin with a full backup. The *xtrabackup* binary writes a file called `xtrabackup_checkpoints` into the backup's target directory. This file contains a line showing the `to_lsn`, which is the database's LSN at the end of the backup. First you need to create a full backup with the following command:

```
$ xtrabackup --backup --target-dir=/data/backups/base \  
--keyring-file-data=/var/lib/mysql-keyring/keyring
```



In order to prepare the backup, you must make a copy of the keyring file yourself. *xtrabackup* does not copy the keyring file into the backup directory. Restoring the backup after the keyring has been changed causes errors like `ERROR 3185 (HY000): Can't find master key from keyring, please check keyring plugin is loaded.` when the restore process tries accessing an encrypted table.

If you look at the `xtrabackup_checkpoints` file, you should see the output similar to the following:

Expected output

```
backup_type = full-backupped
from_lsn = 0
to_lsn = 7666625
last_lsn = 7666634
compact = 0
recover_binlog_info = 1
```

Now that you have a full backup, you can make an incremental backup based on it. Use a command such as the following:

```
$ xtrabackup --backup --target-dir=/data/backups/inc1 \
--incremental-basedir=/data/backups/base \
--keyring-file-data=/var/lib/mysql-keyring/keyring
```

To prepare the backup, you must copy the keyring file manually. xtrabackup does not copy the keyring file into the backup directory.

If the keyring has not been rotated you can use the same as the one you've backed-up with the base backup. If the keyring has been rotated, or you have upgraded the plugin to a component, you must back up the keyring file. Otherwise, you cannot prepare the backup.

The `/data/backups/inc1/` directory should now contain delta files, such as `ibdata1.delta` and `test/table1.ibd.delta`. These represent the changes since the LSN 7666625. If you examine the `xtrabackup_checkpoints` file in this directory, you should see the output similar to the following:

Expected output

```
backup_type = incremental
from_lsn = 7666625
to_lsn = 8873920
last_lsn = 8873929
compact = 0
recover_binlog_info = 1
```

Use this directory as the base for yet another incremental backup:

```
$ xtrabackup --backup --target-dir=/data/backups/inc2 \
--incremental-basedir=/data/backups/inc1 \
--keyring-file-data=/var/lib/mysql-keyring/keyring
```

Prepare incremental backups

The `--prepare` step for incremental backups is not the same as for normal backups. In normal backups, two types of operations are performed to make the database consistent: committed transactions are replayed from the log file against the data files, and uncommitted transactions are rolled back. You must skip the rollback of uncommitted transactions when preparing a backup, because transactions that were uncommitted at the time of your backup may be in progress, and it's likely that they will be committed in the next incremental backup. You should use the `--apply-log-only` option to prevent the rollback phase.

If you do not use the `--apply-log-only` option to prevent the rollback phase, then your incremental backups are useless. After transactions have been rolled back, further incremental backups cannot be applied.

Beginning with the full backup you created, you can prepare it and then apply the incremental differences to it. Recall that you have the following backups:

```
/data/backups/base
/data/backups/inc1
/data/backups/inc2
```

To prepare the base backup, you need to run `--prepare` as usual, but prevent the rollback phase:

```
$ xtrabackup --prepare --apply-log-only --target-dir=/data/backups/base \
--keyring-file-data=/var/lib/mysql-keyring/keyring
```



Expected output



```
InnoDB: Shutdown completed; log sequence number 7666643
InnoDB: Number of pools: 1
160401 12:31:11 completed OK!
```

To apply the first incremental backup to the full backup, you should use the following command:

```
$ xtrabackup --prepare --apply-log-only --target-dir=/data/backups/base \
--incremental-dir=/data/backups/inc1 \
--keyring-file-data=/var/lib/mysql-keyring/keyring
```



Note

If you have many InnoDB Data (IBD) files, speed up the prepare phase for incremental backups by using the `--parallel` option. This option lets you process multiple delta files simultaneously. When using `--parallel` in the prepare phase, always specify a numeric value. The recommended minimum value is 4 (for example, `--parallel=4`).

An example of the backup command using the `--parallel` option:

```
xtrabackup --prepare --apply-log-only --target-dir=/data/backups/base \  
--incremental-dir=/data/backups/inc1 \  
--keyring-file-data=/var/lib/mysql-keyring/keyring --parallel=4
```

Prepare the backup using the same keyring file and type that were used when the backup was created. If the keyring has changed or you upgraded from a plugin to a component between the base and incremental backup, use the keyring from when the first incremental backup was made. If the original keyring is missing or has changed, recover or replace it before restoring. If you cannot recover the keyring, restore from a backup that matches the most recent available keyring.

Preparing the second incremental backup is a similar process: apply the deltas to the (modified) base backup, and you will roll the base backup's data forward in time to the point of the second incremental backup:

```
xtrabackup --prepare --target-dir=/data/backups/base \  
--incremental-dir=/data/backups/inc2 \  
--keyring-file-data=/var/lib/mysql-keyring/keyring
```

You can also use the `--parallel` option here to speed up the process:

```
xtrabackup --prepare --target-dir=/data/backups/base \  
--incremental-dir=/data/backups/inc2 \  
--keyring-file-data=/var/lib/mysql-keyring/keyring --parallel=4
```

Use `--apply-log-only` when merging all incremental backups except the last one. That's why the previous line does not contain the `--apply-log-only` option. Even if the `--apply-log-only` was used on the last step, backup would still be consistent but in that case server would perform the rollback phase.

The backup is now prepared and can be restored with `--copy-back` option. In case the keyring has been rotated you'll need to restore the keyring which was used to take and prepare the backup.

Use the Key Management Interoperability Protocol (KMIP)

Percona XtraBackup 8.0.27-19 adds support for the Key Management Interoperability Protocol (KMIP) which enables the communication between the key management system and encrypted database server.

Percona XtraBackup has no special requirements for backing up a database that contains encrypted InnoDB tablespaces.

Percona XtraBackup performs the following actions:

1. Connects to the MySQL server
2. Pulls the configuration
3. Connects to the KMIP server
4. Fetches the necessary keys from the KMIP server
5. Stores the KMIP server configuration settings in the `xtrabackup_component_keyring_kmip.cnf` file in the backup directory

When preparing the backup, Percona XtraBackup connects to the KMIP server with the settings from the `xtrabackup_component_keyring_kmip.cnf` file.

Restore a backup when the keyring is not available

While the described restore method works, this method requires access to the same keyring that the server is using. It may not be possible if the backup is prepared on a different server or at a much later time, when keys in the keyring are purged, or, in the case of a malfunction, when the keyring vault server is not available at all.

The `--transition-key=<passphrase>` option should be used to make it possible for *xtrabackup* to process the backup without access to the keyring vault server. In this case, *xtrabackup* derives the AES encryption key from the specified passphrase and will use it to encrypt tablespace keys of tablespaces that are being backed up.

Create a backup with a passphrase

The following example illustrates how the backup can be created in this case:

```
xtrabackup --backup --user=root -p --target-dir=/data/backup \  
--transition-key=MySecretKey
```



If `--transition-key` is specified without a value, *xtrabackup* will ask for it.

xtrabackup scrapes `--transition-key` so that its value is not visible in the `ps` command output.

Prepare a backup with a passphrase

The same passphrase should be specified for the prepare command:

```
xtrabackup --prepare --target-dir=/data/backup \  
--transition-key=MySecretKey
```

There are no `--keyring-vault...`, `--keyring-file...`, or `--component-keyring-file-config` options here, because *xtrabackup* does not talk to the keyring in this case.

Restore a backup with a generated key

When restoring a backup you will need to generate a new master key. Here is the example for `keyring_file` plugin or component:

```
xtrabackup --copy-back --target-dir=/data/backup --datadir=/data/mysql \  
--transition-key=MySecretKey --generate-new-master-key \  
--keyring-file-data=/var/lib/mysql-keyring/keyring
```

In case of `keyring_vault`, it will look like this:

```
xtrabackup --copy-back --target-dir=/data/backup --datadir=/data/mysql \  
--transition-key=MySecretKey --generate-new-master-key \  
--keyring-vault-config=/etc/vault.cnf
```

xtrabackup will generate a new master key, store it in the target keyring vault server and re-encrypt the tablespace keys using this key.

Make a backup with a stored transition key

Finally, there is an option to store a transition key in the keyring. In this case, *xtrabackup* will need to access the same keyring file or vault server during prepare and copy-back but does not depend on whether the server keys have been purged.

In this scenario, the three stages of the backup process look as follows.

- Backup

```
xtrabackup --backup --user=root -p --target-dir=/data/backup \  
--generate-transition-key
```

- Prepare

- `keyring_file` variant:

```
xtrabackup --prepare --target-dir=/data/backup \  
--keyring-file-data=/var/lib/mysql-keyring/keyring
```



- `keyring_vault` variant:

```
xtrabackup --prepare --target-dir=/data/backup \  
--keyring-vault-config=/etc/vault.cnf
```



- Copy-back

- `keyring_file` variant:

```
xtrabackup --copy-back --target-dir=/data/backup --datadir=/data/mysql \  
--generate-new-master-key --keyring-file-data=/var/lib/mysql-keyring/keyring
```



- `keyring_vault` variant:

```
xtrabackup --copy-back --target-dir=/data/backup --datadir=/data/mysql \  
--generate-new-master-key --keyring-vault-config=/etc/vault.cnf
```



FLUSH TABLES WITH READ LOCK option

The `FLUSH TABLES WITH READ LOCK` option does the following with a global read lock:

- Closes all open tables
- Locks all tables for all databases

Release the lock with `UNLOCK TABLES`.

Note

`FLUSH TABLES WITH READ LOCK` does not prevent inserting rows into the log tables.

To ensure consistent backups, use the `FLUSH TABLES WITH READ LOCK` option before taking a non-InnoDB file backup. The option does not affect long-running queries.

Long-running queries with `FLUSH TABLES WITH READ LOCK` enabled can leave the server in a read-only mode until the queries finish. Killing the `FLUSH TABLES WITH READ LOCK` does not help if the database is in either the `Waiting for table flush` or `Waiting for master to send event` state. To return to normal operation, you must kill any long-running queries.

Note

All described in this section has no effect when backup locks are used. *Percona XtraBackup* will use [Backup locks](#) where available as a lightweight alternative to `FLUSH TABLES WITH READ LOCK`. This feature is available in *Percona Server for MySQL 5.6+*. *Percona XtraBackup* uses this automatically to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.

In order to prevent this from happening two things have been implemented:

- *xtrabackup* waits for a good moment to issue the global lock
- *xtrabackup* kills all queries or only the `SELECT` queries which prevent the global lock from being acquired

Wait for queries to finish

You should issue a global lock when no long queries are running. Waiting to issue the global lock for extended period of time is not a good method. The wait can extend the time needed for backup to take place. The `[-ftwrl-wait-timeout]` option can limit the waiting time. If it cannot issue the lock during this time, *xtrabackup* stops the option, exits with an error message, and backup is not be taken.

The default value for this option is zero (0) value which turns off the option.

Another possibility is to specify the type of query to wait on. In this case `-ftwrl-wait-query-type`. Possible values are `all` and `update`. When `all` is used *xtrabackup* will wait for all long running queries (execution time longer than allowed by `-ftwrl-wait-threshold`) to finish before running the `FLUSH TABLES WITH READ LOCK`. When `update` is used *xtrabackup* will wait on `UPDATE/ALTER/REPLACE/INSERT` queries to finish.

The time needed for a specific query to complete is hard to predict. We assume that the long-running queries will not finish in a timely manner. Other queries which run for a short time finish quickly. *xtrabackup* uses the value of the `[-ftwrl-wait-threshold]` option to specify the long-running queries and will block a global lock. In order to use this option *xtrabackup* user should have `PROCESS` and `SUPER` privileges.

Kill the blocking queries

The second option is to kill all the queries which prevent from acquiring the global lock. In this case, all queries which run longer than `FLUSH TABLES WITH READ LOCK` are potential blockers. Although all queries can be killed, additional time can be specified for the short running queries to finish using the `-kill-long-queries-timeout` option. This option specifies the time for queries to complete, after the value is reached, all the running queries will be killed. The default value is zero, which turns this feature off.

The `-kill-long-query-type` option can be used to specify all or only `SELECT` queries that are preventing global lock from being acquired. In order to use this option *xtrabackup* user should have `PROCESS` and `SUPER` privileges.

Options summary

- `-ftwrl-wait-timeout` (seconds) - how long to wait for a good moment. Default is 0, not to wait.
- `-ftwrl-wait-query-type` - which long queries should be finished before `FLUSH TABLES WITH READ LOCK` is run. Default is `all`.
- `-ftwrl-wait-threshold` (seconds) - how long query should be running before we consider it long running and potential blocker of global lock.

- [--kill-long-queries-timeout](#) (seconds) - how many time we give for queries to complete after `FLUSH TABLES WITH READ LOCK` is issued before start to kill. Default if `0`, not to kill.
- [--kill-long-query-type](#) - which queries should be killed once [--kill-long-queries-timeout](#) has expired.

Example

Running the *xtrabackup* with the following options will cause *xtrabackup* to spend no longer than 3 minutes waiting for all queries older than 40 seconds to complete.

```
$ xtrabackup --backup --ftwrl-wait-threshold=40 \  
--ftwrl-wait-query-type=all --ftwrl-wait-timeout=180 \  
--kill-long-queries-timeout=20 --kill-long-query-type=all \  
--target-dir=/data/backups/
```



After `FLUSH TABLES WITH READ LOCK` is issued, *xtrabackup* will wait for 20 seconds for lock to be acquired. If lock is still not acquired after 20 seconds, it will kill all queries which are running longer than the `FLUSH TABLES WITH READ LOCK`.

Improved log statements

Percona XtraBackup is an open-source command-line utility. Command-line tools have limited interaction with the background operations and the logs provide the progress of an operation or more information about errors.

In earlier versions, the error logs did not have a standard structure. Notice in the following examples the variance in the log statements:

- The backup log statement header has the name of the module, `xtrabackup`, which generated the statement but no timestamp:

```
$ xtrabackup: recognized client arguments: --parallel=4 --target-dir=/data/backups/ --backup=1
```

The output should be similar to the following:

Expected output

```
./bin/xtrabackup version 8.0.27-19 based on MySQL server 8.0.27 Linux (x86_64) (revision id: b0f75188ca3)
```

- The copy-back log statement has a timestamp but no module name. The timestamp is a mix of UTC and the local timezone.

```
220322 19:05:13 [01] Copying undo_001 to /data/backups/undo_001
```

- The following prepare log statements do not have header information, which makes diagnosing an issue more difficult.

```
Completed space ID check of 1008 files.  
Initializing buffer pool, total size = 128.000000M, instances = 1, chunk size  
=128.000000M  
Completed initialization of buffer pool  
If the mysqld execution user is authorized, page cleaner thread priority can be  
changed. See the man page of setpriority().
```

Log statement structure

Starting in **Percona XtraBackup** 8.0.28-20, changes have been made to improve the log statements. The improved log structure is displayed in the backup, prepare, move-back/copy-back error logs.

Each log statement has the following attributes:

- **Timestamp** - a timestamp for when the event occurred in a UTC format.
- **Severity** - the severity level of a statement indicates the importance of an event.
- **ID** - this identifier is currently not used but may be used in future versions.
- **Context** - the name of the module that issued the log statement, such as **XtraBackup**, **InnoDB**, or **Server**.
- **Message** - a description of the event generated by the module.

An example of a `prepare` log that is generated with the improved structure. The uniformity of the headers makes it easier to follow an operation's progress or review the log to diagnose issues.

Expected output

```
2022-03-22T19:15:36.142247+05:30 0 [Note] [MY-011825] [Xtrabackup] This target seems to be
not prepared yet.
2022-03-22T19:15:36.142792+05:30 0 [Note] [MY-013251] [InnoDB] Number of pools: 1
2022-03-22T19:15:36.149212+05:30 0 [Note] [MY-011825] [Xtrabackup] xtrabackup_logfile
detected: size=8388608, start_lsn=(33311656)
2022-03-22T19:15:36.149998+05:30 0 [Note] [MY-011825] [Xtrabackup] using the following InnoDB
configuration for recovery:
2022-03-22T19:15:36.150023+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_data_home_dir = .
2022-03-22T19:15:36.150036+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_data_file_path =
ibdata1:12M:autoextend
2022-03-22T19:15:36.150078+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_log_group_home_dir
= .
2022-03-22T19:15:36.150095+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_log_files_in_group
= 1
2022-03-22T19:15:36.150111+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_log_file_size =
8388608
2022-03-22T19:15:36.151667+05:30 0 [Note] [MY-011825] [Xtrabackup]
inititalize_service_handles succeeded
2022-03-22T19:15:36.151903+05:30 0 [Note] [MY-011825] [Xtrabackup] using the following InnoDB
configuration for recovery:
2022-03-22T19:15:36.151926+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_data_home_dir = .
2022-03-22T19:15:36.151954+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_data_file_path =
ibdata1:12M:autoextend
2022-03-22T19:15:36.151976+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_log_group_home_dir
= .
2022-03-22T19:15:36.151991+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_log_files_in_group
= 1
2022-03-22T19:15:36.152004+05:30 0 [Note] [MY-011825] [Xtrabackup] innodb_log_file_size =
8388608
2022-03-22T19:15:36.152021+05:30 0 [Note] [MY-011825] [Xtrabackup] Starting InnoDB instance
for recovery.
2022-03-22T19:15:36.152035+05:30 0 [Note] [MY-011825] [Xtrabackup] Using 104857600 bytes for
buffer pool (set by --use-memory parameter)
```

lock-ddl-per-table option improvements

Important

The lock-ddl-per-table option is deprecated in Percona Server for MySQL 8.0. Use `--lock-ddl` instead of this variable.

To block DDL statements on an instance, Percona Server implemented LOCK TABLES FOR BACKUP. Percona XtraBackup uses this lock for the duration of the backup. This lock does not affect DML statements.

Percona XtraBackup has also implemented [--lock-ddl-per-table](#), which blocks DDL statements by using metadata locks (MDL).

The following procedures describe a simplified backup operation when using [--lock-ddl-per-table](#):

1. Parse and copy all redo logs after the checkpoint mark
2. Fork a dedicated thread to continue following new redo log entries
3. List the tablespaces required to copy
4. Iterate through the list. The following steps occur with each listed tablespace:
5. Query INFORMATION_SCHEMA.INNODB_TABLES to find which tables belong to the tablespace ID and take an MDL on the underlying table or tables in case there is a shared tablespace.
6. Copy the tablespace `.ibd` files.

The backup process may encounter a redo log event, generated by the bulk load operations, which notifies backup tools that data file changes have been omitted from the redo log. This event is an `MLOG_INDEX_LOAD`. If this event is found by the redo follow thread, the backup continues and assumes the backup is safe because the MDL protects tablespaces already copied and the `MLOG_INDEX_LOAD` event is for a tablespace that is not copied.

These assumptions may not be correct and may lead to inconsistent backups.

[--lock-ddl-per-table](#) redesign

Implemented in Percona XtraBackup version 8.0.22-15.0, the [--lock-ddl-per-table](#) has been redesigned to minimize inconsistent backups.

The following procedure reorders the steps:

- The MDL lock acquired at the beginning of the backup
- Scan the redo logs. An `MLOG_INDEX_LOAD` event may be recorded if a `CREATE INDEX` statement has occurred before the backup starts. At this time, the backup process is safe and can parse and accept the event.
- After the first scan has completed, the follow redo log thread is initiated. This thread stops the backup process if an `MLOG_INDEX_LOAD` event is found.
- Gather the tablespace list to copy
- Copy the `.ibd` files.

Other improvements

The following improvements have been added:

- If the `.ibd` file belongs to a temporary table, the `SELECT` query is skipped.
- For a FullText Index, an MDL is acquired on the base table.
- A `SELECT` query that acquires an MDL does not retrieve any data.

InnoDB redo log archiving

A physical backup contains raw copies of the database contents and any related files, such as configuration files or logs. During crash recovery, InnoDB corrects the data from incomplete transactions using the redo log. A redo log contains the entries necessary to bring the database back to a consistent state. An LSN value represents the data in the redo log.

Data in the redo log is always appended, and the movement of the checkpoint truncates old data. When a file is full, a new file is created and the movement of the checkpoint truncates the old files. If the MySQL server has significant activity during a backup operation and the redo log file storage media operates faster than the backup storage media, the backup operation may fail to keep pace with the redo log generation and cause data loss.

Before the redo log archiving feature, the solutions were:

- Increasing the redo log size
- Checking for I/O congestion if the read speed is slow
- Checking for I/O or network congestion if the write speed is slow
- Taking the backup at an off-peak time

redo log archiving

MySQL 8.0.17 added redo log archiving. To enable this feature, the following actions are required:

- Create a directory for the archiving logs
- Call the function to start the archiving process

It would help if you created a directory for the redo log archives. The system user running `mysqld` must have access to this directory.

You can verify if MySQL has configured the archive directory with the following command. The result provides the path or returns NULL, meaning you have not set the archiving directory.

```
$ mkdir -p /var/lib/mysql-redo-archive/backup1
$ chown mysql. -R /var/lib/mysql-redo-archive
$ chmod -R 700 /var/lib/mysql-redo-archive/
```



The `innodb_redo_log_archive_dirs` global variable contains labeled directories. The format is `label:directory-path`. Multiple objects can be listed with a semi-colon between each set.

```
mysql> SET PERSIST innodb_redo_log_archive_dirs='backup1:/var/lib/mysql-redo-  
archive/';
```

Verify that the archive was created with the following command:

```
mysql> SHOW GLOBAL VARIABLES LIKE 'innodb_redo_log_ar%';
```

Expected output

```
+-----+-----+  
| Variable_name | Value |  
+-----+-----+  
| innodb_redo_log_archive_dirs | backup1:/var/lib/mysql-redo-archive/ |  
+-----+-----+  
1 row in set (0.0025 sec)
```

Users with the [INNODB REDO LOG ARCHIVE](#) privilege can call the `innodb_redo_log_archive_start()` or `innodb_redo_log_archive_stop()` functions. to either start or stop archiving.

```
mysql> SELECT innodb_redo_log_archive_start('backup1','backup1');
```

Expected output

```
+-----+-----+  
| innodb_redo_log_archive_start('backup1','backup1') |  
+-----+-----+  
| 0 |  
+-----+-----+  
1 row in set (0.00225 sec)
```

Percona XtraBackup and redo log archiving

Percona XtraBackup 8.0.34-29 added the [redo-log-arch-dir](#) option. If the archiving directory is not configured in MySQL, Percona XtraBackup creates a temporary directory for the archive file. To use this option, you must run Percona XtraBackup as the same owner as `mysqld`. If Percona XtraBackup does not have the appropriate permissions, the action generates a “Permission denied” error, and archiving is not used.

Example of reading a log

```
$ sudo -H -u mysql bash
$ xtrabackup --no-lock=1 --compress --parallel=4 --host=localhost --user=root --
password='password_string' --backup=1 --target-dir=/Backup/13oct 2> /tmp/b0-with-
redo-archiving-as-mysql-os-user.log
```

Verify that the archiving is used:

```
$ cat /tmp/b0-with-redo-archiving-as-mysql-os-user.log
```

Expected output

```
[Note] [MY-011825] [Xtrabackup] recognized server arguments: --datadir=/var/lib/mysql
[Note] [MY-011825] [Xtrabackup] recognized client arguments: --no-lock=1 --compress --
parallel=4 --host=localhost --user=root --password=* --backup=1 --target-dir=/Backup/22Aug
xtrabackup version 8.0.34-29 based on MySQL server 8.0.34 Linux (x86_64) (revision id:
5ba706ee) 230822 13:36:02 version_check Connecting to MySQL server with DSN
'dbi:mysql;;mysql_read_default_group=xtrabackup;host=localhost' as 'root' (using password:
YES). 230822 13:36:02 version_check Connected to MySQL server 230822 13:36:02 version_check
Executing a version check against the server... 230822 13:36:02 version_check Done.
[Note] [MY-011825] [Xtrabackup] Connecting to MySQL server host: localhost, user: root,
password: set, port: not set, socket: not set
[Note] [MY-011825] [Xtrabackup] Using server version 8.0.34
....
....
[Note] [MY-011825] [Xtrabackup] Connecting to MySQL server host: localhost, user: root,
password: set, port: not set, socket: not set
[Note] [MY-011825] [Xtrabackup] xtrabackup redo_log_arch_dir is set to
backup1:/var/lib/mysql-redo-archive/
[Note] [MY-011825] [Xtrabackup] Waiting for archive file '/var/lib/mysql-redo-
archive//1692711362986/archive.94f4ab58-3d1c-11ee-9e4f-0af1fd7c44c9.000001.log'
[Note] [MY-011825] [Xtrabackup] >> log scanned up to (19630018)
....
....
[Note] [MY-011825] [Xtrabackup] Done: Compressing file /Backup/22Aug/xtrabackup_info.zst
[Note] [MY-011825] [Xtrabackup] Transaction log of lsn (19630018) to (19630028) was copied.
[Note] [MY-011825] [Xtrabackup] completed OK!
```

Smart memory estimation

The Smart memory estimation is [tech preview](#) feature. Before using Smart memory estimation in production, we recommend that you test restoring production from physical backups in your environment and also use the alternative backup method for redundancy.

[Percona XtraBackup 8.0.30-23](#) adds support for the Smart memory estimation feature. With this feature, Percona XtraBackup computes the memory required for `prepare` phase, while copying redo log entries during the `backup` phase. Percona XtraBackup also considers the number of InnoDB pages to be fetched from the disk.

Percona XtraBackup performs the backup procedure in two steps:

- Creates a backup

To create a backup, Percona XtraBackup copies your InnoDB data files. While copying the files, Percona XtraBackup runs a background process that watches the InnoDB redo log, also called the transaction log, and copies changes from it.

- Prepares a backup

During the `prepare` phase, Percona XtraBackup performs crash recovery against the copied data files using the copied transaction log file. Percona XtraBackup reads all the redo log entries into memory, categorizes them by space id and page id, reads the relevant pages into memory, and checks the log sequence number (LSN) on the page and on the redo log record. If the redo log LSN is more recent than the page LSN, Percona XtraBackup applies the redo log changes to the page.

To `prepare` a backup, Percona Xtrabackup uses InnoDB Buffer Pool memory. Percona Xtrabackup reserves memory to load 256 pages into the buffer pool. The remaining memory is used for hashing/categorizing the redo log entries.

The available memory is controlled by the [-use-memory](#) option. If the available memory on the buffer pool is insufficient, the work is performed in multiple batches. After the batch is processed, the memory is freed to release space for the next batch. This process greatly impacts performance as an InnoDB page holds data from multiple rows. If a change on a page happens in different batches, that page is fetched and evicted numerous times.

How does Smart memory estimation work

In the `prepare` phase, Percona XtraBackup checks the server's available free memory and uses that memory up to the limit specified in the `--use-free-memory-pct` option to run `--prepare`. Due to backward compatibility, the default value for the `--use-free-memory-pct` option is 0 (zero), which defines the option as disabled. For example, if you set `--use-free-memory-pct=50`, then 50% of the free memory is used to `prepare` a backup.

Starting with Percona XtraBackup 8.0.32-26, you can enable or disable the memory estimation during the `backup` phase with the `--estimate-memory` option. The default value is `OFF`. Enable the memory estimation with `--estimate-memory=ON`:

```
$ xtrabackup --backup --estimate-memory=ON --target-dir=/data/backups/
```

In the `prepare` phase, enable the `--use-free-memory-pct` option by specifying the percentage of free memory to be used to `prepare` a backup. The `--use-free-memory-pct` value must be larger than 0.

For example:

```
$ xtrabackup --prepare --use-free-memory-pct=50 --target-dir=/data/backups/
```

Example of Smart memory estimation usage

The examples of how Smart memory estimation can improve the time spent on `prepare` in different versions of Percona XtraBackup:

In Percona XtraBackup 8.0.32-26 and higher

We back up 16, 32, and 64 tables using sysbench. Each set contains 1M rows. In the `backup` phase, we enable Smart memory estimation with `--estimate-memory=ON`. In the `prepare` phase, we set `--use-free-memory-pct=50`, and Percona XtraBackup uses 50% of the free memory to prepare a backup. The backup is run on an ec2 c4.8xlarge instance (36 vCPUs / 60GB memory / General Purpose SSD (gp2)).

In Percona XtraBackup 8.0.30-23 and higher

We back up 16, 32, and 64 tables using sysbench. Each set contains 1M rows. In the `prepare` phase, we set `--use-free-memory-pct=50`, and Percona XtraBackup uses 50% of the free memory to prepare a backup. The backup is run on an ec2 c4.8xlarge instance (36 vCPUs / 60GB memory / General Purpose SSD (gp2)).

During each `--backup`, the following sysbench is run:

```
sysbench --db-driver=mysql --db-ps-mode=disable --mysql-user=sysbench --mysql-  
password=sysbench --table_size=1000000 --tables=${NUM_OF_TABLES} --threads=24 --  
time=0 --report-interval=1 /usr/share/sysbench/oltp_write_only.lua run
```



The following table shows the backup details (all measurements are in Gigabytes):

	Used memory	Size of XtraBackup log	Size of backup
16 tables	3.375	0.7	4.7
32 tables	8.625	2.6	11
64 tables	18.5	5.6	22


- **Used memory** - the amount of memory required by Percona XtraBackup with `--use-free-memory-pct=50`
- **Size of XtraBackup log** - the size of Percona XtraBackup log file (redo log entries copied during the backup)
- **Size of backup** - the size of the resulting backup folder

`Prepare` executed without Smart memory estimation uses the default of 128MB for the buffer pool.

The results are the following:

Note

The following results are based on tests in a specific environment. Your results may vary.



- **16 tables result** - prepare time dropped to ~5.7% of the original time. An improvement in recovery time of about 17x.
- **32 tables result** - prepare time dropped to ~8,2% of the original time. An improvement in recovery time of about 12x.
- **64 tables result** - prepare time dropped to ~9.9% of the original time. An improvement in recovery time of about 10x.

Work with binary logs

The `xtrabackup` binary integrates with the `log_status` table. This integration enables `xtrabackup` to print out the backup's corresponding binary log position, so that you can use this binary log position to provision a new replica or perform point-in-time recovery.

Find the binary log position

You can find the binary log position corresponding to a backup after the backup has been taken. If your backup is from a server with binary logging enabled, `xtrabackup` creates a file named `xtrabackup_binlog_info` in the target directory. This file contains the binary log file name and position of the exact point when the backup was taken.

Expected output during the backup stage

```
210715 14:14:59 Backup created in directory '/backup/'
MySQL binlog position: filename 'binlog.000002', position '156'
. . .
210715 14:15:00 completed OK!
```

Note

As of Percona XtraBackup 8.0.26-18.0, `xtrabackup` no longer creates the `xtrabackup_binlog_pos_innodb` file. This change is because MySQL and Percona Server no longer update the binary log information on global transaction system section of `ibdata`. You should rely on `xtrabackup_binlog_info` regardless of the storage engine in use.

Point-in-time recovery

To perform a point-in-time recovery from an `xtrabackup` backup, you should prepare and restore the backup, and then replay binary logs from the point shown in the `xtrabackup_binlog_info` file.

Find a more detailed procedure in the [Point-in-time recovery](#) document.

Set up a new replication replica

To set up a new replica, you should prepare the backup, and restore it to the data directory of your new replication replica. If you are using version 8.0.22 or earlier, in your `CHANGE MASTER TO` command, use the binary log filename and position shown in the `xtrabackup_binlog_info` file to start replication.

If you are using 8.0.23 or later, use the [CHANGE_REPLICATION_SOURCE_TO](#) and the [appropriate options](#).
`CHANGE_MASTER_TO` is deprecated.

A more detailed procedure is found in [How to setup a replica for replication in 6 simple steps with Percona XtraBackup](#).

Analyze table statistics

Note

As of *Percona XtraBackup* 8.0.30 the `--stats` option is deprecated and may be removed in future releases. We no longer support using the `--stats` option in *Percona XtraBackup* 8.0.29 and older versions.

The *xtrabackup* binary can analyze InnoDB data files in read-only mode to give statistics about them. To do this, you should use the `--stats` option. You can combine this with the `--tables` option to limit the files to examine. It also uses the `--use-memory` option.

You can perform the analysis on a running server, with some chance of errors due to the data being changed during analysis. Or, you can analyze a backup copy of the database. Either way, to use the statistics feature, you need a clean copy of the database including correctly sized log files, so you need to execute with `--prepare` twice to use this functionality on a backup.

The result of running on a backup might look like the following:

Expected output

```
<INDEX STATISTICS>
table: test/table1, index: PRIMARY, space id: 12, root page 3
estimated statistics in dictionary:
  key vals: 25265338, leaf pages 497839, size pages 498304
real statistics:
  level 2 pages: pages=1, data=5395 bytes, data/pages=32%
  level 1 pages: pages=415, data=6471907 bytes, data/pages=95%
    leaf pages: recs=25958413, pages=497839, data=7492026403 bytes, data/pages=91%
```

This can be interpreted as follows:

- The first line simply shows the table and index name and its internal identifiers. If you see an index named `GEN_CLUST_INDEX`, that is the table's clustered index, automatically created because you did not explicitly create a `PRIMARY KEY`.
- The estimated statistics in dictionary information is similar to the data that's gathered through `ANALYZE TABLE` inside of *InnoDB* to be stored as estimated cardinality statistics and passed to the query optimizer.
- The real statistics information is the result of scanning the data pages and computing exact information about the index.

- The `level <X> pages` : output means that the line shows information about pages at that level in the index tree. The larger `<X>` is, the farther it is from the leaf pages, which are level 0. The first line is the root page.
- The `leaf pages` output shows the leaf pages, of course. This is where the table's data is stored.
- The `external pages` : output (not shown) shows large external pages that hold values too long to fit in the row itself, such as long `BLOB` and `TEXT` values.
- The `recs` is the real number of records (rows) in leaf pages.
- The `pages` is the page count.
- The `data` is the total size of the data in the pages, in bytes.
- The `data/pages` is calculated as $(\text{data} / (\text{pages} * \text{PAGE_SIZE})) * 100\%$. It will never reach 100% because of space reserved for page headers and footers.

A more detailed example is posted as a MySQL Performance Blog [post](#).

Script to format output

The following script can be used to summarize and tabulate the output of the statistics information:

Expected output

```
tabulate-xtrabackup-stats.pl

#!/usr/bin/env perl
use strict;
use warnings FATAL => 'all';
my $script_version = "0.1";

my $PG_SIZE = 16_384; # InnoDB defaults to 16k pages, change if needed.
my ($cur_idx, $cur_tbl);
my (%idx_stats, %tbl_stats);
my ($max_tbl_len, $max_idx_len) = (0, 0);
while ( my $line = <> ) {
    if ( my ($t, $i) = $line =~ m/table: (.*), index: (.*), space id:/ ) {
        $t =~ s/!/.!;
        $cur_tbl = $t;
        $cur_idx = $i;
        if ( length($i) > $max_idx_len ) {
            $max_idx_len = length($i);
        }
        if ( length($t) > $max_tbl_len ) {
            $max_tbl_len = length($t);
        }
    }
    elsif ( my ($kv, $lp, $sp) = $line =~ m/key vals: (\d+), \D*(\d+), \D*(\d+)/ ) {
        @{$idx_stats{$cur_tbl}->{$cur_idx}}{qw(est_kv est_lp est_sp)} = ($kv, $lp, $sp);
        $tbl_stats{$cur_tbl}->{est_kv} += $kv;
        $tbl_stats{$cur_tbl}->{est_lp} += $lp;
        $tbl_stats{$cur_tbl}->{est_sp} += $sp;
    }
    elsif ( my ($l, $pages, $bytes) = $line =~ m/(?::level (\d+)|leaf) pages:.*pages=(\d+),
data=(\d+) bytes/ ) {
        $l ||= 0;
        $idx_stats{$cur_tbl}->{$cur_idx}->{real_pages} += $pages;
        $idx_stats{$cur_tbl}->{$cur_idx}->{real_bytes} += $bytes;
        $tbl_stats{$cur_tbl}->{real_pages} += $pages;
        $tbl_stats{$cur_tbl}->{real_bytes} += $bytes;
    }
}

my $hdr_fmt = "%${max_tbl_len}s ${max_idx_len}s %9s %10s %10s\n";
my @headers = qw(TABLE INDEX TOT_PAGES FREE_PAGES PCT_FULL);
printf $hdr_fmt, @headers;

my $row_fmt = "%${max_tbl_len}s ${max_idx_len}s %9d %10d %9.1f%%\n";
foreach my $t ( sort keys %tbl_stats ) {
    my $tbl = $tbl_stats{$t};
    printf $row_fmt, $t, "", $tbl->{est_sp}, $tbl->{est_sp} - $tbl->{real_pages},
        $tbl->{real_bytes} / ($tbl->{real_pages} * $PG_SIZE) * 100;
    foreach my $i ( sort keys %{$idx_stats{$t}} ) {
        my $idx = $idx_stats{$t}->{$i};
        printf $row_fmt, $t, $i, $idx->{est_sp}, $idx->{est_sp} - $idx->{real_pages},
            $idx->{real_bytes} / ($idx->{real_pages} * $PG_SIZE) * 100;
    }
}
```

Sample script output

The output of the above Perl script, when run against the sample shown in the previously mentioned blog post, will appear as follows:

Expected output

TABLE	INDEX	TOT_PAGES	FREE_PAGES	PCT_FULL
art.link_out104		832383	38561	86.8%
art.link_out104	PRIMARY	498304	49	91.9%
art.link_out104	domain_id	49600	6230	76.9%
art.link_out104	domain_id_2	26495	3339	89.1%
art.link_out104	from_message_id	28160	142	96.3%
art.link_out104	from_site_id	38848	4874	79.4%
art.link_out104	revert_domain	153984	19276	71.4%
art.link_out104	site_message	36992	4651	83.4%

The columns are the table and index, followed by the total number of pages in that index, the number of pages not actually occupied by data, and the number of bytes of real data as a percentage of the total size of the pages of real data. The first line in the above output, in which the INDEX column is empty, is a summary of the entire table.

xbcloud features

FIFO data sink

The feature is in [tech preview](#).

[Percona XtraBackup 8.0.33-28](#) implements a new data sink that uses the first in, first out (FIFO) method. With the `FIFO` data sink feature, users with a streaming capacity of 10Gbps (typically on a Local Area Network (LAN)) can benefit from faster backups by streaming data in parallel to an object storage.

FIFO data sink options

[Percona XtraBackup 8.0.33-28](#) implements the following options:

- `--fifo-streams=#` - specifies the number of FIFO files to use for parallel data stream. To disable FIFO data sink and send stream to STDOUT, set `--fifo-streams=1`. The default value is `1` (STDOUT) to ensure the backward compatibility. The `--fifo-streams` value must match on both the XtraBackup and xbccloud sides.
- `--fifo-dir=name` - specifies a directory to write Named Pipe.
- `--fifo-timeout=#` - specifies the number of seconds to wait for the other end to open the stream for reading. The default value is `60` seconds.

How to enable FIFO data sink

To use FIFO data sink, you can either run two commands in separate terminal sessions or run xtrabackup in the background.

For example, run the following commands in separate terminal sessions:

```
$ xtrabackup --backup --stream --fifo-streams=2 --fifo-dir=/tmp/fifo
```

```
$ xbccloud put --fifo-streams=2 --fifo-dir=/tmp/fifo full
```

Run xtrabackup in the background with the following commands:


```
$ xtrabackup --backup --stream --fifo-streams=2 --fifo-dir=/tmp/fifo &  
$ xbccloud put --fifo-streams=2 --fifo-dir=/tmp/fifo full
```

Stream to an object storage

When taking a backup, you can save the files locally or stream the files to either a different server or an object storage.

Before [Percona XtraBackup 8.0.33-28](#), XtraBackup streams data to an object storage writing to STDOUT (a pipe) and using [xbcloud](#) to read from STDIN (standard input device).


XtraBackup spawns multiple copy threads and each copy thread reads a data chunk from a specific file. Then each copy thread writes the data chunks to a pipe (STDOUT). An xbcloud read thread reads each chunk of STDIN data. Then xbcloud uploads the chunks to an object storage using an async request, and adds a callback to an event handler list. The xbcloud event handler executes the callback depending on the response from the object storage (success or failure).



The streaming capacity for XtraBackup using STDOUT is 1.8G. This capacity is sufficient for streaming data using Wide Area Network (WAN) into, for example, Amazon Web Services or Google Cloud Platform.

Starting with [Percona XtraBackup 8.0.33-28](#), when you stream backups to Amazon S3 compatible storage using LAN with a streaming capacity of 10Gbps, XtraBackup can use multiple FIFO streams to stream the backups faster.

XtraBackup spawns multiple copy threads and each copy thread reads a data chunk from a specific file. Then multiple FIFO files are created to store the data from XtraBackup. Each XtraBackup copy thread writes the data chunks to a specific FIFO file. Xbcloud reads from the FIFO streams and uploads data to an object storage using an async request. The xbcloud event handler executes the callback depending on the response from the object storage (success or failure).



Performance improvement

Consider an example of using a FIFO data sink compared to the STDOUT method.

The database has 1TB of data in multiple tables. The link speed between the source server and destination server using MinIO is ~ 9.2 Gbps.

Both STDOUT and FIFO data sink scenarios push 1TB of data from two servers.

For the FIFO data sink we configure 8 parallel streams with `--fifo-streams=8` both for XtraBackup and xbcloud.

The results are the following:

- The **STDOUT** method takes 01:25:24 to push 1TB of data using 239 MBps (1.8 Gbps).
- The **FIFO** method, using 8 streams, takes 00:16:01 to push 1TB of data using 1.15 GBps (9.2 Gbps).

Exponential backoff

This feature was implemented in [Percona XtraBackup 8.0.26-18.0](#) in the xbcloud binary.

Exponential backoff increases the chances for the completion of a backup or a restore operation. For example, a chunk upload or download may fail if you have an unstable network connection or other network issues. This feature adds an exponential backoff, or sleep, time and then retries the upload or download.

When a chunk upload or download operation fails, xbcloud checks the reason for the failure. This failure can be a CURL error or an HTTP error, or a client-specific error. If the error is listed in the [Retriable errors](#) list, xbcloud pauses for a calculated time before retrying the operation until that time reaches the `--max-backoff` value.

The operation is retried until the `--max-retries` value is reached. If the chunk operation fails on the last retry, xbcloud aborts the process.

The default values are the following:

- `--max-backoff` = 300000 (5 minutes)
- `--max-retries` = 10

You can adjust the number of retries by adding the `--max-retries` parameter and adjust the maximum length of time between retries by adding the `--max-backoff` parameter to an xbcloud command.

Since xbcloud does multiple asynchronous requests in parallel, a calculated value, measured in milliseconds, is used for `max-backoff`. This algorithm calculates how many milliseconds to sleep before the next retry. A number generated is based on the combining the power of two (2), the number of retries already attempted and adds a random number between 1 and 1000. This number avoids network congestion if multiple chunks have the same backoff value. If the default values are used, the final retry attempt to be approximately 17 minutes after the first try. The number is no longer calculated when the milliseconds reach the `--max-backoff` setting. At that point, the retries continue by using the `--max-backoff` setting until the `max-retries` parameter is reached.

Retriable errors

We retry for the following CURL operations:

- `CURLE_GOT_NOTHING`
- `CURLE_OPERATION_TIMEOUT`

- `CURLE_RECV_ERROR`
- `CURLE_SEND_ERROR`
- `CURLE_SEND_FAIL_REWIND`
- `CURLE_PARTIAL_FILE`
- `CURLE_SSL_CONNECT_ERROR`

We retry for the following HTTP operation status codes:

- 503
- 500
- 504
- 408

Each cloud provider may return a different CURL error or an HTTP error, depending on the issue. Add new errors by setting the following variables `--curl-retriable-errors` or `--http-retriable-errors` on the command line or in `my.cnf` or in a custom configuration file under the `[xbcloud]` section. You can add multiple errors using a comma-separated list of codes.

The error handling is enhanced when using the `--verbose` output. This output specifies which error caused xcloud to fail and what parameter a user must add to retry on this error.

The following is an example of a verbose output:


Expected output


```

210701 14:34:23 /work/pxb/ins/8.0/bin/xbcloud: Operation failed. Error: Server returned
nothing (no headers, no data)
210701 14:34:23 /work/pxb/ins/8.0/bin/xbcloud: Curl error (52) Server returned nothing (no
headers, no data) is not configured as retriable. You can allow it by adding --curl-
retriable-errors=52 parameter

```

Example

The following example adjusts the maximum number of retries and the maximum time between retries.

```
xbcloud [options] --max-retries=5 --max-backoff=10000
```



The following list describes the process using `--max-backoff=10000`:

- The chunk `xtrabackup_logfile.00000000000000000006` fails to upload the first time and sleeps for 2384 milliseconds.
- The same chunk fails for the second time and the sleep time is increased to 4387 milliseconds.
- The same chunk fails for the third time and the sleep time is increased to 8691 milliseconds.
- The same chunk fails for the fourth time. The `max-backoff` parameter has been reached. All retries sleep the same amount of time after reaching the parameter.
- The same chunk is successfully uploaded.



An example of the output for this setting



```
210702 10:07:05 /work/pxb/ins/8.0/bin/xbcloud: Operation failed. Error: Server returned
nothing (no headers, no data)
210702 10:07:05 /work/pxb/ins/8.0/bin/xbcloud: Sleeping for 2384 ms before retrying
backup3/xtrabackup_logfile.00000000000000000006
. . .
210702 10:07:23 /work/pxb/ins/8.0/bin/xbcloud: Operation failed. Error: Server returned
nothing (no headers, no data)
210702 10:07:23 /work/pxb/ins/8.0/bin/xbcloud: Sleeping for 4387 ms before retrying
backup3/xtrabackup_logfile.00000000000000000006
. . .
210702 10:07:52 /work/pxb/ins/8.0/bin/xbcloud: Operation failed. Error: Failed sending data
to the peer
210702 10:07:52 /work/pxb/ins/8.0/bin/xbcloud: Sleeping for 8691 ms before retrying
backup3/xtrabackup_logfile.00000000000000000006
. . .
210702 10:08:47 /work/pxb/ins/8.0/bin/xbcloud: Operation failed. Error: Failed sending data
to the peer
210702 10:08:47 /work/pxb/ins/8.0/bin/xbcloud: Sleeping for 10000 ms before retrying
backup3/xtrabackup_logfile.00000000000000000006
. . .
210702 10:10:12 /work/pxb/ins/8.0/bin/xbcloud: successfully uploaded chunk:
backup3/xtrabackup_logfile.00000000000000000006, size: 8388660
```

Binaries

Percona XtraBackup binaries overview

Percona XtraBackup is a set of the following binaries:

- [xtrabackup](#) - a compiled C binary that provides functionality to back up a *MySQL* database instance with *MyISAM*, *InnoDB*, and *XtraDB* tables.
- [xbcrypt](#) - a utility used for encrypting and decrypting backup files.
- [xbstream](#) - a utility that allows streaming and extracting files to / from the *xbstream* format.
- [xbcloud](#) - a utility used for downloading and uploading full or part of *xbstream* archive from / to cloud.

The recommended way to take the backup is by using the *xtrabackup* script. For more information on script options, see [xtrabackup](#).

The xtrabackup binary

The xtrabackup binary overview

The *xtrabackup* binary is a compiled C program that is linked with the *InnoDB* libraries and the standard *MySQL* client libraries.

For details on the xtrabackup command-line options, see [xtrabackup command-line options](#).

xtrabackup enables point-in-time backups of *InnoDB* / *XtraDB* tables together with the schema definitions, *MyISAM* tables, and other portions of the server.

The *InnoDB* libraries provide the functionality to apply a log to data files. The *MySQL* client libraries are used to parse command-line options and configuration file.

The tool runs in either `--backup` or `--prepare` mode, corresponding to the two main functions it performs. There are several variations on these functions to accomplish different tasks, and there are two less commonly used modes, `--stats` and `--print-param`.

xtrabackup implementation details

This page contains notes on various internal aspects of the *xtrabackup* tool's operation.

File permissions

xtrabackup opens the source data files in read-write mode, although it does not modify the files. This means that you must run *xtrabackup* as a user who has permission to write the data files. The reason for opening the files in read-write mode is that *xtrabackup* uses the embedded *InnoDB* libraries to open and read the files, and *InnoDB* opens them in read-write mode because it normally assumes it is going to write to them.

Tune the OS buffers

Because *xtrabackup* reads large amounts of data from the filesystem, it uses `posix_fadvise()` where possible, to instruct the operating system not to try to cache the blocks it reads from disk. Without this hint, the operating system would prefer to cache the blocks, assuming that *xtrabackup* is likely to need them again, which is not the case. Caching such large files can place pressure on the operating system's virtual memory and cause other processes, such as the database server, to be swapped out. The *xtrabackup* tool avoids this with the following hint on both the source and destination files:

```
posix_fadvise(file, 0, 0, POSIX_FADV_DONTNEED)
```



In addition, *xtrabackup* asks the operating system to perform more aggressive read-ahead optimizations on the source files:

```
posix_fadvise(file, 0, 0, POSIX_FADV_SEQUENTIAL)
```



Copy data files

When copying the data files to the target directory, *xtrabackup* reads and writes 1 MB of data at a time. This is not configurable. When copying the log file, *xtrabackup* reads and writes 512 bytes at a time. This is also not possible to configure, and matches *InnoDB*'s behavior (workaround exists in *Percona Server for MySQL* because it has an option to tune `innodb_log_block_size` for *XtraDB*, and in that case *Percona XtraBackup* will match the tuning).

After reading from the files, `xtrabackup` iterates over the 1MB buffer a page at a time, and checks for page corruption on each page with InnoDB's `buf_page_is_corrupted()` function. If the page is corrupt, it re-reads and retries up to 10 times for each page. It skips this check on the doublewrite buffer.

The xtrabackup command-line options

Here you can find all of the command-line options for the xtrabackup binary.

Modes of operation

You invoke xtrabackup in one of the following modes:

- `--backup` mode to make a backup in a target directory
- `--prepare` mode to restore data from a backup (created in `--backup` mode)
- `--copy-back` to copy data from a backup to the location that contained the original data; to move data instead of copying use the alternate `--move-back` mode.
- `--stats` mode to scan the specified data files and print out index statistics.

When you intend to run xtrabackup in any of these modes, use the following syntax:

```
$ xtrabackup [--defaults-file=#] --backup|--prepare|--copy-back|--stats [OPTIONS]
```

For example, the `--prepare` mode is applied as follows:

```
$ xtrabackup --prepare --target-dir=/data/backup/mysql/
```

For all modes, the default options are read from the xtrabackup and mysqld configuration groups from the following files in the given order:

1. `/etc/my.cnf`
2. `/etc/mysql/my.cnf`
3. `/usr/etc/my.cnf`
4. `~/ .my.cnf`.

As the first parameter to xtrabackup (in place of the `--defaults-file`, you may supply one of the following:

- `--print-defaults` to have xtrabackup print the argument list and exit.
- `--no-defaults` to forbid reading options from any file but the login file.
- `--defaults-file` to read the default options from the given file.

- `--defaults-extra-file` to read the specified additional file after the global files have been read.
- `--defaults-group-suffix` to read the configuration groups with the given suffix. The effective group name is constructed by concatenating the default configuration groups (xtrabackup and mysqld) with the given suffix.
- `--login-path` to read the given path from the login file.

InnoDB options

A large group of InnoDB options are generally read from the `my.cnf` configuration file, so xtrabackup boots up its embedded InnoDB in the same configuration as your current server. You usually do not need to specify them explicitly. These options have the same behavior in InnoDB and XtraDB. See `--innodb-miscellaneous` for more information.

Options

apply-log-only

Usage: `--apply-log-only`

This option causes only the redo stage to be performed when preparing a backup. It is very important for incremental backups.

backup

Usage: `--backup`

Make a backup and place it in `--target-dir`. See [Create a full backup](#).

backup-lock-timeout

Usage: `--backup-lock-timeout`

The timeout in seconds for attempts to acquire metadata locks.

backup-lock-retry-count

Usage: `--backup-lock-retry-count`

The number of attempts to acquire metadata locks.

backup-locks

Usage: `--backup-locks`

This option controls if backup locks should be used instead of `FLUSH TABLES WITH READ LOCK` on the backup stage. The option has no effect when the server does not support backup locks. This option is enabled by default, disable with `--no-backup-locks`.

check-privileges

Usage: `check-privileges`

This option checks if Percona XtraBackup has all required privileges. If a required privilege is missing for the current operation, the operation terminates and prints an error message. If a privilege is not needed for the current operation, but is missing and may be necessary for another XtraBackup operation, the operation is not aborted, and prints a warning.

Example output

```
```{.text .no-copy}
xtrabackup: Error: missing required privilege LOCK TABLES on *.*
xtrabackup: Warning: missing required privilege REPLICATION CLIENT on *.*
```
```

close-files

Usage: `--close-files`

Do not keep files opened. When xtrabackup opens a tablespace, xtrabackup normally doesn't close its file handle. This operation allows xtrabackup to handle the DDL operations correctly. However, if the number of tablespaces is huge and can not fit into any limit, there is an option to close file handles once they are no longer accessed. Percona XtraBackup can produce inconsistent backups with this option enabled. Use at your own risk.

compress

Usage: `--compress`

xtrabackup compresses all output data, including the transaction log and metadata files.

| Version updates | Option changes |
|--|--|
| Percona
XtraBackup
8.0.34-29 | <ul style="list-style-type: none">* The <code>qpress/QuickLZ</code> compression algorithm is no longer supported for compress operations.* The <code>ZSTD</code> compression algorithm is moved to General Availability. This version makes <code>ZSTD</code> the default compression method for the <code>--compress</code> option.* <code>--compress</code> produces <code>*.zst</code> files. You can specify <code>ZSTD</code> compression level with the <code>--compress-zstd-level=#</code> option.* <code>--compress=lz4</code> produces <code>*.lz4</code> files. You can extract the contents of these files by using <code>lz4</code> program. <p>You can extract the contents of the files by using the <code>--decompress</code> option. To decompress backups taken by older versions of Percona XtraBackup that used a QuickLZ compression algorithm, the <code>--decompress</code> option still supports <code>qpress</code> for backward compatibility.</p> |
| Percona
XtraBackup
8.0.33-28 and
lower versions | <p>The <code>--compress</code> option uses either the <code>quicklz</code>, <code>lz4</code>, or <code>ZSTD</code> compression algorithm to compress all output data. The <code>quicklz</code> algorithm is chosen by default. When using <code>--compress=quicklz</code> or <code>--compress</code>, the resulting files have the <code>qpress</code> archive format. Every <code>*.qp</code> file produced by <code>xtrabackup</code> is essentially a one-file <code>qpress</code> archive and can be extracted and uncompressed by the <code>qpress</code> file archiver.</p> <ul style="list-style-type: none">* <code>--compress=zstd</code> produces <code>*.zst</code> files. You can specify <code>ZSTD</code> compression level with the <code>[--compress-zstd-level=#](#compress-zstd-level)</code> option.* <code>--compress=lz4</code> produces <code>*.lz4</code> files. You can extract the contents of these files by using <code>lz4</code> program. <p>You can extract the contents of the files by using the <code>--decompress</code> option.</p> |
| Percona
XtraBackup
8.0.30-23 | <p>Adds the Zstandard (<code>ZSTD</code>) compression algorithm in tech preview. The <code>ZSTD</code> algorithm is a fast lossless compression algorithm that targets real-time compression scenarios and better compression ratios.</p> <p>To compress files using the <code>ZSTD</code> compression algorithm, set the <code>--compress</code> option to <code>zstd</code>. The <code>--compress=zstd</code> option produces <code>*.zst</code> files. You can extract the contents of these files with the <code>--decompress</code> option. Also, you can specify the <code>ZSTD</code> compression level with the <code>--compress-zstd-level=#</code> option.</p> |

compress-chunk-size

Usage: `--compress-chunk-size=#`

Size of working buffer(s) for compression threads in bytes. The default value is 64K.

compress-threads

Usage: `--compress-threads=#` This option specifies the number of worker threads xtrabackup uses for parallel data compression. This option defaults to `1`. Parallel compression (`--compress-threads`) can be used with [`--parallel`](#).

For example, `--parallel=4 --compress --compress-threads=2` creates 4 I/O threads that read and then pipe that data to 2 compression threads.

compress-zstd-level

Usage: `--compress-zstd-level=#`

This option specifies the `ZSTD` compression level. Compression levels provide a trade-off between the compression speed and the size of the compressed file. A lower compression level provides faster compression speed but larger file sizes. A higher compression level provides lower compression speed but smaller file sizes. For example, set level 1 if the compression speed is the most important for you. Set level 19 if the size of the compressed files is the most important.

The default value is 1. The allowed range is from 1 to 19.

Percona XtraBackup 8.0.30-22 implemented this option.

copy-back

Usage: `--copy-back`

Copy all the files in a previously made backup from the backup directory to their original locations. This option will not copy over existing files unless [`--force-non-empty-directories`](#) option is specified.

core-file

Usage: `--core-file`

Write core on fatal signals.

databases

Usage: `--databases=#`

This option specifies which databases to back up.

Mutual exclusion

`--tables` and `--databases` are mutually exclusive. If you use both options in the same command, XtraBackup ignores `--databases` and only uses `--tables`. Use only one of these options per backup operation.

Accepted syntax:

This option accepts a space-separated list of database names. To include all tables in a database, add `.*` after the database name (for example, `mydb.*`).

Regular expressions are not supported.

Examples:

- `--databases='mysql performance_schema sys'` - Backs up the entire mysql, performance_schema, and sys databases
- `--databases='mydb.*'` - Backs up every table in the mydb database

databases-exclude

Usage: `--databases-exclude=name`

Databases are excluded based on name. This option operates the same way as `--databases`, but excludes the matched names from the backup.

This option accepts a space-separated list of database names.

This option has a higher priority than `--databases`.

databases-file

Usage: `--databases-file=#`

This option specifies the path to the file containing the list of databases and tables that should be backed up. The file can contain the list elements of the form `dbname1[.table_name1]`, one element per line.

datadir

Usage: `--datadir=DIRECTORY`

The source directory for the backup. This directory should be the same as the `datadir` for your MySQL server, so it should be read from `my.cnf` if that exists; otherwise, you must specify it on the command line.

When combined with the `--copy-back` or `--move-back` option, `--datadir` refers to the destination directory.

Once connected to the server, to perform a backup, you will need `READ` and `EXECUTE` permissions at a filesystem level in the server's `datadir`.

debug-sleep-before-unlock

Usage: `--debug-sleep-before-unlock=#`

This is a debug-only option used by the xtrabackup test suite.

debug-sync

Usage: `--debug-sync=name`

The debug sync point. This option is only used by the xtrabackup test suite.

decompress

Usage: `--decompress`

Decompresses all files in a backup previously made with the `--compress` option. The `--parallel` option will allow multiple files to be decrypted simultaneously. To decompress files compressed with `quicklz` compression algorithm, install the `qpress` utility. It must be accessible within the path. Percona XtraBackup does not automatically remove the compressed files. To clean up the backup directory users should use `--remove-original` option.

The `--decompress` option may be used with `xbstream` to decompress individual `qpress` files.

decompress-threads

Usage: `--decompress-threads=#`

Force `xbstream` to use the specified number of threads for decompressing.

decrypt

Usage: `--decrypt=ENCRYPTION-ALGORITHM`

Decrypts all files with the `.xbcrypt` extension in a backup previously made with `--encrypt` option. The `--parallel` option will allow multiple files to be decrypted simultaneously. Percona XtraBackup doesn't automatically remove the encrypted files. To clean up the backup directory users should use `--remove-original` option.

defaults-extra-file

Usage: `--defaults-extra-file=[MY.CNF]`

Read this file after the global files are read. This option must be written as the first option on the command line.

defaults-file

Usage: `--defaults-file=[MY.CNF]`

Only read default options from the given file. Must be given as the first option on the command-line. Must be a real file; it cannot be a symbolic link.

defaults-group

Usage: `--defaults-group=GROUP-NAME` This option sets the group that should be read from the configuration file. This option is used by xtrabackup if you use the `--defaults-group` option. It is needed for `mysqld_multi` deployments.

defaults-group-suffix

Usage: `--defaults-group-suffix=#`

Also reads groups with `concat(group, suffix)`.

dump-innodb-buffer-pool

Usage: `--dump-innodb-buffer-pool`

This option controls whether or not a new dump of buffer pool content should be done.

With `--dump-innodb-buffer-pool`, xtrabackup makes a request to the server to start the buffer pool dump (it takes some time to complete and is done in the background) at the beginning of a backup provided the status variable `innodb_buffer_pool_dump_status` reports that the dump has been completed.

```
$ xtrabackup --backup --dump-innodb-buffer-pool --target-dir=/home/user/backup
```



By default, this option is set to OFF.

If `innodb_buffer_pool_dump_status` reports that there is running dump of the buffer pool, xtrabackup waits for the dump to complete using the value of `--dump-innodb-buffer-pool-timeout`

The file `ib_buffer_pool` stores tablespace ID and page ID data to warm up the buffer pool sooner.

dump-innodb-buffer-pool-timeout

Usage: `--dump-innodb-buffer-pool-timeout`

This option contains the number of seconds that xtrabackup should monitor the value of `innodb_buffer_pool_dump_status` to determine if buffer pool dump has completed.

This option is used in combination with `--dump-innodb-buffer-pool`. By default, it is set to 10 seconds.

dump-innodb-buffer-pool-pct

Usage: `--dump-innodb-buffer-pool-pct`

This option contains the percentage of the most recently used buffer pool pages to dump.

This option is effective if `--dump-innodb-buffer-pool` option is set to ON. If this option contains a value, xtrabackup sets the MySQL system variable `innodb_buffer_pool_dump_pct`. As soon as the buffer pool dump completes or it is stopped (see `--dump-innodb-buffer-pool-timeout`), the value of the MySQL system variable is restored.

encrypt

Usage: `--encrypt=ENCRYPTION_ALGORITHM`

This option instructs xtrabackup to encrypt backup copies of InnoDB data files using the algorithm specified in the ENCRYPTION_ALGORITHM. Currently supported algorithms are: `AES128`, `AES192` and `AES256`

encrypt-key

Usage: `--encrypt-key=ENCRYPTION_KEY`

A proper length encryption key to use. It is not recommended to use this option where there is uncontrolled access to the machine as the command line and thus the key can be viewed as part of the process info.

encrypt-key-file

Usage: `--encrypt-key-file=ENCRYPTION_KEY_FILE`

The name of a file where the raw key of the appropriate length can be read from. The file must be a simple binary (or text) file that contains exactly the key to be used.

It is passed directly to the xtrabackup child process. See the xtrabackup documentation for more details.

encrypt-threads

Usage: `--encrypt-threads=#`

This option specifies the number of worker threads that will be used for parallel encryption/decryption. See the xtrabackup documentation for more details.

encrypt-chunk-size

Usage: `--encrypt-chunk-size=#`

This option specifies the size of the internal working buffer for each encryption thread, measured in bytes. It is passed directly to the xtrabackup child process.

To adjust the chunk size for encrypted files, use `--read-buffer-size` and `--encrypt-chunk-size`.

estimate-memory

Usage: `--estimate-memory=#`

This option is in [tech preview](#). Before using this option in production, we recommend that you test restoring production from physical backups in your environment, and also use the alternative backup method for redundancy.

Implemented in Percona XtraBackup 8.0.32-26, the option lets you enable or disable the [Smart memory estimation](#) feature. The default value is OFF. Enable the feature by setting `--estimate-memory=ON` in the backup phase and setting the `--use-free-memory-pct` option in the `--prepare` phase. If the `--estimate-memory` setting is disabled, the `--use-free-memory-pct` setting is ignored.

An example of how to enable the Smart memory estimation feature:

```
$ xtrabackup --backup --estimate-memory=ON --target-dir=/data/backups/
```



```
$ xtrabackup --prepare --use-free-memory-pct=50 --target-dir=/data/backups/
```



export

Usage: `--export`

Create files necessary for exporting tables. See [restore individual tables](#).

extra-lsdir

Usage: `--extra-lsdir=DIRECTORY`

(for `--backup`): save an extra copy of the `xtrabackup_checkpoints` and `xtrabackup_info` files in this directory.

force-non-empty-directories

Usage: `--force-non-empty-directories`

When specified, it makes `--copy-back` and `--move-back` option transfer files to non-empty directories. No existing files will be overwritten. If files that need to be copied/moved from the backup directory already exist in the destination directory, it will still fail with an error.

ftwrl-wait-timeout

Usage: `--ftwrl-wait-timeout=SECONDS`

This option specifies time in seconds that xtrabackup should wait for queries that would block `FLUSH TABLES WITH READ LOCK` before running it. If there are still such queries when the timeout expires, xtrabackup terminates with an error. Default is `0`, in which case it does not wait for queries to complete and starts `FLUSH TABLES WITH READ LOCK` immediately. Where supported xtrabackup will automatically use [Backup Locks](#) as a lightweight alternative to `FLUSH TABLES WITH READ LOCK` to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.

ftwrl-wait-threshold

Usage: `--ftwrl-wait-threshold=SECONDS`

This option specifies the query run time threshold which is used by xtrabackup to detect long-running queries with a non-zero value of `--ftwrl-wait-timeout`. `FLUSH TABLES WITH READ LOCK` is not started until such long-running queries exist. This option has no effect if `--ftwrl-wait-timeout` is `0`. Default value is `60` seconds. Where supported xtrabackup will automatically use [Backup Locks](#) as a lightweight alternative to `FLUSH TABLES WITH READ LOCK` to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.

ftwrl-wait-query-type

Usage: `--ftwrl-wait-query-type=all|update`

This option specifies which types of queries are allowed to complete before xtrabackup will issue the global lock. Default is `all`.

galera-info

Usage: `--galera-info`

This option creates the `xtrabackup_galera_info` file which contains the local node state at the time of the backup. Option should be used when performing the backup of Percona XtraDB Cluster. It has no effect when backup locks are used to create the backup.

generate-new-master-key

Usage: `--generate-new-master-key`

Generate a new master key when doing a copy-back.

generate-transition-key

Usage: `--generate-transition-key`

xtrabackup needs to access the same keyring file or vault server during prepare and copy-back but it should not depend on whether the server keys have been purged.

`--generate-transition-key` creates and adds to the keyring a transition key for xtrabackup to use if the master key used for encryption is not found because it has been rotated and purged.

get-server-public-key

Usage: `--get-server-public-key`

Get the server public key

help

Usage: `--help`

When run with this option or without any options xtrabackup displays information about how to run the program on the command line along with all supported options and variables with default values where appropriate.

history

Usage: `--history=NAME`

This option enables the tracking of backup history in the `PERCONA_SCHEMA.xtrabackup_history` table. An optional history series name may be specified that will be placed with the history record for the current backup being taken.

host

Usage: `--host=HOST`

This option accepts a string argument that specifies the host to use when connecting to the database server with TCP/IP. It is passed to the mysql child process without alteration. See `mysql --help` for details.

incremental-basedir

Usage: `--incremental-basedir=DIRECTORY`

When creating an incremental backup, this is the directory containing the full backup that is the base dataset for the incremental backups.

incremental-dir

Usage: `--incremental-dir=DIRECTORY`

When preparing an incremental backup, this is the directory where the incremental backup is combined with the full backup to make a new full backup.

incremental-force-scan

Usage: `--incremental-force-scan`

When creating an incremental backup, force a full scan of the data pages in that instance.

incremental-history-name

Usage: `--incremental-history-name=name`

This option specifies the name of the backup series stored in the `PERCONA_SCHEMA.xtrabackup_history` history record to base an incremental backup on. xtrabackup will search the history table looking for the most recent (highest `innodb_to_lsn`), successful backup in the series and take the `to_lsn` value to use as the starting `lsn` for the incremental backup. This will be mutually exclusive with `--incremental-history-uuid`, `--incremental-basedir` and `--incremental-lsn`. If no valid `lsn` can be found (no series by that name, no successful backups by that name) xtrabackup will return with an error. It is used with the `--incremental` option.

incremental-history-uuid

Usage: `--incremental-history-uuid=name`

This option specifies the UUID of the specific history record stored in the `PERCONA_SCHEMA.xtrabackup_history` to base an incremental backup on. `--incremental-history-name`, `--incremental-basedir` and `--incremental-lsn`. If no valid `lsn` can be found (no success record with that UUID) xtrabackup will return with an error. It is used with the `--incremental` option.

incremental-lsn

Usage: `--incremental-lsn=LSN`

When creating an incremental backup, you can specify the log sequence number (LSN) instead of specifying `--incremental-basedir`. For databases created in 5.1 and later, specify the LSN as a single 64-bit integer.

Important

If a wrong LSN value is specified (a user error which Percona XtraBackup does not detect), the backup is unusable. Be careful!

innodb

Usage: `--innodb=name`

This option is ignored for MySQL option compatibility.

innodb-miscellaneous

Usage: `--innodb-miscellaneous`

There is a large group of InnoDB options that are normally read from the `my.cnf` configuration file, so that xtrabackup boots up its embedded InnoDB in the same configuration as your current server. You normally do not need to specify these explicitly. These options have the same behavior in InnoDB and XtraDB:

keyring-file-data

Usage: `--keyring-file-data=FILENAME`

The path to the keyring file. Combine this option with `--xtrabackup-plugin-dir`.

kill-long-queries-timeout

Usage: `--kill-long-queries-timeout=SECONDS`

This option specifies the number of seconds xtrabackup waits between starting `FLUSH TABLES WITH READ LOCK` and killing those queries that block it. Default is 0 seconds, which means xtrabackup will not attempt to kill any queries. In order to use this option xtrabackup user should have the `PROCESS` and `SUPER` privileges. Where supported, xtrabackup automatically uses [Backup locks](#) as a lightweight alternative to `FLUSH TABLES WITH READ LOCK` to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables.

kill-long-query-type

Usage: `--kill-long-query-type=all|select`

This option specifies which types of queries should be killed to unblock the global lock. Default is "select".

lock-ddl

Usage: `--lock-ddl`

Issue `LOCK TABLES FOR BACKUP` if it is supported by server (otherwise use `LOCK INSTANCE FOR BACKUP`) at the beginning of the backup to block all DDL operations.

Note

Prior to Percona XtraBackup 8.0.22-15.0, using a `safe-slave-backup` stops the SQL replica thread after the InnoDB tables and before the non-InnoDB tables are backed up.

As of Percona XtraBackup 8.0.22-15.0, using a `safe-slave-backup` option stops the SQL replica thread before copying the InnoDB files.

lock-ddl-per-table

Usage: `--lock-ddl-per-table`

Lock DDL for each table before xtrabackup starts to copy it and until the backup is completed.

Note

The `[-lock-ddl-per-table]` option is deprecated. Use the `[-lock-ddl]` option instead.

lock-ddl-timeout

Usage: `--lock-ddl-timeout`

If `LOCK TABLES FOR BACKUP` or `LOCK INSTANCE FOR BACKUP` does not return within given timeout, abort the backup.

log

Usage: `--log`

This option is ignored for MySQL

log-bin

Usage: `--log-bin`

The base name for the log sequence.

log-bin-index

Usage: `--log-bin-index=name`

File that holds the names for binary log files.

log-copy-interval

Usage: `--log-copy-interval=#`

This option specifies the time interval between checks done by the log copying thread in milliseconds (default is 1 second).

login-path

Usage: `--login-path`

Read the given path from the login file.

move-back

Usage: `--move-back`

Move all the files in a previously made backup from the backup directory to their original locations. As this option removes backup files, it must be used with caution.

no-backup-locks

Usage: `--no-backup-locks`

Explicitly disables the `--backup-locks` option which is enabled by default.

no-defaults

Usage: `--no-defaults`

The default options are only read from the login file.

no-lock

Usage: `--no-lock`

Disables table lock with `FLUSH TABLES WITH READ LOCK`. Use it only if all your tables are InnoDB and you do not care about the binary log position of the backup. This option shouldn't be used if there are any `DDL` statements being executed or if any updates are happening on non-InnoDB tables (this includes the system MyISAM tables in the mysql database), otherwise it could lead to an inconsistent backup. Where supported xtrabackup will automatically use [Backup locks](#) as a lightweight alternative to `FLUSH TABLES WITH READ LOCK` to copy non-InnoDB data to avoid blocking DML queries that modify InnoDB tables. If you are considering to use this because your backups are failing to acquire the lock, this could be because of incoming replication events are preventing the lock from succeeding. Please try using `--safe-slave-backup` to momentarily stop the replication replica thread, this may help the backup to succeed and you do not need to use this option.

no-server-version-check

Usage: `--no-server-version-check`

Implemented in Percona XtraBackup 8.0.21.

The `--no-server-version-check` option disables the server version check.

The default behavior runs a check that compares the source system version to the Percona XtraBackup version. If the source system version is higher than the XtraBackup version, the backup is aborted with a message.

Adding the option overrides this check, and the backup proceeds, but there may be issues with the backup.

See Server Version and Backup Version Comparison for more information.

no-version-check

Usage: `--no-version-check`

This option disables the version check. If you do not pass this option, the automatic version check is enabled implicitly when xtrabackup runs in the `--backup` mode. To disable the version check, you should pass explicitly the `--no-version-check` option when invoking xtrabackup.

When the automatic version check is enabled, xtrabackup performs a version check against the server on the backup stage after creating a server connection. xtrabackup sends the following information to the server:

- MySQL flavor and version
- Operating system name
- Percona Toolkit version
- Perl version

Each piece of information has a unique identifier. This is a MD5 hash value that Percona Toolkit uses to obtain statistics about how it is used. This is a random UUID; no client information is either collected or stored.

open-files-limit

Usage: `--open-files-limit=#`

The maximum number of file descriptors to reserve with `setrlimit()`.

parallel

Usage: `--parallel=#`

The `--parallel` option specifies the number of threads to use to copy multiple data files concurrently when creating a backup. The default value is 1 (that is, no concurrent transfer). In Percona XtraBackup 2.3.10 and newer, the `--parallel` option can be used with the `--copy-back` option to copy the user data files in parallel (redo logs and system tablespaces are copied in the main thread).

Before Percona XtraBackup 8.0.35-33 and 8.4.0-3, the `--parallel` option didn't have any effect on the prepare phase.

Starting with [Percona XtraBackup 8.0.35-33](#) and 8.4.0-3, using `--parallel=X` has effect on the prepare phase. It will now use X threads to apply the changes from `.delta` files to the IBD files. This option processes multiple delta files simultaneously, improving storage performance and accelerating incremental backups, particularly with numerous InnoDB Data (IBD) files. The option remains effective even if IBD files are unchanged between backups and efficiently handles empty delta files.

When using `--parallel` in the prepare phase, always specify a numeric value. The recommended minimum value is 4 (for example, `--parallel=4`).

Note that each thread operates on a single file. If you have a large delta file, there is still only one thread that processes that `.delta` file. Parallelization occurs at the file level, not within individual files.

password

Usage: `--password=PASSWORD`

This option specifies the password to use when connecting to the database. It accepts a string argument. See `mysql --help` for details.

plugin-load

Usage: `--plugin-load`

List of plugins to load.

port

Usage: `--port=PORT`

This option accepts a string argument that specifies the port to use when connecting to the database server with TCP/IP. It is passed to the `mysql` child process without alteration. See `mysql --help` for details.

prepare

Usage: `--prepare`

Makes `xtrabackup` perform a recovery on a backup created with `--backup`, so that it is ready to use. See preparing a backup.

print-defaults

Usage: `--print-defaults`

Print the program argument list and exit. Must be given as the first option on the command-line.

print-param

Usage: `--print-param`

Makes `xtrabackup` print out parameters that can be used for copying the data files back to their original locations to restore them.

read-buffer-size

Usage: `--read-buffer-size`

Sets the read buffer size. The given value is scaled up to page size. The default size is 10MB. Use this option to increase the xbcrypt/xbstream chunk size from the default size. To adjust the chunk size for encrypted files, use `--read-buffer-size` and `--encrypt-chunk-size`.

redo-log-arch-dir

Usage: `--redo-log-arch-dir=name`

This option sets the redo log archive directory if this directory is not already set on the server. To use this option, you must run Percona XtraBackup as the owner of `mysqld`. Additionally, the user must have at least one of the following privileges: `SUPER` or `SYSTEM_VARIABLES_ADMIN`.

Implemented in [Percona XtraBackup 8.0.34-29](#).

register-redo-log-consumer

Usage: `--register-redo-log-consumer`

The `--register-redo-log-consumer` parameter is disabled by default. When enabled, this parameter lets Percona XtraBackup register as a redo log consumer at the start of the backup. The server does not remove a redo log that Percona XtraBackup (the consumer) has not yet copied. The consumer reads the redo log and manually advances the log sequence number (LSN). The server blocks the writes during the process. Based on the redo log consumption, the server determines when it can purge the log.

Implemented in [Percona XtraBackup 8.0.30-23](#).

remove-original

Usage: `--remove-original`

Implemented in Percona XtraBackup 2.4.6, this option when specified will remove `.qp`, `.xbcrypt` and `.qp.xbcrypt` files after decryption and decompression.

rocksdb-datadir

Usage: `--rocksdb-datadir`

RocksDB data directory

rocksdb-wal-dir

Usage: `--rocksdb-wal-dir`

RocksDB WAL directory.

rocksdb-checkpoint-max-age

Usage: `--rocksdb-checkpoint-max-age`

When backing up MyRocks datadir, Percona XtraBackup creates a checkpoint, measured in seconds, captures binlog coordinates, and then copies the `.SST` files into the backup. Copying may take a while, especially for large datasets. Then, if you want to set up a new replica or apply binary logs on top of the backup, the size of binary logs may become large because the binlog coordinates were captured at the start of the backup, before the file copying began.

This feature provides a way to minimize the size of binary logs and to move binary coordinates closer to the moment when file copying completes. The process works by creating the checkpoint and copying `.SST` files. If the copy takes a while, Percona XtraBackup creates the next checkpoint and copies only new `.SST` files. The second copy should take less time. Percona XtraBackup repeats the process until the copy time falls within an acceptable range, determined by this parameter.

Use `--rocksdb-checkpoint-max-age` to specify the maximum age (in seconds) that a checkpoint can be when the backup completes. If the checkpoint is older than this threshold, Percona XtraBackup creates another checkpoint and copies the additional `.SST` files that were created since the previous checkpoint.

Implemented in [Percona XtraBackup 8.0.8](#).

rocksdb-checkpoint-max-count

Usage: `--rocksdb-checkpoint-max-count`

This option works together with `--rocksdb-checkpoint-max-age` to control the MyRocks checkpoint renewal process. If the checkpoint age requirement specified by `--rocksdb-checkpoint-max-age` has not been met after this number of checkpoint iterations, Percona XtraBackup completes the backup regardless of the checkpoint age. This action prevents the backup from continuing indefinitely if the system is under heavy write load.

Implemented in [Percona XtraBackup 8.0.8](#).

rollback-prepared-trx

Usage: `--rollback-prepared-trx`

Force rollback prepared InnoDB transactions.

rsync

Usage: `--rsync`

Uses the rsync utility to optimize local file transfers. When this option is specified, xtrabackup uses rsync to copy all non-InnoDB files instead of spawning a separate cp for each file, which can be much faster for servers with a large number of databases or tables. This option cannot be used together with `--stream`.

safe-slave-backup

Usage: `--safe-slave-backup`

When specified, xtrabackup will stop the replica SQL thread just before running `FLUSH TABLES WITH READ LOCK` and wait to start backup until `Slave_open_temp_tables` in `SHOW STATUS` is zero. If there are no open temporary tables, the backup will take place, otherwise the SQL thread will be started and stopped until there are no open temporary tables. The backup will fail if `Slave_open_temp_tables` does not become zero after `--safe-slave-backup-timeout` seconds. The replication SQL thread will be restarted when the backup finishes. This option is implemented in order to deal with [replicating temporary tables](#) and isn't necessary with Row-Based-Replication.

Note

Prior to Percona XtraBackup 8.0.22-15.0, using a `safe-slave-backup` stops the SQL replica thread after the InnoDB tables and before the non-InnoDB tables are backed up.

As of Percona XtraBackup 8.0.22-15.0, using a `safe-slave-backup` option stops the SQL replica thread before copying the InnoDB files.

safe-slave-backup-timeout

Usage: `--safe-slave-backup-timeout=SECONDS`

How many seconds `--safe-slave-backup` should wait for `Slave_open_temp_tables` to become zero. Defaults to 300 seconds.

secure-auth

Usage: `--secure-auth`

Refuse client connecting to server if it uses old (pre-4.1.1) protocol. (Enabled by default; use `--skip-secure-auth` to disable.)

server-id

Usage: `--server-id=#`

The server instance being backed up.

server-public-key-path

Usage: `--server-public-key-path`

The file path to the server public RSA key in the PEM format.

skip-tables-compatibility-check

Usage: `--skip-tables-compatibility-check`

See `--tables-compatibility-check`.

slave-info

Usage: `--slave-info`

This option is useful when backing up a replication replica server. It prints the binary log position of the source server. It also writes the binary log coordinates to the `xtrabackup_slave_info` file as a `CHANGE MASTER` command. A new replica for this source can be set up by starting a replica server on this backup and issuing a `CHANGE MASTER` command with the binary log position saved in the `xtrabackup_slave_info` file.

socket

Usage: `--socket`

This option accepts a string argument that specifies the socket to use when connecting to the local database server with a UNIX domain socket. It is passed to the `mysql` child process without alteration. See `mysql --help` for details.

ssl

Usage: `--ssl`

Enable secure connection.

ssl-ca

Usage: `--ssl-ca`

Path of the file which contains list of trusted SSL CAs.

ssl-capath

Usage: `--ssl-capath`

Directory path that contains trusted SSL CA certificates in PEM format.

ssl-cert

Usage: `--ssl-cert`

Path of the file which contains X509 certificate in PEM format.

ssl-cipher

Usage: `--ssl-cipher`

List of permitted ciphers to use for connection encryption.

ssl-crl

Usage: `--ssl-crl`

Path of the file that contains certificate revocation lists. MySQL server documentation.

ssl-crlpath

Usage: `--ssl-crlpath`

Path of directory that contains certificate revocation list files.

ssl-fips-mode

Usage: `--ssl-fips-mode`

SSL FIPS mode (applies only for OpenSSL); permitted values are: OFF, ON, STRICT.

ssl-key

Usage: `--ssl-key`

Path of file that contains X509 key in PEM format.

ssl-mode

Usage: `--ssl-mode`

Security state of connection to server.

ssl-verify-server-cert

Usage: `--ssl-verify-server-cert`

Verify server certificate Common Name value against host name used when connecting to server.

stats

Usage: `--stats`

Causes xtrabackup to scan the specified data files and print out index statistics.

stream

Usage: `--stream=FORMAT`

Stream all backup files to the standard output in the specified format. Currently, this option only supports the xstream format.

strict

Usage: `--strict`

If this option is specified, xtrabackup fails with an error when invalid parameters are passed.

tables

Usage: `--tables=name`

This option filters tables to be backed up based on their fully qualified names.

Mutual exclusion

`--tables` and `--databases` are mutually exclusive. Supplying both in the same command line causes XtraBackup to ignore one of them (currently `--databases`). Use only one of the options per backup operation.

Accepted syntax:

This option accepts either:

- A comma-separated list of fully qualified table names in the format `database.table` (for example, `db1.t1,db1.t2,db2.t3`)
- A POSIX regular expression surrounded by single quotes (for example, `'^(db1\.t1|db2\.t.*)$'`)

Examples:

- List format: `--tables=db1.t1,db1.t2,db2.t3` - Backs up specific tables from multiple databases
- Regex format: `--tables='^(db1\.t1|db2\.t.*)$'` - Uses a regular expression to match tables

See [Create a partial backup](#).

tables-compatibility-check

Usage: `--tables-compatibility-check`

Enables the engine compatibility warning. The default value is ON. To disable the engine compatibility warning use `--skip-tables-compatibility-check`.

tables-exclude

Usage: `--tables-exclude=name`

Filtering by regexp for table names. Operates the same way as `--tables`, but matched names are excluded from backup. Note that this option has a higher priority than `--tables`.

Backing up selected tables together with whole system databases

When you need to back up specific tables from user databases along with entire system databases (such as `mysql`, `performance_schema`, or `sys`), you must use the `--tables` option with a regular expression. This is the only way to combine specific table selection with full-database inclusion, because `--tables` and `--databases` are mutually exclusive.

Example 1: Back up specific tables plus entire mysql database

```
$ xtrabackup --backup --tables='^(mydb\.(t1|t2)|mysql\.)' --target-dir=/data/backup/
```

This command backs up tables `t1` and `t2` from `mydb`, plus all tables in the `mysql` database.

Example 2: Back up tables from multiple databases plus system databases

```
$ xtrabackup --backup --tables='^(db1\.t1|db2\.t.*|mysql\.|performance_schema\.|sys\.)' --target-dir=/data/backup/
```

This command backs up `t1` from `db1`, all tables starting with `t` from `db2`, plus the entire `mysql`, `performance_schema`, and `sys` databases.

Example 3: Back up all tables from a database plus specific tables from another

```
$ xtrabackup --backup --tables='^(mydb\.|otherdb\.(important_table|critical_table))' --target-dir=/data/backup/
```

This command backs up all tables from `mydb` plus specific tables from `otherdb`.

Common pitfalls

The following table lists common mistakes when using `--tables` and `--databases` options and what happens:

| Incorrect usage | What happens |
|--|--|
| Using both <code>--tables</code> and <code>--databases</code> together | XtraBackup ignores <code>--databases</code> and only processes <code>--tables</code> . Your backup may not include the databases you expected. |
| Using <code>--databases</code> with regular expressions | Regular expressions are not supported by <code>--databases</code> . XtraBackup treats the regular expression as a literal database name. If that database name does not exist, the backup will be empty or fail. |

| Incorrect usage | What happens |
|---|---|
| Using <code>--tables</code> without quotes for regex patterns | Without quotes, the shell may interpret special regex characters. This causes incorrect pattern matching or syntax errors. |
| Using <code>--databases=mydb</code> expecting to include all tables | Without the <code>.*</code> wildcard, <code>--databases=mydb</code> may not include all tables. Use <code>--databases=mydb.*</code> to include all tables in the database. |
| Mixing comma-separated lists with regex in <code>--tables</code> | The <code>--tables</code> option accepts either a comma-separated list OR a quoted regular expression, but not both at the same time. Mixing both syntaxes will not work as expected. |

tables-file

Usage: `--tables-file=name`

A file containing one table name per line, in `dbname.tablename` format. The backup will be limited to the specified tables.

target-dir

Usage: `--target-dir=DIRECTORY`

This option specifies the destination directory for the backup. If the directory does not exist, xtrabackup creates it. If the directory does exist and is empty, xtrabackup will succeed. xtrabackup will not overwrite existing files, however; it will fail with operating system error 17, `file exists`.

If this option is a relative path, it is interpreted as being relative to the current working directory from which xtrabackup is executed.

In order to perform a backup, you need `READ`, `WRITE`, and `EXECUTE` permissions at a filesystem level for the directory that you supply as the value of `--target-dir`.

innodb-temp-tablespaces-dir

Usage: `--innodb-temp-tablespaces-dir=DIRECTORY`

Directory where temp tablespace files live, this path can be absolute.

throttle

Usage: `--throttle=#`

This option limits the number of chunks copied per second. The chunk size is 10 MB.

To limit the bandwidth to 10 MB/s, set the option to 1.

tls-ciphersuites

Usage: `--tls-ciphersuites`

TLS v1.3 cipher to use.

tls-version

Usage: `--tls-version`

TLS version to use, permitted values are: TLSv1, TLSv1.1, TLSv1.2, TLSv1.3.

tmpdir

Usage: `--tmpdir=name`

Specify the directory that will be used to store temporary files during the backup

transition-key

Usage: `--transition-key=name`

This option is used to enable processing the backup without accessing the keyring vault server. In this case, xtrabackup derives the AES encryption key from the specified passphrase and uses it to encrypt tablespace keys of tablespaces being backed up.

If `--transition-key` does not have any value, xtrabackup will ask for it. The same passphrase should be specified for the `--prepare` command.

use-free-memory-pct

Usage: `--use-free-memory-pct`

The `--use-free-memory-pct` is a [tech preview](#) option. Before using this option in production, we recommend that you test restoring production from physical backups in your environment, and also use the alternative backup method for redundancy.

Implemented in Percona XtraBackup 8.0.30-23, this option lets you configure the [Smart memory estimation](#) feature. The option controls the amount of free memory that can be used to `--prepare` a backup. The default value is 0 (zero) which defines the option as disabled. For example, if you set `--use-free-memory-pct=50`, then 50% of the free memory is used to `prepare` a backup. The maximum allowed value is 100.

This option works, only if `--estimate-memory` option is enabled. If the `--estimate-memory` option is disabled, the `--use-free-memory-pct` setting is ignored.

An example of how to enable the Smart memory estimation feature:

```
$ xtrabackup --backup --estimate-memory=ON --target-dir=/data/backups/
```



```
$ xtrabackup --prepare --use-free-memory-pct=50 --target-dir=/data/backups/
```



use-memory

Usage: `--use-memory`

This option affects how much memory is allocated and is similar to `innodb_buffer_pool_size`. This option is only relevant in the `--prepare` phase or when analyzing statistics with `--stats`. The default value is 100MB. The recommended value is between 1GB to 2GB. Multiple values are supported if you provide the unit (for example, 1MB, 1M, 1GB, 1G).

user

Usage: `--user=USERNAME`

This option specifies the MySQL username used when connecting to the server, if that's not the current user. The option accepts a string argument. See `mysql --help` for details.

version

Usage: `--version`

This option prints xtrabackup version and exits.

xtrabackup-plugin-dir

Usage: `--xtrabackup-plugin-dir=DIRNAME`

The absolute path to the directory that contains the `keyring` plugin.

Configure xtrabackup

All the xtrabackup configuration is done through options, which behave exactly like standard MySQL program options: they can be specified either at the command-line, or through a file such as `/etc/my.cnf`.

The xtrabackup binary reads the `mysqld` and `xtrabackup` sections from any configuration files, in that order. That is so that it can read its options from your existing MySQL installation, such as the `datadir` or some InnoDB options. If you want to override these options specify them in the `xtrabackup` section. These options are read at a later time and take precedence.

You don't need to put any configuration in your `my.cnf` if you don't want to. You can simply specify the options on the command-line. Normally, the only thing you might find convenient to place in the `xtrabackup` section of your `my.cnf` file is the `target_dir` option to default the directory in which the backups will be placed, for example:

```
[xtrabackup]
target_dir = /data/backups/mysql/
```



This manual will assume that you do not have any file-based configuration for xtrabackup, so it will always show command-line options being used explicitly. Please see the option and variable reference for details on all the configuration options.

The xtrabackup binary does not accept exactly the same syntax in the `my.cnf` file as the `mysqld` server binary does. For historical reasons, the `mysqld` server binary accepts parameters with a `--set-variable=<variable>=<value>` syntax, which xtrabackup does not understand. If your `my.cnf` file has such configuration directives, you should rewrite them in the `--variable=value` syntax.

System configuration and NFS volumes

The xtrabackup tool requires no special configuration on most systems. However, the storage where the `-target-dir` is located must behave properly when `fsync()` is called. In particular, we have noticed that if you mount the NFS volume without the `sync` option the NFS volume does not sync the data. As a result, if you back up to an NFS volume mounted with the `async` option, and then try to prepare the backup from a different server that also mounts that volume, the data might appear to be corrupt. You can use the `sync` mount option to avoid this problem.

The xbcrypt binary

The xbcrypt binary overview

To support encryption and decryption of the backups, a new tool `xbcrypt` was introduced to *Percona XtraBackup*.

Percona XtraBackup 8.0.28-20 implements the `XBCRYPT_ENCRYPTION_KEY` environment variable. The variable is only used in place of the `--encrypt_key=name` option. You can use the environment variable or command line option. If you use both, the command line option takes precedence over the value specified in the environment variable.

This utility has been modeled after The [xbstream binary](#) to perform encryption and decryption outside of *Percona XtraBackup*.

For more details on the command-line options, see [xbcrypt command-line options](#)

The xbcrypt command-line options

Usage:

```
$ xbcrypt[OPTIONS]
```

The `xbcrypt` binary has the following command line options:

decrypt

usage: `-d` `--decrypt`

Decrypt data input to output.

encrypt-algo

usage: `-a=name` `--encrypt-algo=name`

Defines the name of the encryption algorithm.

encrypt-chunk-size

usage: `-s=#` `--encrypt-chunk-size=#`

Defines the size of the working buffer for encryption in bytes. The default value is `64000`.

encrypt-key

usage: `-k=name` `--encrypt-key=name`

The name of the encryption key.

encrypt-key-file

usage: `-f=name` `--encrypt-key-file=name`

The name of the file that contains the encryption key.

encrypt-threads

usage: `--encrypt-threads=#`

This option specifies the number of worker threads used for parallel encryption/decryption. The default value is 1.

input

usage: `-i=name` `--input=name`

Defines the name of the optional input file. If the name is not specified, the input reads from the standard input.

output

usage: `-o=name` `--output=name`

Defines the name of the optional output file. If this name is not specified, the output is written to the standard output.

read-buffer-size

usage: `--read-buffer-size=#`

Read the buffer size. The default value is 10MB.

verbose

usage: `-v` `--verbose`

Display status in verbose mode.

The xbstream binary

The xbstream binary overview

To support simultaneous compression and streaming, a new custom streaming format called xbstream was introduced to Percona XtraBackup in addition to the TAR format. That was required to overcome some limitations of traditional archive formats such as tar, cpio and others which did not allow streaming dynamically generated files, for example dynamically compressed files. Other advantages of xbstream over traditional streaming/archive format include ability to stream multiple files concurrently (so it is possible to use streaming in the xbstream format together with the `-parallel` option) and more compact data storage.

For details on the command-line options, see [xbcloud command-line options](#).

The utility tries to minimize its impact on the OS page cache by using the appropriate `posix_fadvise()` calls when available.

When compression is enabled with xtrabackup all data is compressed, including the transaction log file and meta data files, using the specified compression algorithm. Read more about supported compression algorithms in the [Create a compressed backup](#) document.

The xbstream command-line options

```
$ xbstream -c [OPTIONS]
$ xbstream -x [OPTIONS]
```

This utility has a tar-like interface.

The xbstream binary has the following options:

absolute-names

Usage: `-absolute-names`

Do not strip the leading `//` (slashes) from the file names when creating archives.

c

Usage: `-c`

Streams files specified on the command line to its standard output.

decompress

Usage: `--decompress`

Decompress the individual backup files

decompress-threads

Usage: `--decompress-threads=#`

The number of threads for parallel data decompression. The default value is 1.

decrypt

Usage: `--decrypt=name`

Specifies that xbstream automatically decrypts encrypted files when extracting input stream. The supported values are: `AES128`, `AES192`, and `AES256`.

You can specify either `--encrypt-key` or `--encrypt-key-file` to provide the encryption key, but do not use both options.

directory

Usage: `--directory=name`

Change the current directory to this named directory before streaming or extracting.

encrypt-key

Usage: `--encrypt-key=name`

Specify the encryption key used. Do not use this option with `--encrypt-key-file`; these options are mutually exclusive.

encrypt-key-file

Usage: `--encrypt-key-file=name`

Specify the file that contains the encryption key. Do not use this option with `--encrypt-key`; these options are mutually exclusive.

encrypt-threads

Usage: `--encrypt-threads-=#`

Specify the number of threads for parallel data encryption. The default value is `1`.

extract

Usage: `--extract`

Extract to disk the files from the stream to the standard input.

fifo-dir

Usage: `--fifo-dir=name`

The directory used to read or write named pipes.

fifo-streams

Usage: `--fifo-streams=#`

The number of FIFO files (named pipes) to use for parallel data file streaming. To disable the option and send the stream to STDOUT, set the value to 1.

fifo-timeout

Usage: `--fifo-timeout=#`

The number of seconds to wait for the other end to open the stream. The default value is 60.

parallel

Usage: `--parallel=#`

Defines the number of worker threads for reading or writing.

x

Usage: `-x`

Extracts files from the stream read from its standard input to the current directory unless specified otherwise with the `-c` option. Support for parallel extraction with the `--parallel` option

verbose

Usage: `--verbose`

Print verbose output

The xbccloud binary

The xbcloud binary overview

The purpose of xbcloud is to download from the cloud and upload to the cloud the full or part of an xstream archive. xbcloud will not overwrite the backup with the same name. xbcloud accepts input via a pipe from xstream so that it can be invoked as a pipeline with xtrabackup to stream directly to the cloud without needing a local storage.

Note

In a Bash shell, the `$?` parameter returns the exit code from the last binary. If you use pipes, the `${PIPESTATUS[x]}` array parameter returns the exit code for each binary in the pipe string.

```
xtrabackup --backup --stream=xstream --target-dir=/storage/backups/ | xbcloud put [options]
full_backup
...
$ ${PIPESTATUS[x]}
0 0
$ true | false
$ echo $?
1

# with PIPESTATUS
$ true | false
$ echo ${PIPESTATUS[0]} ${PIPESTATUS[1]}
0 1
```

The xbcloud binary stores each chunk as a separate object with a name `backup_name/database/table.ibd.NNN...`, where `NNN...` is a 0-padded serial number of chunk within a file. The size of chunk produced by xtrabackup and xstream changed to 10M.

To adjust the chunk size use `--read-buffer-size`. To adjust the chunk size for encrypted files, use `--read-buffer-size` and `--encrypt-chunk-size`.

xbcloud has three essential operations: *put*, *get*, and *delete*. With these operations, backups are created, stored, retrieved, restored, and deleted. xbcloud operations clearly map to similar operations within the AWS Amazon S3 API.

For details on the command-line options, see the [xbcloud command line options](#).

The [Exponential Backoff](#) feature was implemented in Percona XtraBackup 8.0.26-18. During the backup, a chunk fails to upload or download operation; this feature adds an exponential backoff, a sleep time, and retries the operation. This feature increases the chances of completing a backup or a restore operation.

The [FIFO datasink](#) feature was implemented in Percona XtraBackup 8.0.33-28. With the `FIFO` data sink feature, users with a streaming capacity of 10Gbps (typically on a Local Area Network (LAN)) can benefit from faster backups by streaming data in parallel to object storage.

Important

To prevent intermittent backup failures, [update the curl utility](#).

Supported cloud storage types

The following cloud storage types are supported:

- OpenStack Object Storage (Swift) - see [Using the xbcloud binary with Swift](#)
- Amazon Simple Storage (S3) - see [Using the xbcloud binary with Amazon S3](#)
- Azure Cloud Storage - see [Using the xbcloud binary with Microsoft Azure Cloud Storage](#)
- Google Cloud Storage (gcs) - see [Using the xbcloud binary with Google Cloud Storage](#)
- MinIO - see [Using the xbcloud binary with MinIO](#)

In addition to OpenStack Object Storage (Swift), which has been the only option for storing backups in a cloud storage until Percona XtraBackup 2.4.14, xbcloud supports Amazon S3, MinIO, and Google Cloud Storage. Other Amazon S3-compatible storages, such as Wasabi or Digital Ocean Spaces, are also supported.

See also

[OpenStack Object Storage\("Swift"\)](#)

[Amazon Simple Storage Service](#)

[MinIO](#)

[Google Cloud Storage](#)

[Wasabi](#)

[Digital Ocean Spaces](#)

Usage

The following sample command creates a full backup:

```
xtrabackup --backup --stream=xbstream --target-dir=/storage/backups/ --extra-  
lsndir=/storage/backups/ | xbccloud \  
put [options] full_backup
```

An incremental backup only includes the changes since the last backup. The last backup can be either a full or incremental backup.

The following sample command creates an incremental backup:

```
xtrabackup --backup --stream=xbstream --incremental-basedir=/storage/backups \  
--target-dir=/storage/inc-backup | xbccloud put [options] inc_backup
```

To prepare an incremental backup, you must first download the full backup with the following command:

```
xbccloud get [options] full_backup | xbstream -xv -C /tmp/full-backup
```

You must prepare the full backup:

```
xtrabackup --prepare --apply-log-only --target-dir=/tmp/full-backup
```

After the full backup has been prepared, download the incremental backup:

```
xbccloud get [options] inc_backup | xbstream -xv -C /tmp/inc-backup
```

The downloaded backup is prepared by running the following command:

```
xtrabackup --prepare --target-dir=/tmp/full-backup --incremental-dir=/tmp/inc-  
backup
```

You do not need the full backup to restore only a specific database. You can specify only the tables to be restored:

```
xbccloud get [options] ibdata1 sakila/payment.ibd > /tmp/partial.xbs
```

An example of the code:

```
xbstream -xv -C /tmp/partial < /tmp/partial/partial.xbs
```

Supplying parameters

Each storage type has mandatory parameters that you can supply on the command line, in a configuration file, and via environment variables.

Configuration files

The parameters the values of which do not change frequently can be stored in `my.cnf` or in a custom configuration file. The following example is a template of configuration options under the `[xbcloud]` group:

```
[xbcloud]
storage=s3
s3-endpoint=http://localhost:9000/
s3-access-key=minio
s3-secret-key=minio123
s3-bucket=backupsx
s3-bucket-lookup=path
s3-api-version=4
```



Note

If you explicitly use a parameter on the command line and in a configuration file, xbcloud uses the value provided on the command line.

Environment variables

If you explicitly use a parameter on the command line, in a configuration file, and the corresponding environment variable contains a value, xbcloud uses the value provided on the command line or in the configuration file.

Check the available options and corresponding environment variables for each supported cloud storage type below:

- OpenStack Object Storage (Swift) - see [Using the xbcloud binary with Swift](#)
- Amazon Simple Storage (S3) - see [Using the xbcloud binary with Amazon S3](#)
- Azure Cloud Storage - see [Using the xbcloud binary with Microsoft Azure Cloud Storage](#)
- Google Cloud Storage (gcs) - see [Using the xbcloud binary with Google Cloud Storage](#)
- MinIO - see [Using the xbcloud binary with MinIO](#)

Shortcuts

For all operations (put, get, and delete), you can use a shortcut to specify the storage type, bucket name, and backup name as one parameter instead of using three distinct parameters (`--storage`, `--s3-bucket`, and backup name per se).

Note

Use the following format: `storage-type://bucket-name/backup-name`

In this example `s3` refers to a storage type, `operator-testing` is a bucket name, and `bak22` is the backup name.

```
xbcloud get s3://operator-testing/bak22 ...
```

This shortcut expands as follows:

```
xbcloud get --storage=s3 --s3-bucket=operator-testing bak22 ...
```

You can supply the mandatory parameters on the command line, configuration files, and in environment variables.

Additional parameters

`xbcloud` accepts additional parameters that you can use with any storage type. The `--md5` parameter computes the MD5 hash value of the backup chunks. The result is stored in files that following the `backup_name.md5` pattern.

```
xtrabackup --backup --stream=xbstream \  
--parallel=8 2>backup.log | xbcloud put s3://operator-testing/bak22 \  
--parallel=8 --md5 2>upload.log
```

You may use the `--header` parameter to pass an additional HTTP header with the server side encryption while specifying a customer key.

An example of using the `--header` for AES256 encryption.

```
xtrabackup --backup --stream=xbstream --parallel=4 | \  
xbcloud put s3://operator-testing/bak-enc/ \  
--header="X-Amz-Server-Side-Encryption-Customer-Algorithm: AES256" \  
--header="X-Amz-Server-Side-Encryption-Customer-Key: CuStoMerKey=" \  
--header="X-Amz-Server-Side-Encryption-Customer-Key-MD5: CuStoMerKeyMd5==" \  
--parallel=8
```

The `--header` parameter is also useful to set the access control list (ACL) permissions: `--header="x-amz-acl: bucket-owner-full-control`

Incremental backups

First, you need to make the full backup on which the incremental one is going to be based:

```
xtrabackup --backup --stream=xbstream --extra-lsdir=/storage/backups/ \
--target-dir=/storage/backups/ | xbccloud put \
--storage=swift --swift-container=test_backup \
--swift-auth-version=2.0 --swift-user=admin \
--swift-tenant=admin --swift-password=xoxoxoxo \
--swift-auth-url=http://127.0.0.1:35357/ --parallel=10 \
full_backup
```

Then you can make the incremental backup:

```
xtrabackup --backup --incremental-basedir=/storage/backups \
--stream=xbstream --target-dir=/storage/inc_backup | xbccloud put \
--storage=swift --swift-container=test_backup \
--swift-auth-version=2.0 --swift-user=admin \
--swift-tenant=admin --swift-password=xoxoxoxo \
--swift-auth-url=http://127.0.0.1:35357/ --parallel=10 \
inc_backup
```

Preparing incremental backups

To prepare a backup you first need to download the full backup:

```
xbccloud get --swift-container=test_backup \
--swift-auth-version=2.0 --swift-user=admin \
--swift-tenant=admin --swift-password=xoxoxoxo \
--swift-auth-url=http://127.0.0.1:35357/ --parallel=10 \
full_backup | xbstream -xv -C /storage/downloaded_full
```

Once you download the full backup, the full backup should be prepared:

```
xtrabackup --prepare --apply-log-only --target-dir=/storage/downloaded_full
```

After the full backup has been prepared you can download the incremental backup:

```
xbccloud get --swift-container=test_backup \
--swift-auth-version=2.0 --swift-user=admin \
--swift-tenant=admin --swift-password=xoxoxoxo \
--swift-auth-url=http://127.0.0.1:35357/ --parallel=10 \
inc_backup | xbstream -xv -C /storage/downloaded_inc
```

Once the incremental backup has been downloaded, you can prepare the incremental backup by running:

```
xtrabackup --prepare --apply-log-only \  
--target-dir=/storage/downloaded_full \  
--incremental-dir=/storage/downloaded_inc
```

and

```
xtrabackup --prepare --target-dir=/storage/downloaded_full
```

Note

Accelerate the prepare phase for incremental backups with many InnoDB Data (IBD) files by using the `--parallel` option to process delta files concurrently. When using `--parallel` in the prepare phase, always specify a numeric value. The recommended minimum value is 4 (for example, `--parallel=4`).

```
xtrabackup --prepare --apply-log-only \  
--target-dir=/storage/downloaded_full \  
--incremental-dir=/storage/downloaded_inc --parallel=4  
  
xtrabackup --prepare --target-dir=/storage/downloaded_full --parallel=4
```

Partial download of the cloud backup

If you do not want to download the entire backup to restore the specific database you can specify only the tables you want to restore:

```
xbcloud get --swift-container=test_backup  
--swift-auth-version=2.0 --swift-user=admin \  
--swift-tenant=admin --swift-password=xoxoxoxo \  
--swift-auth-url=http://127.0.0.1:35357/ full_backup \  
ibdata1 sakila/payment.ibd \  
> /storage/partial/partial.xbs  
  
xbstream -xv -C /storage/partial < /storage/partial/partial.xbs
```

The xbccloud command-line options

Usage:

```
$ xbccloud put [OPTIONS]
$ xbccloud get [OPTIONS]
$ xbccloud delete [OPTIONS]
```

This document contains information on general options and the options available when you select the [Swift authentication version using `swift-auth-version`](#).

The xbccloud binary has the following general command line options:

azure-access-key

Usage: `--azure-access-key=name`

The name of the Azure access-key

azure-container-name

Usage: `--azure-container-name=name`

The name of the Azure container

azure-development-storage

Usage: `--azure-development-storage=name`

If you run the [Azurite emulator](#), use this option, and it works with the default credentials provided by Azurite. You can overwrite these default credentials with other options.

azure-endpoint

Usage: `--azure-endpoint=name`

The name of the Azure endpoint

azure-storage-account

Usage: `--azure-storage-account=name`

The name of the Azure storage account

azure-tier-class

Usage: `--azure-tier-class=name`

The name of the Azure tier-class. The possible values are:

- Hot
- Cool
- Archive

cacert

Usage: `--cacert=name`

The path to the file with Certificate Authority (CA) certificates.

curl-retriable-errors

Usage: `--curl-retriable-errors=name`

Add a new cURL error code as retriable. For multiple codes, use a comma-separated list.

fifo-dir

Usage: `--fifo-dir=name`

The directory used to read or write named pipes. In the put mode, xbcloud reads from the named pipes. In the get mode, xbcloud writes to the named pipes.

fifo-streams

Usage: `--fifo-streams=#`

The number of parallel FIFO stream threads.

fifo-timeout

Usage: `--fifo-timeout=#`

The number of seconds to wait for the other end to open the stream. The default value is 60.

google-access-key

Usage: `--google-access-key=name`

The Google Cloud storage access key

google-bucket-name

Usage: `--google-bucket-name=name`

The Google Cloud storage bucket name

google-endpoint

Usage: `--google-endpoint=name`

The Google Cloud storage endpoint

google-region

Usage: `--google-region=name`

The Google Cloud storage region

google-secret-key

Usage: `--google-secret-key=name`

The Google Cloud storage secret key

google-session-token

Usage: `--google-session-token=name`

The Google Cloud storage session-token

google-storage-class

Usage: `--google-storage-class=name`

The Google Cloud storage storage class

The possible values are the following:

- STANDARD
- NEARLINE
- COLDLINE
- ARCHIVE

header

Usage: `--header=name`

Extra header

http-retriable-errors

Usage: `--http-retriable-errors=name`

Add a new http error code as retriable. For multiple codes, use a comma-separated list.

insecure

Usage: `--insecure`

Do not verify the certificate of the server.

max-backoff

Usage: `--max-backoff=#`

The maximum backoff delay, in milliseconds, between chunk upload or chunk download retries. The default value is 300000.

max-retries

Usage: `--max-retries=#`

The number of retries for chunk uploads or downloads after a failure. The default value is 10.

md5

Usage: `--md5=name`

Uploads an MD5 file into the backup directory.

parallel

Usage: `--parallel=#`

Defines the maximum number of concurrent upload/download requests. The default value is 1.

s3-access-key

Usage: `--s3-access-key=name`

The name of the s3 access key

s3-api-version

Usage: `--s3-api-version=name`

The name of the s3 API version

s3-bucket

Usage: `--s3-bucket=name`

The name of the s3 bucket

s3-bucket-lookup

Usage: `--s3-bucket-lookup=name`

The name of the s3 bucket lookup method

s3-endpoint

Usage: `--s3-endpoint=name`

The name of the s3 endpoint

s3-region

Usage: `--s3-region=name`

The name of the s3 region

s3-secret-key

Usage: `--s3-secret-key=name`

The name of the s3 secret key

s3-session-token

Usage: `--s3-session-token=name`

The name of the s3 session token

s3-storage-class

Usage: `--s3-storage-class=name`

The name of the s3 storage class and is used to pass custom storage class names provided by the other s3 implementations, such as MinIO.

The possible values are:

- STANDARD
- STANDARD_ID
- GLACIER

storage

Usage: `--storage=[S3|SWIFT|GOOGLE|AZURE]`

Defines the Cloud storage option. xcloud supports Swift, MinIO, Google, Azure, and AWS S3. The default value is `swift`.

swift-auth-url

Usage: `--swift-auth-url=name`

The Base URL of the Swift authentication service.

swift-container

Usage: `--swift-container=name`

The Swift container used to store backups.

swift-key

Usage: `--swift-key`

The Swift key/password (X-Auth-Key)

swift-storage-url

Usage: `--swift-storage-url=name`

If a name is specified, the xcloud attempts to get the object-store endpoint for a given region from the keystone response. This option overrides that value.

swift-user

Usage: `--swift-user=name`

The Swift user name (X-Auth-User).

timeout

Usage: `--timeout=#`

The number of seconds to wait for activity on the TCP connection. The default value is 120. If the value is 0 (zero), there is no timeout.

verbose

Usage: `--verbose`

Turns on the cURL tracing.

Swift authentication options

The Swift specification describes several [authentication options](#). The *xbcloud* tool can authenticate against keystone with API version 2 and 3.

swift-auth-version

Usage: `--swift-auth-version=name`

Defines the swift authentication version used.

The possible values are the following:

- `1.0` - TempAuth
- `2.0` - Keystone v2.0
- `3` - Keystone v3.

The default value is `1.0`.

If `--swift-auth-version=2`, the additional options are:

swift-password

Usage: `--swift-password=name`

Swift password for the user

swift-region

Usage: `--swift-region=name`

Swift region object-store endpoint

swift-tenant

Usage: `--swift-tenant=name`

The Swift tenant name

Either of the `--swift-tenant` or `--swift-tenant-id` can be defined, but you should not use both options together. Both of these options are optional.

swift-tenant-id

Usage: `--swift-tenant-id=name`

The Swift tenant ID

Either of the `--swift-tenant` or `--swift-tenant-id` can be defined, but you should not use both options together. Both of these options are optional.

If `--swift-auth-version=3`, the additional options are:

swift-domain

Usage: `--swift-domain=name`

The Swift domain name

swift-domain-id

Usage: `--swift-domain-id=name`

The Swift domain ID

swift-project

Usage: `--swift-project=name`

The Swift project name

swift-project-domain

Usage: `--swift-project-domain=name`

The Swift domain name

swift-project-domain-id

Usage: `--swift-project-domain-id=name`

The Swift domain ID

swift-project-id

Usage: `--swift-project-id=name`

The Swift project ID

swift-user-id

Usage: `--swift-user-id=name`

The Swift user ID

Use the xbccloud binary

Use the xbcloud binary with Amazon S3

Create a full backup with Amazon S3

```
$ xtrabackup --backup --stream=xbstream --extra-lsdir=/tmp --target-dir=/tmp | \
xcloud put --storage=s3 \
--s3-endpoint='s3.amazonaws.com' \
--s3-access-key='YOUR-ACCESSKEYID' \
--s3-secret-key='YOUR-SECRETACCESSKEY' \
--s3-bucket='mysql_backups'
--parallel=10 \
$(date +%Y-%m-%d)-full_backup
```

The following options are available when using Amazon S3:

| Option | Details |
|---------------------|--|
| -s3-access-key | Use to supply the AWS access key ID |
| -s3-secret-key | Use to supply the AWS secret access key |
| -s3-bucket | Use supply the AWS bucket name |
| -s3-region | Use to specify the AWS region. The default value is us-east-1 |
| -s3-api-version = | Select the signing algorithm. The default value is AUTO. In this case, xbcloud will probe. |
| -s3-bucket-lookup = | Specify whether to use bucket.endpoint.com or endpoint.com/bucket*style requests. The default value is AUTO. In this case, xbcloud will probe. |

| Option | Details |
|---------------------------------|--|
| <code>-s3-storage-class=</code> | <p>Specify the S3 storage class. The default storage class depends on the provider. The name options are the following:</p> <ul style="list-style-type: none"> • STANDARD • STANDARD_IA • GLACIER <p>NOTE If you use the GLACIER storage class, the object must be restored to S3 before restoring the backup. Also supports using custom S3 implementations such as MinIO or CephRadosGW.</p> |

Permissions setup

Following the principle of “least-privilege”, these are the minimum bucket permissions needed for xbccloud to write backups to S3: ***LIST/PUT/GET/DELETE***.

The following example shows the policy definition for writing to the `xbccloud-testing` bucket on the AWS S3 storage.

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "*"
      },
      "Action": [
        "s3:ListBucket"
      ],
      "Resource": "arn:aws:s3:::xbcloud-testing"
    },
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "*"
      },
      "Action": [
        "s3:PutObject",
        "s3:PutObjectAcl",
        "s3:GetObject",
        "s3:GetObjectAcl",
        "s3:DeleteObject"
      ],
      "Resource": "arn:aws:s3:::xbcloud-testing/*"
    }
  ]
}

```

Environment variables

The following environment variables are recognized. xbcloud maps them automatically to corresponding parameters applicable to the selected storage.

- AWS_ACCESS_KEY_ID (or ACCESS_KEY_ID)
- AWS_SECRET_ACCESS_KEY (or SECRET_ACCESS_KEY)
- AWS_DEFAULT_REGION (or DEFAULT_REGION)
- AWS_ENDPOINT (or ENDPOINT)
- AWS_CA_BUNDLE

Restore with S3

```
$ xbccloud get s3://operator-testing/bak22 \  
--s3-endpoint=https://storage.googleapis.com/ \  
--parallel=10 2>download.log | xbstream -x -C restore --parallel=8
```



Use the xbccloud binary with an IAM instance profile

This feature is [tech preview](#). Before using this feature in production, we recommend that you test restoring from physical backups in your environment and also use the alternative backup method for redundancy.

[Percona XtraBackup 8.0.31-24](#) adds the ability to use the IAM instance profile when running xbccloud from an EC2 instance.

An authentication system has two elements:

- Who am I?
- What can I do?

A role defines “what can I do.” A role provides a method to define a collection of permissions. Roles are assigned to users, services and EC2 instances, the “who am I” element.

The IAM instance profile is the “who” for an EC2 instance and assumes the IAM role, which has permissions. The instance profile has the same name as the IAM role.

An IAM instance profile does not need the `--s3-secret-key` nor the `--s3-access-key` if they are running `xbccloud` from an Amazon EC2 instance. To configure or attach an instance metadata to an EC2 instance, see [How can I grant my Amazon EC2 instance access to an Amazon S3 bucket](#).

An example of the command:

```
$ xtrabackup ... | xbccloud put --storage=s3 --s3-bucket=bucket-name backup-name
```

The xbccloud binary outputs a connect message when successful.

Expected output

```
221121 13:16:26 Using instance metadata for access and secret key
221121 13:16:26 xbccloud: Successfully connected.
```

An important consideration is that the instance metadata has a time to live (TTL) of 6 hours. A backup that takes more than that time contains Expired token errors. Use [Exponential Backoff](#) to retry the upload after fetching new keys from the instance metadata.

Output when keys have expired



```
221121 13:04:52 xbcloud: S3 error message: The provided token has expired.
221121 13:04:52 xbcloud: Sleeping for 2384 ms before retrying
test/mysql.ibd.00000000000000000002 [1]
221121 13:04:55 xbcloud: S3 error message: The provided token has expired.
221121 13:04:55 xbcloud: Sleeping for 2887 ms before retrying
test/mysql.ibd.00000000000000000003 [1]
221121 13:04:58 xbcloud: S3 error message: The provided token has expired.
221121 13:04:58 xbcloud: Sleeping for 2778 ms before retrying
test/undo_002.00000000000000000000 [1]
221121 13:05:00 xbcloud: S3 error message: The provided token has expired.
221121 13:05:00 xbcloud: Sleeping for 2916 ms before retrying
test/undo_002.00000000000000000001 [1]
221121 13:05:03 xbcloud: S3 error message: The provided token has expired.
221121 13:05:03 xbcloud: Sleeping for 2794 ms before retrying
test/undo_002.00000000000000000002 [1]
221121 13:05:06 xbcloud: S3 error message: The provided token has expired.
221121 13:05:06 xbcloud: Sleeping for 2336 ms before retrying
test/undo_001.00000000000000000000 [1]
221121 13:05:09 xbcloud: successfully uploaded chunk: test/mysql.ibd.00000000000000000002,
size: 5242923
221121 13:05:09 xbcloud: successfully uploaded chunk: test/mysql.ibd.00000000000000000003,
size: 23
221121 13:05:09 xbcloud: successfully uploaded chunk: test/undo_002.00000000000000000000,
size: 10485802
221121 13:05:09 xbcloud: successfully uploaded chunk: test/undo_002.00000000000000000001,
size: 6291498
221121 13:05:09 xbcloud: successfully uploaded chunk: test/undo_002.00000000000000000002,
size: 22
221121 13:05:09 xbcloud: successfully uploaded chunk: test/undo_001.00000000000000000000,
size: 10485802
221121 13:05:10 xbcloud: successfully uploaded chunk: test/undo_001.00000000000000000001,
size: 6291498
221121 13:05:10 xbcloud: successfully uploaded chunk: test/undo_001.00000000000000000002,
size: 22
. . .
221121 13:05:18 xbcloud: successfully uploaded chunk:
test/xtrabackup_tablespace.00000000000000000001, size: 36
221121 13:05:19 xbcloud: Upload completed.
```

Use the xbcloud binary with Swift

Create a full backup with Swift

For details on the command-line options, see the [xbcloud command-line options](#)

The following example shows how to make a full backup and upload it to Swift.

```
$ xtrabackup --backup --stream=xbstream --extra-lsdir=/tmp --target-dir=/tmp | \
xbcloud put --storage=swift \
--swift-container=test \
--swift-user=test:tester \
--swift-auth-url=http://192.168.8.80:8080/ \
--swift-key=testing \
--parallel=10 \
full_backup
```

The following [OpenStack](#) environment variables are also recognized and mapped automatically to the corresponding swift parameters (`--storage=swift`):

- OS_AUTH_URL
- OS_TENANT_NAME
- OS_TENANT_ID
- OS_USERNAME
- OS_PASSWORD
- OS_USER_DOMAIN
- OS_USER_DOMAIN_ID
- OS_PROJECT_DOMAIN
- OS_PROJECT_DOMAIN_ID
- OS_REGION_NAME
- OS_STORAGE_URL
- OS_CACERT

Restore with Swift

```
$ xbcloud get [options] <name> [<list-of-files>] | xbstream -x
```

The following example shows how to fetch and restore the backup from Swift:

```
$ xbccloud get --storage=swift \  
--swift-container=test \  
--swift-user=test:tester \  
--swift-auth-url=http://192.168.8.80:8080/ \  
--swift-key=testing \  
full_backup | xbstream -xv -C /tmp/downloaded_full  
  
$ xbccloud delete --storage=swift --swift-user=xtrabackup \  
--swift-password=xtrabackup123! --swift-auth-version=3 \  
--swift-auth-url=http://openstack.ci.percona.com:5000/ \  
--swift-container=mybackup1 --swift-domain=Default
```



Use the xbccloud binary with Google Cloud Storage

Create a full backup with Google Cloud Storage

The support for Google Cloud Storage is implemented using the interoperability mode. This mode was especially designed to interact with cloud services compatible with Amazon S3.

See also

[Cloud Storage Interoperability](#)

```
$ xtrabackup --backup --stream=xbstream --extra-lsdir=/tmp --target-dir=/tmp | \
xbcloud put --storage=google \
--google-endpoint='storage.googleapis.com' \
--google-access-key='YOUR-ACCESSKEYID' \
--google-secret-key='YOUR-SECRETACCESSKEY' \
--google-bucket='mysql_backups' \
--parallel=10 \
$(date -I)-full_backup
```

The following options are available when using Google Cloud Storage:

- `-google-access-key =`
- `-google-secret-key =`
- `-google-bucket =`
- `-google-storage-class=name`

Note

The Google storage class name options are the following:

- STANDARD
- NEARLINE
- COLDLINE
- ARCHIVE



See also

[Google storage classes - the default Google storage class depends on the storage class of the bucket](#)

Use the xbcloud binary with Microsoft Azure Cloud Storage

Implemented in Percona XtraBackup 8.0.27-19, the **xbcloud** binary adds support for the Microsoft Azure Cloud Storage using the REST API.

Options

The following are the options, environment variables, and descriptions for uploading a backup to Azure using the REST API. The environment variables are recognized by **xbcloud**, which maps them automatically to the corresponding parameters:

| Option name | Environment variables | Description |
|--|------------------------------------|--|
| <code>-azure-storage-account=name</code> | <code>AZURE_STORAGE_ACCOUNT</code> | An Azure storage account is a unique namespace to access and store your Azure data objects. |
| <code>-azure-container-name=name</code> | <code>AZURE_CONTAINER_NAME</code> | A container name is a valid DNS name that conforms to the Azure naming rules |
| <code>-azure-access-key=name</code> | <code>AZURE_ACCESS_KEY</code> | A generated key that can be used to authorize access to data in your account using the Shared Key authorization. |
| <code>-azure-endpoint=name</code> | <code>AZURE_ENDPOINT</code> | The endpoint allows clients to securely access data |
| <code>-azure-tier-class=name</code> | <code>AZURE_STORAGE_CLASS</code> | <p>Cloud tier can decrease the local storage required while maintaining the performance. When enabled, this feature has the following categories:</p> <p>Hot - Frequently accessed or modified data</p> <p>Cool - Infrequently accessed or modified data</p> <p>Archive - Rarely accessed or modified data</p> |

Test your Azure applications with the [Azurite open-source emulator](#). For testing purposes, the **xbcloud** binary adds the `--azure-development-storage` option that uses the default `access_key` and `storage account` of azurite and `testcontainer` for the container. You can overwrite these options, if needed.

Usage

All the available options for **xbcloud**, such as `parallel`, `max-retries`, and others, can be used. For more information, see the [xbcloud binary overview](#).

Examples

An example of an **xbcloud** backup.

```
$ xtrabackup --backup --stream=xbstream |  
xbcloud put backup_name --azure-storage-account=pxbtesting --azure-access-  
key=$AZURE_KEY --azure-container-name=test --storage=azure
```



An example of restoring a backup from **xbcloud**.

```
$ xbcloud get backup_name --azure-storage-account=pxbtesting  
--azure-access-key=$AZURE_KEY --azure-container-name=test --storage=azure --  
parallel=10 2>download.log | xbstream -x -C restore
```



An example of deleting a backup from **xbcloud**.

```
$ xbcloud delete backup_name --azure-storage-account=pxbtesting  
--azure-access-key=$AZURE_KEY --azure-container-name=test --storage=azure
```



An example of using a shortcut restore.

```
$ xbcloud get azure://operator-testing/bak22 ...
```



Use the xbccloud binary with MinIO

Create a full backup with MinIO

```
$ xtrabackup --backup --stream=xbstream --extra-lsdir=/tmp --target-dir=/tmp | \
xbcloud put --storage=s3 \
--s3-endpoint='play.minio.io:9000' \
--s3-access-key='YOUR-ACCESSKEYID' \
--s3-secret-key='YOUR-SECRETACCESSKEY' \
--s3-bucket='mysql_backups'
--parallel=10 \
$(date -I)-full_backup
```

Dependency compatibility and limitations

xbcloud depends on several system libraries for correct operation. Some versions of these dependencies are known to cause unexpected terminations or unstable behavior and are therefore not supported.

If you are experiencing unexpected terminations or intermittent failures, verify that your environment does not use any of the versions listed below.

libcurl

Certain versions of `libcurl` are known to cause unexpected terminations or connection reuse issues in `xbcloud`. The following versions are not supported.

Unsupported versions

`libcurl 7.64.0` the default curl version in Debian 10

- Known to mishandle reuse of closed connections
- May cause intermittent backup failures in `xbcloud`

For more details, see [curl #3750](#) or [curl #3763](#).

`libcurl versions ≥ 8.11.1 and < 8.12.0`

- Affected by double close of an eventfd file descriptor, [CVE-2025-0665](#)
- May cause unexpected terminations when `xbcloud` uses `libcurl`

Upgrade or downgrade options

You should always upgrade `xbcloud` to the latest version, as it includes fixes for all known issues and compatibility problems.

If the installed `libcurl` version is not supported, update it depending on the operating system and available repositories.

- Recommended - Upgrade to `libcurl 8.12.0 or later` if these versions are compatible with your operating system.

- If upgrading to the latest version `≥ 8.12.0` is not possible, downgrade to a supported version earlier than 8.11.1 to avoid the vulnerability associated with [CVE-2025-0665](#). Ensure that the downgrade is compatible with other system dependencies.

Check the installed libcurl version

To check the installed libcurl version, run:

```
curl --version
```



Upgrade libcurl – Recommended

The following steps show how to update the libcurl to version 8.12.0 on Debian 12. Administrator privileges are necessary for installing packages and configuring system services.

- 1 Edit the `/etc/apt/sources.list` to add the following:

```
deb http://deb.debian.org/debian bookworm-backports main
```



- 2 Refresh the `apt` sources:

```
sudo apt update
```



- 3 Install the version from `bookworm-backports`:

```
sudo apt install -t bookworm-backports libcurl4 curl
```



- 4 Verify the version number:

```
curl --version
```



Expected output



```
curl 8.12.0 (x86_64-pc-linux-gnu) libcurl/8.12.0
```

Downgrade libcurl

The following steps show how to downgrade the libcurl to the default Debian 12 version 7.88.1. Administrator privileges are necessary for installing packages and configuring system services.

1 Update the package list:

```
sudo apt update
```

2 Reinstall the default system version of libcurl and curl:

```
sudo apt install --reinstall libcurl4 curl
```

3 Verify the version number:

```
curl --version
```



Expected output



```
curl 7.88.1 (x86_64-pc-linux-gnu) libcurl/7.88.1 OpenSSL/3.0.9 zlib/1.2.13 libidn2/2.3.4  
libssh/0.10.5  
Release-Date: 2023-08-02  
Protocols: dict file ftp ftps gopher http https imap imaps ldap ldaps mqtt pop3 pop3s  
rtsp smb smbs smtp smtps telnet tftp  
Features: AsynchDNS IDN IPv6 Largefile NTLM NTLM_WB SSL libz TLS-SRP HTTP2 HSTS HTTPS-  
proxy Metalink
```

Install

Install Percona XtraBackup 8.0 overview

We recommend that you install Percona XtraBackup 8.0 from the official Percona software repositories using the appropriate package manager for your system:

- [Use an APT repo to install Percona XtraBackup](#)
- [Use a YUM repo to install Percona XtraBackup](#)

Installation alternatives

Percona also provides the following methods:

- [Use DEB downloaded packages to install Percona XtraBackup](#)
- [Use RPM downloaded packages to install Percona XtraBackup](#)
- [Install Percona XtraBackup from a Binary Tarball](#) with all the required files and binaries for a manual installation
- [Compile and install Percona XtraBackup from source code](#)
- [Run Percona XtraBackup in a Docker container](#)

Before installing Percona XtraBackup with any of these methods, we recommend that you review the [release notes](#) and [Server version and backup version comparison](#) and, if you are using Percona XtraBackup 8.0.34 or later, the latest list of Supported versions.

Before you start

Connection and privileges needed

Percona XtraBackup needs to be able to connect to the database server and perform operations on the server and the datadir when creating a backup, when preparing in some scenarios and when restoring it. In order to do so, there are privileges and permission requirements on its execution that must be fulfilled.

Privilege refers to the operations that a system user is permitted to do in the database server. **They are set at the database server and only apply to users in the database server.**

Permissions are those which permits a user to perform operations on the system, like reading, writing or executing on a certain directory or start/stop a system service. **They are set at a system level and only apply to system users.**

When *xtrabackup* is used, there are two actors involved: the user invoking the program - *a system user* - and the user performing action in the database server - *a database user*. Note that these are different users in different places, even though they may have the same username.

All the invocations of *xtrabackup* in this documentation assume that the system user has the appropriate permissions, and you are providing the relevant options for connecting the database server - besides the options for the action to be performed - and the database user has adequate privileges.

Connect to the server

The database user used to connect to the server and its password are specified by the `--user` and `--password` option:

```
$ xtrabackup --user=DVADER --password=14MY0URF4TH3R --backup \
--target-dir=/data/bkps/
```

If you don't use the `--user` option, *Percona XtraBackup* will assume the database user whose name is the system user executing it.

Other connection options

According to your system, you may need to specify one or more of the following options to connect to the server:

| Option | Description |
|--------|---|
| -port | Use this port when connecting to the database with TCP/IP |

| Option | Description |
|----------------------|--|
| <code>-socket</code> | Use this socket when connecting to the local database. |
| <code>-host</code> | Use this host when connecting to the database server with TCP/IP |

These options are passed to the **mysql** child process without alteration, see `mysql --help` for details.

Note

In case of multiple server instances, the correct connection parameters (port, socket, host) must be specified in order for *xtrabackup* to talk to the correct server.

Privileges needed

Once connected to the server, in order to perform a backup you need `READ` and `EXECUTE` permissions at a filesystem level in the server's datadir.

The database user needs the following privileges to back up tables or databases:

- `RELOAD` and `LOCK TABLES` (unless the `--no-lock` option is specified) in order to run `FLUSH TABLES WITH READ LOCK` and `FLUSH ENGINE LOGS` prior to start copying the files, and requires this privilege when [Backup Locks](#) are used
- `BACKUP_ADMIN` privilege is needed to query the `performance_schema.log_status` table, and run `LOCK INSTANCE FOR BACKUP`, `LOCK BINLOG FOR BACKUP`, or `LOCK TABLES FOR BACKUP`.
- `REPLICATION CLIENT` in order to obtain the binary log position,
- `CREATE TABLESPACE` in order to import tables (see Restoring Individual Tables),
- `PROCESS` in order to run `SHOW ENGINE INNODB STATUS` (which is mandatory), and optionally to see all threads which are running on the server (see `FLUSH TABLES WITH READ LOCK` option),
- `SUPER` in order to start/stop the replication threads in a replication environment, use [XtraDB Changed Page Tracking](#) for Incremental Backups and for handling `FLUSH TABLES WITH READ LOCK`,
- `CREATE` privilege in order to create the `PERCONA_SCHEMA.xtrabackup_history` database and table,
- `ALTER` privilege in order to upgrade the `PERCONA_SCHEMA.xtrabackup_history` database and table,
- `INSERT` privilege in order to add history records to the `PERCONA_SCHEMA.xtrabackup_history` table,

- `SELECT` privilege in order to use `--incremental-history-name` or `--incremental-history-uuid` in order for the feature to look up the `innodb_to_lsn` values in the `PERCONA_SCHEMA.xtrabackup_history` table.
- `SELECT` privilege on the [keyring_component_status table](#) to view the attributes and status of the installed keyring component when in use.
- `SELECT` privilege on the [replication_group_members table](#) to validate if the instance is part of group replication cluster.

A SQL example of creating a database user with the minimum privileges required to take full backups would be:

```
mysql> CREATE USER 'bkpuser'@'localhost' IDENTIFIED BY 's3cr%T';  
mysql> GRANT BACKUP_ADMIN, PROCESS, RELOAD, LOCK TABLES, REPLICATION CLIENT ON *.*  
TO 'bkpuser'@'localhost';  
mysql> GRANT SELECT ON performance_schema.log_status TO 'bkpuser'@'localhost';  
mysql> GRANT SELECT ON performance_schema.keyring_component_status TO  
bkpuser@'localhost';  
mysql> GRANT SELECT ON performance_schema.replication_group_members TO  
bkpuser@'localhost';  
mysql> FLUSH PRIVILEGES;
```

Query the privileges

To query the privileges that your database user has been granted at the console of the server execute:

```
mysql> SHOW GRANTS;
```

or for a particular user with:

```
mysql> SHOW GRANTS FOR 'db-user'@'host';
```

It will display the privileges using the same format as for the [GRANT statement](#).

Note that privileges may vary across versions of the server. To list the exact list of privileges that your server support (and a brief description of them) execute:

```
mysql> SHOW PRIVILEGES;
```



See also

[Permissions needed](#)

Permissions needed

We will be referring to *permissions* to the ability of a user to access and perform changes on the relevant parts of the host's filesystem, starting/stopping services and installing software.

By *privileges*, we refer to the abilities of a database user to perform different kinds of actions on the database server.

There are many ways for checking the permission on a file or directory. For example, `ls -ls /path/to/file` or `stat /path/to/file | grep Access` will do the job:

```
$ stat /etc/mysql | grep Access
```

The result could look like this:

Expected output

```
Access: (0755/drwxr-xr-x)  Uid: (    0/    root)   Gid: (    0/    root)
Access: 2011-05-12 21:19:07.129850437 -0300
$ ls -ld /etc/mysql/my.cnf
-rw-r--r-- 1 root root 4703 Apr  5 06:26 /etc/mysql/my.cnf
```

As in this example, `my.cnf` is owned by `root` and not writable for anyone else. Assuming that you do not have `root`'s password, you can check what permissions you have on these types of files with `sudo -l`:

```
$ sudo -l
```

The results could look like this:

Expected output

```
Password:
You may run the following commands on this host:
(root) /usr/bin/
(root) NOPASSWD: /etc/init.d/mysqld
(root) NOPASSWD: /bin/vi /etc/mysql/my.cnf
(root) NOPASSWD: /usr/local/bin/top
(root) NOPASSWD: /usr/bin/ls
(root) /bin/tail
```

Being able to execute with `sudo` scripts in `/etc/init.d/`, `/etc/rc.d/` or `/sbin/service` is the ability to start and stop services.

Also, If you can execute the package manager of your distribution, you can install or remove software with it. If not, having `rwX` permission over a directory will let you do a local installation of software by compiling it there. This is a typical situation in many hosting companies' services.

There are other ways for managing permissions, such as using *PolicyKit*, * Extended ACLs or SELinux*, which may be preventing or allowing your access. You should check them in that case.

 See also

[Connection and privileges needed](#)

Percona Software Download instructions

Select the software

Do the following steps to select the software:

1. Open [Software Downloads](#)
2. Locate the Percona Software, for example, Percona XtraBackup
3. In `Select Product`, select which product, for example, Percona XtraBackup 8.0
4. In `Select Product Version`, select the version, for example, PERCONA-XTRABACKUP-8.0.35-34
5. In `Select Platform`, select the operating system, for example, RED HAT ENTERPRISE LINUX / CENTOS/ ORACLE LINUX 9

The easiest method is to download all packages.

Download to a local computer

In `Package Download Options`, select a specific package or choose the `DOWNLOAD ALL PACKAGES` button. This action downloads the selected packages to the local computer.

Download to another computer

In `Package Download Options`, you can do either of the following:

- Select a specific package and right-click on the arrow next to `DOWNLOAD`, and in the drop-down menu, choose `Copy Link Address`
- Select `DOWNLOAD ALL PACKAGES` and right-click on the arrow, and in the drop-down menu, choose `Copy Link Address`

Paste the link in your terminal prompt to download the selected package.

Use APT

Install Percona XtraBackup on Debian-based systems

Ready-to-use packages are available from the Percona XtraBackup software repositories and the [download page](#).

Specific information on the supported platforms, products, and versions is described in [Percona Release Lifecycle Overview](#).

To prevent intermittent backup failures, [update the curl utility in Debian 10](#).

This guide walks you through the process of installing Percona XtraBackup 8.0 on Debian-based systems using the Percona-release tool.

Check out our [Get more help](#) for ways that we can work with you.

Version changes

Starting with [Percona XtraBackup 8.0.35-31](#), the APT builds for the following platforms include ARM packages with the `arm64.deb` extension:

- Debian 12
- Debian 11
- Ubuntu 24.04
- Ubuntu 22.04
- Ubuntu 20.04

This extension means Percona XtraBackup is now available for users on ARM-based systems.

We recommend downloading Percona XtraBackup for the same platform as your MySQL-compatible server. For instance, if your server runs on an ARM64 platform, you should download and use the Percona XtraBackup with the `arm64.deb` extension for that operating system.

Prerequisites

- Sudo or root access
- Internet connection

Installation steps

The following procedure installs Percona XtraBackup 8.0.

1 Update the system package list

```
$ sudo apt update
```

Example output

```
Get:1 http://ports.ubuntu.com/ubuntu-ports jammy InRelease [270 kB]
Get:2 http://ports.ubuntu.com/ubuntu-ports jammy-updates InRelease [128 kB]
Get:3 http://ports.ubuntu.com/ubuntu-ports jammy-backports InRelease [127 kB]
Get:4 http://ports.ubuntu.com/ubuntu-ports jammy-security InRelease [129 kB]
Get:5 http://ports.ubuntu.com/ubuntu-ports jammy/restricted arm64 Packages [24.2 kB]
Get:6 http://ports.ubuntu.com/ubuntu-ports jammy/multiverse arm64 Packages [224 kB]
Get:7 http://ports.ubuntu.com/ubuntu-ports jammy/universe arm64 Packages [17.2 MB]
Get:8 http://ports.ubuntu.com/ubuntu-ports jammy/main arm64 Packages [1758 kB]
Get:9 http://ports.ubuntu.com/ubuntu-ports jammy-updates/restricted arm64 Packages [2400
kB]
Get:10 http://ports.ubuntu.com/ubuntu-ports jammy-updates/multiverse arm64 Packages
[29.5 kB]
Get:11 http://ports.ubuntu.com/ubuntu-ports jammy-updates/universe arm64 Packages [1379
kB]
Get:12 http://ports.ubuntu.com/ubuntu-ports jammy-updates/main arm64 Packages [2217 kB]
Get:13 http://ports.ubuntu.com/ubuntu-ports jammy-backports/universe arm64 Packages
[31.8 kB]
Get:14 http://ports.ubuntu.com/ubuntu-ports jammy-backports/main arm64 Packages [81.0
kB]
Get:15 http://ports.ubuntu.com/ubuntu-ports jammy-security/main arm64 Packages [1949 kB]
Get:16 http://ports.ubuntu.com/ubuntu-ports jammy-security/restricted arm64 Packages
[2330 kB]
Get:17 http://ports.ubuntu.com/ubuntu-ports jammy-security/multiverse arm64 Packages
[24.1 kB]
Get:18 http://ports.ubuntu.com/ubuntu-ports jammy-security/universe arm64 Packages [1096
kB]
Fetched 31.4 MB in 7s (4264 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
9 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

2 Install the curl download utility if it's not installed:

```
$ sudo apt install curl
```

Example output

```
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  ca-certificates libbrotli1 libcurl4 libldap-2.5-0 libldap-common libnghttp2-14 libpsl5
  librtmp1
Suggested packages:
  ...
Updating certificates in /etc/ssl/certs...
0 added, 0 removed;
done.#####...]
Running hooks in /etc/ca-certificates/update.d...
done.
```

3 Download the percona-release the repository package:

```
$ curl -O https://repo.percona.com/apt/percona-release_latest.generic_all.deb
```

Example output

| % Total | % Received | % Xferd | Average Speed | Time | Time | Time | Current |
|-----------|------------|---------|---------------|-------|----------|----------|---------|
| | | | Dload Upload | Total | Spent | Left | Speed |
| 100 16510 | 100 16510 | 0 | 0 31433 | 0 | --:--:-- | --:--:-- | 31447 |

4 Install the downloaded package and its dependencies using apt:

```
$ sudo apt install gnupg2 lsb-release ./percona-release_latest.generic_all.deb
```

Example output

```
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
Note, selecting 'percona-release' instead of './percona-release_latest.generic_all.deb'
The following additional packages will be installed:
```

5 Refresh the local cache to update the package information:

```
$ sudo apt update
```



Example output



```
Hit:1 http://repo.percona.com/pmm2-client/apt jammy InRelease
Hit:2 http://ports.ubuntu.com/ubuntu-ports jammy InRelease
Hit:3 http://repo.percona.com/prel/apt jammy InRelease
...
```

6 Enable the repository:

```
$ sudo percona-release setup pxb-80
```



Example output



```
* Disabling all Percona Repositories
* Enabling the Percona XtraBackup 8.0 repository
Hit:1 http://repo.percona.com/prel/apt jammy InRelease
Get:2 http://repo.percona.com/pxb-80/apt jammy InRelease [18.7 kB]
Hit:3 http://ports.ubuntu.com/ubuntu-ports jammy InRelease
Get:4 http://ports.ubuntu.com/ubuntu-ports jammy-updates InRelease [128 kB]
Hit:5 http://repo.percona.com/telemetry/apt jammy InRelease
Get:6 http://repo.percona.com/pxb-80/apt jammy/main Sources [3393 B]
Get:7 http://ports.ubuntu.com/ubuntu-ports jammy-backports InRelease [127 kB]
Get:8 http://ports.ubuntu.com/ubuntu-ports jammy-security InRelease [129 kB]
Fetched 406 kB in 1s (274 kB/s)
Reading package lists... Done
```

7 Install the package:

```
$ sudo apt install percona-xtrabackup-80
```





Example output



```
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  libaio1 libcurl4-openssl-dev libdbd-mysql-perl libdbi-perl libev4 libgdbm-compat4
  libgdbm6
  libmysqlclient21 libperl5.34 libpopt0 mysql-common netbase perl perl-modules-5.34
  rsync zstd
...
Setting up libdbd-mysql-perl:amd64 (4.050-5ubuntu0.22.04.1)
...#####.....]
Setting up percona-xtrabackup-80 (8.0.35-31-1.jammy)
...#####.....]
Processing triggers for libc-bin (2.35-0ubuntu3.8)
...#####.....]
```

8 Verify the installation:

```
$ xtrabackup --version
```



Example output



```
xtrabackup version 8.0.35-31 based on MySQL server 8.0.35 Linux (x86_64) (revision id:
55ec21d7)
```

9 Install optional compression packages if you plan to use LZ4 or ZSTD compression:

Install the `lz4` package

```
$ sudo apt install lz4
```



Install the `zstd` package

```
$ sudo apt install zstd
```



Troubleshooting

If you encounter any errors during the installation process, check your system's error logs:

```
$ sudo tail -n 50 /var/log/syslog
```



Ensure your system meets the minimum requirements for Percona XtraBackup.

If you encounter permission errors, ensure you're running the commands with `sudo` or have root access.

For AppArmor profile information, see [Working with AppArmor](#).

Additional resources

To install Percona XtraBackup using downloaded deb packages, see [Install Percona XtraBackup 8.0](#).

To uninstall Percona XtraBackup, see [Uninstall Percona XtraBackup 8.0](#).

Files in the DEB package built for Percona XtraBackup 8.0

The following tables show what data each DEB package contains.

| Package | Contains |
|----------------------------|--|
| percona-xtrabackup-80 | The latest Percona XtraBackup GA binaries and associated files |
| percona-xtrabackup-dbg-80 | The debug symbols for binaries in percona-xtrabackup-80 |
| percona-xtrabackup-test-80 | The test suite for Percona XtraBackup |
| percona-xtrabackup | The older version of the Percona XtraBackup |

Install with DEB packages

Version changes

Starting with [Percona XtraBackup 8.0.35-31](#), the APT builds for the following platforms include ARM packages with the `arm64.deb` extension:

- Debian 12
- Debian 11
- Ubuntu 24.04
- Ubuntu 22.04
- Ubuntu 20.04

This `arm64.deb` extension means Percona XtraBackup is now available for users on ARM-based systems.

We recommend downloading Percona XtraBackup for the same platform as your MySQL-compatible server. For instance, if your server runs on an ARM64 platform, you should download and use the Percona XtraBackup with the `arm64.deb` extension for that operating system.

Download and install DEB packages

Download `DEB` packages of the desired series for your architecture from the [Percona Software Downloads](#). If needed, [Percona Software Download instructions](#) are available.

This method requires you to resolve all dependencies and install any missing packages. Always back up your data before making significant system changes.

The following example downloads Percona XtraBackup 8.0 release package for Ubuntu 20.04. Run the following commands as root or use the `sudo` command.

- 1 Use `wget` to download the `DEB` package

```
wget https://downloads.percona.com/downloads/Percona-XtraBackup-LATEST/Percona-XtraBackup-8.0.26-18/binary/debian/focal/x86_64/percona-xtrabackup-80_8.0.26-18-1.focal_amd64.deb
```



- 2 Install the package

```
sudo dpkg -i percona-xtrabackup-80_8.0.26-18-1.focal_amd64.deb
```



Common dependencies for Percona XtraBackup

| Library Name | Description |
|----------------|--|
| libmysqlclient | A client library for MySQL, providing essential functions for connecting to and communicating with MySQL databases. |
| libssl | A cryptographic library used for implementing SSL/TLS encryption, necessary for secure data transmission. |
| libcurl | A library that enables data transfer via various protocols such as HTTP, FTP, and others, commonly used for handling web requests. |
| libev | A high-performance event-loop library, often used in network applications for handling asynchronous events efficiently. |
| libgcrypt | A general-purpose cryptographic library that provides encryption, decryption, and cryptographic hashing algorithms. |
| zlib | A compression library used for data compression and decompression, supporting the popular DEFLATE compression algorithm. |

Identify and add missing dependencies

If there are missing dependencies, dpkg shows error messages. Install any missing packages with the following command:

```
sudo apt-get install -f
```



Retry the Percona XtraBackup installation.

Verify the installation

The following command returns the version information.

Troubleshoot the installation

| Issue | Description |
|----------------------------------|---|
| Permission denied | Ensure you're using sudo or have root access for installation commands. |
| Repository issues | If packages are not found, check your yum repository configuration. |
| Incompatible OS version | Double-check your OS version matches the package requirements. |
| Conflicts with existing packages | Consider removing conflicting packages or use a separate environment. |

If you encounter persistent issues, consult the Percona XtraBackup documentation or [Get help from Percona](#).

Pin packages in Debian-based systems

You may need to keep a specific package at a certain version and prevent it from being automatically updated by your package manager. This action is called “pinning” a package.

1 Create a preference file, `00percona.pref`, in the `/etc/apt/preferences.d/` directory.

2 Set the Pinning priority

To pin a package, set the `Pin-Priority` to a higher number than the default priority of the packages. This setting makes the packages higher priority over the other sources.

For example, add the following lines to the preference file:

```
Package: <package_name>
Pin: release o=Percona Development Team
Pin-Priority: 1001
```



Replace the `<package_name>` with the package's name.

3 Save the `00percona.pref` file.

You can pin multiple packages by adding separate entries to the file.

Remove the entry from the file if you no longer want to pin the package.

For more information about the pinning, check the official [debian wiki](#).

Work with AppArmor

The Linux Security Module implements mandatory access controls (MAC) with AppArmor. Debian and Ubuntu systems install AppArmor by default. AppArmor uses profiles which define which files and permissions are needed for application.

Percona XtraBackup does not have a profile and is not confined by AppArmor.

For a list of common AppArmor commands, see [Percona Server for MySQL - AppArmor](#)

Develop a profile

Download the profile from:

<https://github.com/percona/percona-xtrabackup/tree/8.0/packaging/percona/apparmor/apparmor.d>

The following profile sections should be updated with your system information, such as location of the backup destination directory.

Expected output

```
# enable storing backups only in /backups directory
# /backups/** rwk,

# enable storing backups anywhere in caller user home directory
/ @{HOME}/** rwk,

# enable storing backups only in /backups directory
# /backups/** rwk,

# enable storing backups anywhere in caller user home directory
/ @{HOME}/** rwk,
}

# enable storing backups only in /backups directory
# /backups/** rwk,

# enable storing backups anywhere in caller user home directory
/ @{HOME}/** rwk,
}
```

Move the updated file:

```
$ sudo mv usr.sbin.xtrabackup /etc/apparmor.d/
```



Install the profile with the following command:

```
$ sudo apparmor_parser -r -T -W /etc/apparmor.d/usr.sbin.xtrabackup
```



Run the backup as usual.

No additional AppArmor-related actions are required to restore a backup.

Uninstall Percona XtraBackup 8.0 on Debian-based systems

Remember to back up any important data before uninstalling software.

To completely uninstall Percona XtraBackup, remove all the installed packages:

```
$ sudo apt remove percona-xtrabackup-80
```



Run the following command to remove all configuration files along with the packages:

```
$ sudo apt purge percona-xtrabackup-80
```



Remove any unused dependencies:

```
$ sudo apt autoremove
```



To verify the removal, check the version:

```
$ xtrabackup --version
```



This should return a `command not found`.

You must remove any remaining directories manually:

```
$ sudo rm -rf /var/lib/xtrabackup
```



You can also remove any added Percona repositories from your system manager's configuration.

Use YUM

Install Percona XtraBackup on Red Hat-based systems

Ready-to-use packages are available from the Percona XtraBackup software repositories and the [download page](#). The Percona yum repository supports popular RPM-based operating systems, including the Amazon Linux AMI, which has reached its end of life. Amazon Linux 2 is not supported at this time.

The easiest way to install the Percona Yum repository is to install an RPM that configures yum and installs the [Percona GPG key](#).

Specific information on the supported platforms, products, and versions is described in [Percona Software and Platform Lifecycle](#).

Check out our [Get more help](#) for ways that we can work with you.

Version changes

Starting with [Percona XtraBackup 8.0.35-31](#), the RPM builds for RHEL 8 and RHEL 9 include ARM packages with the `aarch64.rpm` extension. This extension means Percona XtraBackup is now available for users on ARM-based systems.

We recommend downloading Percona XtraBackup for the same platform as your MySQL-compatible server. For instance, if your server runs on an ARM64 platform, you should download and use the Percona XtraBackup with the `aarch64.rpm` extension for that operating system.

Make sure that you have the `libev` package installed before installing Percona XtraBackup on CentOS 6. For this operating system, the `libev` package is available from [EPEL](#).

Prerequisites

- Sudo or root access
- Internet connection

Install from Percona YUM repository

The following procedure installs Percona XtraBackup 8.0.

- 1 Install the Percona yum repository by running the following command as the `root` user or with `sudo`:

```
$ sudo yum install \
https://repo.percona.com/yum/percona-release-latest.\
noarch.rpm
```

Example output

```
Oracle Linux 9 BaseOS Latest (x86_64)                6.6 MB/s | 35 MB
00:05
Oracle Linux 9 Application Stream Packages (x86_64)    8.5 MB/s | 41 MB
00:04
Last metadata expiration check: 0:00:06 ago on Mon Sep 23 09:43:33 2024.
percona-release-latest.noarch.rpm                    49 kB/s | 27 kB
00:00
Dependencies resolved.
...
For more information, please visit:
https://docs.percona.com/percona-software-repositories/percona-release.html

Verifying      : percona-release-1.0-29.noarch
                  1/1

Installed:
percona-release-1.0-29.noarch

Complete!
```

2 Enable the repository:

```
$ sudo percona-release enable pxb-80
```

Example output

```
```.text .no-copy}
* Enabling the Percona XtraBackup 8.0 repository
<*> All done!
...

```

## 3 Install *Percona XtraBackup* by running:

```
$ sudo yum install percona-xtrabackup-80
```



### Example output



```
* Enabling the Percona XtraBackup 8.0 repository
<*> All done!
[root@e1dc549c0c65 /]# yum install percona-xtrabackup-80
PMM2 Client release/x86_64 YUM repository 13 kB/s | 6.4 kB
00:00
Percona Release release/noarch YUM repository 4.3 kB/s | 2.0 kB
00:00
Percona XtraBackup 8.0 release/x86_64 YUM repository 363 kB/s | 416 kB
00:01
Percona Telemetry release/x86_64 YUM repository 3.1 kB/s | 1.5 kB
00:00
Dependencies resolved.
...
Complete!
```

- 4 To decompress backups made using LZ4 or ZSTD compression algorithm, install the corresponding package:

#### Install the lz4 package

```
$ sudo yum install lz4
```



#### Install the zstd package

```
$ sudo yum install zstd
```



## Verify the installation

Run Percona XtraBackup:

```
$ xtrabackup --version
```



If this command fails, the installation is incomplete or incorrect.

## Troubleshoot

To troubleshoot a Percona XtraBackup 8.0.x installation on a Red Hat-based system, follow these steps:

- 1 This command checks for missing libraries.

```
$ ldd $(which xtrabackup) | grep "not found"
```



**2** Install missing dependencies:

```
$ sudo yum install <missing-library-name>
```



**3** Examine log files:

```
$ sudo journalctl -xe | grep xtrabackup
```



**4** Ensure the Percona repository is correctly set up:

```
$ cat /etc/yum.repos.d/percona-release.repo
```



**5** Update the system to ensure all packages are current:

```
$ sudo yum update
```



**6** If issues persist, try a clean reinstall:

```
$ sudo yum remove percona-xtrabackup-80
$ sudo yum clean all
$ sudo yum install percona-xtrabackup-80
```



**7** SELinux might block XtraBackup. Check its status:

```
$ getenforce
```



**8** Set SELinux to permissive temporarily:

```
$ sudo setenforce 0
```



**9** Create a backup to verify functionality:

```
$ xtrabackup --backup --target-dir=/tmp/test_backup
```



If problems continue, gather all error messages and logs, then consult Percona's documentation or support forums for further assistance.



#### See also

To install Percona XtraBackup using downloaded rpm packages, see [Install with package manager](#).

To uninstall Percona XtraBackup, see [Uninstall Percona XtraBackup](#)

# Files in the RPM package built for Percona XtraBackup 8.0

The following tables show what data each RPM package contains:

Package	Contains
percona-xtrabackup-80	The latest Percona XtraBackup GA binaries and associated files
percona-xtrabackup-80-debuginfo	The debug symbols for binaries in percona-xtrabackup-80
percona-xtrabackup-test-80	The test suite for Percona XtraBackup
percona-xtrabackup	The older version of the Percona XtraBackup

# Install with RPM packages

Download `RPM` packages of the desired series for your architecture from the [Percona Software Downloads](#). If needed, [Percona Software Download instructions](#) are available.

This method requires you to resolve all dependencies and install any missing packages. Always back up your data before making significant system changes.

The following example downloads Percona XtraBackup 8.0 release package for *RHEL* 9. Run the following commands as root or use the `sudo` command.

- 1 Use `wget` to download the `RPM` package

```
wget https://downloads.percona.com/downloads/Percona-XtraBackup-LATEST/Percona-XtraBackup-8.0.32-26/binary/redhat/9/x86_64/percona-xtrabackup-80-8.0.32-26.1.el9.x86_64.rpm
```

- 2 Install the package

```
sudo dnf install ./percona-xtrabackup-80-8.0.32-26.1.el9.x86_64.rpm
```

## Common dependencies

Library name	Description
libev	A high-performance event-loop library used in asynchronous network applications.
libgcrypt	A general-purpose cryptographic library providing encryption, decryption, and hashing algorithms.
openssl	A robust library for implementing SSL/TLS encryption and cryptographic functions.
zlib	A compression library that supports data compression and decompression using the DEFLATE algorithm.
libaio	A library for asynchronous I/O operations, providing non-blocking I/O functionality.



If dnf reports missing dependencies, it may offer to install them automatically. If not, you can install these dependencies manually using the following example.

```
sudo dnf install libev libgcrypt openssl zlib libaio
```

## Verify the installation

After installation, verify it by checking the version:

```
xtrabackup --version
```

## Troubleshoot the installation

Issue	Description
Permission denied	Ensure you’re using sudo or have root access for installation commands.
Repository issues	If packages are not found, check your yum repository configuration.
Incompatible OS version	Double-check your OS version matches the package requirements.
Conflicts with existing packages	Consider removing conflicting packages or use a separate environment.

If you encounter persistent issues, consult the Percona XtraBackup documentation or [Get help from Percona](#).

# Work with SELinux

Percona XtraBackup is installed as an unconfined process running in an undefined domain. SELinux allows unconfined processes almost all access and the processes only use Discretionary Access Control (DAC) rules.

You find the current state of the Percona XtraBackup file with the following command:

```
$ ls -Z /usr/bin | grep xtrabackup
```



## Expected output



```
-rwxr-xr-x. root root system_u:object_r:bin_t:s0 xtrabackup
```

The SELinux context is the following:

- user (root)
- role (object\_r)
- type (bin\_t)
- level (s0)

The unconfined domain supports the network-facing services, which are protected by SELinux. These domains are not exposed. In this configuration, SELinux protects against remote intrusions but local intrusions, which require local access, are not confined.

*Percona XtraBackup* works locally. The service is not network-facing and cannot be exploited externally. The service interacts only with the local user, who provides the parameters. *Percona XtraBackup* requires access to the `target-dir` location.

## Confine XtraBackup

You can modify your security configuration to confine *Percona XtraBackup*. The first question is where to store the backup files. The service requires read and write access to the selected location.

You can use either of the following methods:

- Allow *Percona XtraBackup* to write to any location. The user provides any path to the `target-dir` parameter.

- Allow *Percona XtraBackup* to write to a specific location, such as `/backups` or the user's home directory.

The first option opens the entire system to read and write. Select the second option to harden your security.

## Install SELinux tools

To work with policies, you must install the SELinux tools. To find which package provides the `semanage` command and install the package. The following is an example on CentOS 7.

```
$ yum provides *bin/semanage
```

The result should list the packages.

### Expected output

```
...
policycoreutils-python-2.5-34.el7.x86_64 : SELinux policy core python utilities
...
```

To install missing packages, run the following:

```
$ sudo yum install -y policycoreutils-python
```

The following is an example on CentOS 8:

```
$ yum provides *bin/semanage
```

The result should list the missing packages.

### Expected output

```
...
policycoreutils-python-utils-2.8-16.1.el8.noarch : SELinux policy core python utilities
```

Run the following to install the missing packages:

```
$ sudo yum install -y policycoreutils-python-utils
```

## Create a policy

Use a modular approach to create an SELinux policy. Create a policy module to manage XtraBackup. You must create a `.te` file for type enforcement, and an optional `.fc` file for the file contexts.

Use `$ ps -efZ | grep xtrabackup` to verify the service is not confined by SELinux.

Create the `xtrabackup.fc` file and add content. This file defines the security contexts.

```
/usr/bin/xtrabackup -- gen_context(system_u:object_r:xtrabackup_exec_t,s0)
/usr/bin/xbcrypt -- gen_context(system_u:object_r:xtrabackup_exec_t,s0)
/usr/bin/xbstream -- gen_context(system_u:object_r:xtrabackup_exec_t,s0)
/usr/bin/xbcloud -- gen_context(system_u:object_r:xtrabackup_exec_t,s0)
/backups(/.*)? system_u:object_r:xtrabackup_data_t:s0
```

### Note

If you are using the `/backups` directory you must have the last line. If you are storing the backups in the user's home directory, you can omit this line.

Download the `xtrabackup.te` file from the following location:

<https://github.com/percona/percona-xtrabackup/tree/8.0/packaging/percona/selinux>

### Note

In the file, the sections in bold should be modified for your system. The `fc` file can also be downloaded from the same location.

Compile the policy module:

```
$ make -f /usr/share/selinux/devel/Makefile xtrabackup.pp
```

Install the module:

```
$ semodule -i xtrabackup.pp
```

Tag the PXB binaries with the proper SELinux tags, such as `xtrabackup_exec_t`.

```
$ restorecon -v /usr/bin/*
```

If you store your backups at `/backups`, restore the tag in that location:

```
$ restorecon -v /backups
```



 **Note**

Remember to add the standard Linux DAC permissions for this directory.

Perform the backup in the standard way.

# Uninstall Percona XtraBackup 8.0 on Red Hat-based systems

Remember to back up any important data before uninstalling software.

Remove the XtraBackup packages:

```
$ sudo yum remove percona-xtrabackup-80
```



To remove all configuration files and packages, run the following command:

```
$ sudo yum remove percona-xtrabackup-80 --remove-leaves
```



Remove any unused dependencies:

```
$ sudo yum autoremove
```



To verify the removal, check the version:

```
$ xtrabackup --version
```



This should return a `command not found`.

You must remove any remaining directories manually:

```
$ sudo rm -rf /var/lib/xtrabackup
```



You can also remove any added Percona repositories from your system manager's configuration.

# Use binary tarballs

# Install Percona XtraBackup 8.0 from a binary tarball

Binary tarballs provide a ready-to-use application since they contain precompiled executables, libraries, and dependencies. Installing from a tarball is straightforward — extract the files to your desired location. Compiling or building from source is unnecessary, which saves time and effort. You avoid the compilation process, which can be time-consuming and error-prone.

Binary tarballs provide only the necessary files, reducing disk space usage. A binary tarball installation allows customization since you can choose installation paths, configure options, and tailor the application to your requirements. Binary tarballs often come with version numbers, letting you track and update the installed version as needed.

Be aware of the following when installing from a binary tarball:

- Binary tarballs bundle dependencies within themselves. If your system already has specific libraries or versions installed, you may have conflicts.
- Binary tarball installation bypasses the package manager. Package managers handle updates, dependencies, and conflict resolution. Manual management of updates becomes necessary.
- Binary tarballs may not receive security patches automatically. You must actively monitor and apply updates.
- Binary tarball custom installation can scatter files. It may not adhere to standard paths, which makes maintenance challenging.

## Download the tarball

You can download the binary tarballs from the `Linux - Generic` section on the [Percona Product Downloads](#). For more information, see the [download instructions](#). For information on the available binary tarball names, see [binary tarball names](#).

The following command is an example of downloading a tarball for Linux/Generic:

```
$ wget https://downloads.percona.com/downloads/Percona-XtraBackup-8.0/Percona-XtraBackup-8.0.35-30/binary/tarball/percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17.tar.gz
```





### Expected result

```
--2024-03-18 11:34:50-- https://downloads.percona.com/downloads/Percona-XtraBackup-8.0/Percona-XtraBackup-8.0.35-30/binary/tarball/percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17.tar.gz
Resolving downloads.percona.com (downloads.percona.com)... 147.135.54.159,
2604:2dc0:200:69f::2
Connecting to downloads.percona.com (downloads.percona.com)|147.135.54.159|:443... connected.
...
percona-xtrabackup-8.0 100%[=====] 448.02M 2.56MB/s in 5m 5s
2024-03-18 11:39:56 (1.47 MB/s) - 'percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17.tar.gz' saved [469780148/469780148]
```

## Extract the tarball

After downloading the tarball to the selected location, use the following command to extract the tarball:

```
$ tar -xvf percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17.tar.gz
```

### Expected result

```
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin/
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin/xbcloud_osenv
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin/xtrabackup
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin/xbcrypt
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin/xbstream
percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin/xbcloud
...
```

After the extraction, navigate to the `bin` directory to find the binary files.

```
$ cd percona-xtrabackup-8.0.35-30-Linux-x86_64.glibc2.17/bin
```

In the bin directory, run one of the binary files.

```
$./xtrabackup --version
```



### Expected result



```
./xtrabackup version 8.0.35-30 based on MySQL server 8.0.35 Linux (x86_64) (revision id: 6beb4b49)
```

# Binary tarball file names available

Download the binary tarballs from [Percona Product Downloads](#). For more information, see the [download instructions]. To install Percona XtraBackup, see [install with binary tarball](#).

The following table lists the tarball types available in `Linux - Generic`. Select the Percona XtraBackup 8.0 version, the software or the operating system, and the tarball type for your installation. Binary tarballs support all AMD64 distributions.

After you have downloaded the binary tarballs, extract the tarball in the file location of your choice.

Type	Name	Operating systems	Description
Full	percona-xtrabackup-- Linux.x86_64.glibc2.12.tar.gz	Built for CentOS 6	Contains binary files, libraries, test files, and debug symbols
Minimal	percona-xtrabackup-- Linux.x86_64.glibc2.12-minimal.tar.gz	Built for CentOS 6	Contains binary files, and libraries but does not include test files, or debug symbols
Full	percona-xtrabackup-- Linux.x86_64.glibc2.17.tar.gz	Compatible with any operating system except for CentOS 6	Contains binary files, libraries, test files, and debug symbols
Minimal	percona-xtrabackup-- Linux.x86_64.glibc2.17-minimal.tar.gz	Compatible with any operating system except for CentOS 6	Contains binary files, and libraries but does not include test files, or debug symbols

Select a different software, such as Ubuntu 20.04 (Focal Fossa), for a tarball for that operating system. You can download the packages together or separately.

# Compile from source

# Compile and install Percona XtraBackup 8.0 from source

The following instructions install Percona XtraBackup 8.0.

## 1. Install Percona XtraBackup from the Git Source Tree

Percona uses the [Github](#) revision control system for development. To build the latest Percona Server for MySQL from the source tree, you will need `git` installed on your system.

You can now fetch the latest Percona XtraBackup 8.0 sources:

```
$ git clone https://github.com/percona/percona-xtrabackup.git
$ cd percona-xtrabackup
$ git checkout 8.0
$ git submodule update --init --recursive
```



## 2. Installation prerequisites

The following packages and tools must be installed to compile *Percona XtraBackup* from source. These might vary from system to system.

### Important

To build **Percona XtraBackup 8.0** from source, you must use `cmake` version 3. To check which version is currently installed, run `cmake --version` at a command prompt. If the version is not 3, install `cmake3`.

This `cmake` version may be available in your distribution as a separate package `cmake3`. For more information, see [cmake.org](https://cmake.org)

### Debian or Ubuntu using `apt`

```
sudo apt install bison pkg-config cmake devscripts debconf \
debhelper automake bison ca-certificates libprocps-dev \
libcurl4-openssl-dev cmake debhelper libaio-dev \
libncurses-dev libssl-dev libtool libz-dev libgcrypt-dev libev-dev libprocps-dev \
lsb-release build-essential rsync libdbd-mysql-perl \
libnuma1 socat librtmp-dev libtinfo5 vim-common \
libudev-dev \
liblz4-tool liblz4-1 liblz4-dev zstd python-docutils qpress
```

To install the man pages, install the python3-sphinx package first:

```
$ sudo apt install python3-sphinx
```

### CentOS or Red Hat using `yum`

Percona Xtrabackup requires GCC version 5.3 or higher. If the version of GCC installed on your system is lower then you may need to install and enable [the Developer Toolset](#) on RPM-based distributions to make sure that you use the latest GCC compiler and development tools. Then, install `cmake` and other dependencies:

```
$ sudo yum install cmake openssl-devel libaio libaio-devel automake autoconf \
bison libtool ncurses-devel libgcrypt-devel libev-devel libcurl-devel zlib-devel \
libudev-dev \
zstd vim-common procps-ng-devel
```

To install the man pages, install the python3-sphinx package first:

```
$ sudo yum install python3-sphinx
```

## 3. Generate the build pipeline

At this step, you have `cmake` run the commands in the `CMakeList.txt` file to generate the build pipeline, i.e., a native build environment that will be used to compile the source code).

1. Change to the directory where you cloned the Percona XtraBackup repository

```
$ cd percona-xtrabackup
```

2. Create a directory to store the compiled files and then change to that directory:

```
$ mkdir build
$ cd build
```

3. If you use a regular build, run cmake or cmake3. In either case, the options you need to use are the same.

Note

You can build Percona XtraBackup with man pages but this requires `python-sphinx` package which isn't available from the main repositories for every distribution. If you installed the `python-sphinx` package, remove `-DWITH_MAN_PAGES=OFF` from this command.

```
$ cmake -DWITH_BOOST=PATH-TO-BOOST-LIBRARY -DDOWNLOAD_BOOST=ON \
-DBUILD_CONFIG=xtrabackup_release -DWITH_MAN_PAGES=OFF -B ..
```

Parameter	Description
-DWITH_BOOST	For the <code>-DWITH_BOOST</code> parameter, specify the name of a directory to download the boost library to. This directory is created automatically in your current directory.
-DBUILD_CONFIG	This flag allows you to define a build configuration, such as “debug” or “release”, which can be used to control which parts of the source code are compiled.
-DWITH_MAN_PAGES	To build <b>Percona XtraBackup</b> man pages, use <code>ON</code> or remove this parameter from the command line (it is <code>ON</code> by default). To install the man pages, install the <code>python3-sphinx</code> package first.
-B (-build)	<b>Percona XtraBackup</b> is configured to forbid generating the build pipeline for make in the same directory where you store your sources. The <code>-B</code> parameter refers to the directory that contains the source code. In this example, we use the relative path to the parent directory ( <code>..</code> ).

#### Important

CMake Error at CMakeLists.txt:367 (MESSAGE): Please do not build in-source. Out-of source builds are highly recommended: you can have multiple builds for the same source, and there is an easy way to do cleanup, simply remove the build directory (note that 'make clean' or 'make distclean' does **not** work)

You **can** force in-source build by invoking cmake with `-DFORCE_INSOURCE_BUILD=1`.

## 4. Compile the source code

To compile the source code in your `build` directory, use the `make` command.

1. Change to the `build` directory (created at Step 2: Generating the build pipeline).
2. Run the `make` command. This command may take a long time to complete.

```
$ make
```



To use all CPU threads and make compilation faster please use:

```
$ make -j$(nproc --all)
```



## 5. Install on the target system

The following command installs all *Percona XtraBackup* binaries *xtrabackup* and tests to default location on the target system: `/usr/local/xtrabackup`.

Run `make install` to install *Percona XtraBackup* to the default location.

```
$ sudo make install
```



### Install to a non-default location

You may use the `DESTDIR` parameter with `make install` to install *Percona XtraBackup* to another location. Make sure that the effective user is able to write to the destination you choose.

```
$ sudo make DESTDIR=<DIR_NAME> install
```





In fact, the destination directory is determined by the installation layout ( `-DINSTALL_LAYOUT` ) that `cmake` applies (see Step 2: Generating the build pipeline). In addition to the installation directory, this parameter controls a number of other destinations that you can adjust for your system.

By default, this parameter is set to `STANDALONE`, which implies the installation directory to be `/usr/local/xtrabackup`.

 **See also**

[MySQL Documentation: -DINSTALL\\_LAYOUT](#)

## 6. Run Percona XtraBackup

After *Percona XtraBackup* is installed on your system, you may run it by using the full path to the `xtrabackup` command:

```
$ /usr/local/xtrabackup/bin/xtrabackup
```



Update your `PATH` environment variable if you would like to use the command on the command line directly.

```
$# Setting $PATH on the command line
$ PATH=$PATH:/usr/local/xtrabackup/bin/xtrabackup

$# Run xtrabackup directly
$ xtrabackup
```



Alternatively, you may consider placing a soft link (using `ln -s`) to one of the locations listed in your `PATH` environment variable.

To view the documentation with `man`, update the `MANPATH` variable.

# Docker

# Run in a Docker container

Docker allows you to run Percona XtraBackup and Percona Server for MySQL in separate containers, keeping your host system clean and organized. Each container is fully equipped with everything needed to run the application, creating an isolated environment for Percona XtraBackup without requiring direct installation. By managing your database server and backups in separate containers, you can maintain a tidy separation of services, enhancing your workflow and efficiency.

When we say “container” here, we mean the Docker container itself. The “instance” refers specifically to the database server running inside that container.

Percona maintains an official Docker image on [Docker Hub](#), making it easy to get started. You have two options for the image version:

- Use the `latest` tag to get the most recent release
- Pick a specific version through the Docker tag filter

A critical first step: ensure you have the current version of Docker installed. The package manager versions (APT, YUM) often lag behind and can cause issues. Instead, follow Docker’s official installation guide to get the most recent release.

Check out our [Get more help](#) for ways that we can work with you.

## Install the Docker Engine

You need Docker to run Percona XtraBackup in a container. The instructions are subject to change, so always refer to the [Docker documentation](#) for the most up-to-date information.

## Debian/Ubuntu

1. Remove any existing Docker version to prevent conflicts:

```
$ sudo apt-get remove docker docker-engine docker.io containerd runc
```

2. Update package lists:

```
$ sudo apt-get update
```

3. Install prerequisites to allow the use of repositories over HTTPS:

```
$ sudo apt-get install apt-transport-https ca-certificates curl gnupg lsb-release
```

4. Add Docker's official GPG key:

```
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-archive-keyring.gpg
```

5. Set up the Docker repository:

```
$ echo \
"deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

6. Install the Docker Engine:

```
$ sudo apt-get update
$ sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-plugin
```

7. Verify the installation:

```
$ sudo docker --version
```

## Red Hat Enterprise Linux (RHEL) 7 or below

1. Remove any existing Docker version to prevent conflicts:

```
$ sudo yum remove docker docker-client docker-client-latest docker-common
docker-latest docker-latest-logrotate docker-logrotate docker-engine
```

2. Install the required packages:

```
$ sudo yum install -y yum-utils
```

3. Set up the Docker repository:

```
$ sudo yum-config-manager --add-repo
https://download.docker.com/linux/rhel/docker-ce.repo
```

4. Install the Docker Engine:

```
$ sudo yum install docker-ce docker-ce-cli containerd.io docker-buildx-plugin
docker-compose-plugin
```

5. Start the Docker service:

```
$ sudo systemctl start docker
```

6. Verify the installation:

```
$ sudo docker --version
```

## Red Hat Enterprise Linux (RHEL) 8+

1. Remove any existing Docker version to prevent conflicts:

```
$ sudo dnf remove docker docker-client docker-client-latest docker-common
docker-latest docker-latest-logrotate docker-logrotate docker-engine
```

2. Install the required packages:

```
$ sudo dnf install dnf-plugins-core
```

3. Set up the Docker repository:

```
$ sudo dnf config-manager --add-repo
https://download.docker.com/linux/rhel/docker-ce.repo
```



#### 4. Install the Docker Engine:

```
$ sudo dnf install docker-ce docker-ce-cli containerd.io docker-buildx-plugin
docker-compose-plugin
```



#### 5. Start the Docker service:

```
$ sudo systemctl enable --now docker
```



#### 6. Verify the installation:

```
$ sudo docker --version
```



## Docker images

Docker images are pre-packaged software that contains everything needed to run an application. They include the application code, runtime, libraries, environment variables, and configuration files. You can download images from Docker Hub, a public repository of Docker images.

1. Go to Docker Hub at <https://hub.docker.com>
2. Type “percona xtrabackup” in the search bar
3. Look for the official Percona repository ( `percona/percona-xtrabackup` )

The repository name shows as `percona/percona-xtrabackup` where:

Name	Description
percona	Organization name
percona-xtrabackup	Specific image name

## Image tags

Each Docker image has tags that tell you its version. For Percona XtraBackup, you’ll see tags like:

```
percona/percona-xtrabackup:8.0
percona/percona-xtrabackup:8.0.35
percona/percona-xtrabackup:latest
```



- `8.0` tracks the latest 8.0 release
- `8.0.35` points to that specific version
- `latest` always gets you the newest version

## Pull images

To download an image, use `docker pull`. Here's how:

```
Pull the latest version
docker pull percona/percona-xtrabackup:latest

Pull a specific version
docker pull percona/percona-xtrabackup:8.0.35
```



## What's in a Percona XtraBackup image

A Percona XtraBackup image contains: - The XtraBackup backup tool - Required system libraries - Configuration files - Runtime dependencies

Think of it as a pre-packed toolkit with everything XtraBackup needs to run.

## Check images

See what images you've downloaded:

```
List all images
docker images

Look for specific Percona images
docker images | grep percona
```



## Use an image

Run a basic test to make sure your image works:

```
docker run --rm percona/percona-xtrabackup:latest xtrabackup --version
```



This command shows the XtraBackup version and confirms your setup works.

## 1. Run Percona Server for MySQL container

Percona XtraBackup works in combination with a MySQL-compatible database, Percona Server for MySQL, in our example. To take a backup of Percona Server for MySQL, run Percona Server for MySQL in a Docker container and create a database with data. Read the instructions on how to [Run Percona Server for MySQL in a Docker Container](#).

As soon as Percona Server for MySQL runs and the database contains data, you are ready to connect the Percona XtraBackup container to the Percona Server for MySQL container and take a backup.

## 2. Create a container and connect to Percona Server for MySQL container

You have two options to create a Docker container from a Percona XtraBackup image:

- `docker create`: Makes a container ready to start later
- `docker run`: Creates and starts the container immediately

When you first create a container, Docker downloads the Percona XtraBackup 8.0 image from Docker Hub. For future containers, Docker uses your local copy.

### Important security note

When connecting your XtraBackup container to a MySQL server container, always use the `--user=root` option in your Docker command. This ensures proper backup permissions.

### Container creation

Here's how to create a basic backup container:



```
Create a container for backing up MySQL data
docker create \
 --name percona-xtrabackup \
 --volumes-from percona-server-mysql:8.0 \
 percona/percona-xtrabackup:8.0 \
 xtrabackup --backup \
 --datadir=/var/lib/mysql/ \
 --target-dir=/backup \
 --user=root \
 --password=root
```



This command does the following:

- `--name percona-xtrabackup`: Names your container for easy reference
- `--volumes-from percona-server-mysql:8.0`: Connects to your MySQL container's data
- `percona/percona-xtrabackup:8.0`: Uses the Percona XtraBackup 8.0 image
- The remaining options configure the backup:
- `--backup`: Tells XtraBackup to perform a backup
- `--datadir`: Points to MySQL's data directory
- `--target-dir`: Sets where the backup will be stored
- `--user` and `--password`: MySQL connection credentials

## Next steps

After creating your container, you can:

- Start the container with `docker start percona-xtrabackup`
- Check its status with `docker ps -a`
- View logs with `docker logs percona-xtrabackup`

## Common issues

If you see permission errors:

- Double-check your MySQL user credentials
- Verify the `--volumes-from` connection
- Ensure your target directory is writable

## Remove the container

Need to make changes? You can remove the container and create a new one:

```
docker rm percona-xtrabackup
Then run the create command again with your changes
```



## Percona XtraBackup ARM64

Starting with Percona XtraBackup 8.0.35-31, you can run the Docker ARM64 version of Percona XtraBackup. You can find the version and architecture on [Docker Hub Percona/Percona-Server Tags](#). The Docker images for ARM64 version contain `aarch64` in the tag name, for example, `8.0.35-aarch64`.

We recommend that the architectures match. If the server is ARM64, you should also use Percona XtraBackup in ARM64.

## Start the Percona XtraBackup container

Start the Percona XtraBackup container with exactly the same parameters that were used when the container was created. The following command starts the `percona-xtrabackup` Docker container, attaches to its output, and keeps STDIN open for interactive input. This is useful if you want to monitor the container's logs and interact with it directly.

```
$ sudo docker start -ai percona-xtrabackup
```



The command `sudo docker start -ai percona-xtrabackup` does the following:

- `sudo`: Runs the command with superuser (administrator) privileges.
- `docker`: Refers to the Docker command-line interface.
- `start`: Starts a stopped Docker container.
- `-a`: Attaches to the container's output, so you can see the logs and interact with the container in real time.
- `-i`: Keeps the container's STDIN open, allowing you to provide input to the container.
- `percona-xtrabackup`: Specifies the name of the Docker container you want to start.

## Take a backup

This command creates and starts a new Docker container named `percona-xtrabackup`, mounts the volumes from the existing `percona-server-mysql:8.0` container, and performs a backup of the MySQL data using `xtrabackup`. The backup is stored in the `/backup` directory inside the new container..

```
$ sudo docker run --name percona-xtrabackup --volumes-from percona-server-mysql:8.0 \
percona/percona-xtrabackup:8.0
xtrabackup --backup --datadir=/var/lib/mysql --target-dir=/backup --user=root --
password=root
```

- `sudo`: Runs the command with superuser (administrator) privileges.
- `docker run`: Creates and starts a new Docker container.
- `--name percona-xtrabackup`: Names the new container `percona-xtrabackup`.
- `--volumes-from percona-server-mysql:8.0`: Mounts volumes from the existing container `percona-server-mysql:8.0` into the new container. This allows the new container to access the data and directories of the existing container.
- `percona/percona-xtrabackup:8.0`: Specifies the Docker image to use for creating the container. In this case, it's the `percona/percona-xtrabackup:8.0` image.
- `xtrabackup --backup --datadir=/var/lib/mysql --target-dir=/backup --user=root --password=root`: Executes the `xtrabackup` command within the container to perform a backup. The options provided to the `xtrabackup` command are:
  - `--backup`: Tells `xtrabackup` to perform a backup.
  - `--datadir=/var/lib/mysql`: Specifies the data directory of the MySQL server.
  - `--target-dir=/backup`: Specifies the directory where the backup will be stored.
  - `--user=root`: Uses the root user to perform the backup.
  - `--password=root`: Uses the root password for authentication.

## Restore a backup

To restore a backup in Docker, you need to create a new container from the Percona XtraBackup image and use the `xtrabackup` command to restore the backup data to the MySQL data directory. The following steps outline the process of restoring a backup in Docker:

1. Create a Docker volume for the restored data.

```
docker volume create restorevol
```

2. Run a new Percona Server for MySQL container to restore the backup data:

```
docker run -d -p 3307:3306 --name psmysql2 -e MYSQL_ROOT_PASSWORD=secret -v restorevol:/var/lib/mysql percona/percona-server:8.0.34
```

3. Stop the new Percona Server for MySQL container:

```
docker stop psmysql2
```

4. Remove all files in the new volume to prepare for restoration:

```
docker run --rm --volumes-from psmysql2 -it --user root percona/percona-xtrabackup:8.0.34 /bin/bash -c "rm -rf /var/lib/mysql/*"
```

5. Restore the backup data to the new volume:

```
docker run --rm --volumes-from psmysql2 -it --user root -v restorevol:/backup percona/percona-xtrabackup:8.0.34 xtrabackup --copy-back --target-dir=/backup
```

6. Change the ownership of the restored data:

```
docker run --rm --volumes-from psmysql2 -it --user root percona/percona-xtrabackup:8.0.34 chown -R mysql:mysql /var/lib/mysql
```

7. Start the new Percona Server for MySQL container:

```
docker start psmysql2
```

## Troubleshoot issues with creating the container or during the backup

If you encounter issues when creating the container or during the backup process, consider the following troubleshooting steps:

- Check Docker daemon: Ensure the Docker daemon runs and is properly configured.
- Verify Docker image
  - Confirm you use the correct Docker image.
  - Check available images with the `docker images` command.
- Inspect container logs

- View logs using the `docker logs <container_name>` command.
- Identify any errors or warnings.
- Check volume permissions
  - Ensure the volumes have the correct permissions.
  - Verify that the container can access the volumes.
- Review configuration files: Double-check the configuration files for syntax errors or misconfigurations.
- Resource allocation
  - Ensure the host system has sufficient resources (CPU, memory, disk space).
  - Monitor resource usage during the backup process.
- Network issues
  - Verify there are no network connectivity issues between the Docker containers.
  - Check firewall settings and network configurations.
- Compatibility: Ensure the versions of Percona Server and Percona XtraBackup are compatible.
- Error messages
  - Look for specific error messages in the logs.
  - Search for solutions online or in the Percona documentation.

## For more information

Review the [Docker Docs](#)

# How to

**Take a full backup**

# Create a full backup

To create a backup, run **xtrabackup** with the `--backup` option. You also need to specify the `--target-dir` option, which is where the backup will be stored, if the *InnoDB* data or log files are not stored in the same directory, you might need to specify the location of those, too. If the target directory does not exist, *xtrabackup* creates it. If the directory does exist and is empty, *xtrabackup* will succeed.

*xtrabackup* will not overwrite existing files, it will fail with operating system error 17, `file exists`.

To start the backup process run:

```
$ xtrabackup --backup --target-dir=/data/backups/
```

This will store the backup at `/data/backups/`. If you specify a relative path, the target directory will be relative to the current directory.

During the backup process, you should see a lot of output showing the data files being copied, as well as the log file thread repeatedly scanning the log files and copying from it. Here is an example that shows the log thread scanning the log in the background, and a file copying thread working on the `ibdata1` file:

## Expected output

```
160906 10:19:17 Finished backing up non-InnoDB tables and files
160906 10:19:17 Executing FLUSH NO_WRITE_TO_BINLOG ENGINE LOGS...
xtrabackup: The latest check point (for incremental): '62988944'
xtrabackup: Stopping log copying thread.
.160906 10:19:18 >> log scanned up to (137343534)
160906 10:19:18 Executing UNLOCK TABLES
160906 10:19:18 All tables unlocked
160906 10:19:18 Backup created in directory '/data/backups/'
160906 10:19:18 [00] Writing backup-my.cnf
160906 10:19:18 [00] ...done
160906 10:19:18 [00] Writing xtrabackup_info
160906 10:19:18 [00] ...done
xtrabackup: Transaction log of lsn (26970807) to (137343534) was copied.
160906 10:19:18 completed OK!
```

The last thing you should see is something like the following, where the value of the `<LSN>` will be a number that depends on your system:

```
$ xtrabackup: Transaction log of lsn (<LSN>) to (<LSN>) was copied.
```



#### Note

Log copying thread checks the transactional log every second to see if there were any new log records written that need to be copied, but there is a chance that the log copying thread might not be able to keep up with the amount of writes that go to the transactional logs, and will hit an error when the log records are overwritten before they could be read.

After the backup is finished, the target directory will contain files such as the following, assuming you have a single InnoDB table `test.tb11` and you are using MySQL's `innodb_file_per_table` option:

```
$ ls -lh /data/backups/
```



The result should look like this:

#### Expected output

```
total 182M
drwx----- 7 root root 4.0K Sep 6 10:19 .
drwxrwxrwt 11 root root 4.0K Sep 6 11:05 ..
-rw-r----- 1 root root 387 Sep 6 10:19 backup-my.cnf
-rw-r----- 1 root root 76M Sep 6 10:19 ibdata1
drwx----- 2 root root 4.0K Sep 6 10:19 mysql
drwx----- 2 root root 4.0K Sep 6 10:19 performance_schema
drwx----- 2 root root 4.0K Sep 6 10:19 sbtest
drwx----- 2 root root 4.0K Sep 6 10:19 test
drwx----- 2 root root 4.0K Sep 6 10:19 world2
-rw-r----- 1 root root 116 Sep 6 10:19 xtrabackup_checkpoints
-rw-r----- 1 root root 433 Sep 6 10:19 xtrabackup_info
-rw-r----- 1 root root 106M Sep 6 10:19 xtrabackup_logfile
```

The backup can take a long time, depending on how large the database is. It is safe to cancel at any time, because *xtrabackup* does not modify the database.

The next step is to [prepare](#) the backup in order to restore it.

# Prepare a full backup

After making a backup with the `--backup` option, you need to prepare it in order to [restore](#) it. Data files are not point-in-time consistent until they are *prepared*, because they were copied at different times as the program ran, and they might have been changed while this was happening.

If you try to start InnoDB with these data files, it will detect corruption and stop working to avoid running on damaged data. The `--prepare` step makes the files perfectly consistent at a single instant in time, so you can run *InnoDB* on them.

You can run the *prepare* operation on any machine; it does not need to be on the originating server or the server to which you intend to restore. You can copy the backup to a utility server and prepare it there.

Note that *Percona XtraBackup* 8.0 can only prepare backups of *MySQL 8.0*, *Percona Server for MySQL 8.0*, and *Percona XtraDB Cluster 8.0* databases. Releases prior to 8.0 are not supported.

During the *prepare* operation, *xtrabackup* boots up a kind of modified embedded InnoDB (the libraries *xtrabackup* was linked against). The modifications are necessary to disable InnoDB standard safety checks, such as complaining about the log file not being the right size. This warning is not appropriate for working with backups. These modifications are only for the *xtrabackup* binary; you do not need a modified *InnoDB* to use *xtrabackup* for your backups.

The *prepare* step uses this “embedded InnoDB” to perform crash recovery on the copied data files, using the copied log file. The `prepare` step is very simple to use: you simply run *xtrabackup* with the `--prepare` option and tell it which directory to prepare, for example, to prepare the previously taken backup run:

```
$ xtrabackup --prepare --target-dir=/data/backups/
```

When this finishes, you should see an `InnoDB shutdown` with a message such as the following, where again the value of LSN will depend on your system:

## Expected output

```
InnoDB: Shutdown completed; log sequence number 137345046
160906 11:21:01 completed OK!
```

All following prepares will not change the already prepared data files, you'll see that output says:



#### Expected output



```
xtrabackup: This target seems to be already prepared.
xtrabackup: notice: xtrabackup_logfile was already used to '--prepare'.
```

It is not recommended to interrupt xtrabackup process while preparing backup because it may cause data files corruption and backup will become unusable. Backup validity is not guaranteed if prepare process was interrupted.



#### Note

If you intend the backup to be the basis for further incremental backups, you should use the `--apply-log-only` option when preparing the backup, or you will not be able to apply incremental backups to it. See the documentation on preparing incremental backups for more details.

After preparing the backup, you can [restore](#) the backup.

# **Take an incremental backup**

# Create an incremental backup

*xtrabackup* supports incremental backups, which means that they can copy only all the data that has changed since the last backup.

## Note

Incremental backups on the MyRocks storage engine do not determine if an earlier full backup or incremental backup contains the same files. **Percona XtraBackup** copies all the MyRocks files each time it takes a backup.

You can perform many incremental backups between each full backup, so you can set up a backup process such as a full backup once a week and an incremental backup every day, or full backups every day and incremental backups every hour.

Incremental backups work because each *InnoDB* page contains a log sequence number, or LSN. The LSN is the system version number for the entire database. Each page's LSN shows how recently it was changed.

An incremental backup copies each page which LSN is newer than the previous incremental or full backup's LSN. An algorithm finds the pages that match the criteria. The algorithm reads the data pages and checks the page LSN.

## Use the changed page tracking algorithm

*Percona XtraBackup* 8.0.30 removes the algorithm that used the [changed page tracking](#) feature in *Percona Server for MySQL*. *Percona Server for MySQL* 8.0.27 deprecated the changed page tracking feature.

With PXB 8.0.27 or earlier, another algorithm enabled the *Percona Server for MySQL* changed page tracking feature. The algorithm generates a bitmap file. The *xtrabackup* binary uses that bitmap file to read only those pages needed for the incremental backup. This method potentially saves resources. The backup enables the algorithm by default if the *xtrabackup* binary discovers the bitmap file. You can override the algorithm with `--incremental-force-scan` which forces a read of all pages even if the bitmap file is available.

## Create an incremental backup

To make an incremental backup, begin with a full backup as usual. The *xtrabackup* binary writes a file called `xtrabackup_checkpoints` into the backup's target directory. This file contains a line showing the `to_lsn`, which is the database's LSN at the end of the backup. Create the full backup with a following command:

```
$ xtrabackup --backup --target-dir=/data/backups/base
```



If you look at the `xtrabackup_checkpoints` file, you should see similar content depending on your LSN number:

#### Expected output



```
backup_type = full-backupped
from_lsn = 0
to_lsn = 1626007
last_lsn = 1626007
compact = 0
recover_binlog_info = 1
```

Now that you have a full backup, you can make an incremental backup based on it. Use the following command:

```
$ xtrabackup --backup --target-dir=/data/backups/inc1 \
--incremental-basedir=/data/backups/base
```



The `/data/backups/inc1/` directory should now contain delta files, such as `ibdata1.delta` and `test/table1.ibd.delta`. These represent the changes since the LSN 1626007. If you examine the `xtrabackup_checkpoints` file in this directory, you should see similar content to the following:

#### Expected output



```
backup_type = incremental
from_lsn = 1626007
to_lsn = 4124244
last_lsn = 4124244
compact = 0
recover_binlog_info = 1
```

`from_lsn` is the starting LSN of the backup and for incremental it has to be the same as `to_lsn` (if it is the last checkpoint) of the previous/base backup.

It's now possible to use this directory as the base for yet another incremental backup:

```
$ xtrabackup --backup --target-dir=/data/backups/inc2 \
--incremental-basedir=/data/backups/inc1
```



This folder also contains the `xtrabackup_checkpoints`:



### Expected output



```
backup_type = incremental
from_lsn = 4124244
to_lsn = 6938371
last_lsn = 7110572
compact = 0
recover_binlog_info = 1
```



### Note

In this case you can see that there is a difference between the `to_lsn` (last checkpoint LSN) and `last_lsn` (last copied LSN), this means that there was some traffic on the server during the backup process.

The next step is to [prepare](#) the backup in order to restore it.

# Prepare an incremental backup

The `--prepare` step for incremental backups differs from full backups. For full backups, committed transactions are replayed from the log file to the data files, and uncommitted transactions are rolled back to ensure consistency. When preparing an incremental backup, you must skip the rollback of uncommitted transactions, as these may still be in progress and could be committed in a subsequent incremental backup. Use the `--apply-log-only` option when preparing the first full backup to prevent the rollback phase from occurring.

## Warning

If you do not use the `--apply-log-only` option to prevent the rollback phase, then your incremental backups are unusable. After transactions have been rolled back, further incremental backups cannot be applied.

Start by preparing the full backup, then apply the incremental differences to that backup.

For example, you could have the following backups:

```
/data/backups/base
/data/backups/inc1
/data/backups/inc2
```



To prepare the base backup, you need to run `--prepare` as usual, but prevent the rollback phase:

```
$ xtrabackup --prepare --apply-log-only --target-dir=/data/backups/base
```



The log sequence number should match the `to_lsn` of the base backup. The output should end with text similar to the following:

## Expected output



```
InnoDB: Shutdown completed; log sequence number 1626007
161011 12:41:04 completed OK!
```



### Warning

This backup is actually safe to restore as-is now, even though the rollback phase has been skipped. If you restore the backup and start the server, InnoDB detects that the rollback phase was not performed, and completes the rollback in the background, as InnoDB usually does for a crash recovery. InnoDB notifies you that the database was not shut down normally.

To apply the first incremental backup to the full backup, run the following command:

```
$ xtrabackup --prepare --apply-log-only --target-dir=/data/backups/base \
--incremental-dir=/data/backups/inc1
```

This command applies the delta files to the files in `/data/backups/base`, which rolls them forward in time to the time of the incremental backup. The redo log is then applied as usual. The final data is in `/data/backups/base`, not in the incremental directory.

You should see an output similar to:

### Expected output

```
incremental backup from 1626007 is enabled.
xtrabackup: cd to /data/backups/base
xtrabackup: This target seems to be already prepared with --apply-log-only.
xtrabackup: xtrabackup_logfile detected: size=2097152, start_lsn=(4124244)
...
xtrabackup: page size for /tmp/backups/inc1/ibdata1.delta is 16384 bytes
Applying /tmp/backups/inc1/ibdata1.delta to ./ibdata1...
...
161011 12:45:56 completed OK!
```

Again, the LSN should match what you saw from your earlier inspection of the first incremental backup. If you restore the files from `/data/backups/base`, you should see the state of the database as of the first incremental backup.

## Faster prepare step with `--parallel`

For incremental backups with many InnoDB Data (IBD) files, you can significantly reduce prepare time by using the `--parallel` option. The `--parallel` option enables the concurrent processing of multiple delta files, thereby maximizing storage bandwidth. The `--parallel` option is especially beneficial when there are many IBD files, even if the IBD files didn't change between backups, as empty delta files are processed quickly in parallel.

### Version history

Before Percona XtraBackup 8.0.35-33 and 8.4.0-3, the `--parallel` option didn't have any effect on the prepare phase.

Starting with Percona XtraBackup 8.0.35-33 and 8.4.0-3, using `--parallel=X` has effect on the prepare phase. It will now use X threads to apply the changes from `.delta` files to the IBD files. When using `--parallel` in the prepare phase, always specify a numeric value. The recommended minimum value is 4 (for example, `--parallel=4`).

Note that each thread operates on a single file. If you have a large delta file, there is still only one thread that processes that `.delta` file. Parallelization occurs at the file level, not within individual files.

An example command with the `--parallel` option:

```
$ xtrabackup --prepare --parallel=4 --apply-log-only --target-dir=/data/backups/base \
--incremental-dir=/data/backups/inc1
```



## Prepare a second incremental backup

Percona XtraBackup does not support using the same incremental backup directory to prepare two copies of backup. Do not run `--prepare` with the same incremental backup directory (the value of `--incremental-dir`) more than once.

Preparing the second incremental backup is a similar process: apply the deltas to the (modified) base backup, and you will roll the base backup's data forward in time to the point of the second incremental backup:

```
$ xtrabackup --prepare --target-dir=/data/backups/base \
--incremental-dir=/data/backups/inc2
```



You can also use the `--parallel` option here to speed up the process:

```
$ xtrabackup --prepare --parallel=4 --target-dir=/data/backups/base \
--incremental-dir=/data/backups/inc2
```



### Note

Use `--apply-log-only` when merging the incremental backups except for the last one. This is why the previous command does not include the `--apply-log-only` option. If `--apply-log-only` is used on the last step, backup remains consistent but the server performs the rollback phase.

Once prepared incremental backups are the same as the full backups, and they can be [restored](#) in the same way.

# Take an incremental backup using page tracking

Percona XtraBackup 8.0.27 adds support for page tracking functionality for an [incremental backup](#).

To create an incremental backup with page tracking, Percona XtraBackup uses the MySQL `mysqlbackup` component. This component provides a list of pages modified since the last backup, and Percona XtraBackup copies only those pages. This operation removes the need to scan the pages in the database. If the majority of pages have not been modified, the page tracking feature can improve the speed of incremental backups.

## Install the component

To start using the page tracking functionality, do the following:

1. Install the `mysqlbackup` component and enable it on the server:

```
$ INSTALL COMPONENT "file://component_mysqlbackup";
```



2. Check whether the `mysqlbackup` component is installed successfully:

```
$ SELECT COUNT(1) FROM mysql.component WHERE
component_urn='file://component_mysqlbackup';
```



## Use page tracking

You can enable the page tracking functionality for the full and incremental backups with the `--page-tracking` option.

The option has the following benefits:

- Resets page tracking to the start of the backup. This reset allows the next incremental backup to use page tracking.
- Allows the use of page tracking for an incremental backup if the page tracking data is available from the backup's start checkpoint LSN.

*Percona XtraBackup* processes a list of all the tracked pages in memory. If *Percona XtraBackup* does not have enough available memory to process this list, the process throws an error and exits. For example, if an incremental backup uses 200GB, *Percona XtraBackup* can use an additional 100MB of memory to process and store the page tracking data.

The examples of creating full and incremental backups using the `--page-tracking` option:

### Full backup

```
$ xtrabackup --backup --target-dir=$FULL_BACK --page-tracking
```



### Incremental backup

```
$ xtrabackup --backup --target-dir=$INC_BACKUP
--incremental-basedir=$FULL_BACKUP --page-tracking
```



After enabling the functionality, the next incremental backup finds changed pages using page tracking.

## Limitations

- When creating the first full backup using page tracking, *Percona XtraBackup* may have a delay. The following is an example of the message:

```
xtrabackup: pagetracking: Sleeping for 1 second, waiting for checkpoint lsn
17852922 /
to reach to page tracking start lsn 21353759
```

Enable page tracking before creating the first backup to avoid this delay. This method ensures that the page tracking log sequence number (LSN) is higher than the checkpoint LSN of the server.

- If an Oracle MySQL Server uses page tracking and runs `LOCK=EXCLUSIVE` and `ALGORITHM=INPLACE` DDL operations, *Percona XtraBackup* may take an inconsistent backup. Oracle MySQL Server does not write these DDL operations (Oracle MySQL [bug 106163](#)) to the redo log, and *Percona XtraBackup* only detects the inconsistency during the `--prepare` phase.

This bug does not influence *Percona XtraBackup* taking backup of *Percona Server* for MySQL because the bug is fixed in *Percona Server* for MySQL 8.0.27.

- The *xtrabackup* page tracking requires the system tablespace to be a single file (ibdata1 only). The page tracking feature does not work if the system tablespace is a multi-file tablespace.

## Start page tracking manually

After the mysqlbackup component is loaded and active on the server, you can start page tracking manually with the following option:

```
$ SELECT mysqlbackup_page_track_set(true);
```



## Check the LSN value

Check the LSN value starting from which changed pages are tracked with the following option:

```
$ SELECT mysqlbackup_page_track_get_start_lsn();
```



## Stop page tracking

To stop page tracking, use the following command:

```
$ SELECT mysqlbackup_page_track_set(false);
```



## Purge page tracking data

When you start page tracking, it creates a file under the server's datadir to collect data about changed pages. This file grows until you stop the page tracking. If you stop the server and then restart it, page tracking creates a new file but also keeps the old one. The old file continues to grow until you stop the page tracking explicitly.

If you purge the page tracking data, you should create a full backup afterward. To purge the page tracking data, do the following steps:

```
$ SELECT mysqlbackup_page_track_set(false);
$ SELECT mysqlbackup_page_track_purge_up_to(9223372036854775807);
/* Specify the LSN up to which you want to purge page tracking data. /
9223372036854775807 is the highest possible LSN which purges all page tracking
files.*/
$ SELECT mysqlbackup_page_track_set(true);
```



## Known issue

If the index is built in place using an exclusive algorithm and then is added to a table after the last LSN checkpoint, you may generate a bad incremental backup using page tracking. For more details see [PS-8032](#).

## Uninstall the mysqlbackup component

To uninstall the mysqlbackup component, use the following statement:

```
$ UNINSTALL COMPONENT "file://component_mysqlbackup"
```



# **Take a compressed backup**



# Create a compressed backup

Percona XtraBackup supports compressed backups. A local or streaming backup can be compressed or decompressed with *xbstream*.

To make a compressed backup, use the `--compress` option along with the `--backup` and `--target-dir` options.

## Version updates

## Percona XtraBackup 8.0.34-29 and later



Percona XtraBackup removes `qpress/QuickLZ` and moves the `ZSTD` compression method to [General Availability](#). With this version `ZSTD` becomes the default compression method for the `--compress` option.

The following compression algorithms are supported:

### Zstandard (ZSTD)

`ZSTD` is a fast lossless compression algorithm that targets real-time compression scenarios and better compression ratios. `ZSTD` is the default compression method for the `--compress` option.

To compress files using the `ZSTD` compression algorithm, run the following command:

```
$ xtrabackup --backup --compress --target-dir=/data/backup
```



You can specify `ZSTD` compression level with the `--compress-zstd-level(=#)` option. The default value is `1`.

```
$ xtrabackup --backup --compress --compress-zstd-level=1 --target-dir=/data/backup
```



If you want to speed up the compression you can use the parallel compression, which can be enabled with `--compress-threads` option. Following example will use four threads for compression:

```
$ xtrabackup --backup --compress --compress-threads=4 \
--target-dir=/data/compressed/
```



The expected output:

```
...
170223 13:00:38 [01] Compressing ./test/sbtest1.frm to /tmp/compressed/test/sbtest1.frm.qp
170223 13:00:38 [01] ...done
170223 13:00:38 [01] Compressing ./test/sbtest2.frm to /tmp/compressed/test/sbtest2.frm.qp
170223 13:00:38 [01] ...done
...
170223 13:00:39 [00] Compressing xtrabackup_info
170223 13:00:39 [00] ...done
xtrabackup: Transaction log of lsn (9291934) to (9291934) was copied.
170223 13:00:39 completed OK!
```

### lz4

To compress files using the `lz4` compression algorithm, set `--compress` option to `lz4`:

```
$ xtrabackup --backup --compress=lz4 --target-dir=/data/backup
```



To decompress files, use the `--decompress` option.



#### Percona XtraBackup 8.0.31-24 to 8.0.34-29



Using qpress/QuickLZ to compress backups is deprecated and may be removed in future versions. We recommend using either LZ4 or Zstandard (ZSTD) compression algorithms. ZSTD compression algorithm is in [tech preview](#).

## Before Percona XtraBackup 8.0.31-24

By default, the `--compress` option uses the `qpress` tool that you can install with the `percona-release` package configuration tool as follows:

```
$ sudo percona-release enable tools
$ sudo apt update
$ sudo apt install qpress
```

### Note

Enable the repository: `percona-release enable-only tools release`.

If *Percona XtraBackup* is intended to be used in combination with the upstream MySQL Server, you only need to enable the `tools` repository: `percona-release enable-only tools`.

The following compression algorithms are supported:

`quicklz`

To compress files using the `quicklz` compression algorithm, use `--compress` option:

```
$ xtrabackup --backup --compress --target-dir=/data/backup
```

`lz4`

To compress files using the `lz4` compression algorithm, set `--compress` option to `lz4`:

```
$ xtrabackup --backup --compress=lz4 --target-dir=/data/backup
```

`Zstandard (ZSTD)`

The Zstandard (ZSTD) compression algorithm is a [tech preview](#) feature. Before using ZSTD in production, we recommend that you test restoring production from physical backups in your environment, and also use the alternative backup method for redundancy.

[Percona XtraBackup 8.0.30-23](#) adds support for the `Zstandard (ZSTD)` compression algorithm. `ZSTD` is a fast lossless compression algorithm that targets real-time compression scenarios and better compression ratios.

To compress files using the `ZSTD` compression algorithm, set `--compress` option to `zstd`:

```
$ xtrabackup --backup --compress=zstd --target-dir=/data/backup
```

You can specify `ZSTD` compression level with the `--compress-zstd-level(=#)` option. The default value is `1`.

```
$ xtrabackup --backup --compress --compress-zstd-level=1 --target-dir=/data/backup
```

If you want to speed up the compression you can use the parallel compression, which can be enabled with `--compress-threads` option. Following example will use four threads for compression:

```
$ xtrabackup --backup --compress --compress-threads=4 \
--target-dir=/data/compressed/
```



The expected output:

```
...
170223 13:00:38 [01] Compressing ./test/sbtest1.frm to /tmp/compressed/test/sbtest1.frm.qp
170223 13:00:38 [01] ...done
170223 13:00:38 [01] Compressing ./test/sbtest2.frm to /tmp/compressed/test/sbtest2.frm.qp
170223 13:00:38 [01] ...done
...
170223 13:00:39 [00] Compressing xtrabackup_info
170223 13:00:39 [00] ...done
xtrabackup: Transaction log of lsn (9291934) to (9291934) was copied.
170223 13:00:39 completed OK!
```

Using `--encrypt` might create larger backups than expected when used with InnoDB Page Compression.

To avoid this issue with compressed backups, use the `--compress` option with the `--xstream` option in Percona XtraBackup 8.0.31-24 and later.

The next step is to [prepare](#) the backup in order to [restore](#) it.

# Decompress and prepare a backup

Before you can prepare the backup you need to decompress all the files.

## Decompress a backup

- To decompress backups made using `LZ4` or `ZSTD` compression algorithm, install the corresponding package for your package management system.

Install on APT systems

### Install the `lz4` package

```
$ sudo apt install lz4
```



### Install the `zstd` package

```
$ sudo apt install zstd
```



Install on YUM systems

### Install the `lz4` package

```
$ sudo yum install lz4
```



### Install the `zstd` package

```
$ sudo yum install zstd
```



- Use the `--decompress` option to decompress the backup.

```
$ xtrabackup --decompress --target-dir=/data/compressed/
```



#### Note

`--parallel` can be used with `--decompress` option to decompress multiple files simultaneously.

Percona XtraBackup does not automatically delete compressed files. To clean up the backup directory, use the `--remove-original` option. If you do not remove the compressed files, these files will not be copied or moved to the datadir when using the `--copy-back` or `--move-back` options.

## Prepare the backup

When the files are decompressed you can prepare the backup with the `--prepare` option.

```
$ xtrabackup --prepare --target-dir=/data/compressed/
```



### Confirmation message



```
InnoDB: Starting shutdown...
InnoDB: Shutdown completed; log sequence number 9293846
170223 13:39:31 completed OK!
```

Now the files in `/data/compressed/` are ready to be used by the server.

The next step is to [Restore a backup](#).

# Take an encrypted backup



# Encrypt backups

Percona XtraBackup supports encrypting and decrypting local and streaming backups with the upstream option, adding another protection layer. The encryption is implemented using the `libgcrypt` library from GnuPG.

## Create encrypted backups

The following options create encrypted backups. The `--encrypt-key` and `--encrypt-key-file` options specify the encryption key and are mutually exclusive. You should select one or the other.

- `--encrypt`
- `--encrypt-key`
- `--encrypt-key-file`

For an encryption key, use a command, such as `openssl rand -base64 24`, to generate a random alphanumeric string.

### The `--encrypt-key` option

An example of the `xtrabackup` command using the `--encrypt-key`:

```
$ xtrabackup --backup --encrypt=AES256 --encrypt-key="{randomly-generated-alphanumeric-string}" --target-dir=/data/backup
```



### The `--encrypt-key-file` option

The recommended method uses the command line: `echo -n "{randomly-generated-alphanumeric-string}" > /data/backups/keyfile` to create the file.

Remember that your text editor can automatically insert a CRLF (end of line) character in the `KEYFILE` when using the `--encrypt-key-file` option. This inserted character invalidates the key because the size is wrong.

An example of using the `--encrypt-key-file` option:

```
$ xtrabackup --backup --encrypt=AES256 --encrypt-key-file=/data/backups/keyfile --target-dir=/data/backup
```



## Optimize the encryption process

Additional encrypted backup options, `--encrypt-threads` and `--encrypt-chunk-size`, can speed up the encryption process.

Use the `--encrypt-threads` option to enable parallel encryption with multiple threads.

The `--encrypt-chunk-size` option specifies the size, in bytes, of the working encryption buffer for each encryption thread. The default size is 64K.

## Decrypt encrypted backups

You can decrypt backups with the `xbcrypt` binary. The following example encrypts a backup.

You can use the `--parallel` option and the `--decrypt` option to decrypt multiple files simultaneously.

```
$ for i in `find . -iname "*\.xbcrypt"`; do xbcrypt -d --encrypt-key-
file=/root/secret_key --encrypt-algo=AES256 < $i > $(dirname $i)/$(basename $i
.xbcrypt) && rm $i; done
```

The following example shows a decryption process.

```
$ xtrabackup --decrypt=AES256 --encrypt-key="{randomly-generated-alphanumeric-
string}" --target-dir=/data/backup/
```

Percona XtraBackup doesn't automatically remove the encrypted files. You must remove the `\*.xbcrypt` files manually.

## Prepare encrypted backups

After decrypting the backups, prepare the backups with the `--prepare` option:

```
$ xtrabackup --prepare --target-dir=/data/backup/
```

## Restore encrypted backups

*xtrabackup* offers the `--copy-back` option to restore a backup to the server's datadir:

```
$ xtrabackup --copy-back --target-dir=/data/backup/
```

The option copies all the data-related files to the server's datadir. The server's `my.cnf` configuration file determines the location.

You should check the last line of the output for a success message:

#### Expected output



```
150318 11:08:13 xtrabackup: completed OK!
```

# Take a partial backup

# Create a partial backup

*xtrabackup* supports taking partial backups when the `innodb_file_per_table` option is enabled. There are three ways to create partial backups:

1. matching the tables names with a regular expression
2. providing a list of table names in a file
3. providing a list of databases

## Warning

Do not copy back the prepared backup.

Restoring partial backups should be done by importing the tables, not by using the `--copy-back` option. It is not recommended to run incremental backups after running a partial backup.

Although there are some scenarios where restoring can be done by copying back the files, this may lead to database inconsistencies in many cases and it is not a recommended way to do it.

For the purposes of this manual page, we will assume that there is a database named `test` which contains tables named `t1` and `t2`.

## Warning

If any of the matched or listed tables is deleted during the backup, *xtrabackup* will fail.

There are multiple ways of specifying which part of the whole data is backed up:

- Use the `--tables` option to list the table names
- Use the `--tables-file` option to list the tables in a file
- Use the `--databases` option to list the databases
- Use the `--databases-file` option to list the databases

## Mutual exclusion

The `--tables` and `--databases` options are mutually exclusive. If you use both options in the same command, XtraBackup ignores `--databases` and only uses `--tables`. Use only one of these options per backup operation.

## The `--tables` option

The first method involves the `xtrabackup --tables` option. This option accepts either: \* A comma-separated list of fully qualified table names in the format `database.table` (for example, `db1.t1,db1.t2,db2.t3`) \* A POSIX regular expression surrounded by single quotes that is matched against the fully-qualified database name and table name using the `databasename.tablename` format

To back up only tables in the `test` database, use the following command:

```
$ xtrabackup --backup --datadir=/var/lib/mysql --target-dir=/data/backups/ \
--tables="^test[.].*" 
```

To back up only the `test.t1` table, use the following command:

```
$ xtrabackup --backup --datadir=/var/lib/mysql --target-dir=/data/backups/ \
--tables="^test[.]t1" 
```

## The `--tables-file` option

The `--tables-file` option specifies a file that can contain multiple table names, one table name per line in the file. Only the tables named in the file will be backed up. Names are matched exactly, case-sensitive, with no pattern or regular expression matching. The table names must be fully-qualified in `databasename.tablename` format.

```
$ echo "mydatabase.mytable" > /tmp/tables.txt
$ xtrabackup --backup --tables-file=/tmp/tables.txt 
```

## The `--databases` and `--databases-file` options

The `--databases` option accepts a space-separated list of database names. To include all tables in a database, add `.*` after the database name (for example, `mydb.*`). Regular expressions are not supported.

In addition to your selected databases, make sure to specify the `mysql`, `sys`, and `performance_schema` databases. These databases are required when restoring the databases using `xtrabackup --copy-back`.

#### Note

Tables processed during the `--prepare` step may also be added to the backup even if they are not explicitly listed by the parameter if they were created after the backup started.

```
$ xtrabackup --backup --databases='mysql sys performance_schema test' --target-dir=/data/backups/ 
```

## The `--databases-file` option

The `--databases-file` option specifies a file that can contain multiple databases and tables in the `dbname[.tablename]` format, one element name per line in the file. Names are matched exactly, case-sensitive, with no pattern or regular expression matching.

#### Note

Tables processed during the `--prepare` step may also be added to the backup even if they are not explicitly listed by the parameter if they were created after the backup started.

The next step is to [prepare](#) the backup in order to restore it.

# Prepare a partial backup

The procedure is analogous to restoring individual tables: apply the logs and use the `--export` option:

```
$ xtrabackup --prepare --export --target-dir=/path/to/partial/backup
```



When you use the `--prepare` option on a partial backup, you will see warnings about tables that don't exist. This is because these tables exist in the data dictionary inside InnoDB, but the corresponding .ibd files don't exist. They were not copied into the backup directory. These tables will be removed from the data dictionary, and when you restore the backup and start InnoDB, they will no longer exist and will not cause any errors or warnings to be printed to the log file.

Could not find any file associated with the tablespace ID: 5

Use `--innodb-directories` to find the tablespace files. If that fails then use `--innodb-force-recovery=1` to ignore this and to permanently lose all changes to the missing tablespace(s).

The next step is to [restore](#) restore the partition from the backup.



# **Back up individual partitions**

# Create an individual partitions backup

Percona XtraBackup lets you back up individual partitions because partitions are regular tables with specially formatted names. The only requirement for this feature is to enable the `innodb_file_per_table` option on the server.

There is one caveat about using this kind of backup: you can not copy back the prepared backup. Restoring partial backups should be done by importing the tables.

There are three ways of specifying which part of the whole data will be backed up: regular expressions (`--tables`), enumerating the tables in a file (`--tables-file`) or providing a list of databases (`--databases`).

## Mutual exclusion

The `--tables` and `--databases` options are mutually exclusive. If you use both options in the same command, XtraBackup ignores `--databases` and only uses `--tables`. Use only one of these options per backup operation.

The regular expression provided to this option will be matched against the fully qualified database name and table name, in the form of `database-name.table-name`.

If the partition 0 is not backed up, Percona XtraBackup cannot generate a `.cfg` file. MySQL 8.0 stores the table metadata in partition 0.

For example, this operation takes a back-up of the partition `p4` from the table `name` located in the database `imdb`:

```
$ xtrabackup --tables=^imdb[.]name#p#p4 --backup
```



If partition 0 is not backed up, the following errors may occur:

## The error message



```
xtrabackup: export option not specified
xtrabackup: error: cannot find dictionary record of table imdb/name#p#p4
```

Note that this option is passed to `xtrabackup --tables` and is matched against each table of each database, the directories of each database will be created even if they are empty.

# Prepare an individual partitions backup

To prepare partial backups, follow a procedure similar to restoring individual tables. Apply the logs and use the `--export` option:

```
$ xtrabackup --prepare --export --target-dir=/mnt/backup/2012-08-28_10-29-09
```



You may see warnings in the output about tables that do not exist. This happens because InnoDB-based engines stores its data dictionary inside the tablespace files. During the prepare phase, xtrabackup removes entries for tables that were not included in the partial backup from the data dictionary, preventing future warnings or errors.

The next step is to [restore](#) the backup.

# Restore a backup

# Restore full, incremental, and compressed backups

## Warning

Backup needs to be prepared before it can be restored.

The restore backup procedure is the same for full, incremental, and compressed backups.

Ensure you have one of the following before proceeding:

- Root access
- `sudo` access — administrator privileges are required to restore files, change ownership, and manage the database service.

For convenience, *xtrabackup* binary has the `--copy-back` option to copy the backup to the datadir of the server:

```
xtrabackup --copy-back --target-dir=/data/backups/
```



If you don't want to save your backup, you can use the `--move-back` option which will move the backed up data to the datadir.

If you don't want to use any of the above options, you can also use **rsync** or **cp** to restore the files.

## Note

The datadir must be empty before restoring the backup. Also, it's important to note that MySQL server needs to be shut down before restore is performed. You cannot restore to a datadir of a running mysqld instance (except when importing a partial backup).

An example of the **rsync** command to restore the backup:

```
rsync -avrP /data/backup/ /var/lib/mysql/
```



You should check that the restored files have the correct ownership and permissions.

As files' attributes will be preserved, in most cases you will need to change the files' ownership to `mysql` before starting the database server, as they will be owned by the user who created the backup:

```
chown -R mysql:mysql /var/lib/mysql
```



## Prepare grastate.dat for Galera cluster (Optional)

When restoring a backup of a Percona XtraDB Cluster, be sure the `grastate.dat` file is present and correctly populated before you start the node. The example demonstrates how to restore a backup on a single node and then bootstrap a new Percona XtraDB Cluster. Once the first node is up, additional nodes can join later via IST (Incremental State Transfer) or SST (State Snapshot Transfer).

This `grastate.dat` file contains the following information:

- `UUID` - identifies the specific Galera cluster the node last belonged to.
- `seqno` - indicates the global transaction number of the last committed write-set on that node.
- `safe_to_bootstrap` flag - indicates whether the node was the last to shut down gracefully. A value of `1` means the node was the last to shut down gracefully and it is safe to use this node to bootstrap a new cluster. A value of `0` means the node was not the last to shut down gracefully and it should not be used to start the cluster.

After preparing and restoring the backup, do the following steps:

1. Extract the Galera state information from the backup metadata file located in the restored datadir (for example, `/var/lib/mysql`). The extracted information is represented in the `uuid:seqno` format.

```
cat /var/lib/mysql/xtrabackup_galera_info
```



### Expected output



```
3f9d1c8e-4a12-11ee-bc3c-8c8590c4abcd:145672
```



2. Create or update the `grastate.dat` file in the datadir using the extracted `UUID` and `seqno`. In this example, the node is used to bootstrap a new cluster, so `safe_to_bootstrap` is set to `1`.



### Note

Set `safe_to_bootstrap` to `1` only on the node that will be used to bootstrap a new cluster. All other nodes must have `safe_to_bootstrap` set to `0`.

**Bootstrap node (`safe_to_bootstrap: 1`):**

```
cat > /var/lib/mysql/grastate.dat <<EOF
GALERA saved state
version: 2.1
uuid: 3f9d1c8e-4a12-11ee-bc3c-8c8590c4abcd
seqno: 145672
safe_to_bootstrap: 1
EOF
```

**All other nodes (`safe_to_bootstrap: 0`):**

```
cat > /var/lib/mysql/grastate.dat <<EOF
GALERA saved state
version: 2.1
uuid: 3f9d1c8e-4a12-11ee-bc3c-8c8590c4abcd
seqno: 145672
safe_to_bootstrap: 0
EOF
```

3. Ensure the file has the correct ownership:

```
chown mysql:mysql /var/lib/mysql/grastate.dat
```

If `grastate.dat` is missing or contains incorrect values, the node may fail to start or attempt an unsafe cluster bootstrap.

Data is now restored, and you can start the server.

# Restore a partial backup

Restoring should be done by restoring individual tables in the partial backup to the server.

It can also be done by copying back the prepared backup to a “clean” datadir (in that case, make sure to include the `mysql` database) to the datadir you are moving the backup to. A system database can be created with the following:

```
$ sudo mysql --initialize --user=mysql
```

Once you start the server, you may see mysql complaining about missing tablespaces:

## Expected output

```
2021-07-19T12:42:11.077200Z 1 [Warning] [MY-012351] [InnoDB] Tablespace 4, name 'test1/t1',
file './d2/test1.ibd' is missing!
2021-07-19T12:42:11.077300Z 1 [Warning] [MY-012351] [InnoDB] Tablespace 4, name 'test1/t1',
file './d2/test1.ibd' is missing!
```

In order to clean the orphan database from the data dictionary, you must manually create the missing database directory and then `DROP` this database from the server.

Example of creating the missing database:

```
$ mkdir /var/lib/mysql/test1/d2
```

Example of dropping the database from the server:

```
mysql> DROP DATABASE d2;
```

## Expected output

```
Query OK, 2 rows affected (0.5 sec)
```



# Restore the partition from the backup

Restoring should be done by importing the tables in the partial backup to the server.

First step is to create new table in which data will be restored.

```
mysql> CREATE TABLE `name_p4` (
 `id` int(11) NOT NULL AUTO_INCREMENT,
 `name` text NOT NULL,
 `imdb_index` varchar(12) DEFAULT NULL,
 `imdb_id` int(11) DEFAULT NULL,
 `name_pcode_cf` varchar(5) DEFAULT NULL,
 `name_pcode_nf` varchar(5) DEFAULT NULL,
 `surname_pcode` varchar(5) DEFAULT NULL,
 PRIMARY KEY (`id`)
) ENGINE=InnoDB AUTO_INCREMENT=2812744 DEFAULT CHARSET=utf8
```

## Note

Generate a [.cfg metadata file](#) by running `FLUSH TABLES ... FOR EXPORT`. The command can only be run on a table, not on the individual table partitions. The file is located in the table schema directory and is used for schema verification when importing the tablespace.

To restore the partition from the backup, the tablespace must be discarded for that table:

```
mysql> ALTER TABLE name_p4 DISCARD TABLESPACE;
```

The next step is to copy the `.ibd` file from the backup to the MySQL data directory:

```
cp /mnt/backup/2012-08-28_10-29-09/imdb/name#P#p4.ibd
/var/lib/mysql/imdb/name_p4.ibd
```

## Note

Make sure that the copied files can be accessed by the user running MySQL.

The last step is to import the tablespace:

```
mysql> ALTER TABLE name_p4 IMPORT TABLESPACE;
```

# Restore single tables between databases

Sometimes, you must restore just one table instead of a whole database. Percona XtraBackup helps you do this. This guide shows you how to restore specific tables between different database systems while keeping your data safe.

This process only works with InnoDB tables that use the `file-per-table` setup. This setup stores each table's data in its `.ibd` file. You can restore these tables to systems running Percona Server for MySQL with XtraDB or MySQL 8.0. Your source database can be any type, but the target database must be XtraDB or MySQL 8.0.

When moving tables, you'll work with several important files. Each file has a specific job in the restoration process:

File	Description
ibdata1	The main system file in InnoDB that holds necessary metadata and sometimes table data. You need this file for any InnoDB restoration, including when recovering just one table.
undo_001 and undo_002	Files that help undo transactions that weren't finished. These files help recover from crashes during the preparation phase. They're an essential part of a complete backup plan.
mysql.ibd	A file that contains table definitions and structure information. This file is necessary when restoring individual tables.
xtrabackup_info	A file with information about the backup, including which server version was used and when the backup was made.
xtrabackup_checkpoints	A file that keeps track of important points in the backup process. These checkpoints help with crash recovery.
xtrabackup_logfile	The transaction log that records all database changes. This log ensures your backup stays consistent throughout the restoration.
xtrabackup_tablespaces	This is a list of tablespaces stored outside the main data directory. During restoration, XtraBackup tries to rebuild file paths to match their original locations.

File	Description
.ibd from all tables you want to restore	The most critical files for single-table restoration. Each InnoDB table in file-per-table mode has its .ibd file containing data and indexes. You must have the .ibd file for each table you want to restore.

The following example shows the export and import process for a table called `export_test`:

```
mysql -u <username> -p -e "CREATE DATABASE create_database_test;"
```

```
mysql> CREATE TABLE export_test (
 a int(11) DEFAULT NULL
) ENGINE=InnoDB DEFAULT CHARSET=latin1;
```

## Save the schema

Before importing any table, you must save the schema definition for all tables you want to restore. Think of the schema as a blueprint for your table. XtraBackup only backs up data files, not these blueprints. Without the original schema, MySQL won't be able to use the restored .ibd files.

When you back up your database, you may not know which tables to restore later. This operation prepares you for any future restoration needs. Also, table structures might change between the time when you create the backup and when you use `mysqldump`. Saving the most current structure helps ensure everything works during restoration.

Review the script in [Dump schema for all tables](#) for an easier way to save schemas. You'll likely need to adjust this script to fit your setup.

By saving schemas ahead of time, you ensure that all the table blueprints are ready when you start restoring from backups, making data recovery go more smoothly.

## Export the table

First, create the table using the `innodb_file_per_table` setting. This setting tells MySQL to create a separate `.ibd` file for each table in your source directory. Run this command to find the created file in your system:

```
$ find /data/backups/mysql/ -name export_test.*
/data/backups/mysql/test/export_test.ibd
```

Next, when you run the `--prepare` phase, add the `--export` option to your command. This option tells XtraBackup to create any extra files needed for moving the table:

```
$ xtrabackup --prepare --export --target-dir=/data/backups/mysql/
```

If you're working with encrypted tables (ones that use transparent data encryption), you'll need extra security settings. You must include the keyring file that has the encryption keys:

```
$ xtrabackup --prepare --export --target-dir=/tmp/table \
--keyring-file-data=/var/lib/mysql-keyring/keyring
```

After these commands finish successfully, you'll need these files for the table import.

```
/data/backups/mysql/test/export_test.ibd
/data/backups/mysql/test/export_test.cfg
/data/backups/mysql/test/export_test.cfp
```

## Restore the table

Now that we've exported the table, we can restore it to the destination system. You need to follow these steps in order on the target server:

- Use the saved schema definition to create an empty table at the destination. This schema provides the framework for your imported data.
- Copy all three important files (.ibd, .cfp, and .cfg) from your source directory to the matching directory in the destination system.

The following is the detailed process:

1. Use the schema and create an empty table copy in the new location on your target server (which must be running Percona Server for MySQL with XtraDB or MySQL 8.0).

```
$ mysql -uroot < ./schema_dumps/test/create_database_test.sql;
$ mysql -uroot < ./schema_dumps/test/export_test.sql;
```

2. Run this SQL command to disconnect the new table from its default tablespace. This command prepares it to receive the imported data:

```
mysql> ALTER TABLE test.export_test DISCARD TABLESPACE;
```

If you see an error like “ERROR 1809 (HY000): Table ‘test/export\_test’ in the system tablespace,” the `innodb_file_per_table` setting is disabled on your target server. Enable this in your server setup, then try the steps again.

3. Copy all the table files (`export_test.ibd`, `export_test.cfp`, and `export_test.cfg`) from the source directory to the matching destination schema directory inside the server data directory.
4. Run this SQL command to attach the restored tablespace to the destination table, finishing the data restore:

```
mysql> ALTER TABLE test.export_test IMPORT TABLESPACE;
```



After these steps are finished successfully, your table restore is complete and the data is ready to use. You can run a `SELECT` query on the newly imported table to check that the imported data is correct.

# Replicate

# How to set up a replica for replication in 6 simple steps with Percona XtraBackup

Data is, by far, the most valuable part of a system. Having a backup done systematically and available for a rapid recovery in case of failure is admittedly essential to a system. However, it is not common practice because of its costs, infrastructure needed or even the boredom associated to the task. Percona XtraBackup is designed to solve this problem.

You can have almost real-time backups in 6 simple steps by setting up a replication environment with Percona XtraBackup.

## All the things you will need

Setting up a replica for replication with *Percona XtraBackup* is a straightforward procedure. In order to keep it simple, here is a list of the things you need to follow the steps without hassles:

### Source

A system with a *MySQL*-based server installed, configured and running. This system will be called **Source**, as it is where your data is stored and the one to be replicated. We will assume the following about this system:

- the *MySQL* server is able to communicate with others by the standard TCP/IP port;
- the *SSH* server is installed and configured;
- you have a user account in the system with the appropriate permissions;
- you have a *MySQL*'s user account with appropriate privileges.
- server has binlogs enabled and server-id set up to 1.

### Replica

Another system, with a *MySQL*-based server installed on it. We will refer to this machine as **Replica** and we will assume the same things we did about **Source**, except that the server-id on **Replica** is 2.

### Percona XtraBackup

The backup tool we will use. It should be installed in both computers for convenience.

#### Note

It is not recommended to mix MySQL variants (Percona Server, MySQL) in your replication setup. This may produce incorrect `xtrabackup_slave_info` file when adding a new replica.

## Version updates

The 8.0.23 version deprecates the `CHANGE_MASTER_TO` command. In that version and later, use the [CHANGE REPLICATION SOURCE TO and the appropriate options](#) instead.

The 8.0.22 version deprecates the `START SLAVE` command. In that version or later, use `START REPLICATION` instead.

## 1. Make a backup on the Source and prepare it

At the `Source`, issue the following to a shell:

```
$ xtrabackup --backup --user=yourDBuser --password=MaGiCdB1 --target-dir=/path/to/backupdir
```

After this is finished you should get:

#### Expected output

```
xtrabackup: completed OK!
```

This will make a copy of your *MySQL* data dir to the `/path/to/backupdir` directory. You have told *Percona XtraBackup* to connect to the database server using your database user and password, and do a hot backup of all your data in it (all *MyISAM*, *InnoDB* tables and indexes in them).

In order for snapshot to be consistent you need to prepare the data on the source:

```
$ xtrabackup --prepare --target-dir=/path/to/backupdir
```

You need to select path where your snapshot has been taken. If everything is ok you should get the same OK message. Now the transaction logs are applied to the data files, and new ones are created: your data files are ready to be used by the *MySQL* server.



*Percona XtraBackup* knows where your data is by reading your `my.cnf` file. If you have your configuration file in a non-standard place, you should use the flag `--defaults-file` `=/location/of/my.cnf`.

If you want to skip writing the username and password every time you want to access *MySQL*, you can set it up in `.mylogin.cnf` as follows:

```
mysql_config_editor set --login-path=client --host=localhost --user=root --
password
```

For more information, see [MySQL Configuration Utility](#).

This statement provides root access to *MySQL*.

## 2. Copy backed up data to the Replica

On the Source, use `rsync` or `scp` to copy the data from the Source to the Replica. If you are syncing the data directly to replica's data directory, we recommend that you stop the `mysqld` there.

```
$ rsync -avP -e ssh /path/to/backupdir Replica:/path/to/mysql/
```

After data has been copied, you can back up the original or previously installed *MySQL* datadir (**NOTE:** Make sure `mysqld` is shut down before you move the contents of its datadir, or move the snapshot into its datadir.). Run the following commands on the Replica:

```
$ mv /path/to/mysql/datadir /path/to/mysql/datadir_bak
```

and move the snapshot from the `Source` in its place:

```
$ xtrabackup --move-back --target-dir=/path/to/mysql/backupdir
```

After you copy data over, make sure the Replica *MySQL* has the proper permissions to access them.

```
$ chown mysql:mysql /path/to/mysql/datadir
```

If the `ibdata` and `iblog` files are located in directories outside the datadir, you must put these files in their proper place after the logs have been applied.

## 3. Configure the Source's MySQL server

On the source, create a user specifically for replication tasks. Use the following command to define the user and set a secure password:

```
mysql> CREATE USER 'repl'@$replicaip IDENTIFIED BY '$replicapass';
```

Grant the replication privileges to the newly created user to allow the replica to connect to the source server:

```
mysql> GRANT REPLICATION SLAVE ON *.* TO 'repl'@$replicaip;
```

Apply the changes:

```
mysql> FLUSH PRIVILEGES;
```

Also make sure that firewall rules are correct and that the **Replica** can connect to the **Source**. Run the following command on the **Replica** to test that you can run the **mysql** client on **Replica**, connect to the **Source**, and authenticate.

```
mysql> mysql --host=Source --user=repl --password=$replicapass
```

Verify the privileges.

```
mysql> SHOW GRANTS;
```

## 4. Configure the Replica's MySQL server

Copy the **my.cnf** file from the **Source** to the **Replica**:

```
$ scp user@Source:/etc/mysql/my.cnf /etc/mysql/my.cnf
```

and change the following options in **/etc/mysql/my.cnf**:

```
server-id=2
```

and start/restart **mysqld** on the **Replica**.

In case you're using **init script** on **Debian-based** system to start **mysqld**, be sure that the password for **debian-sys-maint** user has been updated, and it's the same as that user's password on the **Source**. Password can be seen and updated in **/etc/mysql/debian.cnf**.

## 5. Start the replication

On the `Replica`, review the content of the file `xtrabackup_binlog_info`, it will be something like:

```
$ cat /var/lib/mysql/xtrabackup_binlog_info
```



The results should resemble the following:

### Expected output



```
Source-bin.000001 481
```

[The term `master` is deprecated](#). Do the following on a MySQL console and use the username and password you've set up in STEP 3 :

- Version 8.0.23 or later, use the `CHANGE REPLICATION SOURCE TO` statement
- Before 8.0.23, use the `CHANGE MASTER` statement

```
CHANGE REPLICATION SOURCE TO
 SOURCE_HOST=' $sourceip',
 SOURCE_USER=' repl',
 SOURCE_PASSWORD=' $replicapass',
 SOURCE_LOG_FILE=' Source-bin.000001',
 SOURCE_LOG_POS=481;
```



Start the replica:

```
START REPLICA;
```



The [term `slave` is deprecated](#). Do the following:

- Version 8.0.22 or later, use `START REPLICA`
- Before version 8.0.22, use `START SLAVE`

## 6. Check

On the `Replica`, check that everything went OK with:

```
SHOW REPLICA STATUS\G
```



The result shows the status:

#### Expected output

```
Slave_IO_Running: Yes
Slave_SQL_Running: Yes
Seconds_Behind_Master: 13
```

Both `IO` and `SQL` threads need to be running. The `Seconds_Behind_Master` means the `SQL` currently being executed has a `current_timestamp` of 13 seconds ago. It is an estimation of the lag between the `Source` and the `Replica`. Note that at the beginning, a high value could be shown because the `Replica` has to “catch up” with the `Source`.

## Adding more replicas to the Source

You can use this procedure with slight variation to add new replicas to a source. We will use *Percona XtraBackup* to clone an already configured replica. We will continue using the previous scenario for convenience, but we will add a `NewReplica` to the plot.

At the `Replica`, do a full backup:

```
$ xtrabackup --user=yourDBUser --password=MaGiCiGaM \
 --backup --slave-info --target-dir=/path/to/backupdir
```

By using the `--slave-info` *Percona XtraBackup* creates additional file called `xtrabackup_slave_info`.

Apply the logs:

```
$ xtrabackup --prepare --use-memory=2G --target-dir=/path/to/backupdir/
```

#### Note

In the `prepare` phase, the `--use-memory` parameter speeds up the process if the amount of RAM assigned to the option is available. Use the parameter only in the `prepare` phase. In the other phases the parameter makes the application lazy allocate this memory (reserve) but does not affect database pages.

Copy the directory from the `Replica` to the `NewReplica` (**NOTE:** Make sure `mysqld` is shut down on the `NewReplica` before you copy the contents the snapshot into its `datadir`):

```
rsync -avprP -e ssh /path/to/backupdir NewReplica:/path/to/mysql/datadir
```

For example, to set up a new user, `user2`, you add another grant on the Source:

```
> GRANT REPLICATION SLAVE ON *.* TO 'user2'@'$newreplicaip'
 IDENTIFIED BY '$replicapass';
```

On the `NewReplica`, copy the configuration file from the `Replica`:

```
$ scp user@Replica:/etc/mysql/my.cnf /etc/mysql/my.cnf
```

Make sure you change the `server-id` variable in `/etc/mysql/my.cnf` to 3 and disable the replication on start:

```
skip-slave-start
server-id=3
```

After setting `server_id`, start **mysqld**.

Fetch the `master_log_file` and `master_log_pos` from the file `xtrabackup_slave_info`, execute the statement for setting up the source and the log file for the `NewReplica`:

```
> CHANGE MASTER TO
 MASTER_HOST='$Sourceip',
 MASTER_USER='repl',
 MASTER_PASSWORD='$replicapass',
 MASTER_LOG_FILE='Source-bin.000001',
 MASTER_LOG_POS=481;
```

[The term `master` is deprecated](#). Do the following and then start the replica:

- Version 8.0.23 or later, use the `CHANGE REPLICATION SOURCE TO` statement
- Before 8.0.23, use the `CHANGE MASTER` statement

[The term `slave` is deprecated](#). Do the following:

- Version 8.0.22 or later, use `START REPLICATION`.
- Before version 8.0.22, use `START SLAVE`

```
> START REPLICATION;
```

If both IO and SQL threads are running when you check the `NewReplica`, server is replicating the `Source`.



#### See also

[How to create a new \(or repair a broken\) GTID based slave](#)

# How to create a new (or repair a broken) GTID-based replica

MySQL 5.6 introduced the Global Transaction ID ([GTID](#)) support in replication. *Percona XtraBackup* automatically stores the `GTID` value in the `xtrabackup_binlog_info` when doing the backup of MySQL and *Percona Server for MySQL 5.7* with the `GTID` mode enabled. This information can be used to create a new (or repair a broken) `GTID`-based replica.

## 1. Take a backup from any server on the replication environment, source or replica

The following command takes a backup and saves it in the `/data/backups/$TIMESTAMP` folder:

```
$ xtrabackup --backup --target-dirs=/data/backups/
```

In the destination folder, there will be a file with the name `xtrabackup_binlog_info`. This file contains both binary log coordinates and the `GTID` information.

```
$ cat xtrabackup_binlog_info
```

The result could look like this:

### Expected output

```
mysql-bin.000002 1232 c777888a-b6df-11e2-a604-080027635ef5:1-4
```

That information is also printed by `xtrabackup` after taking the backup:

### Expected output

```
xtrabackup: MySQL binlog position: filename 'mysql-bin.000002', position 1232, GTID of the last change 'c777888a-b6df-11e2-a604-080027635ef5:1-4'
```

## 2. Prepare the backup

The backup will be prepared with the following command on the Source:

```
$ xtrabackup --prepare --target-dir=/data/backup
```



You need to select the path where your snapshot has been taken, for example `/data/backups/2013-05-07_08-33-33`. If everything is ok you should get the same OK message. Now, the transaction logs are applied to the data files, and new ones are created: your data files are ready to be used by the MySQL server.

### 3. Move the backup to the destination server

Use **rsync** or **scp** to copy the data to the destination server. If you are synchronizing the data directly to the already running replica's data directory it is advised to stop the *MySQL* server there.

```
$ rsync -avprP -e ssh /path/to/backupdir/$TIMESTAMP NewSlave:/path/to/mysql/
```



After you copy the data over, make sure *MySQL* has proper permissions to access them.

```
$ chown mysql:mysql /path/to/mysql/datadir
```



### 4. Configure and start replication

Set the `gtid_purged` variable to the `GTID` from `xtrabackup_binlog_info`. Then, update the information about the source node and, finally, start the replica. Run the following commands on the replica if you are using a version before 8.0.22:

```
Using the mysql shell
> SET SESSION wsrep_on = 0;
> RESET MASTER;
> SET SESSION wsrep_on = 1;
> SET GLOBAL gtid_purged='<gtid_string_found_in_xtrabackup_binlog_info>';
> CHANGE MASTER TO
 MASTER_HOST="$masterip",
 MASTER_USER="repl",
 MASTER_PASSWORD="$slavepass",
 MASTER_AUTO_POSITION = 1;
> START SLAVE;
```



If you are using version 8.0.22 or later, use `START REPLICA` instead of `START SLAVE`. `START SLAVE` is deprecated as of that release. If you are using version 8.0.21 or earlier, use `START SLAVE`.

If you are using a version 8.0.23 or later, run the following commands:



```
Using the mysql shell
> SET SESSION wsrep_on = 0;
> RESET MASTER;
> SET SESSION wsrep_on = 1;
> SET GLOBAL gtid_purged='<gtid_string_found_in_xtrabackup_binlog_info>';
> CHANGE REPLICATION SOURCE TO
 SOURCE_HOST="$masterip",
 SOURCE_USER="repl",
 SOURCE_PASSWORD="$slavepass",
 SOURCE_AUTO_POSITION = 1;
> START REPLICA;
```

If you are using version 8.0.23 or later, use [CHANGE\\_REPLICATION\\_SOURCE\\_TO](#) and the [appropriate options](#). `CHANGE_MASTER_TO` is deprecated as of that release.

#### Note

The example above is applicable to Percona XtraDB Cluster. The `wsrep_on` variable is set to 0 before resetting the source (`RESET MASTER`). The reason is that Percona XtraDB Cluster will not allow resetting the source if `wsrep_on=1`.

## 5. Check the replication status

The following command returns the replica status:

```
mysql> SHOW REPLICA STATUS\G
```

The results should be similar to the following:

#### Expected output

```
[...]
Slave_IO_Running: Yes
Slave_SQL_Running: Yes
[...]
Retrieved_Gtid_Set: c777888a-b6df-11e2-a604-080027635ef5:5
Executed_Gtid_Set: c777888a-b6df-11e2-a604-080027635ef5:1-5
```

#### Note

The command [SHOW SLAVE STATUS](#) is deprecated. Use [SHOW REPLICA STATUS](#).

We can see that the replica has retrieved a new transaction with number 5, so transactions from 1 to 5 are already on this slave.

We have created a new replica in our `GTID` based replication environment.

# Make backups in replication environments

There are options specific to back up from a replication replica.

## The `--slave-info` option

This option is useful when backing up a replication replica server. It prints the binary log position and name of the source server. It also writes this information to the `xtrabackup_slave_info` file as a `CHANGE MASTER` statement.

This option is useful for setting up a new replica for this source. You can start a replica server with this backup and issue the statement saved in the `xtrabackup_slave_info` file. More details of this procedure can be found in [How to setup a replica for replication in 6 simple steps with Percona XtraBackup](#).

## The `--safe-slave-backup` option

In order to assure a consistent replication state, this option stops the replication SQL thread and waits to start backing up until `Slave_open_temp_tables` in `SHOW STATUS` is zero. If there are no open temporary tables, the backup will take place, otherwise the SQL thread will be started and stopped until there are no open temporary tables. The backup will fail if `Slave_open_temp_tables` does not become zero after `--safe-slave-backup-timeout` seconds (defaults to 300 seconds). The replication SQL thread will be restarted when the backup finishes.

### Note

Prior to *Percona XtraBackup* 8.0.22-15.0, using a `safe-slave-backup` stops the SQL replica thread after the InnoDB tables and before the non-InnoDB tables are backed up.

As of *Percona XtraBackup* 8.0.22-15.0, using a `safe-slave-backup` option stops the SQL replica thread before copying the InnoDB files.

Using this option is always recommended when taking backups from a replica server.

### Warning

Make sure your replica is a true replica of the source before using it as a source for backup. A good tool to validate a replica is [pt-table-checksum](#).

# Verify backups with replication and pt-checksum

One way to verify if the backup is consistent is by setting up the replication and running [pt-table-checksum](#). This can be used to verify any type of backups, but before setting up replication, backup should be prepared and be able to run (this means that incremental backups should be merged to full backups, encrypted backups decrypted etc.).

## Set up the replication

How to set up a replica for replication in 6 simple steps with Percona XtraBackup guide provides a detailed instructions on how to take the backup and set up the replication.

For checking the backup consistency you can use either the original server where the backup was taken, or another test server created by using a different backup method (such as cold backup, mysqldump or LVM snapshots) as the source server in the replication setup.

## Use pt-table-checksum

This tool is part of the *Percona Toolkit*. It performs an online replication consistency check by executing checksum queries on the source, which produces different results on replicas that are inconsistent with the source.

After you confirmed that replication has been set up successfully, you can [install](#) or download *pt-table-checksum*. This example shows downloading the latest version of *pt-table-checksum*:

```
$ wget percona.com/get/pt-table-checksum
```



### Note

In order for *pt-table-checksum* to work correctly `libdbd-mysql-perl` will need to be installed on *Debian/Ubuntu* systems or `perl-DBD-MySQL` on *RHEL/CentOS*. If you installed the *percona-toolkit* package from the Percona repositories package manager should install those libraries automatically.

After this command has been run, *pt-table-checksum* will be downloaded to your current working directory.

Running the *pt-table-checksum* on the source will create `percona` database with the `checksums` table which will be replicated to the replicas as well. Example of the *pt-table-checksum* will look like this:

```
$./pt-table-checksum
```



The results are similar to the following:

#### Expected output



```
TS ERRORS DIFFS ROWS CHUNKS SKIPPED TIME TABLE
04-30T11:31:50 0 0 633135 8 0 5.400 exampledb.aka_name
04-30T11:31:52 0 0 290859 1 0 2.692 exampledb.aka_title
Checksumming exampledb.user_info: 16% 02:27 remain
Checksumming exampledb.user_info: 34% 01:58 remain
Checksumming exampledb.user_info: 50% 01:29 remain
Checksumming exampledb.user_info: 68% 00:56 remain
Checksumming exampledb.user_info: 86% 00:24 remain
04-30T11:34:38 0 0 22187768 126 0 165.216 exampledb.user_info
04-30T11:38:09 0 0 0 1 0 0.033 mysql.time_zone_name
04-30T11:38:09 0 0 0 1 0 0.052 mysql.time_zone_transition
04-30T11:38:09 0 0 0 1 0 0.054 mysql.time_zone_transition_type
04-30T11:38:09 0 0 8 1 0 0.064 mysql.user
```

If all the values in the `DIFFS` column are 0 that means that backup is consistent with the current setup.

In case backup was not consistent *pt-table-checksum* should spot the difference and point to the table that does not match. Following example shows adding new user on the backed up replica in order to simulate the inconsistent backup:

```
mysql> grant usage on exampledb.* to exampledb@localhost identified by
'thisisNewpassword';
```



If we run the *pt-table-checksum* now difference should be spotted

```
$./pt-table-checksum
```



The results are similar to the following:



## Expected output



```
TS ERRORS DIFFS ROWS CHUNKS SKIPPED TIME TABLE
04-30T11:31:50 0 0 633135 8 0 5.400 exampledb.aka_name
04-30T11:31:52 0 0 290859 1 0 2.692 exampledb.aka_title
Checksumming exampledb.user_info: 16% 02:27 remain
Checksumming exampledb.user_info: 34% 01:58 remain
Checksumming exampledb.user_info: 50% 01:29 remain
Checksumming exampledb.user_info: 68% 00:56 remain
Checksumming exampledb.user_info: 86% 00:24 remain
04-30T11:34:38 0 0 22187768 126 0 165.216 exampledb.user_info
04-30T11:38:09 0 0 0 1 0 0.033 mysql.time_zone_name
04-30T11:38:09 0 0 0 1 0 0.052 mysql.time_zone_transition
04-30T11:38:09 0 0 0 1 0 0.054 mysql.time_zone_transition_type
04-30T11:38:09 1 0 8 1 0 0.064 mysql.user
```

This output shows that source and the replica aren't in consistent state and that the difference is in the `mysql.user` table.

More information on different options that `pt-table-checksum` provides can be found in the *pt-table-checksum* [documentation](#).

# Troubleshoot

# Error Message: Found tables with row versions due to INSTANT ADD/DROP columns

MySQL 8.0.29 extended the support for `ALTER TABLE ... ALGORITHM=INSTANT` to allow users to add a column in any table position or drop any column. As part of this update, the redo log format has changed for all server DML operations. This updated redo log format has a design flaw that can cause data corruption for tables with INSTANT ADD/DROP COLUMNS.

The corruption happens when the crash recovery occurs. InnoDB applies redo logs at startup. Xtrabackup copies the redo log during the backup and applies the log as part of the `--prepare` step to bring the backup to a consistent state.

Percona fixed the corruption issue and several other issues with the INSTANT ADD/DROP column feature in the upcoming Percona Server for MySQL 8.0.29.

For more details, see the following:

- [PS-8291](#),
- [PS-8292](#)
- [PS-8303](#)

Percona XtraBackup 8.0.32-25 allows instant columns in the backup for MySQL 8.0.32. This limitation remains for backing up MySQL 8.0.31 and lower versions. This limitation does not exist for backing up Percona Server for any version.

Percona XtraBackup 8.0.29 is able to take backups of Percona Server for MySQL 8.0.29 with tables that have INSTANT ADD/DROP COLUMNS. However, the issues in MySQL 8.0.29 make this version unsafe to take backups.

It is impossible for XtraBackup to deal with the corrupted redo log generated by MySQL 8.0.29 and, for this reason, XtraBackup 8.0.29 version will not take backups if it detects tables with `INSTANT ADD/DROP COLUMNS` and generates an error listing the affected tables and providing instructions to convert them to regular tables.

Please avoid `ALTER ADD/DROP COLUMN` without an explicit `ALGORITHM=INPLACE`. The default `ALGORITHM` is `INSTANT`, so `ALTER TABLE` without the `ALGORITHM` keyword uses the newly introduced `INSTANT` algorithm.

If you already have such tables (information below on how to find such tables), users are advised to do `OPTIMIZE TABLE` against these tables before taking backups.



To find tables with `INSTANT ADD/DROP COLUMNS` run the following command:

```
mysql> SELECT NAME FROM INFORMATION_SCHEMA.INNODB_TABLES WHERE TOTAL_ROW_VERSIONS
> 0;
+-----+
| NAME |
+-----+
| test/t1 |
| test/t2 |
| test/t3 |
+-----+
3 rows in set (0.02 sec)
```

If this query shows an empty result set, you have no issues. Percona Xtrabackup takes the backup of your MySQL 8.0.29 server. If the results are a list of tables, please run `OPTIMIZE TABLE` on the list before taking the backup.

## Percona XtraBackup Error Message

If Percona XtraBackup detects that MySQL 8.0.29 server has tables with instant add/drop columns, it aborts with the following error message

```
2022-07-01T15:18:35.127689+05:30 0 [ERROR] [MY-011825] [Xtrabackup] Found tables
with row versions due to INSTANT ADD/DROP columns
2022-07-01T15:18:35.127714+05:30 0 [ERROR] [MY-011825] [Xtrabackup] This feature
is not stable and will cause backup corruption.
2022-07-01T15:18:35.127723+05:30 0 [ERROR] [MY-011825] [Xtrabackup] Tables found:
2022-07-01T15:18:35.127730+05:30 0 [ERROR] [MY-011825] [Xtrabackup] test/t1
2022-07-01T15:18:35.127737+05:30 0 [ERROR] [MY-011825] [Xtrabackup] test/t2
2022-07-01T15:18:35.127744+05:30 0 [ERROR] [MY-011825] [Xtrabackup] test/t3
2022-07-01T15:18:35.127752+05:30 0 [ERROR] [MY-011825] [Xtrabackup] Please run
OPTIMIZE TABLE or ALTER TABLE ALGORITHM=COPY on all listed tables to fix this
issue.
```

## Summary

The option, `ALGORITHM=INSTANT` is the default value in MySQL 8.0.29. If you do not specify an algorithm, all `ALTER TABLE ADD/DROP COLUMN` statements use this algorithm.

The `INSTANT` algorithm is unstable at this point.

Percona XtraBackup refuses to take backups from MySQL 8.0.29 tables that have used this algorithm. Percona Server users do not suffer the same limitations as the corruption issues are fixed in the upcoming release of Percona Server for MySQL 8.0.29

# xtrabackup exit codes

The *xtrabackup* binary exits with the traditional success value of 0 after a backup when no error occurs. If an error occurs during the backup, the exit value is 1.

In certain cases, the exit value can be something other than 0 or 1, due to the command-line option code included from the *MySQL* libraries. An unknown command-line option, for example, will cause an exit code of 255.

# Reference

# Topic index

- [About Percona Xtrabackup](#)
- [Accelerate the backup process](#)
- [Analyze table statistics](#)
- [Apt pinning the Percona XtraBackup 8.0 packages](#)
- [Binary tarball file names available](#)
- [Compile and install Percona XtraBackup 8.0 from source](#)
- [Configure xtrabackup](#)
- [Connection and privileges needed](#)
- [Copyright and licensing information](#)
- [Create a compressed backup](#)
- [Create a full backup](#)
- [Create an incremental backup](#)
- [Create an individual partitions backup](#)
- [Create a partial backup](#)
- [Decompress and prepare a compressed backup](#)
- [Dictionary cache](#)
- [Encrypt backups](#)
- [Encrypted InnoDB tablespace backups](#)
- [Error message: Found tables with row versions due to INSTANT ADD/DROP columns](#)
- [Exponential backup](#)
- [FIFO data sink](#)
- [Files in the DEB package built for Percona XtraBackup 8.0](#)
- [Files in the RPM package built for Percona XtraBackup 8.0](#)
- [FLUSH TABLES WITH READ LOCK option](#)
- [Frequently asked questions](#)
- [Glossary](#)
- [Home](#)
- [How Percona XtraBackup works](#)
- [How to create a new \(or repair a broken\) GTID-based replica](#)

- [How to set up a replica for replication in 6 simple steps with Percona XtraBackup](#)
- [Improved log statements](#)
- [Install Percona XtraBackup 8.0 from a binary tarball](#)
- [Install Percona XtraBackup 8.0 overview](#)
- [Install Percona XtraBackup 8.0 using downloaded DEB packages](#)
- [Install Percona XtraBackup 8.0 using downloaded RPM packages](#)
- [Index of files created by Percona XtraBackup](#)
- [lock-ddl-per-table option improvements](#)
- [LRU dump backup](#)
- [Make backups in replication environments](#)
- [Permissions needed](#)
- [Point-in-time recovery](#)
- [Prepare a full backup](#)
- [Prepare an incremental backup](#)
- [Prepare an individual partitions backup](#)
- [Prepare a partial backup](#)
- [Percona Toolkit version checking](#)
- [Percona XtraBackup binaries overview](#)
- [Quickstart Guide for Percona XtraBackup](#)
- [Release notes index](#)
- [Restore a backup](#)
- [Restore a partial backup](#)
- [Restore individual tables](#)
- [Restore the partition from a backup](#)
- [Run Percona XtraBackup 8.0 in a Docker container](#)
- [Server version and backup version comparison](#)
- [Smart memory estimation](#)
- [Store backup history on the server](#)
- [Take a streaming backup](#)
- [Take an incremental backup using page tracking](#)

- [Throttling backups](#)
- [Trademark policy](#)
- [The xtrabackup option reference](#)
- [The xtrabackup binary overview](#)
- [The xbcloud binary overview](#)
- [Understand version numbers](#)
- [Uninstall Percona XtraBackup 8.0 on Debian and Ubuntu](#)
- [Uninstall Percona XtraBackup 8.0 on Red Hat Enterprise Linux and CentOS](#)
- [Update curl utility](#)
- [Use an APT repository to install Percona XtraBackup 8.0](#)
- [Use the xbcloud binary with Amazon S3](#)
- [Use the xbcloud binary with Google Cloud Storage](#)
- [Use the xbcloud binary with an IAM instance profile](#)
- [Use the xbcloud binary with Microsoft Azure Cloud Storage](#)
- [Use the xbcloud binary with MinIO](#)
- [Use the xbcloud binary with Swift](#)
- [Use a YUM repository to install Percona XtraBackup](#)
- [The xbcrypt binary overview](#)
- [The xbstream binary overview](#)
- [Verify backups with replication and pt-checksum](#)
- [Work with AppArmor](#)
- [Work with binary logs](#)
- [Work with SELinux](#)
- [xtrabackup exit codes](#)
- [xtrabackup implementation details](#)

# Index of files created by Percona XtraBackup

- Information related to the backup and the server

File Name	Description
<code>backup-my.cnf</code>	This file contains information to start the mini-instance of InnoDB during the <code>--prepare</code> . This <b>not</b> a backup of the original <code>my.cnf</code> . The InnoDB configuration is read from the file <code>backup-my.cnf</code> created by <b>xtrabackup</b> when the backup was made. The <code>--prepare</code> uses InnoDB configuration from <code>backup-my.cnf</code> by default, or from <code>--defaults-file</code> , if specified. The InnoDB's configuration in this context means server variables that affect dataformat, i.e. <code>innodb_page_size</code> option, <code>innodb_log_block_size</code> , etc. Location-related variables, like <code>innodb_log_group_home_dir</code> or <code>innodb_data_file_path</code> are always ignored by <code>--prepare</code> , so preparing a backup always works with data files from the back directory, rather than any external ones.
<code>xtrabackup_checkpoints</code>	The type of the backup (for example, full or incremental), its state (for example, prepared) and the LSN range contained in it. This information is used for incremental backups. Example of the <code>xtrabackup_checkpoints</code> after taking a full backup: <code>backup_type = full-backup</code> from <code>lsn= 0</code> to <code>lsn = 15188961605</code> last <code>lsn = 15188961605</code> Example of the <code>xtrabackup_checkpoints</code> after taking an incremental backup: <code>backup_type = incremental</code> from <code>lsn = 15188961605</code> to <code>lsn = 15189350111</code> last <code>lsn = 15189350111</code>
<code>xtrabackup_binlog_info</code>	The binary log file used by the server and its position at the moment of the backup. A result of the following query: <code>SELECT server_uuid, local, replication, storage_engines FROM performance_schema.log_status;</code>
<code>xtrabackup_binlog</code>	The <b>xtrabackup</b> binary used in the process.
<code>xtrabackup_logfile</code>	Contains data needed for running the: <code>--prepare</code> . The bigger this file is the <code>--prepare</code> process will take longer to finish.
<code>&lt;table_name&gt;.delta.metadata</code>	This file is going to be created when performing the incremental backup. It contains the per-table delta metadata: page size, size of compressed page (if the value is 0 it means the tablespace isn't compressed) and space id. Example of this file: <code>page_size = 16384</code> <code>zip_size = 0</code> <code>space_id = 0</code>

- Information related to the replication environment (if using the `--slave-info` option):

`xtrabackup_slave_info`

The `CHANGE MASTER` statement needed for setting up a replica.



- Information related to the *Galera* and *Percona XtraDB Cluster* (if using the `--galera-info` option):

`xtrabackup_galera_info`

Contains the values of `wsrep_local_state_uuid` and `wsrep_last_committed` status variables

# Understand version numbers

A version number identifies the product release. The product contains the latest Generally Available (GA) features at the time of that release.

8.0.30	-23
Base version	Minor build version

Percona uses semantic version numbering, which follows the pattern of base version and build version. Percona assigns unique, non-negative integers in increasing order for each version release. The version number combines the base [MySQL 8.0](#) version number and the minor build version.

Note that Percona XtraBackup numbering changed after the 8.0.14 version to align Percona XtraBackup versions with Percona Server and MySQL. Find more information in [Aligning Percona XtraBackup Versions with Percona Server for MySQL](#) blog post.

The version numbers for Percona XtraBackup 8.0.30-23 define the following information:

- Base version - the leftmost numbers indicate the [MySQL 8.0](#) version used as a base. An increase in base version resets the minor build version to 0.
- Minor build version - an internal number that denotes the version of the software. A build version increases by one each time the Percona XtraBackup is released.

Percona XtraBackup 8.0.28-20 and 8.0.28-21 are multiple releases based on MySQL 8.0.28.

# Frequently asked questions

## Does Percona XtraBackup 8.0 support taking backups of databases in versions prior to 8.0?

Percona XtraBackup 8.0 can back up only MySQL, Percona Server for MySQL, or Percona XtraDB Cluster instances currently running version 8.0. The data-dictionary, redo-log, and undo-log formats introduced in MySQL 8.0 are incompatible with pre-8.0 releases, so XtraBackup 8.0 does not work with older servers.

The internal structures are converted after an in-place upgrade to MySQL 8.0 (or the equivalent Percona 8.0 product).

After the in-place upgrade, XtraBackup 8.0 can back up the database without restriction, regardless of its original creation version.

For more information, see [Percona XtraBackup 8.x and MySQL 8.0.20](#)

## Does Percona XtraBackup 8.0 support taking backups of databases in higher versions within the 8.0 series?

Yes, starting with **Percona XtraBackup 8.0.34**, you can back up MySQL-compatible databases of **higher versions within the 8.0 series**.

Check the latest list of all [Supported versions](#).

## Does Percona XtraBackup 8.0 support taking backups of databases in lower versions within the 8.0 series?

Yes, if you use **Percona XtraBackup 8.0.35** or higher, you can back up MySQL-compatible databases starting from version **8.0.34** or higher.

Check the latest list of all [Supported versions](#).

## Why will `innobackupex` not run in Percona XtraBackup 8.0?

`innobackupex` has been removed from *Percona XtraBackup* 8.0 in favor of `xtrabackup`.

## Are you aware of any web-based backup management tools (commercial or not) built around *Percona XtraBackup*?

[Zmanda Recovery Manager](#) is a commercial tool that uses *Percona XtraBackup* for Non-Blocking Backups:

"ZRM provides support for non-blocking backups of MySQL using *Percona XtraBackup*. ZRM with *\*Percona XtraBackup* provides resource utilization management by providing throttling based on the number of IO operations per second. *Percona XtraBackup* based backups also allow for table level recovery even though the backup was done at the database level (needs the recovery database server to be *Percona Server for MySQL* with XtraDB)."\*

## xtrabackup binary fails with a floating point exception

In most of the cases this is due to not having installed the required libraries (and version) by *xtrabackup*. Installing the *GCC* suite with the supporting libraries and recompiling *xtrabackup* solves the issue. See [Compiling and Installing from Source Code](#) for instructions on the procedure.

## How xtrabackup handles the ibdata/ib\_log files on restore if they aren't in mysql datadir?

In case the `ibdata` and `ib_log` files are located in different directories outside the datadir, you will have to put them in their proper place after the logs have been applied.

## Backup fails with Error 24: 'Too many open files'

This usually happens when database being backed up contains large amount of files and *Percona XtraBackup* can't open all of them to create a successful backup. In order to avoid this error the operating system should be configured appropriately so that *Percona XtraBackup* can open all its files. On Linux, this can be done with the `ulimit` command for specific backup session or by editing the `/etc/security/limits.conf` to change it globally (**NOTE**: the maximum possible value that can be set up is `1048576` which is a hard-coded constant in the Linux kernel).

## How to deal with skipping of redo logs for DDL operations?

To prevent creating corrupted backups when running DDL operations, *Percona XtraBackup* aborts if it detects that redo logging is disabled. In this case, the following error is printed:

[FATAL] InnoDB: An optimized (without redo logging) DDL operation has been performed. All modified pages may not have been flushed to the disk yet. Percona XtraBackup will not be able to take a consistent backup. Retry the backup operation.

 **Note**

Redo logging is disabled during a [sorted index build](#). To avoid this error, Percona XtraBackup can use metadata locks on tables while they are copied:

- To block all DDL operations, use the `--lock-ddl` option that issues `LOCK TABLES FOR BACKUP`.
- If `LOCK TABLES FOR BACKUP` is not supported, you can block DDL for each table before XtraBackup starts to copy it and until the backup is completed using the `--lock-ddl-per-table` option.

 **Note**

As of *Percona XtraBackup* 8.0.15, the `--lock-ddl-per-table` option is deprecated. Use the `--lock-ddl` option.

# Glossary

## .CSM

Each table with the **CSV Storage Engine** has `.CSM` file which contains the metadata of it.

## .CSV

Each table with the **CSV Storage** engine has `.CSV` file which contains the data of it (which is a standard Comma Separated Value file).

## .exp

Files with the `.exp` extension are created by **Percona XtraBackup** per each *InnoDB* tablespace when the `-export` option is used on prepare. See [restore individual tables](#).

## .frm

For each table, the server will create a file with the `.frm` extension containing the table definition (for all storage engines).

## General availability (GA)

A finalized version of the product which is made available to the general public. It is the final stage in the software release cycle.

## .ibd

On a multiple tablespace setup ([`innodb_file_per_table`] enabled), *MySQL* will store each newly created table on a file with a `.ibd` extension.

## .MRG

Each table using the **MERGE** storage engine, besides of a `.frm` file, will have `.MRG` file containing the names of the *MyISAM* tables associated with it.

## .MYD

Each *MyISAM* table has `.MYD` (MYData) file which contains the data on it.

## **.MYI**

Each *MyISAM* table has `.MYI` (MYIndex) file which contains the table's indexes.

## **.opt**

*MySQL* stores options of a database (like charset) in a file with a `.opt` extension in the database directory.

## **.par**

Each partitioned table has `.par` file which contains metadata about the partitions.

## **.TRG**

The file contains the triggers associated with a table, for example, `\mytable.TRG`. With the `.TRN` file, they represent all the trigger definitions.

## **.TRN**

The file contains the names of triggers that are associated with a table, for example, `\mytable.TRN`. With the `.TRG` file, they represent all the trigger definitions.

## **backup**

The process of copying data or tables to be stored in a different location.

## **compression**

The method that produces backups in a reduced size.

## **configuration file**

The file that contains the server startup options.

## **crash**

An unexpected shutdown which does not allow the normal server shutdown cleanup activities.

## crash recovery

The actions that occur when MySQL is restarted after a crash.

## data dictionary

The metadata for the tables, indexes, and table columns stored in the InnoDB system tablespace.

## datadir

The directory in which the database server stores its data files. Most Linux distribution use `/var/lib/mysql` by default.

## full backup

A backup that contains the complete source data from an instance.

## ibdata

The default prefix for tablespace files. For example, `ibdata1` is a 10MB auto-extensible file that MySQL creates for a shared tablespace by default.

## incremental backup

A backup stores data from a specific point in time.

## InnoDB

Storage engine which provides ACID-compliant transactions and foreign key support, among others improvements over *MyISAM*. It is the default engine for *MySQL* as of the 8.0 series.

## innodb\_buffer\_pool\_size

The size in bytes of the memory buffer to cache data and indexes of *InnoDB*'s tables. This aims to reduce disk access to provide better performance.

```
[mysqld]
```

```
innodb_buffer_pool_size=8MB
```

## innodb\_data\_home\_dir



The directory (relative to `datadir`) where the database server stores the files in a shared tablespace setup. This option does not affect the location of `innodb\_file\_per\_table`. For example:

```
[mysqld]
innodb_data_home_dir = ./
```

## innodb\_data\_file\_path

Specifies the names, sizes and location of shared tablespace files:

```
[mysqld]
innodb_data_file_path=ibdata1:50M;ibdata2:50M:autoextend
```

## innodb\_file\_per\_table

By default, InnoDB creates tables and indexes in a [file-per-tablespace](#). If the `innodb_file_per_table` variable is disabled, you can enable the variable in your configuration file:

```
[mysqld]
innodb_file_per_table
or
start the server with --innodb_file_per_table.
```

## innodb\_log\_group\_home\_dir

Specifies the location of the *InnoDB* log files:

```
[mysqld]
innodb_log_group_home=/var/lib/mysql
```

## logical backup

A backup which contains a set of SQL statements. The statements can be used to recreate the databases.

## LSN

Each InnoDB page contains a log sequence number(LSN). The LSN is the system version number for the database. Each page's LSN shows how recently it was changed.

## my.cnf

The database server's main configuration file. Most Linux distributions place it as `/etc/mysql/my.cnf` or `/etc/my.cnf`, but the location and name depends on the particular installation. Note that this method is not the only way of configuring the server, some systems rely on the command options.

## MyISAM

The *MySQL* default storage engine until version 5.5. It doesn't fully support transactions but in some scenarios may be faster than *InnoDB*. Each table is stored on disk in 3 files: `.frm`, `.MYD`, `.MYI`.

## physical backup

A backup that copies the data files.

## point in time recovery

This method restores the data into the state it was at any selected point of time.

## prepared backup

A consistent set of backup data that is ready to be restored.

## restore

Copies the database backups taken using the backup command to the original location or a different location. A restore returns data that has been either lost, corrupted, or stolen to the original condition at a specific point in time.

## Tech preview

A tech preview item can be a feature, a variable, or a value within a variable. Before using this feature in production, we recommend that you test restoring production from physical backups in your environment and also use an alternative backup method for redundancy. A tech preview item is included in a release for users to provide feedback. The item is either updated and released as [general availability\(GA\)](#) or removed if not useful. The functionality can change from tech preview to GA.

## xbcrypt

To support the encryption and the decryption of the backups, a new tool `xbcrypt` was introduced to **Percona XtraBackup**. This utility has been modeled after the **xbstream** binary to perform encryption and decryption outside **Percona XtraBackup**.

## **xbstream**

To support simultaneous compression and streaming, **Percona XtraBackup** uses the **xbstream** format. For more information see [xbstream](#)

## **XtraDB**

*Percona XtraDB* is an enhanced version of the InnoDB storage engine, designed to better scale on modern hardware. *Percona XtraDB* includes features which are useful in a high performance environment. It is fully backward-compatible, and is a drop-in replacement for the standard InnoDB storage engine. For more information, see [The Percona XtraDB Storage Engine](#).

## **Zstandard (ZSTD)**

**ZSTD** is a fast lossless compression algorithm that targets real-time compression scenarios and better compression ratios.

**ZSTD** compression algorithm was implemented in [Percona XtraBackup 8.0.30-23](#).

# Dump schema for all tables

This script is designed to dump the schemas of all tables from each database in a MySQL server instance, excluding the system databases, into separate files.

Before you run the script, modify it for your specific use case. For example, you may want to exclude certain tables or databases or specify a different output directory.

```

print_usage() {
 echo "Usage: $0 [--defaults-file=<file>] [--destination-dir=<directory>]"
 echo " --defaults-file=<file> (Optional) Path to MySQL defaults file
(e.g., /etc/my.cnf)"
 echo " --destination-dir=<directory> (Optional) Directory to save the schema
dumps (default: schema_dumps)"
 exit 1
}

USER=""
DEFAULTS_FILE=""
HARDCODED_PASSWORD="test2"
DEST_DIR="schema_dumps"

Check for -h or --help option
if [["$1" == "-h"]] || [["$1" == "--help"]]; then
 print_usage
fi

Parse arguments using a loop to check for --defaults-file and --destination-dir
options
while [["$#" -gt 0]]; do
 case "$1" in
 --defaults-file=*)
 DEFAULTS_FILE="--defaults-file=${1#*=}"
 echo "Using provided defaults file: ${1#*=}"
 ;;
 --destination-dir=*)
 DEST_DIR="${1#*=}"
 echo "Saving schema dumps in: $DEST_DIR"
 ;;
 *)
 echo "Error: Invalid option '$1'"
 print_usage
 ;;
 esac
 shift
done

Default behavior if no defaults-file is provided
if [-z "$DEFAULTS_FILE"]; then
 if [-n "$MYSQL_PWD"]; then
 echo "Using password from environment variable MYSQL_PWD"
 else
 DEFAULTS_FILE="--password=$HARDCODED_PASSWORD"
 echo "Using hardcoded password in the script"
 fi
 USER="-uroot"
 echo "Using hardcoded user: $USER"
fi

```

```

mkdir -p "$DEST_DIR" # Root directory for schema dumps
echo "Schema will be saved in: $DEST_DIR"

Get the list of all databases, excluding system databases
databases=$(mysql $DEFAULTS_FILE $USER -N -e "SHOW DATABASES" | grep -Ev
"^(information_schema|performance_schema|mysql|sys)$")

for db in $databases; do
 echo "Dumping schemas for database: $db"

 mkdir -p "$DEST_DIR/$db" # Create a directory for each database

 # Dump only the database creation statement
 mysqldump $DEFAULTS_FILE $USER --no-data --databases "$db" >
"$DEST_DIR/$db/create_database_${db}.sql"

 # Get the list of tables for the current database
 tables=$(mysql $DEFAULTS_FILE $USER -N -e "SHOW TABLES FROM $db")

 for table in $tables; do
 echo " Dumping schema for table: $table"

 # Ensure `USE <database>;` is at the top of each table's schema file
 echo "USE \`${db}\`;" > "$DEST_DIR/$db/${table}.sql"

 # Append the table schema dump
 mysqldump $DEFAULTS_FILE $USER --no-data --skip-add-drop-table "$db"
"$table" >> "$DEST_DIR/$db/${table}.sql"
 done
done

echo "Schema dumps are saved in the '$DEST_DIR' directory."

```

## Usage

To use this script, follow these steps:

### 1. Save the script

- Copy the script to your local machine.
- Save the file with a `.sh` extension, such as `dump_schemas.sh`.

### 2. Make the script executable

- Run the following command to make the script executable:

```
$ chmod +x dump_schemas.sh
```



### 3. Verify dependencies

- Check if `mysqldump` is installed by running:

```
$ which mysqldump
```



- If `mysqldump` is not installed or not available in your PATH, install it using:

```
$ sudo apt install mysql-client
```



### 4. Use one of the following commands to execute the script:

- See the usage instructions, run:

```
$./dump_schemas.sh --help
```



- Specify a MySQL defaults file (for example, `~/ .my.cnf`) and destination directory for the schema dumps, run:

```
$./dump_schemas.sh --defaults-file=/path/to/my.cnf --destination-dir=/path/to/output
```



- If you do not have a `defaults-file`, use the script defaults of a hardcoded password or an environment variable, if set. Specify only the destination directory:

```
$./dump_schemas.sh --destination-dir=/path/to/output
```



- To use the hardcoded password and default destination directory, `schema_dumps`, use:

```
$./dump_schemas.sh
```



### 5. Inspect the output

- The script creates a `schema_dumps` directory in the current working directory.
- The dumped schemas are organized as follows:

```
$ find .
```



 Expected output 

```
.
./test
./test/t4.sql
./test/t3.sql
./test/create_database_test.sql
./test2
./test2/create_database_test2.sql
./test2/create_database_test2.sql
./test2/t1.sql
...
```

## 6. Review logs for errors or warnings

- After execution, check the terminal output for any errors or warnings that may indicate issues with the backup process.

# Troubleshooting

If you encounter any issues while using this script, please check the following:

- Ensure that `mysqldump` is installed and available in your PATH.
- Verify that the MySQL root user and password are correct.
- Check the permissions of the output directory and ensure that the script has write access.
- Ensure that the MySQL server is running and accessible.
- Ensure that the MySQL server is configured to accept remote connections.



# Supported versions

Percona XtraBackup 8.0.35-36 supports backing up databases compatible with server version 8.0.34 and later. Percona XtraBackup 8.0.35-36 can create backups of older server versions when you override the server version check and enable specific options, such as `lock-ddl=ON`. This flexibility helps you maintain compatibility and perform backups across mixed-version environments.

Percona XtraBackup 8.0.35-36 supports taking backups of the following MySQL-compatible databases:

- Percona Server for MySQL 8.0.34, and MySQL 8.0.34
- Percona Server for MySQL 8.0.35, and MySQL 8.0.35
- Percona Server for MySQL 8.0.36 and MySQL 8.0.36 and any later 8.0.x release

For more information, see [Percona XtraBackup Version Compatibility and Server Version Checks](#).

[For details on the backup process, see How XtraBackup Works.](#)

# Percona Toolkit version checking

Some Percona software contains “version checking” functionality which is a feature that enables Percona software users to be notified of available software updates to improve your environment security and performance. Alongside this, the version check functionality also provides Percona with information relating to which software versions you are running, coupled with the environment confirmation which the software is running within. This helps enable Percona to focus our development effort accordingly based on trends within our customer community.

The purpose of this document is to articulate the information that is collected, as well as to provide guidance on how to disable this functionality if desired.

## Usage

*Version Check* was implemented in *Percona Toolkit* 2.1.4, and was enabled by default in version 2.2.1. Currently, it is supported as a `--[no]version-check` option by a number of tools in Percona Toolkit, *Percona XtraBackup*, and *Percona Monitoring and Management* (PMM).

When launched with Version Check enabled, the tool that supports this feature connects to a Percona’s *version check service* via a secure HTTPS channel. It compares the locally installed version for possible updates, and also checks versions of the following software:

- Operating System
- Percona Monitoring and Management (PMM)
- MySQL
- Perl
- MySQL driver for Perl (DBD::mysql)
- Percona Toolkit

Then it checks for and warns about versions with known problems if they are identified as running in the environment.

Each version check request is logged by the server. Stored information consists of the checked system unique ID followed by the software name and version. The ID is generated either at installation or when the *version checking* query is submitted for the first time.

#### Note

Prior to version 3.0.7 of *Percona Toolkit*, the system ID was calculated as an MD5 hash of a hostname, and starting from *Percona Toolkit* 3.0.7 it is generated as an MD5 hash of a random number. *Percona XtraBackup* continues to use hostname-based MD5 hash.

As a result, the content of the sent query is as follows:

#### Expected output

```
85624f3fb5d2af8816178ea1493ed41a;DBD: :mysql;4.044
c2b6d625ef3409164cbf8af4985c48d3;MySQL;MySQL Community Server (GPL) 5.7.22-log
85624f3fb5d2af8816178ea1493ed41a;OS;Manjaro Linux
85624f3fb5d2af8816178ea1493ed41a;Percona::Toolkit;3.0.11-dev
85624f3fb5d2af8816178ea1493ed41a;Perl;5.26.2
```

## Disabling version check

Although the *version checking* feature does not collect any personal information, you might prefer to disable this feature, either one time or permanently. To disable it one time, use `--no-version-check` option when invoking the tool from a Percona product which supports it. Here is a simple example which shows running [pt-diskstats](#) tool from the *Percona Toolkit* with *version checking* turned off:

```
pt-diskstats --no-version-check
```

Disabling *version checking* permanently can be done by placing `no-version-check` option into the configuration file of a Percona product (see correspondent documentation for exact file name and syntax). For example, in case of *Percona Toolkit* [this can be done](#) in a global configuration file `/etc/percona-toolkit/percona-toolkit.conf`:

```
Disable Version Check for all tools:
no-version-check
```

In case of *Percona XtraBackup* this can be done [in its configuration file](#) in a similar way:

```
[xtrabackup]
no-version-check
```

## Frequently asked questions

## **Why is this functionality enabled by default?**

We believe having this functionality enabled improves security and performance of environments running Percona Software and it is good choice for majority of the users.

## **Why not rely on Operating System's built in functionality for software updates?**

In many environments the Operating Systems repositories may not carry the latest version of software and newer versions of software often installed manually, so not being covered by operating system wide check for updates.

## **Why do you send more information than just the version of software being run as a part of version check service?**

Compatibility problems can be caused by versions of various components in the environment, for example problematic versions of Perl, DBD or MySQL could cause operational problems with Percona Toolkit.

# Trademark policy

This [Trademark Policy](#) is to ensure that users of Percona-branded products or services know that what they receive has really been developed, approved, tested and maintained by Percona. Trademarks help to prevent confusion in the marketplace, by distinguishing one company's or person's products and services from another's.

Percona owns a number of marks, including but not limited to Percona, XtraDB, Percona XtraDB, XtraBackup, Percona XtraBackup, Percona Server, and Percona Live, plus the distinctive visual icons and logos associated with these marks. Both the unregistered and registered marks of Percona are protected.

Use of any Percona trademark in the name, URL, or other identifying characteristic of any product, service, website, or other use is not permitted without Percona's written permission with the following three limited exceptions.

*First*, you may use the appropriate Percona mark when making a nominative fair use reference to a bona fide Percona product.

*Second*, when Percona has released a product under a version of the GNU General Public License ("GPL"), you may use the appropriate Percona mark when distributing a verbatim copy of that product in accordance with the terms and conditions of the GPL.

*Third*, you may use the appropriate Percona mark to refer to a distribution of GPL-released Percona software that has been modified with minor changes for the sole purpose of allowing the software to operate on an operating system or hardware platform for which Percona has not yet released the software, provided that those third party changes do not affect the behavior, functionality, features, design or performance of the software. Users who acquire this Percona-branded software receive substantially exact implementations of the Percona software.

Percona reserves the right to revoke this authorization at any time in its sole discretion. For example, if Percona believes that your modification is beyond the scope of the limited license granted in this Policy or that your use of the Percona mark is detrimental to Percona, Percona will revoke this authorization. Upon revocation, you must immediately cease using the applicable Percona mark. If you do not immediately cease using the Percona mark upon revocation, Percona may take action to protect its rights and interests in the Percona mark. Percona does not grant any license to use any Percona mark for any other modified versions of Percona software; such use will require our prior written permission.

Neither trademark law nor any of the exceptions set forth in this Trademark Policy permit you to truncate, modify or otherwise use any Percona mark as part of your own brand. For example, if XYZ creates a modified version of the Percona Server, XYZ may not brand that modification as "XYZ Percona Server" or "Percona XYZ Server", even if that modification otherwise complies with the third exception noted above.

In all cases, you must comply with applicable law, the underlying license, and this Trademark Policy, as amended from time to time. For instance, any mention of Percona trademarks should include the full trademarked name, with proper spelling and capitalization, along with attribution of ownership to Percona LLC and/or its affiliates. For example, the full proper name for XtraBackup is Percona XtraBackup. However, it is acceptable to omit the word “Percona” for brevity on the second and subsequent uses, where such omission does not cause confusion.

In the event of doubt as to any of the conditions or exceptions outlined in this Trademark Policy, please contact [trademarks@percona.com](mailto:trademarks@percona.com) for assistance and we will do our very best to be helpful.

# Copyright and licensing information

## Documentation licensing

Percona XtraBackup documentation is (C)2009-2026 Percona LLC and/or its affiliates and is distributed under the [Creative Commons Attribution 4.0 International License](#).

# Index

- [About Percona XtraBackup](#)
- [Accelerate the backup process](#)
- [Analyze table statistics](#)
- [Work with AppArmor](#)
- [Backup example](#)
- [Backup process overview](#)
- [Backup strategy and backup policy](#)
- [Work with binary logs](#)
- [Binary tarball file names available](#)
- [Clean up](#)
- [Compile and install Percona XtraBackup 8.0 from source](#)
- [Configure xtrabackup](#)
- [Connection and privileges needed](#)
- [Copyright and licensing information](#)
- [Create a compressed backup](#)
- [Create a full backup](#)
- [Create a partial backup](#)
- [Create an incremental backup](#)
- [Create an individual partitions backup](#)
- [Decompress and prepare a backup](#)
- [Dictionary cache](#)
- [Differences between a logical backup and Percona XtraBackup](#)
- [Dump schema for all tables](#)
- [Encrypt backups](#)
- [Encrypted InnoDB tablespace backups](#)
- [Error Message: Found tables with row versions due to INSTANT ADD/DROP columns](#)
- [Exponential backoff](#)
- [FIFO data sink](#)
- [Files in the DEB package built for Percona XtraBackup 8.0](#)



- [Files in the RPM package built for Percona XtraBackup 8.0](#)
- [FLUSH TABLES WITH READ LOCK option](#)
- [Frequently asked questions](#)
- [Get help from Percona](#)
- [Glossary](#)
- [How Percona XtraBackup works](#)
- [How to create a new \(or repair a broken\) GTID-based replica](#)
- [How to set up a replica for replication in 6 simple steps with Percona XtraBackup](#)
- [Improved log statements](#)
- [Index of files created by Percona XtraBackup](#)
- [InnoDB redo log archiving](#)
- [Install Percona XtraBackup 8.0 from a binary tarball](#)
- [Install Percona XtraBackup 8.0 overview](#)
- [Install Percona XtraBackup on Debian-based systems](#)
- [Install Percona XtraBackup on Red Hat-based systems](#)
- [Install with DEB packages](#)
- [Install with RPM packages](#)
- [lock-ddl-per-table option improvements](#)
- [LRU dump backup](#)
- [Make backups in replication environments](#)
- [Next steps](#)
- [Overview](#)
- [Percona Product Download Instructions](#)
- [Percona Toolkit version checking](#)
- [Percona XtraBackup - Documentation](#)
- [Percona XtraBackup 8.0 release notes index](#)
- [Percona XtraBackup binaries overview](#)
- [Percona XtraBackup Version Compatibility and Server Version Checks](#)
- [Permissions needed](#)
- [Pin packages in Debian-based systems](#)

- [Point-in-time recovery](#)
- [Prepare a full backup](#)
- [Prepare a partial backup](#)
- [Prepare an incremental backup](#)
- [Prepare an individual partitions backup](#)
- [Restore a partial backup](#)
- [Restore full, incremental, and compressed backups](#)
- [Restore single tables between databases](#)
- [Restore the backup](#)
- [Restore the partition from the backup](#)
- [Run in a Docker container](#)
- [Work with SELinux](#)
- [Smart memory estimation](#)
- [Start a Docker container and take a backup](#)
- [Store backup history on the server](#)
- [Supported versions](#)
- [Take a streaming backup](#)
- [Take an incremental backup using page tracking](#)
- [Throttling backups](#)
- [Topic index](#)
- [Trademark policy](#)
- [Understand version numbers](#)
- [Uninstall Percona XtraBackup 8.0 on Debian-based systems](#)
- [Uninstall Percona XtraBackup 8.0 on Red Hat-based systems](#)
- [update-curl-utility](#)
- [Use the xbcloud binary with Amazon S3](#)
- [Use the xbcloud binary with an IAM instance profile](#)
- [Use the xbcloud binary with Google Cloud Storage](#)
- [Use the xbcloud binary with Microsoft Azure Cloud Storage](#)
- [Use the xbcloud binary with MinIO](#)

- [Use the xbcloud binary with Swift](#)
- [Verify backups with replication and pt-checksum](#)
- [The xbcloud binary overview](#)
- [The xbcloud command-line options](#)
- [The xbcrypt binary overview](#)
- [The xbcrypt command-line options](#)
- [The xbstream binary overview](#)
- [The xbstream command-line options](#)
- [The xtrabackup binary overview](#)
- [The xtrabackup command-line options](#)
- [xtrabackup exit codes](#)
- [xtrabackup implementation details](#)